



nVIDIA

CERTIFIED PROFESSIONAL AGENTIC AI



NCP-AAI-BETA

PRACTICE EXAM QUESTIONS

Cloud Certification Store

NVIDIA CERTIFIED PROFESSIONAL AGENTIC AI

NVIDIA Certified Professional Agentic - AI NCP-AAI-BETA (NVIDIA-NCP-AAI-0010)

© 2025 [Cloud Certification Store](https://cloudcertificationstore.com/b/8eYLD) All rights reserved.

NVIDIA® is a registered trademark of NVIDIA Corporation.

This practice set is an original work for educational use and is NOT endorsed by or affiliated with NVIDIA Corporation. “NVIDIA,” “NVIDIA Certified Professional,” “Agentic AI,” and related marks are trademarks of NVIDIA Corporation, used here for identification only.

DISCLAIMER

- **As of October 2025, the certification is a Beta exam, and can only be taken by invitation.**
- **This practice test is based on the NVIDIA study guide for the beta exam** sent by invitation to potential beta exam takers, plus exam materials from related certifications.
- It includes questions **compiled from various exam preparation platforms.**
- **Important: some answers were curated using generative AI with human review** Verify accuracy with official documentation before relying on this material.
- **Users are strongly encouraged to** double-check all content against **official documentation and trusted sources** before using it for exam preparation or making important decisions.
- **The creators of this material assume no responsibility** for any errors, inaccuracies, or outcomes, including exam results, based on the use of this content.
- **Some questions might be duplicated or close** to previous ones. This is done on purpose as a way to show possible scenarios and to re-inforce your learning.
- **Single-user licence only:**
 - Includes one unique Payhip Licence Key per purchase, along with a Product Key.
 - **Redistribution, resale, or public posting is prohibited.** We can trace any file to the purchaser, with the use of the purchased License Key, Product Key and a watermark at the top/left corner of each page containing the email of the purchaser.

This is a Preview Copy

with only 20 out of 520 questions

>> Get the Full Version at

<https://cloudcertificationstore.com/b/8eYLD>

PREVIEW ONLY

NVIDIA Certified Professional Agentic - AI

NCP-AAI-BETA (NVIDIA-NCP-AAI-0010)



Issued by [NVIDIA](https://cloudcertificationstore.com/b/8eYLD)

The NVIDIA Certified Professional Agentic AI certification (in Beta phase), is an intermediate practitioner adept at designing, evaluating, and deploying autonomous AI systems.

Responsibilities include construction of resilient, secure, and trustworthy agentic solutions such as customer/employee assistance, automated meeting companion summaries, and productivity tools (e.g., content generation, analytics).

The candidate must master Agentic Architecture Fundamentals — including reactive, deliberative, and hybrid systems —and excel in goal-oriented reasoning, chain-of-thought prompt engineering, and tool orchestration with robust error handling.



NVIDIA-Certified Professional: Agentic AI - Official Study Guide

Contents

Agent Architecture and Design: Exam Weight 15%

Agent Development: Exam Weight 15%

Evaluation & Tuning: Exam Weight 13%

Deployment and Scaling: Exam Weight 13%

Cognition, Planning & Memory: Weight 10%

Knowledge Integration & Data Handling: Exam Weight 10%

NVIDIA Platform Implementation: Exam Weight 7%

Run, Monitor & Maintain : Exam Weight 7%

Safety, Ethics & Compliance : Exam Weight 5%

Human-AI Interaction & Oversight : Exam Weight 5%

This study guide provides an overview of each topic covered on the NVIDIA Certified Professional-Agent AI certification exam, recommended training, and suggested reading to prepare for the exam.

Job Description

The Agentic AI professional is an intermediate practitioner adept at designing, evaluating, and deploying autonomous AI systems. Responsibilities include construction of resilient, secure, and trustworthy agentic solutions such as customer/employee assistance, automated meeting companion summaries, and productivity tools (e.g., content generation, analytics). The candidate must master Agentic Architecture Fundamentals—including reactive, deliberative, and hybrid systems—and excel in goal-oriented reasoning, chain-of-thought prompt engineering, and tool orchestration with robust error handling.

Expertise includes memory management (short- and long-term context), evaluation frameworks, deployment of containerized workflows, monitoring/logging, ethical safeguards (e.g., bias detection, privacy preservation, guardrails), troubleshooting hallucinations, tracing failures, and managing stateful orchestration for complex, multi-turn tasks. They combine generative AI and reasoning systems knowledge with MLOps workflow deployment and GPU-optimized operations to ensure scalable, ethical implementation.

Key Responsibilities

- End-to-End Agent Development Exposure
- Model and framework selection/integration
- Agent structure and tool creation
- Orchestration of agent workflows
- Assessment, evaluation, and iterative improvement

Recommended Qualifications and Experience

- 1-2 years in AI/ML roles
- Experience with production-level agentic AI projects (e.g., chatbots, workflow automation)

This is a Preview Copy

with only 20 out of 520 questions

>> Get the Full Version at

<https://cloudcertificationstore.com/b/8eYLD>

Certification Topics and References

Agent Architecture and Design: Exam Weight 15%

Foundational structuring and design of agentic systems, focusing on how agents interact, reason, and communicate within their environments.

- 1.1 Design user interfaces for intuitive human-agent interaction.
- 1.2 Implement reasoning and action frameworks (e.g., ReAct).
- 1.3 Configure agent-to-agent communication protocols for collaboration.
- 1.4 Manage short-term and long-term memory for context retention.
- 1.5 Orchestrate multi-agent workflows and coordination.
- 1.6 Apply logic trees, prompt chains, and stateful orchestration for multi-step reasoning.
- 1.7 Integrate knowledge graphs to enable relational reasoning.
- 1.8 Ensure adaptability and scalability of the agent's architecture.

Recommended NVIDIA Course and Suggested Readings

NVIDIA Course

- ☐ [Building Agentic AI Applications with LLMs](#)

Suggested Readings

- ☐ [Agentic AI in the Factory](#)
- ☐ [Building Autonomous AI with NVIDIA Agentic NeMo](#)
- ☐ [Three Building Blocks for Creating AI Virtual Assistants for Customer Service with an NVIDIA NIM Agent Blueprint](#)
- ☐ [Agentic AI: Towards Autonomous Artificial Intelligence Agents](#)
- ☐ [Catch Me If You Can: A Multi-Agent Framework for Financial Fraud Detection](#)
- ☐ [What Are Multi-Agent Systems?](#)

Agent Development: Exam Weight 15%

Practical building, integration, and enhancement of agents.

- 2.1 Engineer prompts and dynamic prompt chains for reliable performance.
- 2.2 Integrate generative and multimodal models (text, vision, audio).
- 2.3 Build and connect custom tools, APIs, and functions for external system interaction.
- 2.4 Implement error handling (retry logic, graceful failure recovery).
- 2.5 Develop dynamic conversation flows with real-time streaming and feedback mechanisms.
- 2.6 Evaluate and refine agent decision-making strategies.

Recommended NVIDIA Course and Suggested Readings

NVIDIA Course

- ☐ [Building RAG Agents with LLMs](#)

Suggested Readings

- ☐ [Optimization — NVIDIA Triton Inference Server](#)
- ☐ [NVIDIA Agent Intelligence Toolkit Overview — NVIDIA Agent Intelligence Toolkit \(1.1.0\)](#)
- ☐ [An Introduction to Large Language Models: Prompt Engineering and P-Tuning | NVIDIA Technical Blog](#)
- ☐ [Building Multimodal AI RAG with LlamaIndex, NVIDIA NIM, and Milvus](#)
- ☐ [Design Considerations of Advanced Agentic AI for Real-World Applications](#)
- ☐ [Transient fault handling - Azure Architecture Center | Microsoft Learn](#)
- ☐ [Circuit Breaker Pattern - Azure Architecture Center | Microsoft Learn](#)
- ☐ [Retry pattern - Azure Architecture Center | Microsoft Learn](#)

Evaluation & Tuning: Exam Weight 13%

Measuring, comparing, and optimizing agent performance.

- 3.1 Implement evaluation pipelines and task benchmarks to measure performance.
- 3.2 Compare agent performance across tasks and datasets.
- 3.3 Collect and integrate structured user feedback for iterative improvements.
- 3.4 Tune model parameters (e.g., accuracy, latency-efficiency trade-offs).
- 3.5 Analyze evaluation results to guide targeted optimization.

Recommended NVIDIA Course and Suggested Readings

NVIDIA Course

- ☐ [Building Agentic AI Applications with LLMs](#)

Suggested Readings

- ☐ [Powering the Next Generation of AI Agents](#)
- ☐ [NVIDIA Agent Intelligence Toolkit](#)
- [Overview](#) ☐ [NVIDIA Agent Intelligence Toolkit Tutorials](#) ☐ [NVIDIA Agent Intelligence Toolkit FAQ](#)
- ☐ [Launching the NVIDIA Agent Intelligence Toolkit API Server and User Interface — NVIDIA Agent Intelligence Toolkit \(1.1\)](#)
- ☐ [NVIDIA/NeMo-Agent-Toolkit - GitHub](#)
- ☐ [Agentic AI: The Next Big Thing in Artificial Intelligence](#)
- ☐ [Agentic AI: The Top 5 Challenges and How to Overcome Them](#)
- ☐ [AI Agents for Beginners – Production Patterns \(Microsoft\)](#)
- ☐ [Navigating the Challenges: 5 Common Pitfalls in Agentic AI Adoption](#)

Deployment and Scaling : Exam Weight 5%

Operationalizing and scaling agentic systems.

- 4.1 Deploy and orchestrate multi-agent systems at production scale.
- 4.2 Apply MLOps practices for CI/CD workflows, monitoring, and governance.
- 4.3 Profile performance and reliability under distributed system loads.
- 4.4 Scale deployments using containerization (Docker, Kubernetes) with load balancing.
- 4.5 Optimize deployment costs while ensuring high availability.

Recommended NVIDIA Courses and Suggested Readings

NVIDIA Courses

- ☐ [Building Agentic AI Applications with LLMs](#)
- ☐ [Building RAG Agents with LLMs](#)

Suggested Readings

- ☐ [Agentic AI in the Factory \(NVIDIA White Paper\)](#)
- ☐ [TensorRT-LLM \(NVIDIA GitHub\)](#)
- ☐ [Measure and Improve AI Workload Performance with NVIDIA DGX Cloud Benchmarking](#)
- ☐ [Kubernetes Glossary \(NVIDIA\)](#)
- ☐ [NVIDIA Nsight Systems](#)
- ☐ [Kube Prometheus for GPU Telemetry \(NVIDIA Docs\)](#)
- ☐ [Scaling LLMs with NVIDIA Triton and TensorRT-LLM Using Kubernetes](#)
- ☐ [TensorRT-LLM Performance Analysis Documentation](#)

Cognition, Planning & Memory: Exam Weight 10%

Core cognitive processes underlying intelligent agent behavior, including reasoning strategies, decision-making, and memory management.

- 5.1 Implement memory mechanisms for short- and long-term context retention.
- 5.2 Apply reasoning frameworks (chain-of-thought, task decomposition).
- 5.3 Engineer planning strategies for sequential and multi-step decision-making.
- 5.4 Manage stateful orchestration to coordinate complex tasks and knowledge retention.
- 5.5 Adapt reasoning strategies based on prior experiences and feedback.

Recommended NVIDIA Course and Suggested Readings

NVIDIA Course

- ☐ [Building Agentic AI Applications with LLMs](#)

Suggested Readings

- ☐ [NVIDIA NeMo](#)
- ☐ [Large Language Models are in Context Learners \(arXiv:2310.10501\)](#)
- ☐ [NeMo RL Documentation](#)
- ☐ [Jamba 1.5 LLMS Leverage Hybrid Architecture](#)
- ☐ [Understanding the Planning of LLM Agents: A Survey \(HTML version\)](#)
- ☐ [AI Agent Memory | IBM](#)
- ☐ [MCP Agent Memory Types, Management, Implementation](#)
- ☐ [Understanding the Planning of LLM Agents: A Survey \(arXiv:2402.02716\)](#)

Knowledge Integration & Data Handling: Exam Weight 10%

Integration of external knowledge and the management of diverse data types.

- 6.1 Implement retrieval pipelines (RAG, embedded search, hybrid approaches).
- 6.2 Configure and optimize vector databases for fast retrieval.
- 6.3 Build ETL pipelines to integrate enterprise or client data sources.
- 6.4 Conduct data quality checks, augmentation, and preprocessing.
- 6.5 Enable real-time access and reasoning over structured/unstructured knowledge.

Recommended NVIDIA Course and Suggested Readings

NVIDIA Course

- ☐ [Building RAG Agents with LLMs](#)

Suggested Readings

- ☐ [How to Make Your LLM More Accurate with RAG & Fine-Tuning | Towards Data Science](#)

NVIDIA Platform Implementation: Exam Weight 7%

Leveraging NVIDIA's AI hardware and software platforms for agentic systems.

- 7.1 Integrate NeMo Guardrails for compliance and safety enforcement.
- 7.2 Deploy NVIDIA NIM microservices for high-performance inference.

7.3 Optimize workflows with NVIDIA Agent Intelligence Toolkit.

7.4 Leverage TensorRT-LLM and Triton Inference Server for latency reduction.

7.5 Manage and optimize multi-modal input pipelines on NVIDIA hardware.

Recommended NVIDIA Course and Suggested Readings

NVIDIA Course

☐ [Building RAG Agents with LLMs](#)

Suggested Readings

☐ [Best Practices — NVIDIA TensorRT](#)

[Documentation](#) ☐ [Batchers — NVIDIA Triton Inference Server](#)

☐ [Triton Inference Server Backend - NVIDIA](#)

[Documentation](#) ☐ [NeMo Guardrails | NVIDIA Developer](#)

☐ [NVIDIA/NeMo-Guardrails - GitHub](#)

☐ [Performance Tuning Guide — NVIDIA NeMo Framework User](#)

[Guide](#) ☐ [Best Practices — NVIDIA NeMo Framework User Guide](#)

☐ [Optimization — NVIDIA Triton Inference](#)

[Server](#) ☐ [NVIDIA NeMo Agent Toolkit](#)

☐ [NVIDIA Agent Intelligence](#)

[toolkit](#) ☐ [NVIDIA/AIQToolkit](#)

☐ [Mastering LLM Techniques: Inference](#)

[Optimization](#) ☐ [Deploy Inference Workloads with NVIDIA NIM](#)

☐ [How to Use the NVIDIA Llama Nemotron API for Advanced AI](#)

[Agents](#) ☐ [How to deploy Llama-3.1-Nemotron-70B-Instruct on a Virtual](#)

...

☐ [AI Agents Blueprint: Designing Foundation Models and Agents for the Next Wave of](#)

[AI](#) ☐ [Improve AI Code Generation Using NVIDIA Agent Intelligence Toolkit](#)

- ☐ [NVIDIA/NeMo: A scalable generative AI framework built for ... - GitHub](#)

Run, Monitor & Maintain : Exam Weight 7%

Ongoing operation, monitoring, and maintenance of agentic systems post-deployment.

- 8.1 Define monitoring dashboards and reliability metrics.
- 8.2 Track logs, errors, and anomalies for root cause diagnosis.
- 8.3 Continuously benchmark deployed agents against prior versions.
- 8.4 Implement automated tuning, retraining, and versioning in production.
- 8.5 Ensure continuous uptime, transparency, and trust in live deployments.

Suggested Readings

- ☐ [What is AI Agent Evaluation?](#)
- ☐ [Log, Trace, and Monitor](#)
- ☐ [Time-Weighted Retriever](#)
- ☐ [Troubleshooting](#)
- ☐ [LangChain Tracing Concepts](#)
- ☐ [LangChain Structured Outputs Concepts](#)
- ☐ [Smith LangChain Model Evaluation: Rate Limiting](#)
- ☐ [A Guide to Monitoring Machine Learning Models in Production](#)
- ☐ [Monitoring Machine Learning Models in Production: How to Track Data Quality and Integrity](#)

Safety, Ethics & Compliance : Exam Weight 5%

Principles and practices that ensure agentic systems operate responsibly, uphold ethical

standards, and comply with legal and regulatory frameworks.

9.1 Design and enforce system security and audit trails.

9.2 Integrate compliance guardrails (privacy, enterprise policy).

9.3 Mitigate bias and toxicity in outputs.

9.4 Deploy layered safety frameworks (filters, escalation protocols).

9.5 Ensure compliance with licensing and regulatory standards.

Suggested Readings

- ☐ [Building safer LLM apps with LangChain templates and NVIDIA NeMo Guardrails](#)
- ☐ [NVIDIA NeMo Guardrails](#)
- ☐ [Artificial Intelligence and Machine Learning in Software as a Medical Device](#)
- ☐ [Proposal for a Regulation laying down harmonised rules on artificial intelligence \(Artificial Intelligence Act\)](#)
- ☐ [Ethically Aligned Design: A Vision for Prioritizing Human Well-being with Autonomous and Intelligent Systems](#)
- ☐ [NVIDIA/NeMo-Guardrails](#)
- ☐ [Securing generative AI deployments with NVIDIA NIM and NVIDIA NeMo Guardrails](#)
- ☐ [Metrics for Agentic AI](#)
- ☐ [AI for Regulatory Compliance](#)
- ☐ [Responsible AI Revisited](#)

Human-AI Interaction & Oversight : Exam Weight 5%

The design and implementation of systems that facilitate effective human oversight and interaction with agents.

10.1 Build intuitive UIs with user-in-the-loop interaction.

10.2 Design structured feedback loops that guide iterative agent improvements.

10.3 Implement transparency mechanisms (explainable reasoning, decision traceability).

10.4 Enable human oversight and intervention for accountability and trust.

Recommended NVIDIA Course and Suggested Readings

NVIDIA Course

☐ [Building Conversational AI Applications](#)

Suggested Readings

☐ [NVIDIA Agent Intelligence](#)

[Toolkit](#) ☐ [NVIDIA Data Flywheel](#)

[Glossary](#)

☐ [AI Agents With Human-in-the-Loop](#)

[\(Medium\)](#) ☐ [Human-in-the-Loop AI \(HologicAI\)](#)

☐ [Human-in-the-Loop Agentic AI Systems](#)

[\(OneReach.ai\)](#) ☐ [Aporia: AI Guardrails](#)

☐ [Improve AI Code Generation Using NVIDIA Agent Intelligence](#)

[Toolkit](#) ☐ [Chain-of-Thought \(CoT\) Prompting | Codecademy](#)

© 2025 NVIDIA Corporation and affiliates. All rights reserved. NVIDIA, the NVIDIA logo, NeMo, NGC, and Triton are trademarks and/or registered trademarks of NVIDIA Corporation and affiliates in the

U.S. and other countries. Other company and product names may be trademarks of the respective owners with which they are associated.

Practice Questions

Question 1

You are designing a customer-support agent that must reason about a user's issue, decide whether to call a troubleshooting API, and then summarize actions taken back to the user. Which design choice most directly supports intertwined reasoning and tool use for this workflow?

- A. Use a pure retrieval pipeline and rely on embeddings for all actions.
- B. Apply the ReAct pattern to interleave chain-of-thought with function calls.
- C. Configure a rule-based system that only executes if/then logic trees.
- D. Depend on a long system prompt that lists all possible tools and their docs.

✓ **Correct answer: B. Apply the ReAct pattern to interleave chain-of-thought with function calls.**

✚ ReAct explicitly couples reasoning traces with actions, enabling the agent to decide, call tools, and reflect in multi-step tasks.

Incorrect answers:

✗ **A. Use a pure retrieval pipeline and rely on embeddings for all actions.** – Retrieval augments knowledge but does not decide or invoke tools by itself. It supplies context, not

procedural control flow. Without an action framework, the agent cannot determine when to escalate to an API. This limits the system to passive lookup rather than active problem solving.

✗ C. Configure a rule-based system that only executes if/then logic trees. – Rules can encode simple flows but struggle with open-ended reasoning and uncertain inputs. They do not inherently blend free-form deliberation with dynamic tool selection. Pure rules are brittle as tasks and tools evolve. A hybrid ReAct approach provides the needed flexibility.

✗ D. Depend on a long system prompt that lists all possible tools and their docs. – A long prompt is not a control framework; it's static guidance. Without an explicit reasoning-and-action loop, the model may hallucinate or ignore tools. Tool discovery alone does not guarantee correct sequencing or error recovery. You still need a method like ReAct to structure decisions and calls.

Question 2

Your design requires multiple specialized agents: one plans, another executes API calls, and a third audits results for safety. What is the best architectural step to ensure they coordinate effectively?

- A. Share one global prompt for all agents to read and write.
- B. Enable agent-to-agent protocols with a task router and shared state.

C. Give each agent the same memory store to maximize overlap.

D. Force all agents to respond sequentially on a fixed timer.

✅ **Correct answer: B. Enable agent-to-agent protocols with a task router and shared state.**

🔗 The guide emphasizes agent-to-agent communication and orchestration for collaborative workflows.

Incorrect answers:

❌ **A. Share one global prompt for all agents to read and write.** – A single shared prompt invites race conditions and prompt collisions. There's no routing or role clarity, so responsibilities blur. It also complicates provenance and auditability of decisions. Purposeful protocols and state management are superior.

❌ **C. Give each agent the same memory store to maximize overlap.** – Memory sharing without partitioning causes context bloat and leakage of irrelevant details. Different roles need different slices of context. Unscoped memory makes debugging and oversight harder. The better approach is coordinated exchange with controlled state.


❌ **D. Force all agents to respond sequentially on a fixed timer.** – Timed sequencing does not equal coordination. It ignores task dependencies and can delay critical actions. It also wastes resources when agents are idle. Coordination should be event-driven with explicit handoffs.

Question 3


A multiturn analyst agent must remember project details for hours and also recall recent turns precisely. What memory approach fits best?

- A. Only short-term working memory with a small token window.
- B. Only a global long-term vector store for all conversations.
- C. Combine short-term context with long-term memory retrieval.
- D. Disable memory and ask the user to repeat details each turn.

 **Correct answer: C. Combine short-term context with long-term memory retrieval.**

 The blueprint calls for managing both short- and long-term memory to retain context over time.

Incorrect answers:

 **A. Only short-term working memory with a small token window.** – This will truncate important details as the dialogue grows. The agent soon forgets earlier constraints or decisions. Token limits force eviction of relevant facts. Long-term recall is required for consistency.

✗ B. Only a global long-term vector store for all conversations. – A single LTM without a fresh working set hinders step-by-step reasoning. The agent must juggle current intent while selectively recalling history. Pure retrieval can be noisy and slow without a focused short-term cache. Pairing both layers yields stable performance.

✗ D. Disable memory and ask the user to repeat details each turn. – This degrades UX and creates unnecessary friction. It also increases token use and response latency. The system appears incompetent at following through. Proper memory mechanisms are core to agent design.

Question 4

A conversational agent occasionally fails when dependent APIs time out. Which engineering improvement most directly increases reliability?

- A. Increase the model temperature so it tries more variations.
- B. Add retry logic with exponential backoff and circuit breaker guards.
- C. Switch to a larger LLM for better comprehension.
- D. Disable tool calls when the API is slow.

✓ Correct answer: B. Add retry logic with exponential backoff and circuit breaker guards.

✚ The exam guide highlights robust error handling with retry and failure recovery patterns.

Incorrect answers:

✗ **A. Increase the model temperature so it tries more variations.** – Temperature controls sampling diversity, not network resiliency. It won't fix transient faults or flaky endpoints. In fact, more variability can reduce determinism in error handling. Reliability needs infrastructure patterns, not randomization.

✗ **C. Switch to a larger LLM for better comprehension.** – Model size doesn't repair HTTP timeouts or upstream failures. Bigger models may even increase latency and costs. Availability problems call for defensive patterns at the integration layer. Reliability patterns are the correct lever.

✗ **D. Disable tool calls when the API is slow.** – Blanket disabling breaks core functionality and user expectations. It avoids the symptom rather than addressing the cause. Intelligent fallbacks and circuit breakers are more nuanced and recoverable. Graceful degradation maintains utility.

Question 5

You must add real-time vision input (screenshots) to an existing chat agent and stream interim tokens as it responds. What is the primary development focus?

- A. Replace the agent with a rules engine that parses images.
- B. Integrate multimodal models and implement streaming output.
- C. Use retrieval only; images can be embedded as text.
- D. Turn off streaming to simplify session state.

✓ **Correct answer: B. Integrate multimodal models and implement streaming output.**

📌 The guide calls out multimodal integration and dynamic conversation flows with streaming.

Incorrect answers:

✗ **A. Replace the agent with a rules engine that parses images.** – Rules engines are not suited for perception tasks and lack generalization. Vision understanding requires learned multimodal capabilities. Replacing the agent discards reasoning and dialog strengths. It undercuts the overall design goals.

✗ **C. Use retrieval only; images can be embedded as text.** – Images are not directly convertible to semantic text without a vision model. Retrieval augments knowledge rather than interpreting pixels. Without multimodal inference, the agent can't "see" the screenshot. Streaming improves responsiveness but still needs proper inputs.

✗ **D. Turn off streaming to simplify session state.** – Disabling streaming harms UX and perceived speed. Streaming is a recommended mechanism for responsive conversations.

Complexity should be managed, not removed, by good orchestration. The priority is adding vision capabilities with proper streaming.

Question 6

A platform team wants to serve high-performance inference via NVIDIA's standardized microservices with minimal custom glue. What should they deploy?

- A. A bespoke Flask app that wraps raw CUDA kernels.
- B. NVIDIA NIM microservices to expose consistent inference endpoints.
- C. A single-node Triton server without any model repository.
- D. A GPUless Kubernetes cluster running CPU-only pods.

✅ **Correct answer: B. NVIDIA NIM microservices to expose consistent inference endpoints.**

✚ The guide lists deploying NVIDIA NIM microservices for high-performance inference as a core platform skill.

Incorrect answers:

❌ **A. A bespoke Flask app that wraps raw CUDA kernels.** – This is fragile and reinvents platform plumbing. It lacks standardized lifecycle and observability. Maintenance and portability suffer over time. NIM provides supported, optimized endpoints out of the box.

✗ C. A single-node Triton server without any model repository. – Triton benefits from a structured model repository and orchestration around it. A single node is a scaling bottleneck. Without proper config, you lose batching, versioning, and health checks. NIM streamlines this in a service form.

✗ D. A GPUless Kubernetes cluster running CPU-only pods. – Many agentic workloads rely on GPUs for latency and throughput. CPU-only nodes undercut performance targets. Platform integration on NVIDIA hardware is part of the exam scope. NIM expects a suitable GPU environment.

Question 7

To raise throughput for a multilingual agent serving many small requests, which Triton feature is most appropriate?

- A. Model ensembling to chain multiple backends.
- B. Per-request temperature scaling at the client.
- C. Dynamic batching via Triton's batchers.
- D. Disabling health checks during peak hours.

✓ Correct answer: C. Dynamic batching via Triton's batchers.

🔥 Triton's batchers improve utilization and reduce per-request overhead for higher throughput.

Incorrect answers:

- ❌ **A. Model ensembling to chain multiple backends.** – Ensembling coordinates multi-stage inference, not primarily throughput for homogeneous traffic. It may add latency with extra stages. The problem here is per-request overhead, which batching targets directly. Ensembling solves a different class of problem.
- ❌ **B. Per-request temperature scaling at the client.** – Temperature changes output diversity, not system-level throughput. It does not affect GPU scheduling or kernel efficiency. Throughput is bounded by server batching and concurrency. Client tweaks won't fix server utilization.
- ❌ **D. Disabling health checks during peak hours.** – Health checks are crucial to reliability and quick failure detection. Turning them off risks cascading failures. It neither improves batching nor memory reuse. Proper batching is the safe performance lever.
-

Question 8

A team needs a rigorous way to compare two decision-making strategies for the same agent across a benchmark suite. What should they implement first?

- A. A/B test prompts live in production without safeguards.
- B. A structured evaluation pipeline with task benchmarks and metrics.

- C. Manual ad-hoc spot checks by senior developers.
- D. A weekly stakeholder survey about perceived quality.

✅ **Correct answer: B. A structured evaluation pipeline with task benchmarks and metrics.**

🔗 The guide stresses evaluation pipelines and benchmarks to measure and compare performance.

Incorrect answers:

❌ **A. A/B test prompts live in production without safeguards.** – Live tests can be helpful but require guardrails, traffic allocation, and rollback. Without a prior offline eval, you risk poor user experiences. Benchmarks give safer, reproducible comparisons. It's better to iterate before exposing users.

❌ **C. Manual ad-hoc spot checks by senior developers.** – Human inspection can miss systematic errors and won't scale. It introduces bias and inconsistency. Results are hard to reproduce and compare. Formal pipelines prevent these pitfalls.

❌ **D. A weekly stakeholder survey about perceived quality.** – Surveys capture subjective sentiment, not ground-truth task success. They lag and lack task coverage. Optimization needs quantifiable, comparable metrics. Benchmarks provide that structure.

Question 9

Your production agents occasionally produce incorrect final answers due to tool chain failures. What monitoring addition best helps root-cause these errors?

- A. Disable logging to reduce storage costs.
- B. Add tracing for tool calls and structured outputs.
- C. Lower max tokens so generations are shorter.
- D. Use a larger context window only.

✓ **Correct answer: B. Add tracing for tool calls and structured outputs.**

📌 The “Run, Monitor & Maintain” scope highlights logs, errors, anomalies, and tracing to diagnose root causes.

Incorrect answers:

✗ **A. Disable logging to reduce storage costs.** – Turning off logs blindly removes critical forensic data. Cost savings should not trump reliability. Without logs, you can’t explain failures or improve. Monitoring is a core operational requirement.

✗ **C. Lower max tokens so generations are shorter.** – Shorter outputs might reduce latency but won’t reveal why tools fail. You need visibility into the calls, parameters, and responses. Length limits change behavior, not observability. Tracing is the correct instrument.

✗ D. Use a larger context window only. – More context can help reasoning but won't diagnose integration faults. When APIs fail, you must trace the path. Context size is orthogonal to tooling observability. Logging and tracing directly target the problem.

Question 10

Which action most directly reduces harmful or policy-violating outputs from a generative agent in a regulated domain?

- A. Rely on model pretraining to have removed all unsafe data.
- B. Add layered safety guardrails with filters and escalation protocols.
- C. Block all user inputs that contain sensitive words.
- D. Let users flag content and fix issues later.

✓ Correct answer: B. Add layered safety guardrails with filters and escalation protocols.

✚ The safety domain emphasizes layered guardrails, privacy, bias mitigation, and escalation.

Incorrect answers:

- ❌ **A. Rely on model pretraining to have removed all unsafe data.** – Pretraining cannot guarantee perfect safety at inference time. Context, prompts, and tools create new risks. Relying solely on pretraining ignores runtime controls. Guardrails are still required.
- ❌ **C. Block all user inputs that contain sensitive words.** – Naïve keyword blocks over-censor legitimate use cases and under-catch subtle risks. Safety requires context-aware filtering and escalation. Excessive blocking harms usability. Layered controls are more precise and adaptable.
- ❌ **D. Let users flag content and fix issues later.** – Reactive flagging is valuable but insufficient. Harm occurs before remediation. Proactive filters and policies reduce risk at generation time. Oversight should be designed in, not retrofitted.
-

This is a Preview Copy

with only 20 out of 520 questions

>> Get the Full Version at

<https://cloudcertificationstore.com/b/8eYLD>

Question 11

In a medical-support workflow, the agent may suggest actions but a clinician must approve them. What pattern best fits this requirement?

- A. Fully autonomous execution with no human intervention.
- B. Human-in-the-loop UI that supports review and override.
- C. Disable suggestions and only allow direct commands.
- D. Let the agent execute and notify the clinician afterward.

✅ **Correct answer: B. Human-in-the-loop UI that supports review and override.**

📌 The exam includes building intuitive UIs with user-in-the-loop and enabling human oversight.

Incorrect answers:

❌ **A. Fully autonomous execution with no human intervention.** – This conflicts with oversight and accountability needs. In regulated settings, autonomy without approval is risky. It fails to provide guardrails. User-in-the-loop is explicitly recommended.

❌ **C. Disable suggestions and only allow direct commands.** – Removing agent initiative wastes its reasoning capabilities. The clinician loses decision support that can improve outcomes. Better is to propose actions and let humans approve. This balances utility and control.

✗ D. Let the agent execute and notify the clinician afterward. – Post-hoc notification undermines safety and compliance. Approvals must occur before action. This pattern invites irrecoverable errors. Human oversight must be upstream.

Question 12

You need to augment an agent's answers with fresh internal documents and database facts. Which integration approach aligns best with the guide?

- A. Fine-tune only; never retrieve at inference time.
- B. Implement a RAG pipeline with embedded search and hybrid retrieval.
- C. Disable external knowledge to avoid latency.
- D. Mirror the entire data warehouse into the system prompt.

✓ Correct answer: B. Implement a RAG pipeline with embedded search and hybrid retrieval.

✚ The knowledge section explicitly calls out RAG, embedded search, and hybrid approaches.

Incorrect answers:

✗ A. Fine-tune only; never retrieve at inference time. – Fine-tuning fixes static knowledge but won't track rapid changes. Retrieval provides freshness without

retraining. A balanced approach yields accuracy and agility. The guide leans on retrieval pipelines for live knowledge.

✗ **C. Disable external knowledge to avoid latency.** – Latency can be optimized, but accuracy collapses without relevant data. A blind model risks hallucinations and outdated facts. Real-time access is a stated goal. Proper retrieval is essential.

✗ **D. Mirror the entire data warehouse into the system prompt.** – Prompts have token limits and cost implications. Unfiltered dumps introduce noise and privacy risks. Retrieval selectively injects what's needed per query. This is the scalable pattern.

Question 13

Your RAG system's latency spikes at the vector database stage. Where should you focus first?

- A. Increase temperature for more creative answers.
- B. Optimize vector DB configuration for fast retrieval.
- C. Add larger prompts to pre-explain the schema.
- D. Disable filtering and return the top 100 chunks.

✓ **Correct answer: B. Optimize vector DB configuration for fast retrieval.**

✦ The guide calls out configuring and optimizing vector databases for retrieval performance.

Incorrect answers:

✗ **A. Increase temperature for more creative answers.** – Creativity doesn't fix retrieval bottlenecks. The latency occurs before the model even generates. Tuning generation parameters won't accelerate ANN searches. Address the database layer directly.

✗ **C. Add larger prompts to pre-explain the schema.** – Longer prompts add generation cost and may not help retrieval. The issue is in the search stage, not instruction clarity. Schema guidance is useful but orthogonal here. Focus on indexes, vectors, and filters.

✗ **D. Disable filtering and return the top 100 chunks.** – Over-retrieval increases token usage and harms answer quality. It can slow the model with excessive context. Smart filtering is key to relevance and speed. Proper DB tuning beats brute force.

Question 14

An enterprise wants to unify CSV exports, CRM APIs, and PDF manuals into one agent's knowledge layer. What is the correct first move?

- A. Build ETL pipelines and perform data quality checks/augmentation.
- B. Paste all documents into the system prompt.

C. Ask users to upload relevant snippets manually per question.

D. Skip preprocessing to preserve “raw truth.”

✅ **Correct answer: A. Build ETL pipelines and perform data quality checks/augmentation.**

📌 The guide specifically includes ETL, data quality, augmentation, and preprocessing.

Incorrect answers:

❌ **B. Paste all documents into the system prompt.** – Prompts cannot hold large, evolving corpora. This causes token bloat and inconsistency. It’s unmaintainable at enterprise scale. ETL and retrieval are the scalable path.

❌ **C. Ask users to upload relevant snippets manually per question.** – This shifts burden onto users and destroys UX. It also leads to inconsistent coverage. Automation via pipelines ensures completeness and freshness. It’s central to enterprise readiness.

❌ **D. Skip preprocessing to preserve “raw truth.”** – Raw data often contains duplicates, errors, and noise. Quality checks improve retrieval precision and trust. Preprocessing is a recommended best practice in the guide. Skipping it undermines reliability.

Question 15

You must scale an agentic service to thousands of concurrent users with high availability.

Which deployment step maps best to the guide?

- A. Run the service on a single GPU VM and vertically scale.
- B. Containerize and orchestrate with Kubernetes and load balancing.
- C. Use cron jobs to restart the service hourly.
- D. Disable autoscaling to avoid unpredictable costs.

✅ **Correct answer: B. Containerize and orchestrate with Kubernetes and load balancing.**

📌 “Deployment & Scaling” emphasizes containerization, Kubernetes, and load balancing for scale.

Incorrect answers:

❌ **A. Run the service on a single GPU VM and vertically scale.** – Vertical scaling hits hardware limits and single-point-of-failure risks. It weakens availability guarantees. Modern patterns encourage horizontal scaling. Kubernetes provides resilience and elasticity.

❌ **C. Use cron jobs to restart the service hourly.** – Restarts don’t equal scaling or reliability. They mask issues and disrupt users. Proper orchestration handles health, rollout, and balancing. Cron jobs are not a deployment strategy.

✗ D. Disable autoscaling to avoid unpredictable costs. – Predictability is important, but disabling autoscale harms availability during spikes. Smarter cost controls are better than rigidity. The guide stresses performance and HA alongside cost optimization. Balance, not blanket disabling, is key.

Question 16

A planning agent must decompose ambiguous goals into ordered steps and adapt mid-plan as new information arrives. Which technique is most appropriate?

- A. Single-turn prompts with no intermediate reasoning.
- B. Task decomposition with chain-of-thought and adaptive planning.
- C. Random sampling until a working plan emerges.
- D. Fixed if/then decision trees only.

✓ Correct answer: B. Task decomposition with chain-of-thought and adaptive planning.

📌 The cognition section lists chain-of-thought, task decomposition, and planning strategies.

Incorrect answers:

✗ **A. Single-turn prompts with no intermediate reasoning.** – Complex tasks need intermediate steps and reflection. Single-shot attempts often fail to capture dependencies. They underutilize the agent’s planning capabilities. Multi-step reasoning is recommended.

✗ **C. Random sampling until a working plan emerges.** – Randomness is inefficient and non-deterministic. It undermines repeatability and auditing. Planning should be principled and traceable. The guide champions structured strategies.

✗ **D. Fixed if/then decision trees only.** – Static trees are brittle in changing environments. They cannot flex to novel constraints. Adaptive planning learns from feedback and updates steps. This is central to agentic cognition.

Question 17

Your multi-turn research assistant forgets earlier constraints halfway through a long workflow. What engineering improvement addresses this most directly?

- A. Increase temperature to make outputs more “creative.”
- B. Implement stateful orchestration to persist and reuse key state.
- C. Force the user to restate all constraints each time.
- D. Use a larger model with more parameters.

✓ **Correct answer: B. Implement stateful orchestration to persist and reuse key state.**

📌 The guide calls out stateful orchestration for coordinating complex tasks and knowledge retention.

Incorrect answers:

✗ **A. Increase temperature to make outputs more “creative.”** – Creativity does not fix memory loss. Higher temperature can worsen consistency. The issue is state management. Orchestration, not sampling, is the remedy.

✗ **C. Force the user to restate all constraints each time.** – This degrades UX and shifts burden to the user. It does not scale for long tasks. Agents should handle memory responsibly. The guide emphasizes proper memory and state strategies.

✗ **D. Use a larger model with more parameters.** – Size alone doesn’t ensure persistence of task constraints. Without orchestration, context still drifts. Larger models may increase cost and latency. Proper state handling is the direct fix.

Question 18

To cut end-to-end latency for a coding-assistant agent, your MLOps team proposes an inference stack upgrade. Which pairing aligns with the guide?

- A. CPU inference with no batching and synchronous I/O.
- B. Triton Inference Server plus TensorRT-LLM optimizations.
- C. A single Python worker using naive fp32 matmuls.
- D. WebSockets only; model stack unchanged.

✅ **Correct answer: B. Triton Inference Server plus TensorRT-LLM optimizations.**

💡 The platform section highlights leveraging Triton and TensorRT-LLM for latency reduction.

Incorrect answers:

❌ **A. CPU inference with no batching and synchronous I/O.** – This combination increases latency and underutilizes hardware. Batching and GPU acceleration are proven levers. The guide focuses on NVIDIA-optimized stacks for performance. CPU-only is misaligned.

❌ **C. A single Python worker using naive fp32 matmuls.** – Naive implementations ignore kernel-level optimizations and scheduling. You'll hit throughput and cost walls quickly. Modern inference servers and kernels are essential. The recommended tools address these gaps.

❌ **D. WebSockets only; model stack unchanged.** – Transport changes alone don't transform inference latency. The bottleneck is computation and scheduling. Without

server and kernel optimization, you see little improvement. Triton and TensorRT-LLM are the direct levers.

Question 19

You're standardizing operational visibility across agent services. Which practice best matches the guide's maintenance focus?

- A. Define dashboards, reliability metrics, and benchmark against prior versions.
- B. Rotate logs weekly to save space and skip error tracking.
- C. Disable versioning to minimize storage usage.
- D. Only monitor CPU and memory, not application behavior.

✓ **Correct answer: A. Define dashboards, reliability metrics, and benchmark against prior versions.**

✚ The “Run, Monitor & Maintain” section lists dashboards, metrics, and continuous benchmarking.

Incorrect answers:

✗ **B. Rotate logs weekly to save space and skip error tracking.** – Space management is fine, but skipping error tracking is not. You lose insight into regressions and outages.

Monitoring includes logs and anomalies for root cause analysis. This is explicitly emphasized.

✗ C. Disable versioning to minimize storage usage. – Versioning is vital for rollback and comparisons. Without it, you cannot benchmark or audit changes. Storage costs should be balanced, not eliminated at the expense of control. The guide highlights versioning in production.

✗ D. Only monitor CPU and memory, not application behavior. – System metrics alone miss application-level failures. Tool errors, retries, and quality metrics matter in agentic systems. Holistic monitoring is required. The guide is explicit on this breadth.

Question 20

Legal asks how your agent complies with corporate policy, privacy rules, and external regulations. What should you prioritize?

- A. Rely on team culture; formal controls are unnecessary.
- B. Integrate compliance guardrails and ensure audit trails across the stack.
- C. Allow unfiltered tool outputs and delete logs nightly.
- D. Disable all data retention to avoid audit exposure.

✓ Correct answer: B. Integrate compliance guardrails and ensure audit trails across the stack.

✚ The safety section covers security, audit trails, compliance guardrails, and regulatory alignment.

Incorrect answers:

✗ **A. Rely on team culture; formal controls are unnecessary.** – Culture helps, but compliance requires enforceable mechanisms. Auditors expect technical controls and evidence. Informality won't satisfy legal obligations. Guardrails and trails are mandatory.

✗ **C. Allow unfiltered tool outputs and delete logs nightly.** – Unfiltered outputs risk policy violations. Deleting logs destroys accountability and traceability. Both moves increase regulatory exposure. The guide advocates layered safety and auditing.

✗ **D. Disable all data retention to avoid audit exposure.** – Zero retention prevents demonstrating compliance and investigating issues. Regulations often require specific retention policies. Thoughtful retention with access controls is better. The guide stresses transparent, trustworthy deployments.

(END OF PREVIEW QUESTIONS)

This is a Preview Copy

with only 20 out of 520 questions

>> Get the Full Version at



<https://cloudcertificationstore.com/b/8eYLD>

Final Review Checklist & Exam Readiness Scorecard

How to Use the Final Review Checklist

This section is meant to **validate your hands-on skills and theoretical readiness** across all exam topics.

Step-by-step:

1. **Print it or load it in a note-taking app** (Notion, Google Docs, OneNote, etc.).
2. Go through each checkbox:
 -  Check it if you **fully understand and can implement** the topic without looking up documentation.
 -  Leave it unchecked if you feel unsure or haven't practiced the task.
3. Prioritize unchecked topics by reviewing:
 - Check the official documentation
 - Practice exams
 - Hands-on labs
4. For each **unchecked item**, write a short action plan or resource link next to it.

How to Use the Exam Readiness Scorecard



This part helps you **self-assess your confidence level** and **focus your revision time** wisely.

Instructions:

1. For each domain (e.g., "Hybrid connectivity and routing"), **rate yourself** from 1 to 5:
 - **1** = No understanding or hands-on practice
 - **3** = Moderate familiarity, but need review
 - **5** = Mastered topic and can apply it in real-world use
2. Add **Notes / Action Items** to explain:
 - Why you scored yourself low
 - What resources you'll use to improve (YouTube, whitepapers, exam guides)
 - Practice test scores if relevant
3. Reassess **2–3 days before your exam**, and compare scores to measure improvement.

Bonus Tips

- Do **timed mock exams** and cross-reference errors with checklist topics
- Use the scorecard to **simulate an exam debrief**: where did you fail? What must you strengthen?

Once all checklist items are  and all categories are at **4–5 stars** and you're consistently scoring **85%+** on full practice exams with confidence in scenario-based reasoning, then 

you're likely **ready to book the real exam**.

Final Review Checklist

Agent Architecture & Design (15%)

- ☐ Design user interfaces for intuitive human-agent interaction
- ☐ Implement reasoning and action frameworks (e.g., ReAct)
- ☐ Configure agent-to-agent collaboration protocols
- ☐ Manage short-term and long-term memory for context
- ☐ Orchestrate multi-agent workflows and coordination
- ☐ Apply logic trees, prompt chains, stateful orchestration
- ☐ Integrate knowledge graphs for relational reasoning
- ☐ Ensure adaptability and scalability of architecture

Agent Development (15%)

- ☐ Engineer prompts and dynamic prompt chains
- ☐ Integrate generative and multimodal models (text, vision, audio)
- ☐ Build and connect tools, APIs, and external functions
- ☐ Implement error handling (retry, graceful recovery)
- ☐ Develop dynamic conversation flows with streaming feedback
- ☐ Evaluate and refine agent decision-making strategies

Evaluation & Tuning (13%)

- ☐ Implement evaluation pipelines and task benchmarks
- ☐ Compare performance across datasets and tasks
- ☐ Collect and integrate structured user feedback
- ☐ Tune model parameters (accuracy, latency trade-offs)
- ☐ Analyze evaluation results for targeted optimization

Deployment & Scaling (13%)

- ☐ Deploy and orchestrate multi-agent systems at scale
- ☐ Apply MLOps practices (CI/CD, monitoring, governance)
- ☐ Profile reliability under distributed system loads
- ☐ Scale with Docker/Kubernetes and load balancing
- ☐ Optimize deployment costs while maintaining availability

Cognition, Planning & Memory (10%)

- ☐ Implement memory for short- and long-term retention
- ☐ Apply reasoning frameworks (CoT, task decomposition)
- ☐ Engineer sequential and multi-step planning strategies
- ☐ Manage stateful orchestration for complex tasks
- ☐ Adapt reasoning strategies from prior experiences

Knowledge Integration & Data Handling (10%)

- ☐ Implement retrieval pipelines (RAG, hybrid, embedded search)
- ☐ Configure and optimize vector databases
- ☐ Build ETL pipelines for enterprise/client data integration
- ☐ Conduct data quality checks and augmentation
- ☐ Enable real-time structured/unstructured reasoning

NVIDIA Platform Implementation (7%)

- ☐ Integrate NeMo Guardrails for safety and compliance
- ☐ Deploy NVIDIA NIM microservices for inference
- ☐ Optimize workflows with Agent Intelligence Toolkit
- ☐ Leverage TensorRT-LLM and Triton for latency reduction
- ☐ Manage multi-modal pipelines on NVIDIA hardware

Run, Monitor & Maintain (7%)

- ☐ Define monitoring dashboards and reliability metrics
- ☐ Track logs, anomalies, and root cause analysis
- ☐ Continuously benchmark deployed agents
- ☐ Automate tuning, retraining, versioning pipelines
- ☐ Ensure uptime, transparency, and trust in live systems











Safety, Ethics & Compliance (5%)

- ☐ Design and enforce system security and audit trails
- ☐ Integrate compliance guardrails (privacy, enterprise policies)
- ☐ Mitigate bias and toxicity in outputs
- ☐ Deploy layered safety frameworks and escalation protocols
- ☐ Ensure regulatory and licensing compliance

Human-AI Interaction & Oversight (5%)

- ☐ Build intuitive UIs with user-in-the-loop mechanisms
- ☐ Design structured feedback loops for improvement
- ☐ Implement transparency (explainable reasoning, traceability)
- ☐ Enable human oversight and intervention

Exam Readiness Scorecard

Domain	Confidence (1-5)	Notes / Action Items
 Agent Architecture & Design (15%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Agent Development (15%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Evaluation & Tuning (13%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Deployment & Scaling (13%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Cognition, Planning & Memory (10%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Knowledge & Data Handling (10%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 NVIDIA Platform Implementation (7%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Run, Monitor & Maintain (7%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Safety, Ethics & Compliance (5%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Human-AI Interaction & Oversight (5%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	

PREVIEW ONLY

Conclusion

🌟 **Congratulations!! You are on the right path to certification. All of our practice exams include 300 + questions. Some, like this one, have over 500 questions.**


**This is a Preview Copy
with only 20 out of 520 questions
Get the Full Version at**

<https://cloudcertificationstore.com/b/8eYLD>

Most of our writers who have taken the exam recently—and the reviewers who purchased these materials—agree that **over 90 %** of the questions matched what they saw on the live test.

Invest in your future: browse the full catalogue of [Cloud practice exams at our store](https://cloudcertificationstore.com)

Featured Collection

 On Sale	 On Sale	 On Sale	 On Sale	 On Sale
2025 Google Cloud Professional Cloud Architect Exam Questions (GOOG-PCA-0010) \$8.89- \$6.22	2025 Google Cloud Professional Cloud Developer Certification Practice Exam Questions (GOOG-PCD-0010) \$8.89- \$6.22	2025 Google Cloud Professional Machine Learning Engineer Practice Exam Questions (GOOG-PMLE-0010) \$8.89- \$6.29	2025 Google Cloud Professional Data Engineer Certification Practice Exam Questions (GOOG-PDE-0010) \$8.89- \$6.22	2025 Google Cloud Associate Cloud Engineer Exam Questions (GOOG-ACE-0010) \$8.89- \$6.22
 On Sale	 On Sale	 On Sale	 On Sale	 On Sale
2025 Google Cloud Certified Professional Cloud Database Engineer Practice Exam Questions (GOOG-PCDE-0010) \$8.89- \$6.22	2025 AWS Certified Cloud Practitioner Practice Exam Questions (AWS-CLF-002-0010) \$8.89- \$6.29	2025 Microsoft Azure Fundamentals AZ-900 Practice Exam Questions (AZ-900-0010) \$8.89- \$6.22	2025 Microsoft Azure Developer Associate AZ-204 Practice Exam Questions (AZ-204-0010) \$8.89- \$6.22	FREE Practice Questions for the Google Cloud Generative AI Leader Exam \$0.00+