

# AWS Certified AI Practitioner (AIF-C01)

## 100 Questions & Answers

*Welcome to your complete AWS Certified AI Practitioner (AIF-C01) practice questions collection. This set is designed not just to quiz you, but to teach, reinforce, and sharpen your **exam readiness using real-world AWS AI scenarios**.*



## Learning Objectives and Expectations

You'll get:

- Realistic, scenario-based questions modeled after the AWS certification exam structure
- Organized in **batches of 10 questions, followed by 10 answers** and explanations
- Clear answer keys with short explanations to reinforce core concepts, AWS services, and decision logic

## AWS AI Practitioner (AIF-C01) Domains

Each domain carries a specific weight. Domains 2 and 3 (Generative AI and Foundation Model Applications) are the largest:

- **Domain 1:** Fundamentals of AI and ML – 20%
- **Domain 2:** Fundamentals of Generative AI – 24%
- **Domain 3:** Applications of Foundation Models – 28%
- **Domain 4:** Guidelines for Responsible AI – 14%
- **Domain 5:** Security, Compliance, and Governance – 14%

## Quick Reminder: How the Exam Works

- **Number of Questions:** 65
- **Format:** Multiple choice, multiple response, matching, and scenario-based
- **Time Limit:** 90 minutes
- **Passing Score:** 700/1000
- **Test Provider:** Pearson VUE (online or onsite)

## Questions By Domain

Domain	Title	Questions Assigned	Question Numbers
<b>Domain 1</b>	Fundamentals of AI and ML (20%)	20 Questions	Q1–2, Q12, Q14, Q27, Q31, Q34–35, Q42, Q52, Q56, Q63, Q66, Q72–73, Q74, Q94, Q97
<b>Domain 2</b>	Fundamentals of Generative AI (24%)	24 Questions	Q3–4, Q10, Q16, Q21, Q26, Q36, Q45–46, Q55, Q58, Q60, Q64, Q70, Q75, Q79–80, Q84–85, Q88, Q98–100
<b>Domain 3</b>	Applications of Foundation Models (28%)	28 Questions	Q5–6, Q9, Q13, Q19–20, Q22–23, Q25, Q30, Q33, Q37, Q47, Q50, Q57, Q59, Q61, Q67, Q69, Q71, Q76–77, Q82, Q86–87, Q90–91, Q95
<b>Domain 4</b>	Responsible AI (14%)	14 Questions	Q7, Q11, Q15, Q28–29, Q32, Q38, Q48–49, Q53–54, Q83, Q89, Q92
<b>Domain 5</b>	Security, Compliance, and Governance (14%)	14 Questions	Q8, Q17–18, Q24, Q39–41, Q43–44, Q51, Q62, Q65, Q68, Q78, Q93

## Remember — You Don't Need to Be Perfect to Pass!

The AWS AI Practitioner passing score is **700/1000**, which means you can **miss 15–20 questions** and still pass. Focus on concepts, AWS AI service selection, and responsible model deployment — trust your understanding and eliminate wrong answers.

# Questions 1–10

**Q1.**

Which AWS service allows you to build, train, and deploy custom machine learning models at scale?

- A) Amazon Lex
  - B) Amazon Bedrock
  - C) Amazon SageMaker
  - D) AWS Lambda
- 

**Q2.**

Which of the following best describes unsupervised learning?

- A) Predicting numeric values using labeled data
  - B) Identifying patterns or clusters in unlabeled data
  - C) Using rewards and penalties to train an agent
  - D) Applying labeled data to train a model
- 

**Q3.**

A retail company wants to summarize customer reviews. Which AWS service is best suited for this task?

- A) Amazon Rekognition
  - B) Amazon Comprehend
  - C) Amazon Translate
  - D) Amazon Polly
- 

**Q4.**

What is one advantage of using Retrieval-Augmented Generation (RAG) with a foundation model?

- A) Reduces latency by skipping the model altogether
  - B) Improves prediction speed by pruning the model
  - C) Provides up-to-date, context-specific answers using external data
  - D) Decreases API cost by caching prompt results
- 

**Q5.**

Which type of model training is used to specialize a foundation model for a specific task or dataset?

- A) Pre-training
  - B) Prompt chaining
  - C) Fine-tuning
  - D) Embedding
- 

**Q6.**

Which AWS service provides access to foundation models like Claude, Jurassic, or Titan via API without managing infrastructure?

- A) Amazon SageMaker
  - B) Amazon Kendra
  - C) Amazon Bedrock
  - D) AWS Lambda
- 

**Q7.**

In the context of AI, what is “bias” most accurately defined as?

- A) An overuse of GPU resources during training
  - B) A measure of model complexity
  - C) Systematic error that unfairly influences model outcomes
  - D) A type of deep learning optimization
- 

**Q8.**

Which service would help you build a multilingual chatbot with speech recognition?

- A) Amazon Translate + Amazon Comprehend
  - B) Amazon Polly + Amazon Rekognition
  - C) Amazon Lex + Amazon Transcribe
  - D) Amazon SageMaker + Amazon Textract
- 

**Q9.**

Which AWS service helps generate speech from text, useful for voice-enabled applications?

- A) Amazon Polly
  - B) Amazon Transcribe
  - C) Amazon Lex
  - D) Amazon Q
-

### Q10.

What is a “token” in the context of large language models?

- A) A multi-factor authentication method
- B) A pre-processed image input
- C) A segment of text used during model inference
- D) An access credential for invoking APIs

## Answers 1–10

### A1.

**Answer:** C) Amazon SageMaker

**Explanation:** SageMaker is AWS’s end-to-end ML platform for building, training, and deploying custom models.

### A2.

**Answer:** B) Identifying patterns or clusters in unlabeled data

**Explanation:** Unsupervised learning finds structure in data without labeled outputs.

### A3.

**Answer:** B) Amazon Comprehend

**Explanation:** Comprehend can perform NLP tasks like sentiment analysis and summarization.

### A4.

**Answer:** C) Provides up-to-date, context-specific answers using external data

**Explanation:** RAG enhances foundation models by retrieving relevant documents before generating responses.

### A5.

**Answer:** C) Fine-tuning

**Explanation:** Fine-tuning adapts a pre-trained model to new tasks using labeled domain-specific data.

### A6.

**Answer:** C) Amazon Bedrock

**Explanation:** Bedrock provides fully managed API access to multiple foundation models.

### A7.

**Answer:** C) Systematic error that unfairly influences model outcomes

**Explanation:** Bias in AI typically refers to unintended, unfair impacts on decisions.

### A8.

**Answer:** C) Amazon Lex + Amazon Transcribe

**Explanation:** Lex handles chatbot logic; Transcribe converts speech to text.

**A9.**

**Answer:** A) Amazon Polly

**Explanation:** Polly converts text to speech in lifelike voices.

**A10.**

**Answer:** C) A segment of text used during model inference

**Explanation:** Tokens are subunits of text (words, subwords) processed by LLMs.

# Questions 11–20

## Q11.

Which AWS service can help identify personal information like credit card numbers or names in a training dataset stored in Amazon S3?

- A) AWS IAM
  - B) Amazon Macie
  - C) Amazon Textract
  - D) AWS CloudTrail
- 

## Q12.

What is the purpose of SageMaker Clarify?

- A) Launching endpoint servers for foundation models
  - B) Evaluating model speed and GPU usage
  - C) Identifying bias and explaining model predictions
  - D) Serving prompts to Amazon Bedrock
- 

## Q13.

A developer needs a vector database to store text embeddings for semantic search. Which AWS service is most appropriate?

- A) Amazon DynamoDB
  - B) Amazon DocumentDB
  - C) Amazon OpenSearch with k-NN
  - D) Amazon Polly
- 

## Q14.

Which of the following is a common risk of generative AI in production?

- A) Deployment via HTTPS
  - B) High recall
  - C) Model hallucination
  - D) Insufficient GPU logging
- 

## Q15.

What is the primary role of Amazon Lex in AI-powered applications?

- A) Sentiment analysis
- B) Speech synthesis
- C) Building conversational interfaces and chatbots
- D) Object detection

**Q16.**

What type of prompting helps models reason more clearly by breaking tasks into logical steps?

- A) Zero-shot prompting
  - B) Prompt chaining
  - C) Chain-of-thought prompting
  - D) Embedding prompting
- 

**Q17.**

Which AWS service allows a user to catalog, govern, and manage AI-related datasets?

- A) Amazon DataZone
  - B) Amazon Inspector
  - C) Amazon S3
  - D) AWS Artifact
- 

**Q18.**

What is a *model card* in SageMaker used for?

- A) Encrypting model input/output
  - B) Documenting model use, data, metrics, and ethical considerations
  - C) Launching auto-scaling endpoints
  - D) Selecting the correct foundation model from Amazon Bedrock
- 

**Q19.**

Which AWS service helps with identifying and logging all API activity related to your AI infrastructure?

- A) Amazon Macie
  - B) AWS Config
  - C) AWS CloudTrail
  - D) Amazon Inspector
-



## Q20.

Which AWS tool helps ensure that only approved users and roles can access an Amazon Bedrock model endpoint?

- A) AWS KMS
- B) AWS IAM
- C) Amazon Polly
- D) Amazon Route 53

# Answers 11–20

## A11.

**Answer:** B) Amazon Macie

**Explanation:** Macie detects and classifies sensitive data in S3, such as PII.

## A12.

**Answer:** C) Identifying bias and explaining model predictions

**Explanation:** SageMaker Clarify supports fairness and explainability through metrics and SHAP values.

## A13.

**Answer:** C) Amazon OpenSearch with k-NN

**Explanation:** OpenSearch supports vector search for embeddings using k-Nearest Neighbors.

## A14.

**Answer:** C) Model hallucination

**Explanation:** A common risk where the model outputs plausible but false or made-up content.

## A15.

**Answer:** C) Building conversational interfaces and chatbots

**Explanation:** Lex provides voice/text-based interfaces powered by NLP.

## A16.

**Answer:** C) Chain-of-thought prompting

**Explanation:** Helps improve reasoning by guiding the model step-by-step.

## A17.

**Answer:** A) Amazon DataZone

**Explanation:** DataZone enables cataloging, governance, and sharing of data assets across teams.

**A18.**

**Answer:** B) Documenting model use, data, metrics, and ethical considerations

**Explanation:** Model cards improve transparency and compliance by summarizing key model details.

**A19.**

**Answer:** C) AWS CloudTrail

**Explanation:** Tracks all API actions across your AWS account, including those for AI services.

**A20.**

**Answer:** B) AWS IAM

**Explanation:** IAM controls access and permissions for Bedrock, SageMaker, and all AWS resources.

## Questions 21–30

### Q21.

A user needs to generate multiple image variations from a single text prompt. Which type of model would be best for this task?

- A) Recurrent neural network
  - B) Diffusion model
  - C) Transformer model
  - D) Logistic regression
- 

### Q22.

Which AWS service helps convert handwritten forms into structured digital data?

- A) Amazon Textract
  - B) Amazon Polly
  - C) Amazon Lex
  - D) AWS Glue
- 

### Q23.

In a generative AI application, which AWS service provides no-code tools to experiment with foundation models?

- A) Amazon Q
  - B) Amazon Bedrock Studio
  - C) Amazon PartyRock
  - D) AWS Lambda
- 

### Q24.

Which of the following is a benefit of using AWS-managed foundation models instead of training your own?

- A) Requires more infrastructure setup
  - B) Lower initial compute costs and faster time-to-value
  - C) Allows unlimited control of the model's internal layers
  - D) Offers higher hallucination risk due to limited access
- 

### Q25.

Which metric is best suited to evaluate the quality of machine translation output?

- A) ROUGE

- B) BLEU
  - C) AUC
  - D) F1 Score
- 

### Q26.

What is one major benefit of in-context learning over fine-tuning?

- A) It updates the model's weights for new tasks
  - B) It costs less and doesn't require retraining
  - C) It guarantees higher accuracy
  - D) It allows GPU scaling
- 

### Q27.

A machine learning model performs well on training data but poorly on new data. What is the most likely issue?

- A) Underfitting
  - B) Data drift
  - C) Overfitting
  - D) High bias
- 

### Q28.

What is the purpose of Amazon SageMaker Feature Store?

- A) Manage IAM roles for training jobs
  - B) Store structured logs for model performance
  - C) Serve consistent features for training and inference
  - D) Replace Glue for visual data transformation
- 

### Q29.

What is a key feature of the AWS Shared Responsibility Model as it applies to AI workloads?

- A) AWS is responsible for securing custom model code
  - B) Customers are responsible for physical data center access
  - C) AWS secures the infrastructure, customers secure configurations and data
  - D) Customers rely on AWS to guarantee model explainability
-

**Q30.**

Which AWS tool enables human reviewers to verify low-confidence model predictions before results are finalized?

- A) Amazon Lex
- B) Amazon A2I
- C) AWS Macie
- D) Amazon Q

---

## Answers 21–30

**A21.**

**Answer:** B) Diffusion model

**Explanation:** Diffusion models generate high-quality images from noise guided by text prompts.

**A22.**

**Answer:** A) Amazon Textract

**Explanation:** Textract extracts text and structure (like tables, forms) from scanned documents.

**A23.**

**Answer:** C) Amazon PartyRock

**Explanation:** PartyRock allows users to experiment with Bedrock-hosted models via a no-code interface.

**A24.**

**Answer:** B) Lower initial compute costs and faster time-to-value

**Explanation:** Managed models save cost and time since training and infra are already done by AWS.

**A25.**

**Answer:** B) BLEU

**Explanation:** BLEU evaluates translation accuracy based on n-gram overlap with reference text.

**A26.**

**Answer:** B) It costs less and doesn't require retraining

**Explanation:** In-context learning uses well-structured prompts without changing the model itself.

**A27.**

**Answer:** C) Overfitting

**Explanation:** Overfitting happens when a model learns training data too well but fails to generalize.

**A28.**

**Answer:** C) Serve consistent features for training and inference

**Explanation:** Feature Store provides a centralized source for ML features used in both stages.

**A29.**

**Answer:** C) AWS secures the infrastructure, customers secure configurations and data

**Explanation:** The Shared Responsibility Model splits infra vs user responsibilities.

**A30.**

**Answer:** B) Amazon A2I

**Explanation:** A2I (Augmented AI) lets humans review ML outputs before they are accepted.

# Questions 31–40

## Q31.

Which AWS service would you use to translate chatbot messages from English to Spanish in real-time?

- A) Amazon Comprehend
  - B) Amazon Polly
  - C) Amazon Translate
  - D) Amazon Lex
- 

## Q32.

What is the main advantage of using Amazon SageMaker JumpStart for a new AI project?

- A) It provides serverless access to third-party models.
  - B) It allows full model interpretability by default.
  - C) It offers pre-built models and solutions for quick deployment.
  - D) It replaces the need for IAM roles in ML workflows.
- 

## Q33.

Which method is commonly used to measure the similarity between user input and stored content in a vector database?

- A) Hash matching
  - B) k-Nearest Neighbors (k-NN)
  - C) Confusion matrix
  - D) ROC curve
- 

## Q34.

What is an *embedding* in the context of AI and vector databases?

- A) A label used for supervised learning
  - B) A data format used for model deployment
  - C) A dense vector representing the meaning of data
  - D) A specific GPU setting for SageMaker endpoints
-

### Q35.

A model always predicts the majority class but fails to detect rare cases. What metric is most likely misleading in this case?

- A) Precision
  - B) F1 Score
  - C) Accuracy
  - D) Recall
- 

### Q36.

What is the purpose of using SHAP values in model evaluation?

- A) To encrypt data before inference
  - B) To calculate how each feature impacts a model's output
  - C) To standardize tokenization of inputs
  - D) To classify model bias as legal or illegal
- 

### Q37.

Which AWS service supports storing AI model outputs securely with version control and access logging?

- A) AWS Lambda
  - B) AWS Glue
  - C) Amazon S3
  - D) AWS CloudTrail
- 

### Q38.

What does Amazon Rekognition primarily analyze?

- A) Document structure
  - B) Audio waveforms
  - C) Image and video content
  - D) Token-based language prompts
- 

### Q39.

How does Amazon Q for Business support enterprise AI use cases?

- A) By running real-time image recognition
- B) By providing GPU acceleration for model training



- C) By answering natural language questions using internal company data
- D) By launching distributed training jobs on SageMaker

#### Q40.

Which of the following defines *prompt injection* in generative AI applications?

- A) Feeding the model with missing values
- B) Tricking the model into executing unintended behavior via crafted input
- C) Replacing GPU memory during inference
- D) Automatically chaining model outputs to data pipelines

## Answers 31–40

#### A31.

**Answer:** C) Amazon Translate

**Explanation:** Translate provides real-time, high-quality language translation.

#### A32.

**Answer:** C) It offers pre-built models and solutions for quick deployment

**Explanation:** JumpStart accelerates AI adoption by providing ready-to-use templates and models.

#### A33.

**Answer:** B) k-Nearest Neighbors (k-NN)

**Explanation:** Vector search compares embeddings using k-NN for similarity ranking.

#### A34.

**Answer:** C) A dense vector representing the meaning of data

**Explanation:** Embeddings convert text/images into vectors for comparison and semantic search.

#### A35.

**Answer:** C) Accuracy

**Explanation:** High accuracy can be misleading when rare classes are missed (i.e., imbalanced data).

#### A36.

**Answer:** B) To calculate how each feature impacts a model's output

**Explanation:** SHAP values explain model predictions by showing feature contribution.

#### A37.

**Answer:** C) Amazon S3

**Explanation:** S3 securely stores data, including models and logs, with versioning and access control.

**A38.**

**Answer:** C) Image and video content

**Explanation:** Rekognition detects labels, faces, and unsafe content in media.

**A39.**

**Answer:** C) By answering natural language questions using internal company data

**Explanation:** Amazon Q for Business is designed for internal Q&A across enterprise documents.

**A40.**

**Answer:** B) Tricking the model into executing unintended behavior via crafted input

**Explanation:** Prompt injection exploits LLMs by injecting commands that override the intended prompt.

# Questions 41–50

## Q41.

Which AWS service can automate data transformations and prep for machine learning with a visual interface?

- A) AWS Glue DataBrew
  - B) Amazon S3
  - C) Amazon Q
  - D) Amazon Polly
- 

## Q42.

What is the purpose of a system prompt in a generative AI application?

- A) To define the format of the output
  - B) To provide instructions that guide the model's role or behavior
  - C) To set GPU configuration
  - D) To compress the input tokens
- 

## Q43.

Which AWS service would help detect data drift or anomalies in live model input and predictions?

- A) Amazon Rekognition
  - B) SageMaker Model Monitor
  - C) AWS Config
  - D) Amazon Kendra
- 

## Q44.

An organization wants to classify documents by language and topic. Which combination of services is best?

- A) Amazon Polly and Amazon Textract
  - B) Amazon Translate and Amazon S3
  - C) Amazon Comprehend and Amazon S3
  - D) Amazon Bedrock and Amazon Q
-

#### Q45.

Which type of model evaluation metric is most suitable for summarization tasks?

- A) ROUGE
  - B) BLEU
  - C) Perplexity
  - D) MAE
- 

#### Q46.

How does Amazon Bedrock simplify access to foundation models?

- A) It auto-trains any model uploaded by the user
  - B) It hosts open-source models in SageMaker endpoints
  - C) It provides API access to multiple foundation models without managing infrastructure
  - D) It requires a local container environment for use
- 

#### Q47.

What is the purpose of a vector database in Retrieval-Augmented Generation (RAG)?

- A) To reduce training time by compressing models
  - B) To generate GPU-friendly model files
  - C) To store embeddings and enable similarity search for grounding responses
  - D) To support IAM token storage
- 

#### Q48.

Which AWS tool enables ongoing audit tracking and policy alignment for AI projects?

- A) Amazon Q
  - B) AWS Artifact
  - C) AWS Audit Manager
  - D) Amazon Comprehend
- 

#### Q49.

What is the main purpose of Amazon Titan in the AWS AI ecosystem?

- A) It offers containerized GPU inference for local deployments
  - B) It's AWS's proprietary foundation model family, accessible via Bedrock
  - C) It provides OCR and handwriting recognition
  - D) It's an Amazon robotics platform for automation
-

**Q50.**

In what situation would using human-in-the-loop (via Amazon A2I) be most appropriate?

- A) To store model parameters in S3
- B) When model outputs must be verified for accuracy before release
- C) When you want to increase token limits in Bedrock
- D) To convert speech to text

## Answers 41–50

**A41.**

**Answer:** A) AWS Glue DataBrew

**Explanation:** DataBrew lets users visually transform, clean, and prepare data for ML without code.

**A42.**

**Answer:** B) To provide instructions that guide the model's role or behavior

**Explanation:** System prompts establish tone, behavior, and rules for the AI's responses.

**A43.**

**Answer:** B) SageMaker Model Monitor

**Explanation:** Model Monitor tracks model inputs and outputs for drift, helping detect issues post-deployment.

**A44.**

**Answer:** C) Amazon Comprehend and Amazon S3

**Explanation:** Comprehend analyzes text for language and topic classification; S3 stores the documents.

**A45.**

**Answer:** A) ROUGE

**Explanation:** ROUGE compares overlaps between generated and reference summaries.

**A46.**

**Answer:** C) It provides API access to multiple foundation models without managing infrastructure

**Explanation:** Bedrock is a managed service that provides simple access to FMs from various providers.

**A47.**

**Answer:** C) To store embeddings and enable similarity search for grounding responses

**Explanation:** Vector databases hold embeddings to support retrieval in RAG workflows.

**A48.**

**Answer:** C) AWS Audit Manager

**Explanation:** Audit Manager helps you map and maintain compliance controls for audit readiness.

**A49.**

**Answer:** B) It's AWS's proprietary foundation model family, accessible via Bedrock

**Explanation:** Titan models support generative AI use cases and are natively integrated with Bedrock.

**A50.**

**Answer:** B) When model outputs must be verified for accuracy before release

**Explanation:** Amazon A2I allows for human intervention to validate or correct AI results.

# Questions 51–60

## Q51.

What is the benefit of using Amazon SageMaker Studio for AI development?

- A) It replaces IAM permissions
  - B) It offers a browser-based IDE with integrated ML tools
  - C) It automatically generates prompts for Bedrock
  - D) It converts images to embeddings
- 

## Q52.

Which AWS service would best support training and deploying a machine learning model with custom Python code?

- A) Amazon Lex
  - B) Amazon SageMaker
  - C) Amazon Translate
  - D) Amazon Polly
- 

## Q53.

A company wants to track changes in model performance over time. What AWS service helps with this?

- A) Amazon Textract
  - B) SageMaker Model Monitor
  - C) Amazon Lex
  - D) Amazon Transcribe
- 

## Q54.

Which of the following AWS services provides an interface to manage and govern data shared across teams?

- A) AWS IAM
  - B) Amazon Macie
  - C) Amazon DataZone
  - D) Amazon Translate
-

### Q55.

What is the main benefit of using Amazon Comprehend's custom classification feature?

- A) It allows training of vision models on image datasets
  - B) It enables model quantization for faster inference
  - C) It lets you categorize documents based on your own business labels
  - D) It encrypts output summaries for compliance
- 

### Q56.

Which of the following best defines "zero-shot prompting"?

- A) The model is fine-tuned before inference
  - B) The model answers a task without seeing any examples
  - C) The model is prompted after receiving labeled training data
  - D) The prompt contains multiple few-shot examples
- 

### Q57.

What is a major advantage of using Amazon Q (Business) over a general LLM interface?

- A) It supports GPU-based training workloads
  - B) It automatically clusters training data
  - C) It grounds answers in enterprise documents and internal knowledge
  - D) It supports image classification natively
- 

### Q58.

Which AWS service provides a knowledge base-backed AI chatbot using Bedrock models and RAG?

- A) Amazon Lex
  - B) Amazon Rekognition
  - C) Amazon Q
  - D) Amazon S3
- 

### Q59.

A security team wants to ensure AI data in S3 is encrypted at rest. Which AWS service helps manage encryption keys?

- A) Amazon A2I
- B) AWS Artifact



- C) AWS KMS  
D) Amazon Bedrock

### Q60.

What is the benefit of using Guardrails in Amazon Bedrock?

- A) To detect and prevent drift in training pipelines  
B) To convert structured data into image format  
C) To restrict harmful, unsafe, or non-compliant model outputs  
D) To perform live speech translation

## Answers 51–60

### A51.

**Answer:** B) It offers a browser-based IDE with integrated ML tools

**Explanation:** SageMaker Studio centralizes data science workflows with Jupyter notebooks and built-in tools.

### A52.

**Answer:** B) Amazon SageMaker

**Explanation:** SageMaker supports bringing your own code and containers for ML model training and deployment.

### A53.

**Answer:** B) SageMaker Model Monitor

**Explanation:** Tracks prediction quality and input distribution over time to identify drift or degradation.

### A54.

**Answer:** C) Amazon DataZone

**Explanation:** DataZone supports cataloging, sharing, and governance of datasets across teams.

### A55.

**Answer:** C) It lets you categorize documents based on your own business labels

**Explanation:** Comprehend custom classification lets you train a classifier on your own label taxonomy.

### A56.

**Answer:** B) The model answers a task without seeing any examples

**Explanation:** Zero-shot prompting relies on pretrained capabilities and instructions alone.

**A57.**

**Answer:** C) It grounds answers in enterprise documents and internal knowledge

**Explanation:** Amazon Q uses RAG to answer based on your company's data.

**A58.**

**Answer:** C) Amazon Q

**Explanation:** Q integrates Bedrock models with internal data sources for grounded, enterprise chat.

**A59.**

**Answer:** C) AWS KMS

**Explanation:** Key Management Service handles encryption key creation, storage, and access.

**A60.**

**Answer:** C) To restrict harmful, unsafe, or non-compliant model outputs

**Explanation:** Bedrock Guardrails filter and control generative model responses based on safety rules.

## Questions 61–70

### Q61.

Which AWS service is specifically designed to help build, test, and deploy AI solutions using an end-to-end managed environment?

- A) Amazon Textract
  - B) Amazon Q
  - C) Amazon SageMaker
  - D) AWS Lambda
- 

### Q62.

What is the main reason organizations use Retrieval-Augmented Generation (RAG) instead of relying only on a foundation model's internal knowledge?

- A) It reduces GPU power requirements
  - B) It improves model explainability
  - C) It enables grounding in up-to-date or proprietary information
  - D) It increases BLEU score in translation tasks
- 

### Q63.

In a classification model, which metric is most useful when trying to reduce false negatives in rare event detection?

- A) Accuracy
  - B) Recall
  - C) BLEU
  - D) MAE
- 

### Q64.

What does a model card typically include? (*Choose TWO.*)

- A) GPU utilization metrics
  - B) Intended use cases
  - C) Dataset details and sources
  - D) Token cost per query
- 

### Q65.

Which AWS service helps prevent unauthorized access to AI endpoints by controlling permissions?

- A) AWS IAM
  - B) Amazon Translate
  - C) Amazon Polly
  - D) Amazon CloudWatch
- 

**Q66.**

What is one reason to use prompt engineering instead of fine-tuning?

- A) It reduces latency
  - B) It allows reuse of the same model without additional training
  - C) It guarantees full explainability
  - D) It converts model outputs into tokens
- 

**Q67.**

Which AWS service is best suited for real-time speech-to-text transcription?

- A) Amazon Transcribe
  - B) Amazon Lex
  - C) Amazon Polly
  - D) Amazon Kendra
- 

**Q68.**

A team wants to perform image analysis, including detecting objects and facial attributes. Which service should they use?

- A) Amazon Bedrock
  - B) Amazon Comprehend
  - C) Amazon Rekognition
  - D) Amazon A2I
- 

**Q69.**

What is the main risk of not monitoring model inputs after deployment?

- A) GPU underuse
  - B) Increased training cost
  - C) Input drift leading to degraded model performance
  - D) Longer prompt tokenization
-

**Q70.**

What AWS feature supports safe handling of sensitive prompts and outputs in generative AI applications?

- A) SageMaker Model Cards
- B) Amazon Macie
- C) Bedrock Guardrails
- D) AWS CodeCommit

## Answers 61–70

**A61.**

**Answer:** C) Amazon SageMaker

**Explanation:** SageMaker provides a complete ML platform for model building, training, and deployment.

**A62.**

**Answer:** C) It enables grounding in up-to-date or proprietary information

**Explanation:** RAG lets the model reference external documents for current or custom data.

**A63.**

**Answer:** B) Recall

**Explanation:** Recall focuses on capturing true positives — critical when false negatives are costly.

**A64.**

**Answer:** B) Intended use cases and C) Dataset details and sources

**Explanation:** Model cards document purpose, data origin, performance, and limitations.

**A65.**

**Answer:** A) AWS IAM

**Explanation:** IAM enforces fine-grained access control to AI resources like Bedrock endpoints.

**A66.**

**Answer:** B) It allows reuse of the same model without additional training

**Explanation:** Prompt engineering lets you tailor outputs with no retraining or cost overhead.

**A67.**

**Answer:** A) Amazon Transcribe

**Explanation:** Transcribe converts live or recorded speech into accurate text.

**A68.**

**Answer:** C) Amazon Rekognition

**Explanation:** Rekognition analyzes images and videos for faces, objects, and text.

**A69.**

**Answer:** C) Input drift leading to degraded model performance

**Explanation:** If input data shifts from what the model was trained on, accuracy may drop.

**A70.**

**Answer:** C) Bedrock Guardrails

**Explanation:** Guardrails help filter and manage prompts and outputs for compliance and safety.

# Questions 71–80

**Q71.**

What AWS service enables access to pre-trained models from Anthropic, Stability AI, and AI21 Labs?

- A) Amazon SageMaker
  - B) Amazon Bedrock
  - C) Amazon Comprehend
  - D) Amazon EC2
- 

**Q72.**

What is a foundational concept of human-centered design in AI systems?

- A) Reducing memory usage in training
  - B) Prioritizing GPU acceleration
  - C) Explaining outputs in a way users can understand
  - D) Minimizing vector database usage
- 

**Q73.**

Which AWS service is primarily used to manage machine learning features and avoid duplication between training and inference?

- A) AWS Glue
  - B) Amazon S3
  - C) SageMaker Feature Store
  - D) Amazon Lex
- 

**Q74.**

An organization wants to control how long ML data is retained and ensure deletion after a specific period. Which feature helps?

- A) S3 Bucket Lifecycle Policies
  - B) Amazon Lex intents
  - C) AWS Lambda
  - D) SageMaker model tuning
-

### Q75.

What is the purpose of Amazon A2I in AI workflows?

- A) Create multi-modal embeddings for large models
  - B) Allow human review of low-confidence or sensitive ML predictions
  - C) Accelerate foundation model training
  - D) Detect personally identifiable information in documents
- 

### Q76.

What does BLEU score evaluate in NLP tasks?

- A) Output tone
  - B) Model inference time
  - C) Translation accuracy based on reference overlap
  - D) Ethical use of AI
- 

### Q77.

Which AWS service provides compliance documentation (e.g., SOC, ISO certifications)?

- A) AWS CloudTrail
  - B) AWS Artifact
  - C) Amazon Rekognition
  - D) Amazon S3
- 

### Q78.

A financial company wants to trace which dataset was used to train a particular ML model. Which governance tool supports this need?

- A) Amazon Polly
  - B) Amazon Translate
  - C) Model Cards
  - D) Amazon Kendra
- 

### Q79.

What is *latent space* in the context of foundation models?

- A) A model's reserved memory area
- B) A range of GPU utilization values



- C) A compressed internal representation of features or concepts
- D) A multi-user chat interface for AI responses

---

### Q80.

Which AWS service would best support a chatbot that needs to understand spoken commands and reply with generated speech?

- A) Amazon Textract
  - B) Amazon Q
  - C) Amazon Transcribe + Amazon Polly
  - D) Amazon Comprehend + Amazon S3
- 

## Answers 71–80

### A71.

**Answer:** B) Amazon Bedrock

**Explanation:** Bedrock provides access to foundation models from multiple providers via API.

### A72.

**Answer:** C) Explaining outputs in a way users can understand

**Explanation:** Human-centered design emphasizes clarity, trust, and user comprehension.

### A73.

**Answer:** C) SageMaker Feature Store

**Explanation:** Stores and serves features consistently across ML pipelines.

### A74.

**Answer:** A) S3 Bucket Lifecycle Policies

**Explanation:** These automate data archiving and deletion for governance and compliance.

### A75.

**Answer:** B) Allow human review of low-confidence or sensitive ML predictions

**Explanation:** A2I adds human-in-the-loop validation to critical model decisions.

### A76.

**Answer:** C) Translation accuracy based on reference overlap

**Explanation:** BLEU measures how closely machine translation matches a human reference.

**A77.**

**Answer:** B) AWS Artifact

**Explanation:** Provides downloadable compliance and security documents.

**A78.**

**Answer:** C) Model Cards

**Explanation:** Document model inputs, training data, use cases, and limitations.

**A79.**

**Answer:** C) A compressed internal representation of features or concepts

**Explanation:** Foundation models use latent space to understand and relate input data.

**A80.**

**Answer:** C) Amazon Transcribe + Amazon Polly

**Explanation:** Transcribe converts speech to text; Polly converts generated replies into audio.

# Questions 81–90

## Q81.

Which AWS service helps you define custom vocabulary and language models for speech recognition tasks?

- A) Amazon Translate
  - B) Amazon Polly
  - C) Amazon Transcribe
  - D) Amazon Kendra
- 

## Q82.

What does the *Shared Responsibility Model* mean in the context of AWS AI services?

- A) AWS and users co-own model weights
  - B) AWS handles physical security; users manage data, configuration, and access
  - C) AWS manages all aspects of AI, including ethics
  - D) Users only manage billing and logs
- 

## Q83.

Which tool would help a company detect bias in both datasets and model predictions?

- A) Amazon Comprehend
  - B) SageMaker Clarify
  - C) AWS Glue
  - D) Amazon Bedrock
- 

## Q84.

Why might an organization use *continuous pretraining* with foundation models?

- A) To delete unused tokens
  - B) To keep the model updated with recent data
  - C) To reset model weights
  - D) To avoid needing GPUs
-

**Q85.**

What is the function of *instruction tuning* in foundation models?

- A) Adjusts latency parameters
  - B) Teaches the model to better follow user commands
  - C) Converts binary to text
  - D) Optimizes GPU memory
- 

**Q86.**

Which AWS service supports chunking text input and adding grounding context for large prompts?

- A) Amazon Lex
  - B) Amazon Comprehend
  - C) Amazon Bedrock
  - D) AWS Lambda
- 

**Q87.**

Which term best describes the risk of a model memorizing and reproducing training data verbatim?

- A) Overfitting
  - B) Latency
  - C) Prompt injection
  - D) Token duplication
- 

**Q88.**

An AI assistant produces different responses each time for the same prompt. What parameter is likely causing this?

- A) Accuracy
  - B) Recall
  - C) Temperature
  - D) Latency
- 

**Q89.**

In responsible AI development, what is one purpose of using *model explainability* tools?

- A) To reduce GPU usage
- B) To measure token frequency

- C) To help understand feature impact on model predictions
- D) To optimize latency in real-time endpoints

### Q90.

A company needs to prevent public access to its SageMaker inference endpoints. What should it use?

- A) Public API Gateway
- B) Amazon Macie
- C) AWS PrivateLink
- D) Amazon Polly

## Answers 81–90

### A81.

**Answer:** C) Amazon Transcribe

**Explanation:** Transcribe supports custom vocabularies and language models for tailored speech recognition.

### A82.

**Answer:** B) AWS handles physical security; users manage data, configuration, and access

**Explanation:** The Shared Responsibility Model divides control between AWS (infra) and the customer (data, permissions, models).

### A83.

**Answer:** B) SageMaker Clarify

**Explanation:** Clarify detects bias and provides SHAP-based explainability during and after training.

### A84.

**Answer:** B) To keep the model updated with recent data

**Explanation:** Continuous pretraining extends a model's knowledge beyond its original training cutoff.

### A85.

**Answer:** B) Teaches the model to better follow user commands

**Explanation:** Instruction tuning helps models learn to respond well to prompts phrased as instructions.

### A86.

**Answer:** C) Amazon Bedrock

**Explanation:** Bedrock supports prompt engineering, chunking, and RAG for grounding context.

**A87.**

**Answer:** A) Overfitting

**Explanation:** Overfitting can cause models to memorize training data instead of generalizing.

**A88.**

**Answer:** C) Temperature

**Explanation:** Higher temperature adds randomness to model output, affecting determinism.

**A89.**

**Answer:** C) To help understand feature impact on model predictions

**Explanation:** Explainability helps interpret why the model made a certain prediction.

**A90.**

**Answer:** C) AWS PrivateLink

**Explanation:** PrivateLink allows private, VPC-internal access to SageMaker and Bedrock endpoints.

# Questions 91–100

## Q91.

What is the role of Amazon Kendra in a Retrieval-Augmented Generation (RAG) pipeline?

- A) Train a new foundation model from scratch
  - B) Manage cloud cost for AI pipelines
  - C) Retrieve relevant documents or passages from a data source
  - D) Generate embeddings from images
- 

## Q92.

Which of the following is a typical component of a RAG system on AWS? (Choose TWO.)

- A) Amazon Bedrock
  - B) Amazon Textract
  - C) Amazon OpenSearch
  - D) AWS CodePipeline
- 

## Q93.

What AWS service provides a no-code interface for creating apps powered by Bedrock foundation models?

- A) Amazon Rekognition
  - B) Amazon Lex
  - C) Amazon Q
  - D) Amazon PartyRock
- 

## Q94.

What type of learning uses labeled input-output pairs to train a model?

- A) Reinforcement learning
  - B) Supervised learning
  - C) Unsupervised learning
  - D) Federated learning
-

### Q95.

Which AWS service allows users to transcribe streaming audio with speaker identification?

- A) Amazon Polly
  - B) Amazon Transcribe
  - C) Amazon Comprehend
  - D) Amazon SageMaker
- 

### Q96.

Why would an AI team use AWS Config in an ML pipeline?

- A) To build no-code AI assistants
  - B) To validate prompts for Bedrock
  - C) To monitor compliance of infrastructure configurations
  - D) To perform model hyperparameter tuning
- 

### Q97.

Which type of bias occurs when the training dataset does not represent the target population?

- A) Algorithmic bias
  - B) Data drift
  - C) Sampling bias
  - D) Tuning bias
- 

### Q98.

What is one limitation of few-shot prompting compared to fine-tuning?

- A) It requires GPU access
  - B) It needs a full labeled dataset
  - C) It has limited task-specific learning retention
  - D) It supports only classification models
- 

### Q99.

A company needs to host a multilingual question-answering chatbot using proprietary documentation. Which combination is best?

- A) Amazon Translate + Amazon S3
- B) Amazon Kendra + Amazon Bedrock



- C) Amazon Polly + AWS Lambda
  - D) AWS Glue + Amazon SageMaker
- 

**Q100.**

Which AWS service allows teams to detect and redact sensitive data from documents or text inputs before model training?

- A) Amazon Bedrock
  - B) Amazon Rekognition
  - C) Amazon Macie
  - D) Amazon Comprehend
- 

## Answers 91–100

**A91.**

**Answer:** C) Retrieve relevant documents or passages from a data source

**Explanation:** Kendra is used to search and retrieve documents based on semantic or keyword relevance.

**A92.**

**Answer:** A) Amazon Bedrock and C) Amazon OpenSearch

**Explanation:** Bedrock provides the generative model; OpenSearch stores embeddings for vector search in RAG.

**A93.**

**Answer:** D) Amazon PartyRock

**Explanation:** PartyRock is a no-code app builder that uses Bedrock models to build generative AI apps.

**A94.**

**Answer:** B) Supervised learning

**Explanation:** Supervised learning uses labeled input-output pairs to learn mappings.

**A95.**

**Answer:** B) Amazon Transcribe

**Explanation:** Transcribe can transcribe speech with features like speaker identification and vocabulary filtering.

**A96.**

**Answer:** C) To monitor compliance of infrastructure configurations

**Explanation:** AWS Config tracks configuration changes and compliance of AWS resources.

**A97.**

**Answer:** C) Sampling bias

**Explanation:** Sampling bias arises when the training data does not represent the broader population.

**A98.**

**Answer:** C) It has limited task-specific learning retention

**Explanation:** Few-shot prompting is flexible but doesn't permanently adapt the model like fine-tuning.

**A99.**

**Answer:** B) Amazon Kendra + Amazon Bedrock

**Explanation:** Kendra retrieves relevant data, Bedrock generates grounded responses using that context.

**A100.**

**Answer:** C) Amazon Macie

**Explanation:** Macie scans for sensitive data (like PII) and helps with redaction and classification.