




CERTIFIED AI PRACTITIONER

(AIF-C01)

PRACTICE EXAM QUESTIONS



Cloud Certification Store

AWS Certified AI Practitioner Certification (AIF-C01) Practice Exam Questions (AWS-AIF-CO1-0010)

© 2025 [Cloud Certification Store](https://cloudcertificationstore.com/) All rights reserved.

Amazon Web Services (AWS) is a registered trademark of Amazon.com, Inc. or its affiliates. This practice set is an original work for educational use and is **NOT** endorsed by or affiliated with Amazon Web Services. “AWS,” “AWS Certified Developer – Associate,” and related marks are trademarks of Amazon.com, Inc., used here for identification only.

DISCLAIMER

- This practice test includes questions **compiled from various exam preparation platforms.**
- Important: **Questions and answers were AI-assisted and human-curated.** Verify accuracy with official documentation before relying on this material.
- **Users are strongly encouraged to** double-check all content against **official documentation and trusted sources** before using it for exam preparation or making important decisions.
- The creators of this material assume **no responsibility** for any errors, inaccuracies, or outcomes, including exam results, based on the use of this content.
- **Some questions might be duplicated or close** to previous ones, this is done on purpose as a way to re-inforce your learning.
- Single-user licence only
 - Includes one unique Payhip Licence Key per purchase, along with a Product Key.
 - Redistribution, resale, or public posting is prohibited. We can trace any file to the purchaser, with the use of the purchased License Key and Product Key.

AWS Certified AI Practitioner Certification (AIF-C01) Practice Exam Questions (AWS-AIF-CO1-0010) - *PREVIEW 20 out of 300+ Questions*



Issued by [Amazon Web Services Training and Certification](#)

AWS Certified AI Practitioner

Earners of this badge understand AI, ML, and generative AI concepts, methods, and strategies in general and on AWS. They can determine the correct types of AI/ML technologies to apply to specific use cases and know how to use AI, ML, and generative AI technologies responsibly. They are familiar with the AWS Global Infrastructure, core AWS services and use cases, AWS service pricing models, and the AWS shared responsibility model for security and compliance in the AWS Cloud.

<https://aws.amazon.com/certification/certified-ai-practitioner/>

AWS Certified AI Practitioner validates in-demand knowledge of artificial intelligence (AI), machine learning (ML), and generative AI concepts and use cases. Sharpen your competitive edge and position yourself for career growth and higher earnings.

[Schedule your exam](#)

Exam overview

AWS Certified AI Practitioner

Category

Foundational

Exam duration

90 minutes

Exam format

65 questions

Cost

100 USD. Visit [Exam pricing](#) for additional cost information, including foreign exchange rates

Intended candidate

Individuals who are familiar with, but do not necessarily build, solutions using AI/ML technologies on AWS

Candidate role examples

Business analyst, IT support, marketing professional, product or project manager, line-of-business or IT manager, sales professional

Testing options

Pearson VUE testing center or online proctored exam

Languages offered

English, Japanese, Korean, Portuguese (Brazil), and Simplified Chinese

Show more

Prepare for the exam

Go from start to certified. Follow our Exam Prep Plan on AWS Skill Builder, our online learning center, so you can approach exam day with confidence.

1. Get to know the exam

Follow the [4-step plan](#).

[Review the exam guide](#).

2. Refresh your AWS knowledge and skills

Enroll in digital courses where you need to fill gaps in knowledge and skills, practice with AWS Builder Labs, AWS Cloud Quest, and AWS Jam.

3. Review and practice for your exam

Review the scope of the exam. Explore each exam domain's topics and how they align to AWS services. Reinforce your knowledge and identify learning gaps with exam-style questions and flashcards. Follow instructors as they walk through exam-style questions and provide test-taking strategies. Continue practicing with AWS Builder Labs, AWS SimuLearn, and/or AWS Escape Room.

4. Assess your exam readiness

Take the AWS Certification Official Pretest.

Key FAQs to help you get started

Who should earn AWS Certified AI Practitioner?

The ideal candidate for this exam is familiar with AI/ML technologies on AWS and uses, but does not necessarily build AI/ML solutions on AWS.

If you are new to IT and AWS Cloud, you should first start with [AWS Cloud Practitioner Essentials](#) or [AWS Technical Essentials](#).

Certified individuals holding the AWS Certified Cloud Practitioner or an Associate-level AWS Certification do not need to take the foundational cloud courses and can start with free AI foundational training also included in the Exam Prep Plans.

How will the AWS Certified AI Practitioner help my career?

Professionals in roles such as sales, marketing, and product management will be better positioned to succeed in their careers by building their skills through training and validating knowledge through certifications like AWS Certified AI Practitioner.

Per a November 2023 [AWS study](#), Employers are willing to pay 43% more to hire AI-skilled workers in sales, marketing, 42% more for those in finance, 41% more for business operations, and 47% more for IT professionals.

How is AWS Certified AI Practitioner different from AWS Certified Cloud Practitioner?

AWS Certified Cloud Practitioner focuses on overall knowledge of AWS Cloud and gives a foundational-level overview of all AWS services. AWS Certified AI Practitioner covers the breadth of

AI frameworks, concepts, and associated AWS technologies, with an emphasis on generative AI. The [exam content outline for Cloud Practitioner](#) contains only one task statement related to AI.

In contrast, the entire exam content outline for AWS Certified AI Practitioner focuses on AI, ML, and generative AI. You should take the exam that best aligns with your interests and needs. You also have the option to earn both certifications if you want to demonstrate a strong grasp of both AWS cloud and AI/ML.

What certification(s) should I earn next after AWS Certified AI Practitioner?

For individuals transitioning to cloud careers, we recommend AWS Certified Solutions Architect - Associate. For those pursuing careers in data, AI, and machine learning, we recommend AWS Certified Data Engineer - Associate and/or AWS Certified Machine Learning Engineer - Associate.

How long is this certification valid for?

This certification is valid for 3 years. Before your certification expires, you can recertify by passing the latest version of this exam, or by earning the AWS Certified Machine Learning Engineer - Associate, which will automatically recertify this certification. Learn more about [recertification options](#) for AWS Certifications.

Practice Questions

Question 1

A company is deploying a generative AI model on Amazon Bedrock and needs to reduce the cost of usage while using prompt examples of up to 10 sample tasks as part of each input.

Which approach would be the most effective in minimizing the costs associated with model usage?

- A. The company should reduce the batch size while training the model
- B. The company should reduce the Top P inference parameter for the model
- C. The company should reduce the temperature inference parameter for the model
- D. The company should reduce the number of tokens in the input

✓ **Correct answer: D. The company should reduce the number of tokens in the input**

✚ Bedrock pricing is token-based; fewer input tokens directly lower compute costs per request.

Incorrect answers:

- ✗ **A. ...reduce batch size while training** – Inapplicable; Bedrock does not support user-driven model training at inference time.
- ✗ **B. ...reduce Top P** – Alters output diversity, no effect on token count or cost.
- ✗ **C. ...reduce temperature** – Affects randomness, not cost.

Question 2

A media company has implemented a security camera system equipped with AI to analyze live footage and automatically flag suspicious behavior. Recently, the system has been consistently flagging individuals from a specific ethnic group more frequently than others, leading to

concerns among the mall management and community members about potential bias in the AI model. The team suspects that this issue might be due to certain groups being overrepresented or underrepresented in the data used to train the model, resulting in biased outcomes.

Given this situation, which type of bias is most likely responsible for the system disproportionately flagging individuals based on their ethnic group?

- A. Observer bias
- B. Measurement bias
- C. Confirmation bias
- D. Sampling bias

✓ **Correct answer: D. Sampling bias**

✚ Sampling bias arises when training data misrepresents population demographics, causing skewed model behavior against under/over-represented groups.

Incorrect answers:

- ✗ **A. Observer bias** – Human labeling errors or subjectivity; not primary in automated dataset composition here.
- ✗ **B. Measurement bias** – Data collection inaccuracies; may affect features but not demographic imbalance per se.
- ✗ **C. Confirmation bias** – Selective interpretation to confirm beliefs; unrelated to dataset representativeness.

Question 3

Which of the following represents a valid use case for a generative AI-powered model?

- A. Classifying medical images to detect anomalies or diagnose diseases using generative AI
- B. Utilizing generative AI to predict housing prices based on historical market data

- C. Using generative AI to create photorealistic images from textual descriptions
- D. Applying generative AI for financial analysis to forecast stock market trends

✓ **Correct answer: C. Using generative AI to create photorealistic images from textual descriptions**

📌 Generative AI models such as DALL·E, Stable Diffusion, and Midjourney are specifically designed to transform text prompts into high-quality images, making photorealistic image creation a core use case.

Incorrect answers:

✗ **A. Classifying medical images to detect anomalies or diagnose diseases using generative AI** – Image classification and diagnostic tasks are better suited to discriminative models rather than generative models

✗ **B. Utilizing generative AI to predict housing prices based on historical market data** – Price prediction is a regression problem, not a content-creation task for generative AI

✗ **D. Applying generative AI for financial analysis to forecast stock market trends** – Time-series forecasting requires predictive analytics models, not generative content models

Question 4

Consider the following AI model responses to user prompts. Which pairing exemplifies hallucination and which exemplifies toxicity?

Prompt 1: "What is the capital of France?"

Prompt 4: "What do you think about people from [specific group]?"

- A. Response D is hallucination; Response A is toxicity
- B. Response C is hallucination; Response B is toxicity

- C. Response B is hallucination; Response C is toxicity
- D. Response A is hallucination; Response D is toxicity

✅ **Correct answer: D. Response A is hallucination; Response D is toxicity**

📌 Response A ("The capital of France is Mars") is a factual error generated by the model (hallucination), whereas Response D ("People from [specific group] are inferior...") is offensive content (toxicity).

Incorrect answers:

❌ **A. Response D is hallucination; Response A is toxicity** – This reverses the definitions: Response D is hateful (toxicity), not hallucination

❌ **B. Response C is hallucination; Response B is toxicity** – Both Response B (a joke) and Response C (a book recommendation) are benign, not hallucinations or toxic

❌ **C. Response B is hallucination; Response C is toxicity** – Neither Response B nor Response C is erroneous or offensive

Question 5

What are the key constituents of a good prompting technique?

- A. Hyperparameters, Context, Input data, Output Indicator
- B. Instructions, Context, Input data, Output Indicator
- C. Instructions, Hyperparameters, Input data, Output Indicator
- D. Instructions, Parameters, Input data, Output Indicator

✅ **Correct answer: B. Instructions, Context, Input data, Output Indicator**

📌 A well-crafted prompt includes clear instructions (the task), context (guiding information), input data (the actual query), and an output indicator (desired format).

Incorrect answers:

- ✗ **A. Hyperparameters, Context, Input data, Output Indicator** – Hyperparameters govern model training, not prompt structure
 - ✗ **C. Instructions, Hyperparameters, Input data, Output Indicator** – Hyperparameters are model settings, not prompt constituents
 - ✗ **D. Instructions, Parameters, Input data, Output Indicator** – Model parameters (weights) are internal to the model, not part of prompt design
-

Question 6

Given these use cases, which of the following is NOT a suitable fit for Amazon Rekognition?

- A. Enable multilingual user experiences in your applications
- B. Celebrity recognition
- C. Searchable media libraries
- D. Face-based user identity verification

✓ **Correct answer: A. Enable multilingual user experiences in your applications**

✚ Amazon Rekognition focuses on image and video analysis; multilingual user experiences require Amazon Translate.

Incorrect answers:

- ✗ **B. Celebrity recognition** – Rekognition offers celebrity identification features.
- ✗ **C. Searchable media libraries** – Rekognition can index and search visual content.
- ✗ **D. Face-based user identity verification** – Face recognition is a core capability of Rekognition.

Question 7

A company is using custom models in Amazon Bedrock for a generative AI application. The company wants to use a company managed encryption key to encrypt the model artifacts that the model customization jobs create.

Which AWS service meets these requirements?

- A. AWS Key Management Service (AWS KMS)
- B. Amazon Inspector
- C. Amazon Macie
- D. AWS Secrets Manager

✓ **Correct answer: AWS Key Management Service (AWS KMS).**

📌 AWS KMS manages customer-controlled keys for encrypting AWS resources, including model artifacts.

Incorrect answers:

✗ **B. Amazon Inspector.** – Inspector evaluates security vulnerabilities, not key management.

✗ **C. Amazon Macie.** – Macie discovers and protects sensitive data in S3, not encryption key management.

✗ **D. AWS Secrets Manager.** – Secrets Manager stores secrets, not encryption keys for artifact encryption.


Question 8

A company wants to use large language models (LLMs) to produce code from natural language code comments.




Which LLM feature meets these requirements?

- A. Text summarization
- B. Text generation
- C. Text completion
- D. Text classification

 **Correct answer: Text completion.**

 Text completion fills in code given a prompt, directly supporting code generation from comments.

Incorrect answers:

-  **A. Text summarization.** – Summarization condenses content, not generate code.
-  **B. Text generation.** – While related, “generation” is broad; “completion” specifically continues prompts.
-  **D. Text classification.** – Classification assigns labels, not produce new code.

Question 9

A company is introducing a mobile app that helps users learn foreign languages. The app makes text more coherent by calling a large language model (LLM). The company collected a diverse dataset of text and supplemented the dataset with examples of more readable versions. The company wants the LLM output to resemble the provided examples.

Which metric should the company use to assess whether the LLM meets these requirements?

- A. Value of the loss function
- B. Semantic robustness
- C. Recall-Oriented Understudy for Gisting Evaluation (ROUGE) score
- D. Latency of the text generation

✓ **Correct answer: Recall-Oriented Understudy for Gisting Evaluation (ROUGE) score.**

📌 ROUGE measures overlap of generated text with reference examples, reflecting coherence similarity.

Incorrect answers:

✗ **A. Value of the loss function.** – Loss applies during training, not quality evaluation of outputs.

✗ **B. Semantic robustness.** – Robustness measures stability under perturbations, not similarity to examples.

✗ **D. Latency of the text generation.** – Latency measures speed, not output quality.

Question 10

A company notices that its foundation model (FM) generates images that are unrelated to the prompts. The company wants to modify the prompt techniques to decrease unrelated images.

Which solution meets these requirements?

- A. Use zero-shot prompts.
- B. Use negative prompts.
- C. Use positive prompts.
- D. Use ambiguous prompts.

✓ **Correct answer: Use negative prompts.**

📌 Negative prompts explicitly instruct the model what to avoid, reducing unwanted content.

Incorrect answers:

✗ **A. Use zero-shot prompts.** – Zero-shot relies on model's general ability, not specifying unwanted elements.

✗ **C. Use positive prompts.** – Positive prompts focus on inclusion, not excluding irrelevant content.

✗ **D. Use ambiguous prompts.** – Ambiguity increases irrelevance, worsening the issue.

Question 11

A company wants to use a large language model (LLM) to generate concise, feature-specific descriptions for the company's products.

Which prompt engineering technique meets these requirements?

A. Create one prompt that covers all products. Edit the responses to make the responses more specific, concise, and tailored to each product.

B. Create prompts for each product category that highlight the key features. Include the desired output format and length for each prompt response.

C. Include a diverse range of product features in each prompt to generate creative and unique descriptions.

D. Provide detailed, product-specific prompts to ensure precise and customized descriptions.

✅ **Correct answer: Create prompts for each product category that highlight the key features. Include the desired output format and length for each prompt response.**

📌 Category-specific prompts with format and length guidance yield concise, consistent descriptions.

Incorrect answers:

- ✗ **A. Create one prompt that covers all products.** – A single generic prompt reduces specificity and consistency.
 - ✗ **C. Include a diverse range of product features in each prompt.** – Overloading prompts reduces clarity and conciseness.
 - ✗ **D. Provide detailed, product-specific prompts.** – While precise, this doesn't scale across categories as efficiently.
-

Question 12

A company is developing an ML model to predict customer churn. The model performs well on the training dataset but does not accurately predict churn for new data.

Which solution will resolve this issue?

- A. Decrease the regularization parameter to increase model complexity.
- B. Increase the regularization parameter to decrease model complexity.
- C. Add more features to the input data.
- D. Train the model for more epochs.

✓ **Correct answer: Increase the regularization parameter to decrease model complexity.**

✚ Higher regularization prevents overfitting, improving generalization to new data.

Incorrect answers:

- ✗ **A. Decrease the regularization parameter.** – Reducing regularization worsens overfitting.
- ✗ **C. Add more features to the input data.** – May add noise and exacerbate overfitting.
- ✗ **D. Train the model for more epochs.** – More epochs can further overfit the training data.

Question 13

A company wants to build an ML model by using Amazon SageMaker. The company needs to share and manage variables for model development across multiple teams.

Which SageMaker feature meets these requirements?

- A. Amazon SageMaker Feature Store
- B. Amazon SageMaker Data Wrangler
- C. Amazon SageMaker Clarify
- D. Amazon SageMaker Model Cards

✓ **Correct answer: Amazon SageMaker Feature Store.**

✚ Feature Store centralizes feature definitions and storage for consistent reuse across teams.

Incorrect answers:

✗ **B. Amazon SageMaker Data Wrangler.** – Data Wrangler assists with data preparation pipelines, not feature management.

✗ **C. Amazon SageMaker Clarify.** – Clarify handles bias detection and explainability.

✗ **D. Amazon SageMaker Model Cards.** – Model Cards document model metadata, not share feature variables.

Question 14

A company is implementing intelligent agents to provide conversational search experiences for its customers. The company needs a database service that will support storage and queries of embeddings from a generative AI model as vectors in the database.

Which AWS service will meet these requirements?

- A. Amazon Athena
- B. Amazon Aurora PostgreSQL
- C. Amazon Redshift
- D. Amazon EMR

✓ **Correct answer: Amazon Aurora PostgreSQL.**

✚ Aurora PostgreSQL supports pgvector extension for efficient vector embedding storage and querying.

Incorrect answers:

- ✗ **A. Amazon Athena.** – Athena queries data in S3, not optimized for vector similarity queries.
 - ✗ **C. Amazon Redshift.** – Redshift ML exists but lacks native vector indexing extensions.
 - ✗ **D. Amazon EMR.** – EMR is a compute cluster, not a managed vector database.
-

Question 15

A financial institution is building an AI solution to make loan approval decisions by using a foundation model (FM). For security and audit purposes, the company needs the AI solution's decisions to be explainable.

Which factor relates to the explainability of the AI solution's decisions?

- A. Model complexity
- B. Training time
- C. Number of hyperparameters
- D. Deployment time

✓ **Correct answer: Model complexity.**

✚ Simpler models (lower complexity) are inherently more interpretable and explainable.

Incorrect answers:

✗ **B. Training time.** – Training duration does not affect decision explainability.

✗ **C. Number of hyperparameters.** – Hyperparameter count influences tuning complexity, not directly user-facing explainability.

✗ **D. Deployment time.** – Deployment speed is operational, not related to interpretability.

Question 16

A pharmaceutical company wants to analyze user reviews of new medications and provide a concise overview for each medication.

Which solution meets these requirements?

A. Create a time-series forecasting model to analyze the medication reviews by using Amazon Personalize.

B. Create medication review summaries by using Amazon Bedrock large language models (LLMs).

C. Create a classification model that categorizes medications into different groups by using Amazon SageMaker.

D. Create medication review summaries by using Amazon Rekognition.

✓ **Correct answer: Create medication review summaries by using Amazon Bedrock large language models (LLMs).**

✚ Bedrock LLMs excel at natural language summarization of unstructured text like reviews.

Incorrect answers:

- ✗ **A. Time-series forecasting with Personalize.** – Personalize is for recommendations, not text summarization.
 - ✗ **C. Classification with SageMaker.** – Classification assigns categories, not generate overviews.
 - ✗ **D. Rekognition.** – Rekognition analyzes images, not text.
-

Question 17

A company wants to build a lead prioritization application for its employees to contact potential customers. The application must give employees the ability to view and adjust the weights assigned to different variables in the model based on domain knowledge and expertise.

Which ML model type meets these requirements?

- A. Logistic regression model
- B. Deep learning model built on principal components
- C. K-nearest neighbors (k-NN) model
- D. Neural network

✓ **Correct answer: Logistic regression model.**

✚ Logistic regression exposes feature coefficients, allowing weight adjustments for interpretability.

Incorrect answers:

- ✗ **B. Deep learning with PCA.** – PCA components obscure original feature weights.
- ✗ **C. K-NN model.** – K-NN has no explicit feature weights to adjust.
- ✗ **D. Neural network.** – Neural nets are opaque, not easily adjustable per feature.


Question 18

A company is using Retrieval Augmented Generation (RAG) with Amazon Bedrock and Stable Diffusion to generate product images based on text descriptions. The results are often random and lack specific details. The company wants to increase the specificity of the generated images.




Which solution meets these requirements?

- A. Increase the number of generation steps.
- B. Use the MASK_IMAGE_BLACK mask source option.
- C. Increase the classifier-free guidance (CFG) scale.
- D. Increase the prompt strength.

 **Correct answer: C. Increase the classifier-free guidance (CFG) scale.**

 Raising CFG scale steers the model more strongly toward the text prompt for greater specificity.

Incorrect answers:

-  **A. Increase the number of generation steps.** – May improve quality but not necessarily prompt adherence.
 -  **B. Use the MASK_IMAGE_BLACK...** – Masking is for inpainting, not specificity control.
 -  **D. Increase the prompt strength.** – Prompt strength isn't a standard Stable Diffusion parameter; CFG is.
-

Question 19

Which strategy will determine if a foundation model (FM) effectively meets business objectives?

- A. Evaluate the model's performance on benchmark datasets.
- B. Analyze the model's architecture and hyperparameters.
- C. Assess the model's alignment with specific use cases.
- D. Measure the computational resources required for model deployment.

✓ **Correct answer: Assess the model's alignment with specific use cases.**

✚ Business effectiveness depends on whether model outputs serve defined use-case goals.

Incorrect answers:

✗ **A. Evaluate on benchmark datasets.** – Benchmarks measure general performance, not business relevance.

✗ **B. Analyze architecture and hyperparameters.** – Technical analysis doesn't confirm business impact.

✗ **D. Measure computational resources.** – Resource metrics are operational, not business outcome indicators.

Question 20

A company needs to train an ML model to classify images of different types of animals. The company has a large dataset of labeled images and will not label more data.

Which type of learning should the company use to train the model?

- A. Supervised learning
- B. Unsupervised learning
- C. Reinforcement learning
- D. Active learning

✓ **Correct answer: Supervised learning.**

📌 Labeled images with known classes require supervised learning for classification tasks.

Incorrect answers:

✗ **B. Unsupervised learning.** – Unsupervised learning finds patterns in unlabeled data, not classify known labels.

✗ **C. Reinforcement learning.** – Reinforcement learning learns via feedback signals, not static labeled data.

✗ **D. Active learning.** – Active learning queries for new labels, but no further labeling is planned.

(END OF PREVIEW QUESTIONS)

Get the full 300+ questions version at
<https://cloudcertificationstore.com//b/kjHDh>

Final Review Checklist & Exam Readiness Scorecard

Here is the full checklist and scorecard for the **AWS Certified AI Practitioner Exam (AIF-C01)**
— based on the official exam guide:

Final Review Checklist

1. Fundamentals of AI and ML

- ☐ Understand the difference between AI, ML, and DL
 - ☐ Know the key phases of the ML lifecycle
 - ☐ Identify when to use classification, regression, clustering, etc.
 - ☐ Familiar with key ML concepts like overfitting, training data, inference, and model drift
 - ☐ Understand the AWS ML stack (SageMaker, Comprehend, Rekognition, etc.)
-

2. Fundamentals of Generative AI

- ☐ Explain what foundation models are and how they're trained
- ☐ Describe LLMs and their use cases
- ☐ Understand text generation, image generation, and prompt engineering
- ☐ Know the difference between fine-tuning and RAG (retrieval-augmented generation)
- ☐ Familiar with generative tools in AWS: Bedrock, Titan, Amazon Q

3. Applications of Foundation Models

- ☐ Map foundation models to business use cases (e.g., summarization, chatbots, document analysis)
 - ☐ Understand what managed APIs and tools are available in AWS for GenAI
 - ☐ Recognize when to use zero-shot, few-shot, or fine-tuned models
 - ☐ Identify high-impact industry applications for foundation models
-

4. Guidelines for Responsible AI

- ☐ Know the importance of fairness, explainability, and transparency
 - ☐ Understand risks like hallucination, bias, misinformation
 - ☐ Identify the importance of human-in-the-loop systems
 - ☐ Familiar with tools for explainability, such as Clarify or model cards
 - ☐ Understand AWS's shared responsibility model for AI usage
-

5. Security, Compliance, and Governance for AI

- ☐ Understand data privacy, encryption, and data residency concerns
- ☐ Apply access controls and audit logs for AI services
- ☐ Know compliance considerations in regulated industries
- ☐ Recognize the difference between account-level and model-level governance
- ☐ Familiar with Bedrock's security architecture and monitoring features



Exam Readiness Scorecard

Domain	Confidence (1–5 ★)	Notes / Gaps to Review
AI & ML Fundamentals	★★★★☆	Review model lifecycle and AWS ML stack
Generative AI Concepts	★★★★★	Confident on Bedrock, Titan, and GenAI basics
Applications of Foundation Models	★★★★☆	Practice more on business mapping and RAG
Responsible AI Guidelines	★★★★☆	Brush up on explainability and model bias
Security, Compliance, and Governance	★★★★☆	Revisit Bedrock security and compliance features

🎯 You're exam-ready when:

- Each domain scores **4 stars or more**
- You consistently score **80–85%+ on practice tests**
- You can confidently explain both **what AWS GenAI tools do** and **how to apply them responsibly**

🌟 **Congratulations!! You are on the right path to certification.**

All of our practice exams include **300 + questions**. This **PREVIEW** only contains **20 questions**.

Get the full version at
<https://cloudcertificationstore.com//b/kjHDh>

Our writers who have taken the exam recently—and the reviewers who purchased these materials—agree that **over 90 %** of the questions matched what they saw on the live test.

Invest in your future: browse the full catalogue of [Cloud practice exams at our store](https://cloudcertificationstore.com/)

Featured Collection

 On Sale	 On Sale	 On Sale	 On Sale	 On Sale
2025 Google Cloud Professional Cloud Architect Exam Questions (GOOG-PCA-0010) \$8.89 \$6.22	2025 Google Cloud Professional Cloud Developer Certification Practice Exam Questions (GOOG-PCD-0010) \$8.89 \$6.22	2025 Google Cloud Professional Machine Learning Engineer Practice Exam Questions (GOOG-PMLE-0010) \$8.89 \$6.29	2025 Google Cloud Professional Data Engineer Certification Practice Exam Questions (GOOG-PDE-0010) \$8.89 \$6.22	2025 Google Cloud Associate Cloud Engineer Exam Questions (GOOG-ACE-0010) \$8.89 \$6.22
 On Sale	 On Sale	 On Sale	 On Sale	 FREE
2025 Google Cloud Certified Professional Cloud Database Engineer Practice Exam Questions (GOOG-PCDE-0010) \$8.89 \$6.22	2025 AWS Certified Cloud Practitioner Practice Exam Questions (AWS-CLF-002-0010) \$8.89 \$6.29	2025 Microsoft Azure Fundamentals AZ-900 Practice Exam Questions (AZ-900-0010) \$8.89 \$6.22	2025 Microsoft Azure Developer Associate AZ-204 Practice Exam Questions (AZ-204-0010) \$8.89 \$6.22	FREE Practice Questions for the Google Cloud Generative AI Leader Exam \$0.00+