



TOP 10

LLM APPLICATIONS
& GENERATIVE AI

LLM AI Cybersecurity & Governance Checklist

English

Version 1.1
April, 2024

Revision History

Revision	Date	Author(s)	Description
0.1	2023-11-01	Sandy Dunn	initial draft
0.5	2023-12-06	SD, Team	public draft
0.9	2023-02-15	SD, Team	pre-release draft
1.0	2024-02-19	SD, Team	public release v 1.0
1.1	2024-04-10	SD, Team	public release v 1.1

The information provided in this document does not, and is not intended to, constitute legal advice. All information is for general informational purposes only.

This document contains links to other third-party websites. Such links are only for convenience and OWASP does not recommend or endorse the contents of the third-party sites.



1	Overview	5
1.1	Responsible and Trustworthy Artificial Intelligence	6
1.2	Who is This For?	7
1.3	Why a Checklist?	7
1.4	Not Comprehensive	7
1.5	Large Language Model Challenges	7
1.6	LLM Threat Categories	8
1.7	Artificial Intelligence Security and Privacy Training	9
1.8	Incorporate LLM Security and governance with Existing, Established Practices and Controls	9
1.9	Fundamental Security Principles	9
1.10	Risk	10
1.11	Vulnerability and Mitigation Taxonomy	10
2	Determining LLM Strategy	11
2.1	Deployment Strategy	13
3	Checklist	14
3.1	Adversarial Risk	14
3.2	Threat Modeling	14
3.3	AI Asset Inventory	14
3.4	AI Security and Privacy Training	15
3.5	Establish Business Cases	15
3.6	Governance	16
3.7	Legal	17
3.8	Regulatory	18
3.9	Using or Implementing Large Language Model Solutions	19
3.10	Testing, Evaluation, Verification, and Validation <i>TEVV</i>	19
3.11	Model Cards and Risk Cards	20
3.12	RAG: Large Language Model Optimization	21
3.13	AI Red Teaming	21

4	Resources	22
A	Team	32

Overview

Every internet user and company should prepare for the upcoming wave of powerful generative artificial intelligence (GenAI) applications. GenAI has enormous promise for innovation, efficiency, and commercial success across a variety of industries. Still, like any powerful early stage technology, it brings its own set of obvious and unexpected challenges.

Artificial intelligence has advanced greatly over the last 50 years, inconspicuously supporting a variety of corporate processes until ChatGPT's public appearance drove the development and use of Large Language Models (LLMs) among both individuals and enterprises. Initially, these technologies were limited to academic study or the execution of certain, but vital, activities within corporations, visible only to a select few. However, recent advances in data availability, computer power, GenAI capabilities, and the release of tools such as Llama 2, ElevenLabs, and Midjourney have raised AI from a niche to general widespread acceptance. These improvements have not only made GenAI technologies more accessible, but they have also highlighted the critical need for enterprises to develop solid strategies for integrating and exploiting AI in their operations, representing a huge step forward in how we use technology.

- **Artificial intelligence (AI)** is a broad term that encompasses all fields of computer science that enable machines to accomplish tasks that would normally require human intelligence. Machine learning and generative AI are two subcategories of AI.
- **Machine learning (ML)** is a subset of AI that focuses on creating algorithms that can learn from data. Machine learning algorithms are trained on a set of data, and then they can use that data to make predictions or decisions about new data.
- **Generative AI** is a type of machine learning that focuses on creating new data. Often, GenAI relies on the use of large language models to perform the tasks needed to create the new data.
- A **large language model (LLM)** is a type of AI model that processes and generates human-like text. In the context of artificial intelligence a "model" refers to a system that is trained to make predictions based on input data. LLMs are specifically trained on large data sets of natural language and the name large language models.

Organizations are entering uncharted territory in securing and overseeing GenAI solutions. The rapid advancement of GenAI also opens doors for adversaries to enhance their attack strategies, introducing a dual challenge of defense and threat escalation.

Businesses use artificial intelligence in many areas, including HR for recruiting, email spam screening, SIEM for behavioral analytics, and managed detection and response applications. However, this document's primary focus is on Large Language Model applications and their function in creating generated content.

Responsible and Trustworthy Artificial Intelligence

As challenges and benefits of Artificial Intelligence emerge - and regulations and laws are passed - the principles and pillars of responsible and trustworthy AI usage are evolving from idealistic objects and concerns to established standards. The OWASP AI Exchange Working Group is monitoring these changes and addressing the broader and more challenging considerations for all aspects of artificial intelligence.

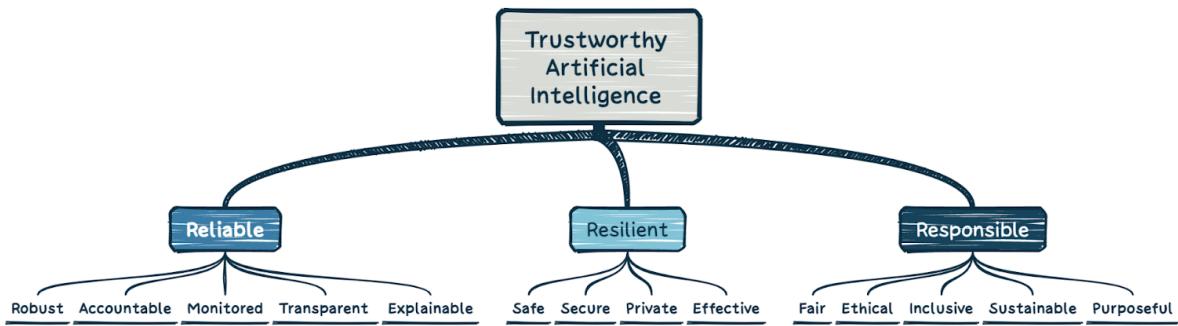


Figure 1.1: Image depicting Pillars of Trustworthy Artificial Intelligence: created from Montreal Ethics Institute Example

Who is This For?

The OWASP Top 10 for LLM Applications Cybersecurity and Governance Checklist is for leaders across executive, tech, cybersecurity, privacy, compliance, and legal areas, DevSecOps, MLSecOps, and Cybersecurity teams and defenders. It is intended for people who are striving to stay ahead in the fast-moving AI world, aiming not just to leverage AI for corporate success but also to protect against the risks of hasty or insecure AI implementations. These leaders and teams must create tactics to grab opportunities, combat challenges, and mitigate risks.

This checklist is intended to help these technology and business leaders quickly understand the risks and benefits of using LLM, allowing them to focus on developing a comprehensive list of critical areas and tasks needed to defend and protect the organization as they develop a Large Language Model strategy.

It is the hope of the OWASP Top 10 for the LLM Applications team that this list will help organizations improve their existing defensive techniques and develop techniques to address the new threats that come from using this exciting technology.

Why a Checklist?

Checklists used to formulate strategies improve accuracy, define objectives, preserve uniformity, and promote focused deliberate work, reducing oversights and missed details. Following a check list not only increases trust in a safe adoption journey, but also encourages future organizations innovations by providing a simple and effective strategy for continuous improvement.

Not Comprehensive

Although this document intends to support organizations in developing an initial LLM strategy in a rapidly changing technical, legal, and regulatory environment, it is not exhaustive and does not cover every use case or obligation. While using this document is Organizations should extend assessments and practices beyond the scope of the provided checklist as required for their use case or jurisdiction.

Large Language Model Challenges

Large Language models face several serious and unique issues. One of the most important is that while working with LLMs, the control and data planes cannot be strictly isolated or separable. Another significant challenge is that LLMs are nondeterministic by design, yielding a different outcome when prompted or requested. LLMs employ semantic search rather than keyword search. The key distinction between the two is that the model's algorithm prioritizes the terms in its response. This is a significant departure from how consumers have previously used technology, and it has an impact on the consistency and reliability of the findings. Hallucinations, emerging from the gaps and training flaws in the data the model is trained on, are the result of this method.

There are methods to improve reliability and reduce the attack surface for jailbreaking, model tricking, and hallucinations, but there is a trade-off between restrictions and utility in both cost and functionality.

LLM use and LLM applications increase an organization's attack surface. Some risks associated

with LLMs are unique, but many are familiar issues, such as the known software bill of materials (SBoM), supply chain, data loss protection (DLP), and authorized access. There are also increased risks not directly related to GenAI, but GenAI increases the efficiency, capability, and effectiveness of attackers who attack and threaten organizations.

Adversaries are increasingly harnessing LLM and Generative AI tools to refine and expedite traditional methods of attacking organizations, individuals, and government systems. LLM facilitates their ability to enhance techniques allowing them to effortlessly craft new malware, potentially embedded with novel zero-day vulnerabilities or designed to evade detection. They can also generate sophisticated, unique, or tailored phishing schemes. The creation of convincing deep fakes, whether video or audio, further promotes their social engineering ploys. Additionally, these tools enable them to execute intrusions and develop innovative hacking capabilities. In the future, more “tailored” and compound use of AI technology by criminal actors will demand specific responses and dedicated solutions for an organization’s appropriate defense and resilience capabilities.

Organizations also face the threat of NOT utilizing the capabilities of LLMs such as a competitive disadvantage, market perception by customers and partners of being outdated, inability to scale personalized communications, innovation stagnation, operational inefficiencies, the higher risk of human error in processes, and inefficient allocation of human resources.

Understanding the different kinds of threats and integrating them with the business strategy will help weigh both the pros and cons of using Large Language Models (LLMs) against not using them, making sure they accelerate rather than hinder the business's meeting business objectives.

LLM Threat Categories



Figure 1.2: Image depicting the types of AI threats: credit sdunn

Artificial Intelligence Security and Privacy Training

Employees throughout organizations benefit from training to understand artificial intelligence, generative artificial intelligence, and the future potential consequences of building, buying, or utilizing LLMs. Training for permissible use and security awareness should target all employees as well as be more specialized for certain positions such as human resources, legal, developers, data teams, and security teams.

Fair use policies and healthy interaction are key aspects that, if incorporated from the very start, will be a cornerstone to the success of future AI cybersecurity awareness campaigns. This will necessarily provide user's with knowledge of the basic rules for interaction as well as the ability to separate good behavior from bad or unethical behavior.

Incorporate LLM Security and governance with Existing, Established Practices and Controls

While AI and generated AI add a new dimension to cybersecurity, resilience, privacy, and meeting legal and regulatory requirements, the best practices that have been around for a long time are still the best way to identify issues, find vulnerabilities, fix them, and mitigate potential security issues.

- Confirm the management of artificial intelligence systems is integrated with existing organizational practices.
- Confirm AIML systems follow existing privacy, governance, and security practices, with AI specific privacy, governance, and security practices implemented when required.

Fundamental Security Principles

LLM capabilities introduce a different type of attack and attack surface. LLMs are vulnerable to complex business logic bugs, such as prompt injection, insecure plugin design, and remote code execution. Existing best practices are the best way to solve these issues. An internal product security team that understands secure software review, architecture, data governance, and third-party assessments. The cybersecurity team should also check how strong the current controls are to find problems that could be made worse by LLM, such as voice cloning, impersonation, or bypassing captchas.

Given recent advancements in machine learning, NLP (Natural Language Processing), NLU (Natural Language Understanding), Deep Learning, and more recently, LLMs (Large Language Models) and Generative AI, it is recommended to include professionals proficient in these areas alongside cybersecurity and devops teams. Their expertise will not only aid in adopting these technologies but also in developing innovative analyses and responses to emerging challenges.

Risk

Reference to risk uses the ISO 31000 definition: Risk = "effect of uncertainty on objectives." LLM risks included in the checklist includes a targeted list of LLM risks that address adversarial, safety, legal, regulatory, reputation, financial, and competitive risks.

Vulnerability and Mitigation Taxonomy

Current systems for classifying vulnerabilities and sharing threat information, like OVAL, STIX, CVE, and CWE, are still developing the ability to monitor and alert defenders about vulnerabilities and threats specific to Large Language Models (LLMs) and Predictive Models. It is expected that organizations will lean on these established and recognized standards, such as CVE for vulnerability classification and STIX for the exchange of cyber threat intelligence (CTI), when vulnerabilities or threats to AI/ML systems and their supply chains are identified.

Determining LLM Strategy

The rapid expansion of Large Language Model (LLM) applications has heightened the attention and examination of all AI/ML systems used in business operations, encompassing both Generative AI and long-established Predictive AI/ML systems. This increased focus exposes potential risks, such as attackers targeting systems that were previously overlooked and governance or legal challenges that may have been disregarded in terms of legal, privacy, liability, or warranty issues. For any organization leveraging AI/ML systems in its operations, it's critical to assess and establish comprehensive policies, governance, security protocols, privacy measures, and accountability standards to ensure these technologies align with business processes securely and ethically.

Attackers, or adversaries, provide the most immediate and harmful threat to enterprises, people, and government agencies. Their goals, which range from financial gain to espionage, push them to steal critical information, disrupt operations, and damage confidence. Furthermore, their ability to harness new technologies such as AI and machine learning increases the speed and sophistication of attacks, making it difficult for defenses to stay ahead of attacks.

The most pressing non-adversary LLM threat for many organizations stem from "Shadow AI": employees using unapproved online AI tools, unsafe browser plugins, and third-party applications that introduce LLM features via updates or upgrades, circumventing standard software approval processes.

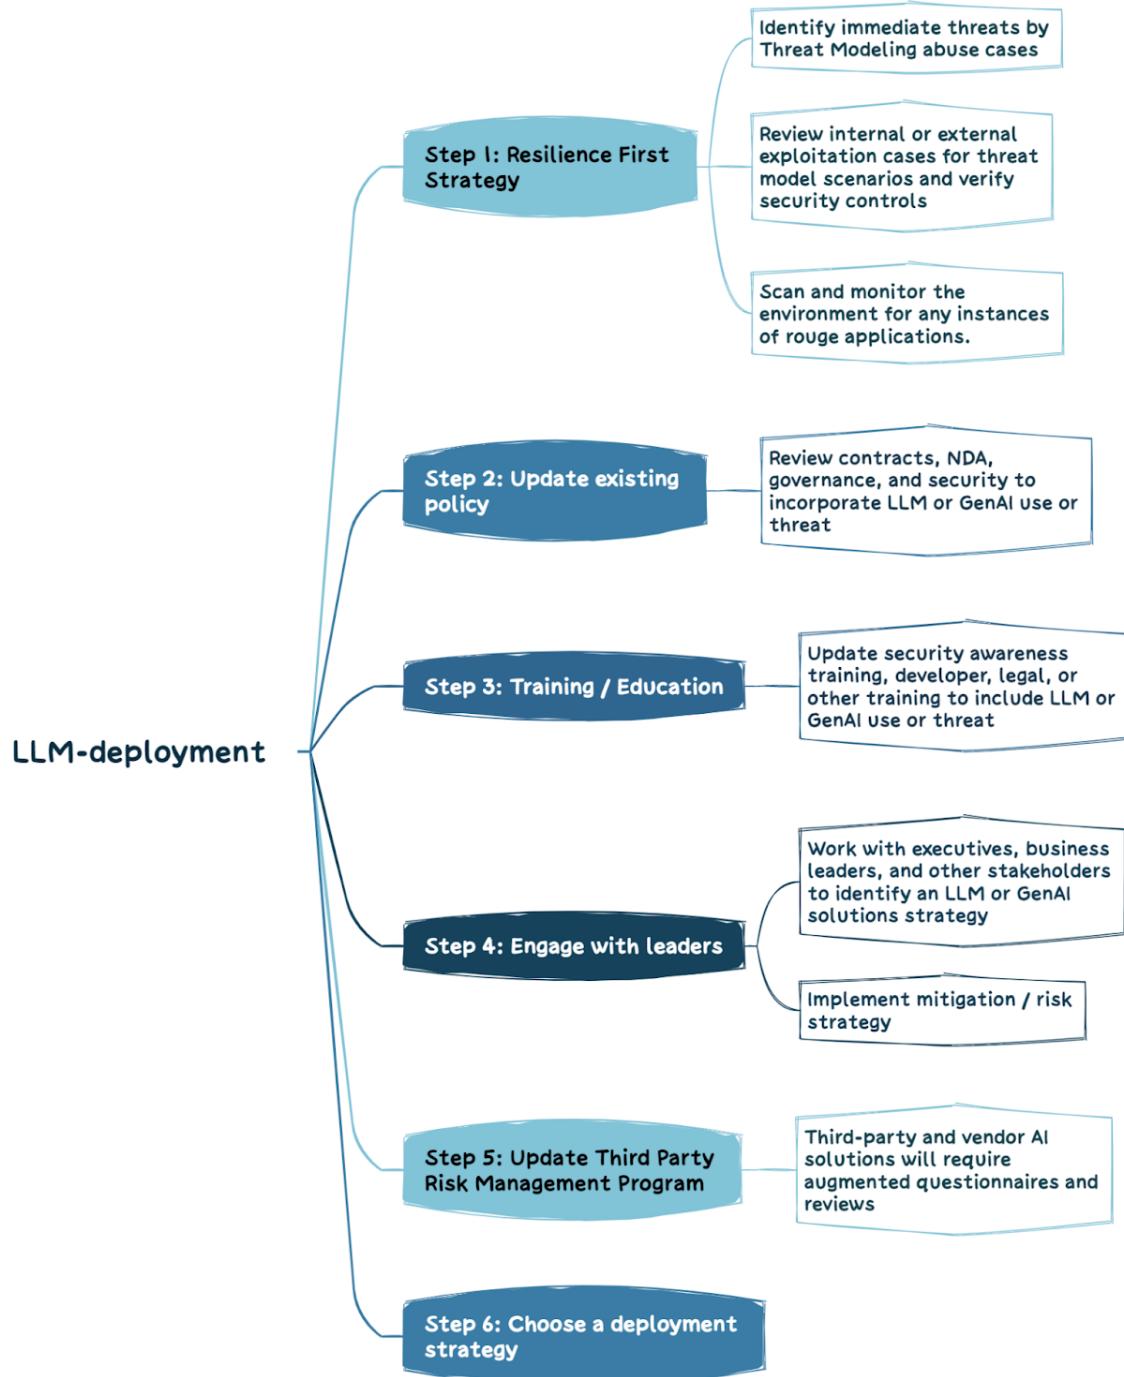


Figure 2.1: Image of options for deployment strategy: credit sdunn

Deployment Strategy

The scopes range from leveraging public consumer applications to training proprietary models on private data. Factors like use case sensitivity, capabilities needed, and resources available help determine the right balance of convenience vs. control. However, understanding these five model types provides a framework for evaluating options.

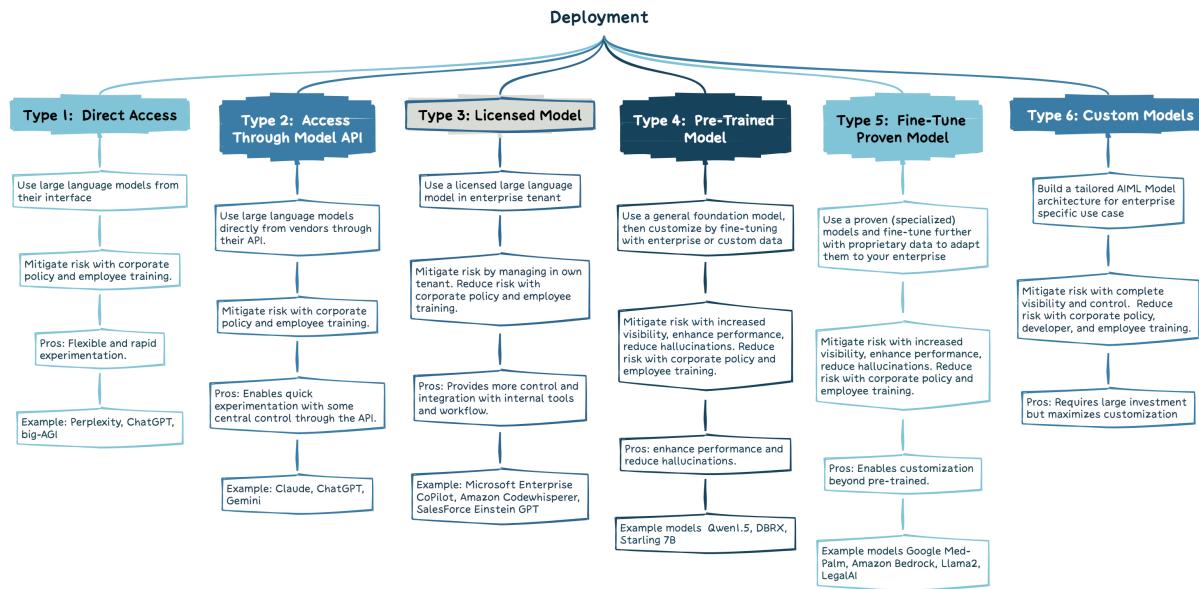


Figure 2.2: Image of options for deployment types: credit sdunn

Checklist

Adversarial Risk

Adversarial Risk includes competitors and attackers.

- Scrutinize how competitors are investing in artificial intelligence. Although there are risks in AI adoption, there are also business benefits that may impact future market positions.
- Investigate the impact of current controls, such as password resets, which use voice recognition which may no longer provide the appropriate defensive security from new GenAI enhanced attacks.
- Update the Incident Response Plan and playbooks for GenAI enhanced attacks and AIML specific incidents.

Threat Modeling

Threat modeling is highly recommended to identify threats and examine processes and security defenses. Threat modeling is a set of systematic, repeatable processes that enable making reasonable security decisions for applications, software, and systems. Threat modeling for GenAI accelerated attacks and before deploying LLMs is the most cost effective way to Identify and mitigate risks, protect data, protect privacy, and ensure a secure, compliant integration within the business.

- How will attackers accelerate exploit attacks against the organization, employees, executives, or users? Organizations should anticipate "hyper-personalized" attacks at scale using Generative AI. LLM-assisted Spear Phishing attacks are now exponentially more effective, targeted, and weaponized for an attack.
- How could GenAI be used for attacks on the business's customers or clients through spoofing or GenAI generated content?
- Can the business detect and neutralize harmful or malicious inputs or queries to LLM solutions?
- Can the business safeguard connections with existing systems and databases with secure integrations at all LLM trust boundaries?
- Does the business have insider threat mitigation to prevent misuse by authorized users?
- Can the business prevent unauthorized access to proprietary models or data to protect Intellectual Property?
- Can the business prevent the generation of harmful or inappropriate content with automated content filtering?

AI Asset Inventory

An AI asset inventory should apply to both internally developed and external or third-party solutions.

- Catalog existing AI services, tools, and owners. Designate a tag in asset management for specific inventory.
- Include AI components in the Software Bill of Material (SBOM), a comprehensive list of all the software components, dependencies, and metadata associated with applications.
- Catalog AI data sources and the sensitivity of the data (protected, confidential, public)
- Establish if pen testing or red teaming of deployed AI solutions is required to determine the current attack surface risk.
- Create an AI solution onboarding process.
- Ensure skilled IT admin staff is available either internally or externally, following SBOM requirements.

AI Security and Privacy Training

- Actively engage with employees to understand and address concerns with planned LLM initiatives.
- Establish a culture of open, and transparent communication on the organization's use of predictive or generative AI within the organization process, systems, employee management and support, and customer engagements and how its use is governed, managed, and risks addressed.
- Train all users on ethics, responsibility, and legal issues such as warranty, license, and copyright.
- Update security awareness training to include GenAI related threats. Voice cloning and image cloning, as well as in anticipation of increased spear phishing attacks
- Any adopted GenAI solutions should include training for both DevOps and cybersecurity for the deployment pipeline to ensure AI safety and security assurances.

Establish Business Cases

Solid business cases are essential to determining the business value of any proposed AI solution, balancing risk and benefits, and evaluating and testing return on investment. There are an enormous number of potential use cases; a few examples are provided.

- Enhance customer experience
- Better operational efficiency
- Better knowledge management
- Enhanced innovation
- Market Research and Competitor Analysis
- Document creation, translation, summarization, and analysis

Governance

Corporate governance in LLM is needed to provide organizations with transparency and accountability. Identifying AI platform or process owners who are potentially familiar with the technology or the selected use cases for the business is not only advised but also necessary to ensure adequate reaction speed that prevents collateral damages to well established enterprise digital processes.

- Establish the organization's AI RACI chart (who is responsible, who is accountable, who should be consulted, and who should be informed)
- Document and assign AI risk, risk assessments, and governance responsibility within the organization.
- Establish data management policies, including technical enforcement, regarding data classification and usage limitations. Models should only leverage data classified for the minimum access level of any user of the system. For example, update the data protection policy to emphasize not to input protected or confidential data into nonbusiness-managed tools.
- Create an AI Policy supported by established policy (e.g., standard of good conduct, data protection, software use)
- Publish an acceptable use matrix for various generative AI tools for employees to use.
- Document the sources and management of any data that the organization uses from the generative LLM models.

Legal

Many of the legal implications of AI are undefined and potentially very costly. An IT, security, and legal partnership is critical to identifying gaps and addressing obscure decisions.

- Confirm product warranties are clear in the product development stream to assign who is responsible for product warranties with AI.
- Review and update existing terms and conditions for any GenAI considerations.
- Review AI EULA agreements. End-user license agreements for GenAI platforms are very different in how they handle user prompts, output rights and ownership, data privacy, compliance, liability, privacy, and limits on how output can be used.
- Organizations EULA for customers, Modify end-user agreements to prevent the organization from incurring liabilities related to plagiarism, bias propagation, or intellectual property infringement through AI-generated content.
- Review existing AI-assisted tools used for code development. A chatbot's ability to write code can threaten a company's ownership rights to its product if a chatbot is used to generate code for the product. For example, it could call into question the status and protection of the generated content and who holds the right to use the generated content.
- Review any risks to intellectual property. Intellectual property generated by a chatbot could be in jeopardy if improperly obtained data was used during the generative process, which is subject to copyright, trademark, or patent protection. If AI products use infringing material, it creates a risk for the outputs of the AI, which may result in intellectual property infringement.
- Review any contracts with indemnification provisions. Indemnification clauses try to put the responsibility for an event that leads to liability on the person who was more at fault for it or who had the best chance of stopping it. Establish guardrails to determine whether the provider of the AI or its user caused the event, giving rise to liability.
- Review liability for potential injury and property damage caused by AI systems.
- Review insurance coverage. Traditional (D&O) liability and commercial general liability insurance policies are likely insufficient to fully protect AI use.
- Identify any copyright issues. Human authorship is required for copyright. An organization may also be liable for plagiarism, propagation of bias, or intellectual property infringement if LLM tools are misused.
- Ensure agreements are in place for contractors and appropriate use of AI for any development or provided services.
- Restrict or prohibit the use of generative AI tools for employees or contractors where enforceable rights may be an issue or where there are IP infringement concerns.
- Assess and AI solutions used for employee management or hiring could result in disparate treatment claims or disparate impact claims.
- Make sure the AI solutions do not collect or share sensitive information without proper consent or authorization.

Regulatory

The EU AI Act is anticipated to be the first comprehensive AI law but will apply in 2025 at the earliest. The EU's General Data Protection Regulation (GDPR) does not specifically address AI but includes rules for data collection, data security, fairness and transparency, accuracy and reliability, and accountability, which can impact GenAI use. In the United States, AI regulation is included within broader consumer privacy laws. Ten US states have passed laws or have laws that will go into effect by the end of 2023. Canada has so far only published a Voluntary Code of Conduct on the Responsible Development and Management of Advanced Generative AI Systems, however, the Artificial Intelligence and Data Act (AIDA) will have stronger requirements. Federal organizations such as the US Equal Employment Opportunity Commission (EEOC), the Consumer Financial Protection Bureau (CFPB), the Federal Trade Commission (FTC), and the US Department of Justice's Civil Rights Division (DOJ) are closely monitoring hiring fairness.

- Determine Country, State, or other Government specific AI compliance requirements.
- Determine compliance requirements for restricting electronic monitoring of employees and employment-related automated decision systems (Vermont, California, Maryland, New York, New Jersey)
- Determine compliance requirements for consent for facial recognition and the AI video analysis required (Illinois, Maryland, Washington, Vermont)
- Review any AI tools in use or being considered for employee hiring or management.
- Confirm the vendor's compliance with applicable AI laws and best practices.
- Ask and document any products using AI during the hiring process. Ask how the model was trained, and how it is monitored, and track any corrections made to avoid discrimination and bias.
- Ask and document what accommodation options are included.
- Ask and document whether the vendor collects confidential data.
- Ask how the vendor or tool stores and deletes data and regulates the use of facial recognition and video analysis tools during pre-employment.
- Review other organization-specific regulatory requirements with AI that may raise compliance issues. The Employee Retirement Income Security Act of 1974, for instance, has fiduciary duty requirements for retirement plans that a chatbot might not be able to meet.

Using or Implementing Large Language Model Solutions

- Threat Model LLM components and architecture trust boundaries.
- Data Security, verify how data is classified and protected based on sensitivity, including personal and proprietary business data. (How are user permissions managed, and what safeguards are in place?)
- Access Control, implement least privilege access controls and implement defense-in-depth measures
- Training Pipeline Security, require rigorous control around training data governance, pipelines, models, and algorithms.
- Input and Output Security, evaluate input validation methods, as well as how outputs are filtered, sanitized, and approved.
- Monitoring and Response, map workflows, monitoring, and responses to understand automation, logging, and auditing. Confirm audit records are secure.
- Include application testing, source code review, vulnerability assessments, and red teaming in the production release process.
- Check for existing vulnerabilities in the LLM model or supply chain.
- Look into the effects of threats and attacks on LLM solutions, such as prompt injection, the release of sensitive information, and process manipulation.
- Investigate the impact of attacks and threats to LLM models, including model poisoning, improper data handling, supply chain attacks, and model theft.
- Supply Chain Security, request third-party audits, penetration testing, and code reviews for third-party providers. (both initially and on an ongoing basis)
- Infrastructure Security, ask how often a vendor performs resilience testing? What are their SLAs in terms of availability, scalability, and performance?
- Update incident response playbooks and include an LLM incident in tabletop exercises.
- Identify or expand metrics to benchmark generative cybersecurity AI against other approaches to measure expected productivity improvements.

Testing, Evaluation, Verification, and Validation *TEVV*

NIST AI Framework recommends a continuous TEVV process throughout the AI lifecycle which includes the AI system operators, domain experts, AI designers, users, product developers, evaluators, and auditors. TEVV includes a range of tasks such as system validation, integration, testing, recalibration, and ongoing monitoring for periodic updates to navigate the risks and changes of the AI system.

- Establish continuous testing, evaluation, verification, and validation throughout the AI model lifecycle.
- Provide regular executive metrics and updates on AI Model functionality, security, reliability, and robustness.

Model Cards and Risk Cards

Model cards and risk cards are foundational elements for increasing the transparency, accountability, and ethical deployment of Large Language Models (LLMs). Model cards help users understand and trust AI systems by providing standardized documentation on their design, capabilities, and constraints, leading them to make educated and safe applications. Risk cards supplement this by openly addressing potential negative consequences, such as biases, privacy problems, and security vulnerabilities, which encourages a proactive approach to harm prevention. These documents are critical for developers, users, regulators, and ethicists equally since they establish a collaborative atmosphere in which AI's social implications are carefully addressed and handled. These cards, developed and maintained by the organizations that created the models, play an important role in ensuring that AI technologies fulfill ethical standards and legal requirements, allowing for responsible research and deployment in the AI ecosystem.

Model cards include key attributes associated with the ML model:

- **Model details:** Basic information about the model, i.e., name, version, and type (neural network, decision tree, etc.), and the intended use case.
- **Model architecture:** Includes a description of the structure of the model, such as the number and type of layers, activation functions, and other key architectural choices.
- **Training data and methodology:** Information about the data used to train the model, such as the size of the dataset, the data sources, and any preprocessing or data augmentation techniques used. It also includes details about the training methodology, such as the optimizer used, the loss function, and any hyperparameters that were tuned.
- **Performance metrics:** Information about the model's performance on various metrics, such as accuracy, precision, recall, and F1 score. It may also include information about how the model performs on different subsets of the data.
- **Potential biases and limitations:** Lists potential biases or limitations of the model, such as imbalanced training data, overfitting, or biases in the model's predictions. It may also include information about the model's limitations, such as its ability to generalize to new data or its suitability for certain use cases.
- **Responsible AI considerations:** Any ethical or responsible AI considerations related to the model, such as privacy concerns, fairness, and transparency, or potential societal impacts of the model's use. It may also include recommendations for further testing, validation, or monitoring of the model.

The precise features contained in a model card may differ based on the model's context and intended usage, but the purpose is to give openness and accountability in the creation and deployment of machine learning models.

- Review a model's model card
- Review risk card if available
- Establish a process to track and maintain model cards for any deployed model including models used through a third party.

RAG: Large Language Model Optimization

Fine tuning, the traditional method for optimizing a pre-trained model, involved retraining an existing model on new, and domain-specific data, modifying it for performance on a task or application. Fine-tuning is expensive but essential to improve performance.

Retrieval-Augmented Generation *RAG* has evolved as a more effective way of optimizing and augmenting the capabilities of large language models by retrieving pertinent data from up to date available knowledge sources. RAG can be customized for specific domains, optimizing the retrieval of domain-specific information and tailoring the generation process to the nuances of specialized fields. RAG is seen as a more efficient and transparent method for LLM optimization, particularly for problems where labeled data is limited or expensive to collect. One of the primary advantages of RAG is its support for continuous learning since new information can be continually updated at the retrieval stage.

The RAG implementation involves several key steps starting from embedding model deployment, indexing the knowledge library, to retrieving the most relevant documents for query processing. Efficient retrieval of the relevant context is made based on vector databases which are used for storage and querying of document embeddings.

RAG Reference

- Retrieval Augmented Generation *RAG* & LLM: Examples
- 12 RAG Pain Points and Proposed Solutions

AI Red Teaming

AI Red Teaming is an adversarial attack test simulation of the AI System to validate there aren't any existing vulnerabilities which can be exploited by an attacker. It is a recommended practice by many regulatory and AI governing bodies including the Biden administration. Red-teaming alone is not a comprehensive solution to validate all real-world harms associated with AI systems and should be included with other forms of testing, evaluation, verification, and validation such as algorithmic impact assessments and external audits.

- Incorporate Red Team testing as a standard practice for AI Models and applications.

Resources

OWASP Top 10 for Large Language Model Applications

OWASP Top 10 for LLM									
LLM01 Prompt Injection This manipulates a large language model (LLM) through crafty inputs, causing unintended actions by the LLM. Direct injections overwrite system prompts, while indirect ones manipulate inputs from external sources.	LLM02 Insecure Output Handling This vulnerability occurs when an LLM output is accepted without scrutiny, exposing backend systems. Misuse may lead to severe consequences like XSS, CSRF, SSRF, privilege escalation, or remote code execution.	LLM03 Training Data Poisoning Training data poisoning refers to manipulating the data or fine-tuning process to introduce vulnerabilities, backdoors or biases that could compromise the model's security, effectiveness or ethical behavior.	LLM04 Model Denial of Service Attackers cause resource-heavy operations on LLMs, leading to service degradation or high costs. The vulnerability is magnified due to the resource-intensive nature of LLMs and unpredictability of user inputs.	LLM05 Supply Chain Vulnerabilities LLM application lifecycle can be compromised by vulnerable components or services, leading to security attacks. Using third-party datasets, pre-trained models, and plugins add vulnerabilities.					
LLM06 Sensitive Information Disclosure LLMs may inadvertently reveal confidential data in its responses, leading to unauthorized data access, privacy violations, and security breaches. Implement data sanitization and strict user policies to mitigate this.	LLM07 Insecure Plugin Design LLM plugins can have insecure inputs and insufficient access control due to lack of application control. Attackers can exploit these vulnerabilities, resulting in severe consequences like remote code execution.	LLM08 Excessive Agency LLM-based systems may undertake actions leading to unintended consequences. The issue arises from excessive functionality, permissions, or autonomy granted to the LLM-based systems.	LLM09 Overreliance Systems or people overly depending on LLMs without oversight may face misinformation, miscommunication, legal issues, and security vulnerabilities due to incorrect or inappropriate content generated by LLMs.	LLM10 Model Theft This involves unauthorized access, copying, or exfiltration of proprietary LLM models. The impact includes economic losses, compromised competitive advantage, and potential access to sensitive information.					

Figure 4.1: Image of OWASP Top 10 for Large Language Model Applications

OWASP Top 10 for Large Language Model Applications Visualized

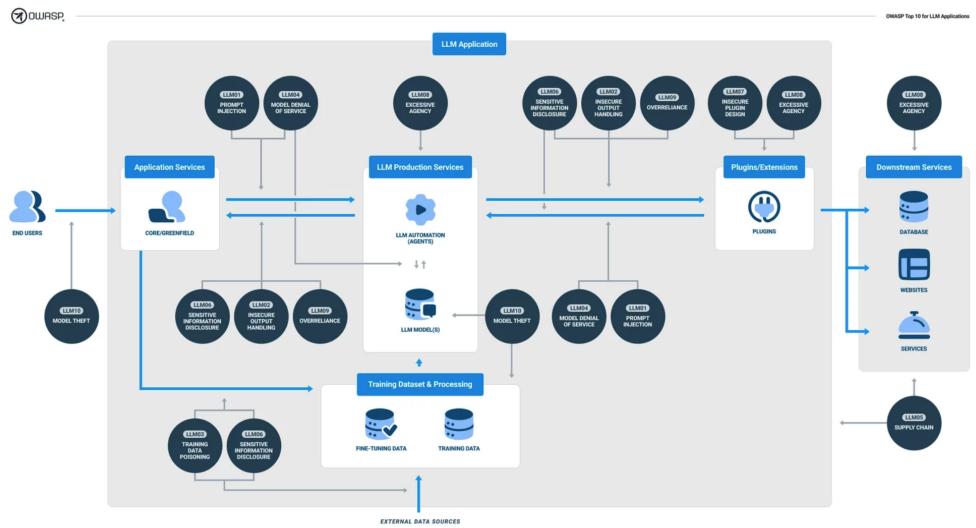


Figure 4.2: Image of OWASP Top 10 for Large Language Model Applications Visualized

OWASP Resources Using LLM solutions expands an organization's attack surface and presents new challenges, requiring special tactics and defenses. It also poses problems that are similar to known issues, and where there are already established cybersecurity procedures and mitigations. Integrating LLM cybersecurity with an organization's established cybersecurity controls, processes, and procedures allows an organization to reduce its vulnerability to threats. How they integrate is available at the OWASP Integration Standards.

OWASP Resource	Description	Why It Is Recommended & Where To Use It
OWASP SAMM	Software Assurance Maturity Model	Provides an effective and measurable way to analyze and improve an organization's secure development lifecycle. SAMM supports the complete software lifecycle. It is iterative and risk-driven, enabling organizations to identify and prioritize gaps in secure software development so resources for improving the process can be dedicated where efforts have the greatest improvement impact.
OWASP AI Security and Privacy Guide	OWASP Project with a goal of connecting worldwide for an exchange on AI security, fostering standards alignment, and driving collaboration.	The OWASP AI Security and Privacy Guide is a comprehensive list of the most important AI security and privacy considerations. It is meant to be a comprehensive resource for developers, security researchers, and security consultants to verify the security and privacy of AI systems.
OWASP AI Exchange	OWASP AI Exchange is the intake method for the OWASP AI Security and Privacy Guide.	The AI Exchange is the primary intake method used by OWASP to drive the direction of the OWASP AI Security and Privacy Guide.

OWASP Resource		Description	Why It Is Recommended & Where To Use It
OWASP Learning Top 10	Machine Security	OWASP Machine Learning Security Top 10 security issues of machine learning systems.	The OWASP Machine Learning Security Top 10 is a community-driven effort to collect and present the most important security issues of machine learning systems in a format that is easy to understand by both a security expert and a data scientist. This project includes the ML Top 10 and is a live working document that provides clear and actionable insights on designing, creating, testing, and procuring secure and privacy-preserving AI systems. It is the best OWASP resource for AI global regulatory and privacy information.
OpenCRE		OpenCRE (Common Requirement Enumeration) is the interactive content-linking platform for uniting security standards and guidelines into one overview.	Use this site to search for standards. You can search by standard name or by control type.
OWASP Threat Modeling		A structured, formal process for threat modeling of an application	Learn everything about Threat Modeling which is a structured representation of all the information that affects the security of an application.
OWASP CycloneDX		OWASP CycloneDX is a full-stack Bill of Materials (BOM) standard that provides advanced supply chain capabilities for cyber risk reduction.	Modern software is assembled using third-party and open source components. They are glued together in complex and unique ways and integrated with original code to achieve the desired functionality. An SBOM provides an accurate inventory of all components which enables organizations to identify risk, allows for greater transparency, and enables rapid impact analysis. EO 14028 provided minimum requirements for SBOM for federal systems.

OWASP Resource	Description	Why It Is Recommended & Where To Use It
OWASP Software Component Verification Standard (SCVS)	A community-driven effort to establish a framework for identifying activities, controls, and best practices can help in identifying and reducing risk in a software supply chain.	Use SCVS to develop a common set of activities, controls, and best-practices that can reduce risk in a software supply chain and identify a baseline and path to mature software supply chain vigilance.
OWASP API Security Project	API Security focuses on strategies and solutions to understand and mitigate the unique vulnerabilities and security risks of Application Programming Interfaces (APIs)	APIs are a foundational element of connecting applications, and mitigating misconfigurations or vulnerabilities is mandatory to protect users and organizations. Use for security testing and red teaming the build and production environments.
OWASP Application Security Verification Standard ASVS	Application Security Verification Standard (ASVS) Project provides a basis for testing web application technical security controls and also provides developers with a list of requirements for secure development.	Cookbook for web application security requirements, security testing, and metrics. Use to establish security user stories and security use case release testing.
OWASP Threat and Safeguard Matrix (TaSM)	An action oriented view to safeguard and enable the business	This matrix allows a company to overlay its major threats with the NIST Cyber Security Framework Functions (Identify, Protect, Detect, Respond, & Recover) to build a robust security plan. Use it as a dashboard to track and report on security across the organization.
Defect Dojo	An open source vulnerability management tool that streamlines the testing process by offering templating, report generation, metrics, and baseline self-service tools.	Use Defect Dojo to reduce the time for logging vulnerabilities with templates for vulnerabilities, imports for common vulnerability scanners, report generation, and metrics.

Table 4.1: OWASP Resources

MITRE Resources The increased frequency of LLM threats emphasizes the value of a resilience-first approach to defending an organization's attack surface. Existing TPPS are combined with new attack surfaces and capabilities in LLM Adversary threats and mitigations. MITRE maintains a well-established and widely accepted mechanism for coordinating opponent tactics and procedures based on real-world observations.

Coordination and mapping of an organization's LLM Security Strategy to MITRE ATT&CK and MITRE ATLAS allows an organization to determine where LLM Security is covered by current processes such as API Security Standards or where security holes exists.

MITRE ATT&CK (Adversarial Tactics, Techniques, and Common Knowledge) is a framework, collection of data matrices, and assessment tool that was made by the MITRE Corporation to help organizations figure out how well their cybersecurity works across their entire digital attack surface and find holes that had not been found before. It is a knowledge repository that is used all over the world. The MITRE ATT&CK matrix contains a collection of strategies used by adversaries to achieve a certain goal. In the ATT&CK Matrix, these objectives are classified as tactics. The objectives are outlined in attack order, beginning with reconnaissance and progressing to the eventual goal of exfiltration or impact.

MITRE ATLAS, which stands for "Adversarial Threat Landscape for Artificial Intelligence Systems," is a knowledge base that is based on real-life examples of attacks on machine learning (ML) systems by bad actors. ATLAS is based on the MITRE ATT&CK architecture, and its tactics and procedures complement those found in ATT&CK.

MITRE Resource	Description	Why It Is Recommended & Where To Use It
MITRE ATT&CK	Knowledge base of adversary tactics and techniques based on real-world observations	The ATT&CK knowledge base is used as a foundation for the development of specific threat models and methodologies. Map existing controls within the organization to adversary tactics and techniques to identify gaps or areas to test.
MITRE Workbench	Create or extend ATT&CK data in a local knowledge base	Host and manage a customized copy of the ATT&CK knowledge base. This local copy of the ATT&CK knowledge base can be extended with new or updated techniques, tactics, mitigation groups, and software that is specific to your organization.

MITRE Resource	Description	Why It Is Recommended & Where To Use It
MITRE ATLAS	MITRE ATLAS (Adversarial Threat Landscape for Artificial-Intelligence Systems) is a knowledge base of adversary tactics, techniques, and case studies for machine learning (ML) systems based on real-world observations, demonstrations from ML red teams and security groups, and the state of the possible from academic research	Use it to map known ML vulnerabilities and map checks and controls for proposed projects or existing systems.
MITRE ATT&CK Powered Suit	ATT&CK Powered Suit is a browser extension that puts the MITRE ATT&CK knowledge base at your fingertips.	Add to your browser to quickly search for tactics, techniques, and more without disrupting your workflow.
The Threat Report ATT&CK Mapper (TRAM)	Automates TTP Identification in CTI Reports	Mapping TTPs found in CTI reports to MITRE ATT&CK is difficult, error prone, and time-consuming. TRAM uses LLMs to automate this process for the 50 most common techniques. Supports Jupyter notebooks.
Attack Flow v2.1.0	Attack Flow is a language for describing how cyber adversaries combine and sequence various offensive techniques to achieve their goals.	Attack Flow helps visualize how an attacker uses a technique, so defenders and leaders understand how adversaries operate and improve their own defensive posture.
MITRE Caldera	A cyber security platform (framework) designed to easily automate adversary emulation, assist manual red-teams, and automate incident response.	Plugins are available for Caldera that help to expand the core capabilities of the framework and provide additional functionality, including agents, reporting, collections of TTPs and others
CALDERA Arsenal plugin:	A plugin developed for adversary emulation of AI-enabled systems.	This plugin provides TTPs defined in MITRE ATLAS to interface with CALDERA.

MITRE Resource	Description	Why It Is Recommended & Where To Use It
Atomic Red Team	Library of tests mapped to the MITRE ATT&CK framework.	Use to validate and test controls in an environment. Security teams can use Atomic Red Team to quickly, portably, and reproducibly test their environments. You can execute atomic tests directly from the command line; no installation is required.
MITRE CTI Blueprints	Automates Cyber Threat Intelligence reporting.	CTI Blueprints helps Cyber Threat Intelligence (CTI) analysts create high-quality, actionable reports more consistently and efficiently.

Table 4.2: MITRE Resources

AI Vulnerability Repositories

Name	Description
AI Incident Database	A repository of articles about different times AI has failed in real-world applications and is maintained by a college research group and crowds sourced.
OECD AI Incidents Monitor (AIM)	Offers an accessible starting point for comprehending the landscape of AI-related challenges.
Three of the leading companies tracking AI Model vulnerabilities	
Huntr Bug Bounty : ProtectAI	Bug bounty platform for AI/ML
AI Vulnerability Database (AVID) : Garak	Database of model vulnerabilities
AI Risk Database: Robust Intelligence	Database of model vulnerabilities

Table 4.3: AI Vulnerability Repositories

AI Procurement Guidance

Name	Description
World Economic Forum: Adopting AI Responsibly: Guidelines for Procurement of AI Solutions by the Private Sector: Insight Report June 2023	The standard benchmarks and assessment criteria for procuring Artificial systems are in early development. The procurement guidelines provide organizations with a baseline of considerations for the end-to-end procurement process. Use this guidance to augment an organization's existing Third Party Risk Supplier and Vendor procurement process.

Table 4.4: AI Procurement Guidance

Team

Thank you to the OWASP Top 10 for LLM Applications Cybersecurity and Governance Checklist Contributors.

Checklist Contributors		
Sandy Dunn	Heather Linn	John Sotiropoulos
Steve Wilson	Fabrizio Cilli	Aubrey King
Bob Simonoff	David Rowe	Rob Vanderveer
Emmanual Guilherme Junior	Andrea Succi	Jason Ross
Talesh Seeparsan	Anthony Glynn	Julie Tao

Table A.1: OWASP LLM AI Security & Governance Checklist Team

This project is licensed under the terms of the Creative Commons Attribution-ShareAlike 4.0 International License