

# The 20 Concepts That Turned My System Design Nightmare into Success

Why 90% of Engineers Bomb System Design Interviews (20 Concepts That Save You)



JAVINPAUL AND SOMA

AUG 31, 2025



19



4

Share

Preparing for system design interviews can feel like climbing a mountain without a map. Unlike **coding interviews** where you can gain confidence by practicing data structures and algorithms on platforms like [AlgoMonster](#), [Exponent](#) and LeetCode, system design questions demand a mix of breadth and depth — architecture principles, scalability patterns, trade-offs, and real-world application.

For me, this part of the interview loop was intimidating at first. I often felt lost in diagrams, unsure which concept to use where, and overwhelmed by the sheer vastness of distributed systems.

The turning point came when I started breaking the subject down into core concepts. Once I understood ideas like [load balancing](#), [caching](#), [database sharding](#), CAP theorem, and [message queues](#), everything else started to click into place.

Instead of memorizing solutions, I began recognizing patterns. That's when I realized system design isn't about giving a "perfect" architecture, but about **reasoning through trade-offs with clarity**.

What really accelerated my learning was leveraging structured resources. Books and visual explanations like [ByteByteGo's System Design Course](#) made the hardest concepts digestible with diagrams and case studies.

I also explored platforms such as [Codemia.io](#) and [Bugfree.ai](#) for hands-on interview prep and [Exponent](#) for mock interviews with engineers from top companies. Each helped me move from feeling clueless to confident, especially when facing open-ended system design questions at FAANG-level interviews.

In this article, I'll share the **20 core concepts that completely changed how I approach system design interviews**. Mastering these will save you from confusion, help you build better mental models, and make those tough whiteboard sessions a lot less scary.

## Stop Failing System Design Interviews: Master These 20 Core Concepts First

Here are the 20 key concepts I learned and master by going through different System Design resources. Once you learn these concepts, half the battle is already one.

### 1. Load Balancing: The Traffic Director

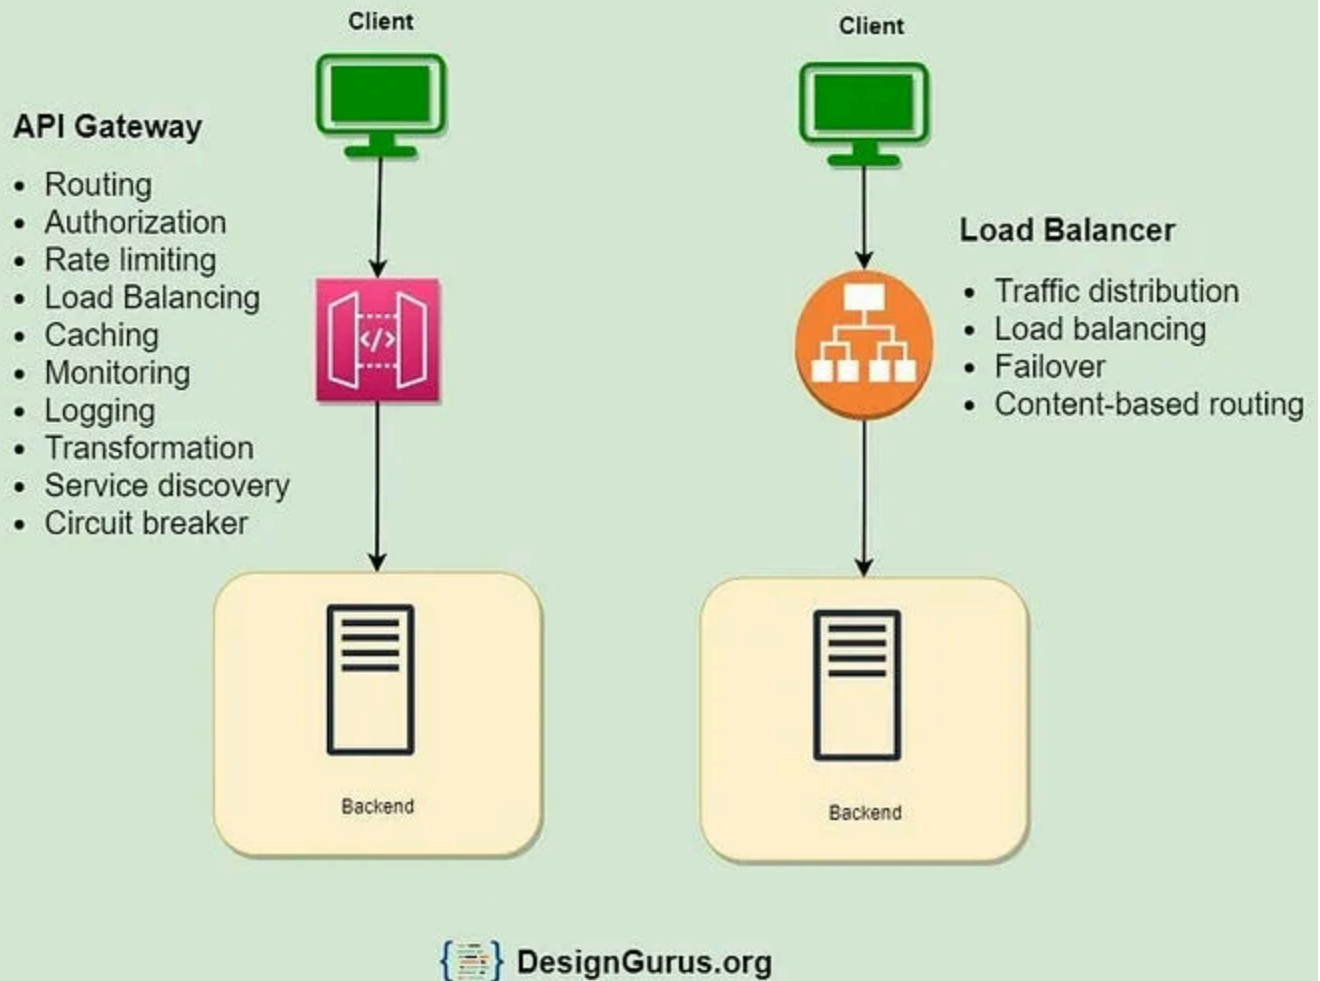
Think of load balancers as smart traffic directors for your application. They distribute incoming requests across multiple servers to prevent any single server from becoming overwhelmed.

**Key insight:** There are different types — Layer 4 (transport layer) and Layer 7 (application layer). Layer 7 load balancers can make routing decisions based on content, while Layer 4 focuses on IP and port information.

**Real-world example:** When you visit Amazon, a load balancer decides which of the thousands of servers will handle your request.

Here is a nice diagram from [designgurus.io](#) which explains the load balancer concept along with API gateway which we will see in a couple of seconds.

## API Gateway vs. Load Balancer




## 2. Horizontal vs Vertical Scaling: The Growth Strategies

- **Vertical Scaling (Scale Up):** Adding more power to existing machines
- **Horizontal Scaling (Scale Out):** Adding more machines to the pool

**Game-changer moment:** Understanding that horizontal scaling is almost always preferred for large systems because it's more cost-effective and provides better fault tolerance.

Here is a visual guide from [ByteByteGo](#) which makes this concept crystal clear

# Vertical Vs Horizontal Scaling



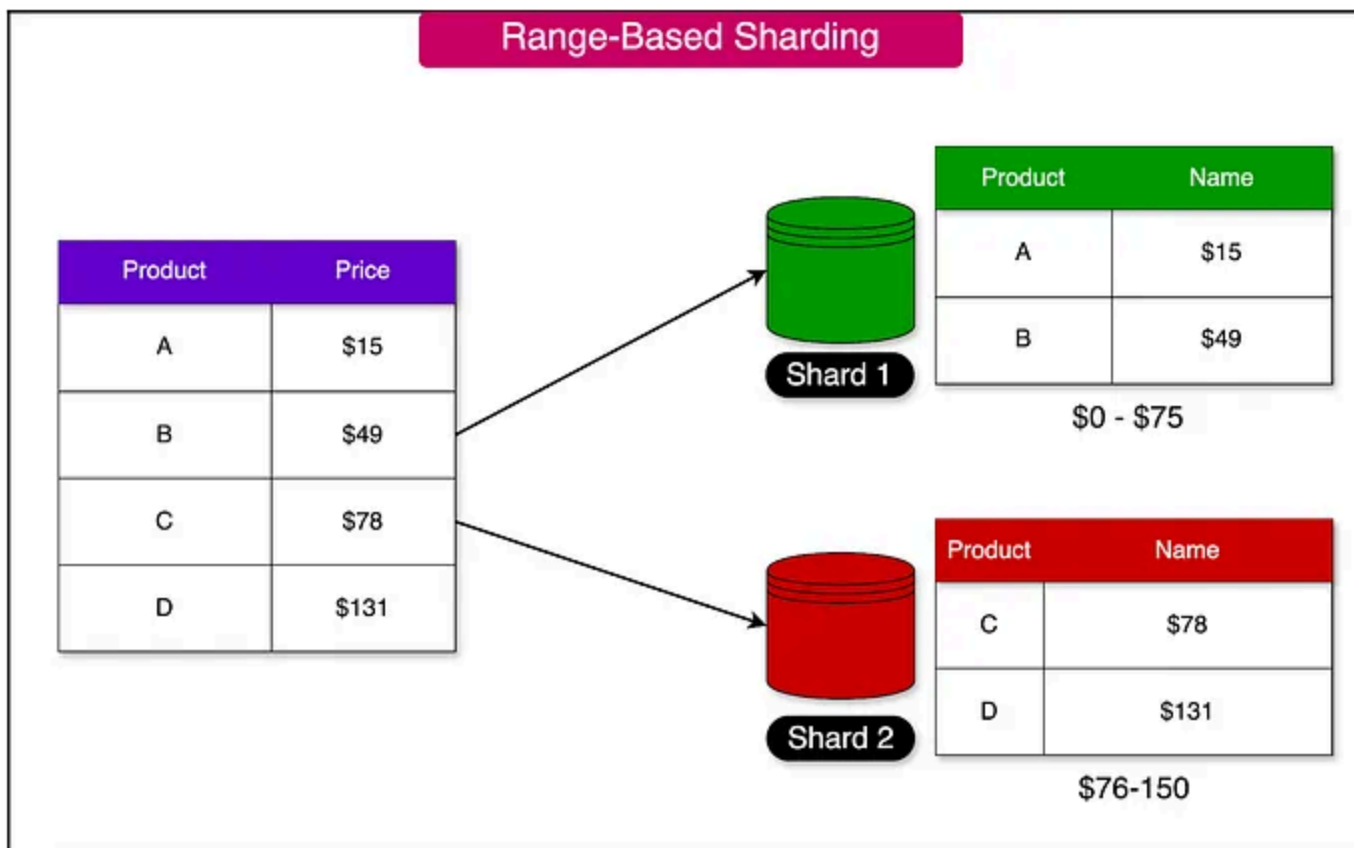
## 3. Database Sharding: Divide and Conquer

**Sharding** splits your database across multiple machines. Each shard contains a subset of your data.

**The breakthrough:** Learning about sharding keys and how poor sharding strategies can create hotspots that defeat the entire purpose.

**Example:** Instagram shards user data based on user ID, ensuring even distribution across databases.

Here is another great visual from [ByteByteGo](#) which explains Range based sharding



## 4. Caching Strategies: The Speed Multiplier

Caching is storing frequently accessed data in fast storage. The key is understanding different caching patterns:


- **Cache-aside (Lazy Loading):** Application manages cache
- **Write-through:** Write to cache and database simultaneously
- **Write-behind:** Write to cache immediately, database later

**Pro tip:** The cache invalidation problem is one of the hardest problems in computer science. Master cache eviction policies (LRU, LFU, FIFO).

A picture is worth thousand words and this visual from ByteByteGo proves that, it nicely explains all the caching strategies a senior developer should be aware of.

If you like visual learning, I highly recommend you to join [ByteByteGo](#) now as they offering **50% discount on their lifetime plan**. I have taken that as its just 2.5 times of annual plan but provides most value.

## Top 5 Caching Strategies

 [blog.bytebytego.com](https://blog.bytebytego.com)

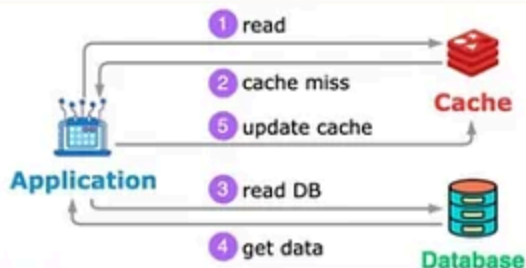
### Cache Aside

#### Pros

1. update logic is on application level, easy to implement
2. cache only contains what the application requests for

#### Cons

1. each cache miss results in 3 trips
2. data may be stale if DB is updated directly




### Read Through

#### Pros

1. application logic is simple
2. can easily scale the reads and only one query hits the DB

#### Cons

- data access logic is in the cache, needs to write a plugin to access DB



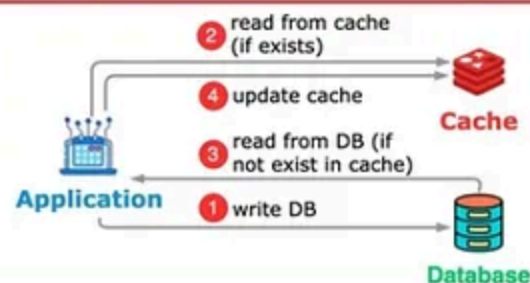
### Write Around

#### Pros

1. the DB is the source of truth
2. lower read latency

#### Cons

1. higher write latency because data is written to DB first
2. the data in the cache may be stale



### Write Back

#### Pros

1. lower write latency
2. lower read latency
3. the cache and DB are eventually consistent

#### Cons

1. there can be data loss if the cache is down
2. infrequent data is also stored in the cache




### Write Through

#### Pros

1. reads have lower latency
2. the cache and DB are in sync

#### Cons

1. writes have higher latency because they need to wait for the DB writes to finish
2. infrequent data is also stored in the cache

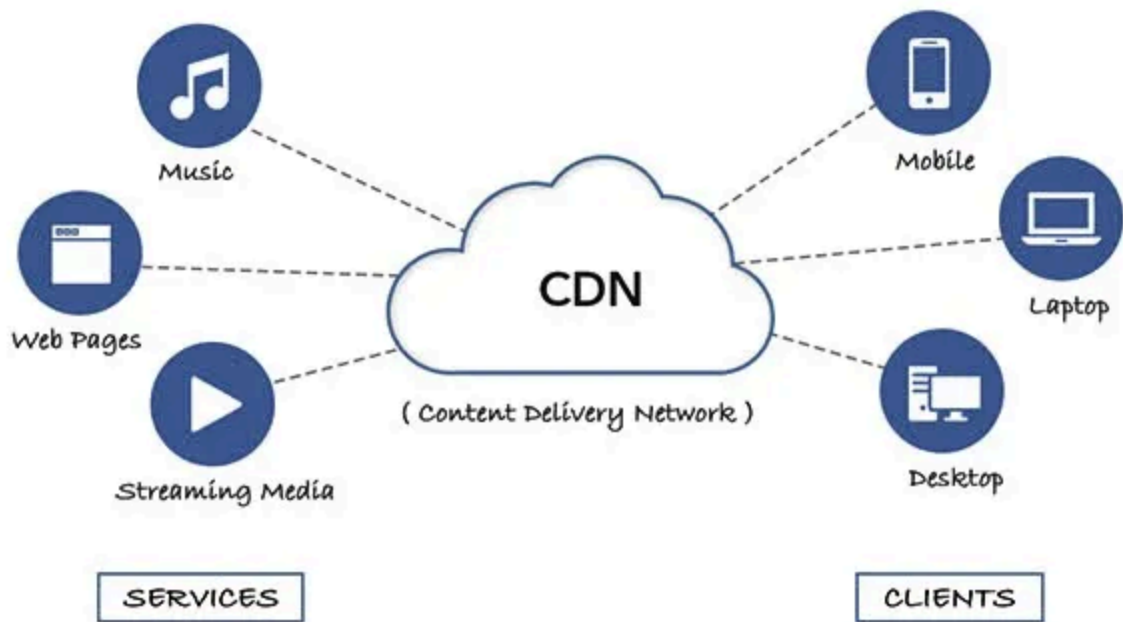




## 5. Content Delivery Networks (CDN): Global Speed

CDNs cache your content at edge locations worldwide, reducing latency for users.

**Aha moment:** Realizing that CDNs don't just cache static content — modern CDNs can cache dynamic content and even run serverless functions at the edge.




## 6. Database Replication: The Backup Plan

- **Master-Slave:** One write node, multiple read nodes
- **Master-Master:** Multiple write nodes (more complex)

**Critical insight:** Understanding eventual consistency and how replication lag can affect your application logic.

# Database Replication



## 7. Consistent Hashing: The Elegant Solution

Regular hashing breaks when you add/remove servers. Consistent hashing minimizes redistribution when the hash table is resized.

**Why it matters:** This is how systems like DynamoDB and Cassandra distribute data across nodes efficiently.

## 8. CAP Theorem: The Fundamental Trade-off

You can only guarantee two out of three:

- **Consistency:** All nodes see the same data simultaneously
- **Availability:** System remains operational
- **Partition Tolerance:** System continues despite network failures



**Real impact:** This guides every distributed system design decision you'll ever make.

## 9. Event-Driven Architecture: The Modern Approach

Systems communicate through events rather than direct calls. This creates loose coupling and better scalability.

**Game-changer:** Understanding that event sourcing can make your system audit-friendly and enable powerful debugging capabilities.

## 10. Message Queues: Asynchronous Communication

**Queues** decouple producers and consumers, enabling asynchronous processing.

**Key patterns:**

- Point-to-point (one consumer)
- Publish-subscribe (multiple consumers)

**Example:** When you upload a video to YouTube, it goes into a queue for processing rather than blocking your upload.

# Message Queue

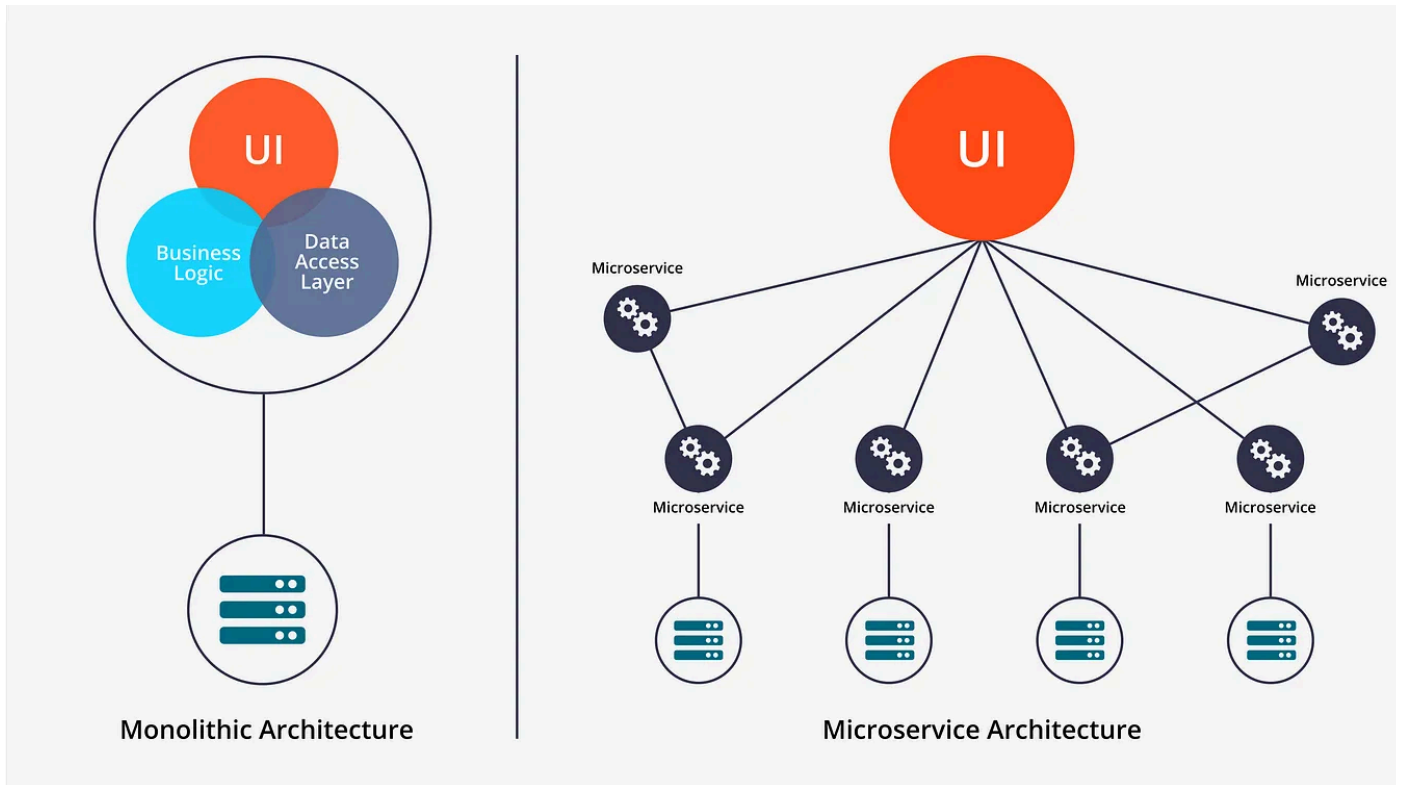


## 11. Microservices vs Monolith: The Architecture Debate

**Monolith advantages:** Simpler deployment, testing, debugging

**Microservices advantages:** Independent scaling, technology diversity, fault isolation


**The insight:** Start with a monolith, extract microservices when you have clear bounded contexts and team structure to support them.



## 12. API Gateway: The Single Entry Point

**API gateways** handle cross-cutting concerns like authentication, rate limiting, and request routing.

**Why crucial:** They prevent every microservice from implementing the same boilerplate code.



## 13. Database Indexing: Query Performance

Indexes are data structures that improve query speed at the cost of storage and write performance.

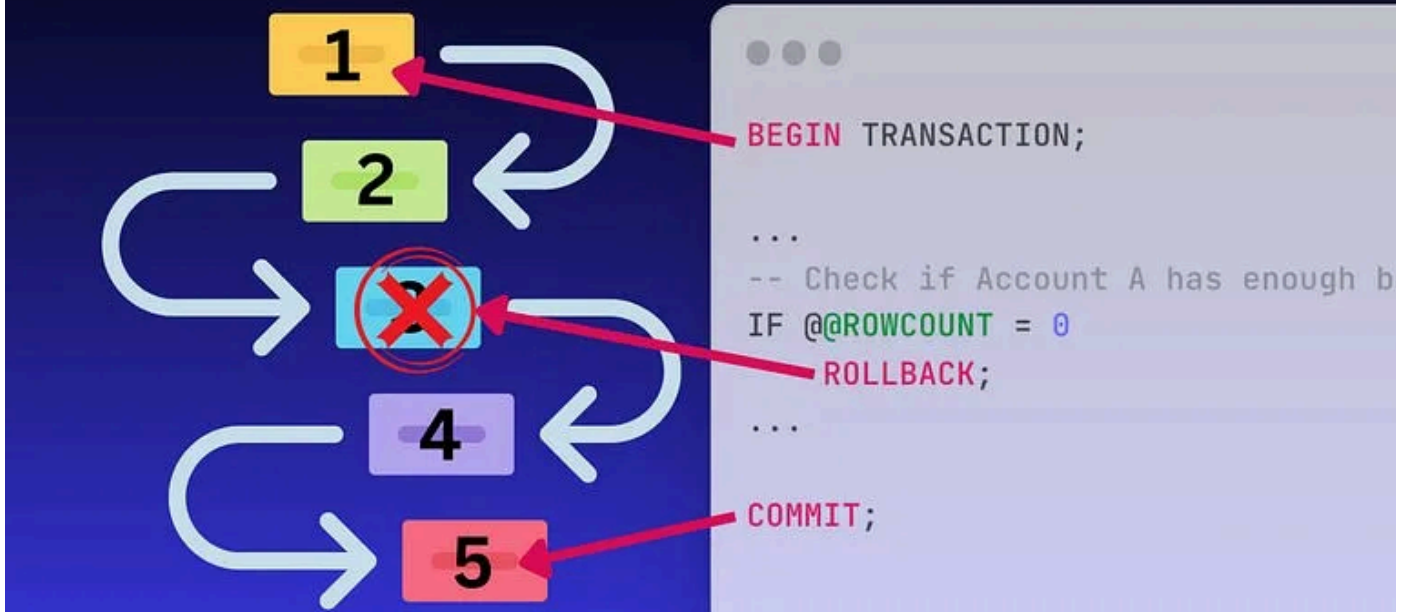
**Advanced concept:** Understand compound indexes, covering indexes, and when indexes actually hurt performance.

## 14. ACID vs BASE: Data Consistency Models

**ACID:** Atomicity, Consistency, Isolation, Durability (SQL databases) **BASE:** Basical Available, Soft state, Eventual consistency (NoSQL)

**The decision framework:** Use ACID for financial transactions, BASE for social media feeds.

# ACID Transactions




## 15. Rate Limiting: Protecting Your System

There are many different **Rate limiting algorithms** for controlling request rates:

- Token bucket
- Leaky bucket
- Fixed window
- Sliding window

**Real-world application:** Twitter's rate limiting prevents spam and ensures fair usage across all users.



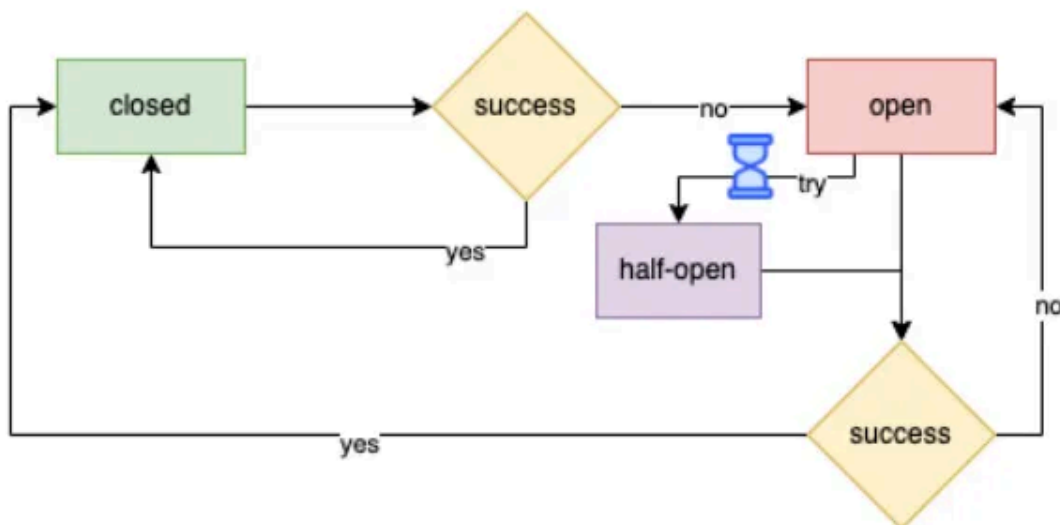
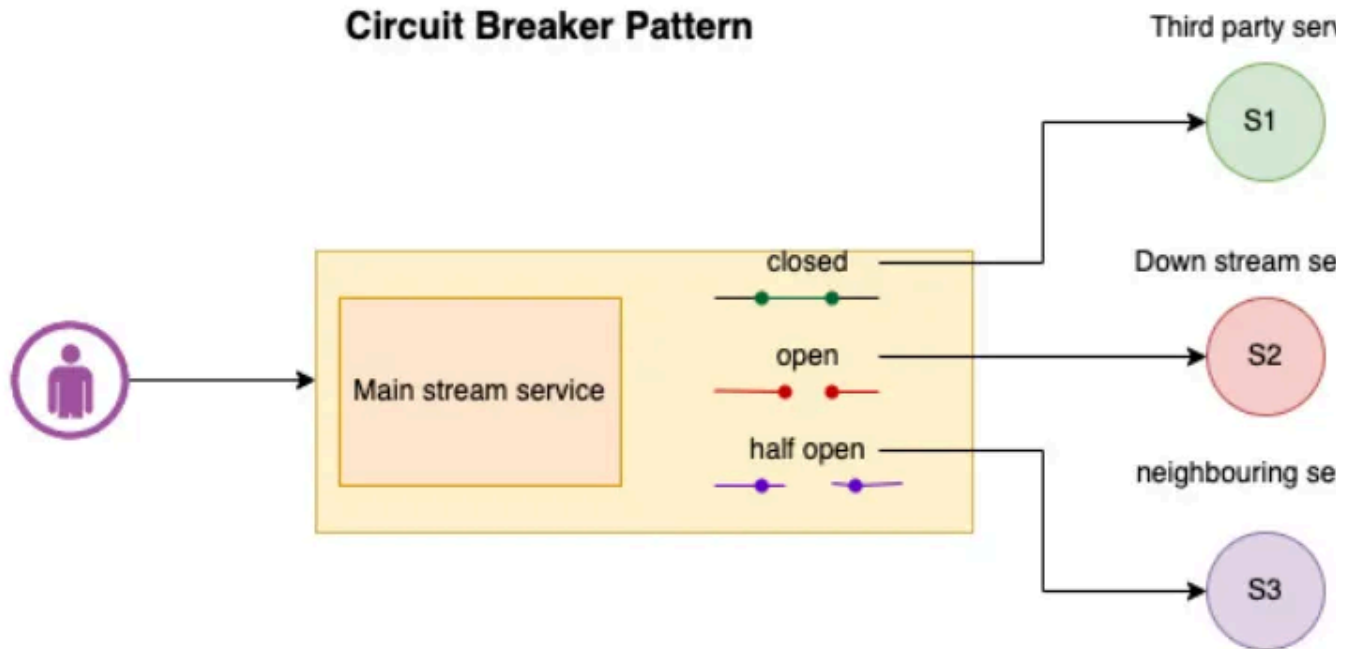
## 16. Circuit Breaker Pattern: Failure Resilience

This is a classic pattern which is asked multiple times on interview. This pattern Prevents cascade failures by temporarily stopping requests to a failing service.

**States:** Closed (normal), Open (failing), Half-open (testing recovery)



## Circuit Breaker Pattern



## 17. Distributed Consensus: Agreement in Chaos

Algorithms like Raft and Paxos help distributed systems agree on a single value even with network partitions and node failures.

**Why it matters:** This is how systems like etcd (Kubernetes) and DynamoDB maintain consistency across replicas.

## 18. Eventual Consistency: The Distributed Reality

In distributed systems, achieving immediate consistency across all nodes is often impossible or impractical.

**Examples:**

- Social media likes (eventual consistency is fine)
- Bank transfers (strong consistency required)

## 19. Bloom Filters: Probabilistic Data Structures

Space-efficient data structure that tells you if an element is “definitely not in a set” “possibly in a set.”

**Use cases:**

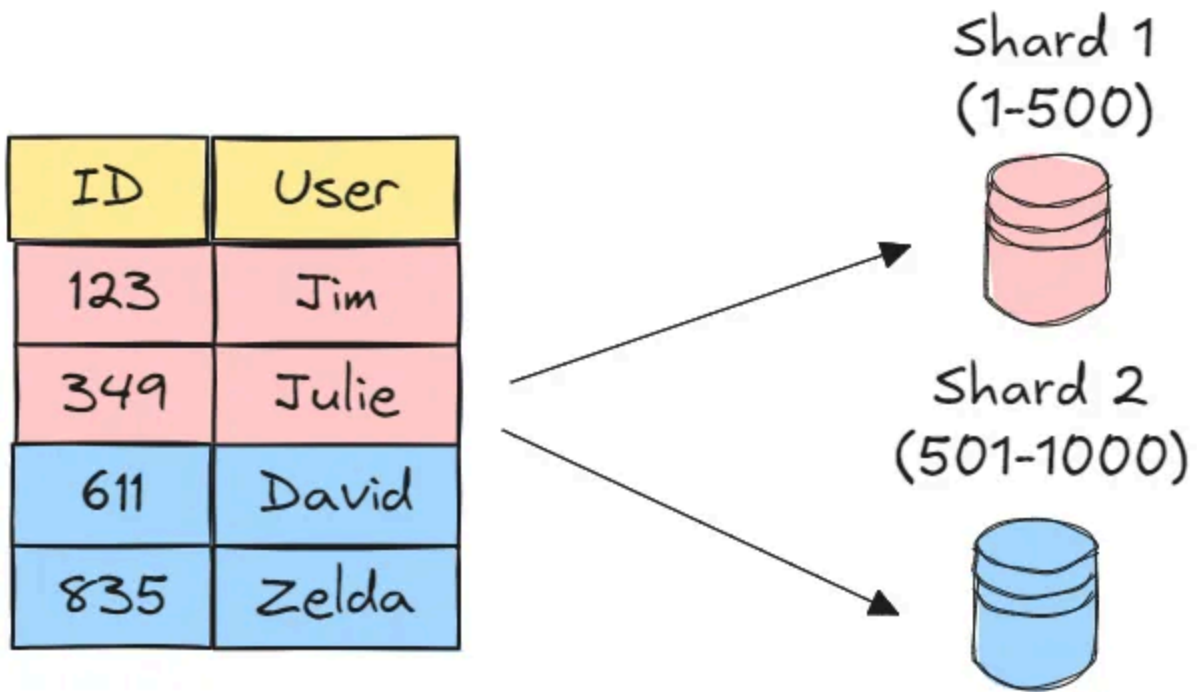
- Web crawlers avoiding duplicate URLs
- Databases checking if data exists before expensive disk reads

## 20. Data Partitioning Strategies

There are three main data partitioning strategies:

- **Vertical:** Split by features/columns
- **Horizontal:** Split by rows
- **Functional:** Split by service boundaries

**Critical insight:** Choosing the wrong partitioning key can create hotspots and uneven load distribution.



## How These Concepts Connected Everything?

The magic happened when I realized these concepts don't exist in isolation.

For example:

- **Netflix's architecture** combines CDNs for content delivery, microservices for different functions, event-driven architecture for recommendations, and sophisticated caching strategies.
- **WhatsApp's messaging** uses consistent hashing for user distribution, message queues for offline message delivery, and database sharding to handle billions of messages.

Once you understand how these concepts work together, you can design systems for any scale.

## My Learning Strategy That Worked

Here's the exact approach I used to master these concepts:

## 1. Start with Fundamentals

I began with the basics using resources like [ByteByteGo](#), which breaks down complex systems into digestible visual explanations. Their system design course was instrumental in building my foundation.

## 2. Practice with Real Examples

Sites like [Codemia](#) and [System Design School](#) provided excellent hands-on practice with real-world system design problems. The interactive approach helped me apply concepts immediately.

## 3. Deep Dive into Patterns

[DesignGuru](#) offered comprehensive coverage of system design patterns. Their [Grokking the System Design Interview](#) course became my bible.

## 4. Mock Interviews

[Exponent](#) and [BugFree.ai](#) provided peer to peer and AI-powered mock interviews that helped me practice explaining my designs clearly and handling follow-up questions.

## 5. Interactive Learning

[Educative](#) offered interactive courses that let me experiment with concepts in a hands-on environment.

## 6. Video Learning

YouTube and [Udemy](#) had comprehensive video courses that I could watch during commutes and lunch breaks.

## 7. Essential Reading

[Designing Data-Intensive Applications](#) became my go-to reference book. This book is gold for understanding distributed systems deeply.

## 8. Open Source Learning

[GitHub Repositories](#) provided real-world examples and system design

## Final Thoughts

Mastering system design doesn't happen overnight, but focusing on these **20 core concepts** will give you a strong foundation to tackle any interview with confidence.

If you want to go even deeper, I highly recommend resources like [ByteByteGo's System Design Interview course](#) and other structured programs that break down real-world problems step by step. They are also offering 50% discount now on their lifetime plan.

With the right preparation and consistent practice, what once felt overwhelming will soon become second nature.

If you like this post, then don't forget to subscribe to [Soma's substack](#) where she shares more tips with System Design, Coding, Java, and React.js for developers.



### React Java

A newsletter about Java, Programming, System Design, and React.js

By Soma

Isantos2000@gmail.com

Subscribe


## Other System Design Articles you may like



### Stop Using HTTP for Everything: The API Protocol Guide That Sav Careers

JAVINPAUL • AUG 30


[Read full story →](#)



### 50 System Design Questions That Actually Prepare You for Real Interviews

JAVINPAUL AND SOMA • AUG 29


[Read full story →](#)



### Why 90% of APIs Fail (And How to Design Ones That Don't)

JAVINPAUL AND HAYK • AUG 26

[Read full story →](#)



### How I Would Learn System Design in 2025 (If I Had To Start Over

JAVINPAUL AND SOMA • AUG 24

[Read full story →](#)

## Subscribe to Javarevisited Newsletter

By javinpaul · Hundreds of paid subscribers

Master Java and System Design Interviews. Level up your Software Engineering career. Subscribe and get copy of my book Grokking the Spring Boot Interview in your inbox

[Upgrade to paid](#)



19 Likes · 4 Restacks



← Previous



A guest post by

**Soma**

Java and React Developer

Subscribe to Son

## Discussion about this post

Comments

Restacks



Write a comment...