



# ANSIBLE DECLARATIVE CONFIGURATION STEP-BY-STEP GUIDE



By DevOps Shack

---

[Click here for DevSecOps & Cloud DevOps Course](#)

## DevOps Shack

# Ansible Declarative Configuration –

## Step-by-Step Guide

### Table of Contents

#### 1. Introduction & Installation

- What is Ansible?
- Declarative vs Imperative configuration
- Installing Ansible on Linux/Mac/Windows
- Verifying installation

#### 2. Project Setup & Inventory

- Understanding Ansible architecture (Control Node, Managed Nodes)
- Creating a project directory
- Writing a static inventory file
- Testing connections via SSH

#### 3. Writing Declarative Playbooks

- YAML syntax basics
- Playbook structure: hosts, tasks, vars
- Using common modules: apt, yum, file, service, etc.
- Ensuring idempotency (declarative behavior)

#### 4. Variables, Templates & Handlers

- Defining and using variables (vars, group\_vars, host\_vars)
- Jinja2 templating in config files
- Using handlers and triggers for reactive changes

#### 5. Organizing with Roles

- 
- Creating and using roles for reusability
  - Role directory structure (tasks/, handlers/, templates/, etc.)
  - Role-based playbook example
  - Using Ansible Galaxy for public roles

## 6. Running, Debugging & Best Practices

- Running playbooks with ansible-playbook
- Dry-run (--check), verbosity (-v)
- Debugging tasks and using ansible-lint
- Recommended folder structure and Git versioning

## Step 1: Introduction & Installation

## ◆ What is Ansible?

Ansible is an open-source automation tool used for **configuration management, application deployment, orchestration, and task automation**. It allows IT administrators and DevOps engineers to define the state of systems in a simple, **declarative way using human-readable YAML files**, known as **playbooks**.

Unlike other tools (like Puppet or Chef), Ansible is **agentless**—it only requires SSH access to target machines and Python installed on them. This simplicity makes Ansible lightweight, easy to adopt, and ideal for both small teams and large-scale infrastructure.

## ◆ Declarative vs. Imperative Configuration

Ansible embraces the **declarative** approach. This means you describe *what* state you want the system to be in (e.g., “nginx should be installed and running”), and Ansible figures out *how* to make it so.

In contrast, **imperative** tools require you to explicitly define every step in sequence (e.g., install nginx, enable nginx, start nginx).

Declarative configurations are:

- Easier to maintain
- More predictable
- Naturally idempotent (no unintended side effects if you run the same playbook multiple times)

## ◆ Installing Ansible

Let's get started by installing Ansible on your local machine (Control Node).

### Requirements:

- Python 3.8+
- pip (Python package manager)
- Internet connection

- 
- Basic understanding of the command line

### Installation on Ubuntu/Debian:

```
sudo apt update
```

```
sudo apt install software-properties-common
```

```
sudo add-apt-repository --yes --update ppa:ansible/ansible
```

```
sudo apt install ansible
```

### Installation on macOS (using Homebrew):

```
brew install ansible
```

### Installation on Windows (via WSL or Python/pip):

For Windows, it's best to use **WSL (Windows Subsystem for Linux)**:

1. Enable WSL and install Ubuntu from the Microsoft Store
2. Open the Ubuntu terminal and follow the Ubuntu steps above

Or install with pip:

```
pip install ansible
```

 Tip: If you plan to use Ansible heavily, prefer WSL2 or a Linux VM for best compatibility.

### ◆ Verifying the Installation

Once Ansible is installed, verify it by checking the version:

```
ansible --version
```

You should see output like:

```
ansible [core 2.16.1]
```

```
config file = /etc/ansible/ansible.cfg
```

```
python version = 3.10.x
```

If you see that, you're all set!

## ◆ Optional: Set Up a Python Virtual Environment

To isolate your Ansible environment:

```
python3 -m venv ansible-env  
source ansible-env/bin/activate  
pip install ansible
```

This is a good practice when working on multiple projects.

## Summary

In this step, you've:

- Learned what Ansible is and how it works
- Understood the difference between declarative and imperative configuration
- Installed Ansible on your system
- Verified that it works

With Ansible installed, you're ready to start creating your first Ansible project, define an inventory, and write your first playbook!

## Step 2: Project Setup & Inventory

Now that Ansible is installed, it's time to set up your first Ansible project and create an **inventory**, which is a key part of how Ansible knows what machines to manage.

### ◆ Understanding Ansible Architecture

Before we jump into the hands-on part, let's quickly understand the core concepts of Ansible's architecture:

- **Control Node**: The machine from which Ansible commands are run (your local machine).
- **Managed Nodes** (or hosts/targets): The remote machines you want to configure.
- **Inventory File**: A file that lists your managed nodes.
- **Playbook**: A YAML file that defines the desired configuration (more on this in the next step).
- **Modules**: The units of work in Ansible (e.g., install package, copy files, start service).

Ansible communicates with managed nodes via SSH (port 22), so make sure passwordless SSH is set up.

### ◆ Setting Up Your First Ansible Project

You can keep your Ansible projects well-organized by using a dedicated folder for each one. Here's how you can start:

```
mkdir my-ansible-project
```

```
cd my-ansible-project
```

Inside this directory, you'll store:

- inventory.ini: your list of managed nodes
- playbook.yml: your Ansible instructions
- group\_vars/ or host\_vars/: variable definitions

- 
- roles/: reusable role directories

#### ◆ Writing the Inventory File

Create a file named inventory.ini in the root of your project:

```
[webservers]
```

```
192.168.1.101
```

```
192.168.1.102
```

```
[dbservers]
```

```
db1.example.com ansible_user=ubuntu ansible_port=22
```

#### Explanation:

- You define **groups** of servers using [groupname].
- You can assign variables like ansible\_user, ansible\_port, etc., on a per-host basis.

You can also use a **YAML-style inventory** (inventory.yml) if you prefer:

```
all:
```

```
  children:
```

```
    webservers:
```

```
      hosts:
```

```
        192.168.1.101:
```

```
        192.168.1.102:
```

```
    dbservers:
```

```
      hosts:
```

```
        db1.example.com:
```

```
          ansible_user: ubuntu
```

```
          ansible_port: 22
```

- 
- ❖ Choose either .ini or .yml format — both work, but .ini is more common for beginners.

- ◆ **SSH Access to Remote Machines**

Make sure your control node (your local machine) can SSH into the target machines without needing a password each time.

**Generate SSH Key (if not already done):**

```
ssh-keygen -t rsa -b 4096
```

**Copy Key to Remote Server:**

```
ssh-copy-id ubuntu@192.168.1.101
```

Repeat for each server.

- ◆ **Test Your Setup**

You can now test Ansible's ability to connect to the nodes in your inventory:

```
ansible all -i inventory.ini -m ping
```

Expected output:

```
192.168.1.101 | SUCCESS => {
```

```
    "changed": false,
```

```
    "ping": "pong"
```

```
}
```

If you get a "**UNREACHABLE**" message, check:

- IP address
- SSH key setup
- ansible\_user value
- Whether the machine is up and reachable

- ◆ **Using ansible.cfg (Optional)**

---

You can simplify command-line usage by creating an `ansible.cfg` file in the project root:

`[defaults]`

`inventory = inventory.ini`

`host_key_checking = False`

`retry_files_enabled = False`

This lets you run commands like:

`ansible all -m ping`

...without needing the `-i inventory.ini` every time.

## Summary

By completing this step, you have:

- Understood how Ansible's architecture works
- Created a clean folder structure for your project
- Written your first inventory file (INI or YAML)
- Set up passwordless SSH between your control node and managed nodes
- Verified your setup using the Ansible ping module

## ❸ Step 3: Writing Declarative Playbooks

Now that you have your Ansible project and inventory set up, it's time to write your first **playbook** — the core of Ansible's declarative automation.

### ◆ What is an Ansible Playbook?

An **Ansible Playbook** is a YAML file that defines a set of tasks to run on managed nodes. Each playbook describes:

- **What** systems to configure (via inventory or group)
- **How** they should be configured (tasks, modules, variables)
- **What state** you want (declaratively)

Playbooks are:

- **Human-readable**
- **Idempotent** (safe to run multiple times without unintended side effects)
- Reusable and modular

### ◆ Basic YAML Syntax (Quick Primer)

Ansible playbooks are written in **YAML**. Here's a quick guide:

- Use spaces, **not tabs**
- Indentation is crucial
- Use : after keys, and prefix lists with -

Example:

---

```
- name: My First Playbook
  hosts: webservers
  become: yes
  tasks:
```

---

```
- name: Install Nginx
```

```
  apt:
```

```
    name: nginx
```

```
    state: present
```

#### ◆ Anatomy of a Playbook

Let's break down each part:

```
---
```

```
- name: Configure Web Server # Play name  
  hosts: webservers      # Target group from inventory  
  become: yes           # Use sudo (for privilege escalation)
```

```
  tasks:                 # List of tasks to execute
```

```
    - name: Ensure Nginx is installed  
      apt:  
        name: nginx  
        state: present  
        update_cache: yes
```

```
    - name: Ensure Nginx is running  
      service:  
        name: nginx  
        state: started  
        enabled: true
```

#### Explanation:

- name: A descriptive label for the play or task

- hosts: Group or host from your inventory
- become: Use sudo to perform privileged tasks
- tasks: A list of operations to apply
- Modules used: apt, service — this is where the **declarative logic** happens

#### ◆ Using Common Ansible Modules Declaratively

Here are some frequently used modules:

Module	Purpose	Example (declarative state)
apt / yum	Install packages	state: present or absent
file	Manage files/directories	state: directory, mode: 0755
copy	Copy local files to remote	src, dest, mode
service	Control services	state: started, enabled: true
user	Manage users	state: present, shell: /bin/bash

All these modules work **declaratively** — you just define the end state.

#### ◆ Running Your Playbook

Once your playbook is ready, run it using the ansible-playbook command:

`ansible-playbook -i inventory.ini playbook.yml`

With ansible.cfg set up (as shown in Step 2), you can simplify:

`ansible-playbook playbook.yml`

Expected output:

`PLAY [Configure Web Server] ***`

`TASK [Ensure Nginx is installed] ***`

`changed: [192.168.1.101]`

`TASK [Ensure Nginx is running] ***`

---

ok: [192.168.1.101]

- ✓ ok means the task was already in the correct state
- ✗ changed means the task made modifications to reach the desired state
- ✗ failed means something went wrong

#### ◆ Dry Run (Check Mode)

You can test a playbook without making changes:

```
ansible-playbook playbook.yml --check
```

Useful for validation before deploying to production.

#### ✓ Summary

In this step, you've:

- Learned how Ansible playbooks work
- Understood YAML syntax and playbook structure
- Used common modules declaratively
- Created and executed your first real playbook
- Verified and validated your configuration

Ansible now ensures your systems reach (and stay in) the exact state you want — with no manual work or surprises.



## Step 4: Variables, Templates & Handlers

As your Ansible projects grow, hardcoding values into your playbooks becomes inefficient and error-prone. This step introduces **variables**, **templates**, and **handlers** — essential tools for writing dynamic and clean configurations.

### ◆ Using Variables in Ansible

Variables allow you to customize your playbooks without rewriting them for different environments or hosts.

#### ◆ Types of Variables:

##### 1. Inline Variables (inside playbooks):

vars:

app\_port: 8080

##### 2. Group Variables (for an entire group): Create a folder named group\_vars/ and file like group\_vars/webservers.yml:

app\_port: 8080

##### 3. Host Variables (for a specific host): Create a folder named host\_vars/ and file like host\_vars/192.168.1.101.yml:

nginx\_enabled: true

##### 4. Extra Variables (CLI-based):

ansible-playbook site.yml -e "app\_port=8080"

### ◆ Using Variables in Tasks

Once a variable is defined, you can use it like this:

- name: Use variable in task

copy:

src: app.conf

dest: /etc/myapp/{{ app\_port }}.conf

---

Ansible uses **Jinja2 templating** to evaluate variables — which brings us to the next topic...

- ◆ **Templating with Jinja2**

Templates allow you to define configuration files (e.g., nginx.conf) dynamically, with variable substitution.

- ◆ **Example:**

Create a template file templates/nginx.conf.j2:

```
server {  
    listen {{ app_port }};  
    server_name {{ server_name }};  
    root /var/www/{{ site_folder }};  
}
```

In your playbook:

```
- name: Deploy Nginx config from template  
  template:  
    src: templates/nginx.conf.j2  
    dest: /etc/nginx/sites-available/default
```

- ◆ **Variables file:**

```
server_name: example.com  
site_folder: html
```

This allows the same template to work across different environments.

 Make sure the destination path has correct permissions or use become: yes.

- ◆ **Handlers (Reactive Tasks)**

---

**Handlers** are special tasks that run only when notified by another task. They're used for tasks that should only happen if something changes — like restarting a service after a config update.

◆ **Example:**

tasks:

```
- name: Deploy nginx config
  template:
    src: templates/nginx.conf.j2
    dest: /etc/nginx/sites-available/default
  notify: Restart Nginx
```

handlers:

```
- name: Restart Nginx
  service:
    name: nginx
    state: restarted
```

If the file changes, Ansible will trigger the Restart Nginx handler. If not, it skips it.

This approach keeps your playbooks efficient and avoids unnecessary restarts.

◆ **Putting It All Together**

Here's a full playbook example that combines everything:

```
---
```

```
- name: Setup web server with variables and templates
  hosts: webservers
  become: yes
  vars:
```

---

```
app_port: 8080
server_name: mysite.com
site_folder: html
```

#### tasks:

```
- name: Install Nginx
  apt:
    name: nginx
    state: present
    update_cache: yes

- name: Deploy Nginx config
  template:
    src: templates/nginx.conf.j2
    dest: /etc/nginx/sites-available/default
  notify: Restart Nginx
```

#### handlers:

```
- name: Restart Nginx
  service:
    name: nginx
    state: restarted
```

## Summary

In this step, you've learned how to:

- Use variables to make your playbooks dynamic

- 
- Create Jinja2 templates for reusable configs
  - Define handlers to respond to changes
  - Build more modular and maintainable infrastructure

## Step 5: Organizing with Roles

As your Ansible playbooks grow in complexity, it becomes necessary to break things down into **roles** — reusable, organized units that can be shared across projects and teams.

Roles help you maintain structure, promote reuse, and reduce duplication.

### ◆ What is an Ansible Role?

A **role** is a collection of related files and directories (tasks, variables, templates, handlers, etc.) that define a reusable automation component.

You can create roles for:

- Installing a web server
- Setting up a database
- Deploying an application
- Configuring a system service

Roles can be reused in multiple playbooks or shared publicly via **Ansible Galaxy**.

### ◆ Role Directory Structure

Here's the standard layout of a role:

```
css
CopyEdit
my-ansible-project/
└── roles/
    └── nginx/
        ├── tasks/
        │   └── main.yml
        └── handlers/
```

```
 | └── main.yml
 └── templates/
    | └── nginx.conf.j2
    └── vars/
        | └── main.yml
        └── defaults/
            | └── main.yml
            └── files/
                └── meta/
                    └── main.yml
```

Each folder serves a purpose:

- `tasks/`: Core tasks for the role
- `handlers/`: Notifications
- `templates/`: Jinja2 template files
- `vars/`: Variables with higher precedence
- `defaults/`: Default variables
- `files/`: Static files to copy
- `meta/`: Metadata about the role (dependencies, author, etc.)

#### ◆ Creating a Role

Use Ansible's built-in generator:

```
ansible-galaxy init roles/nginx
```

This scaffolds a new role named `nginx` inside the `roles/` folder.

#### ◆ Writing Role Content

Let's say you want to create a role for installing and configuring **Nginx**.

---

**Example: roles/nginx/tasks/main.yml**

```
---
```

```
- name: Install nginx

  apt:
    name: nginx
    state: present
    update_cache: yes

- name: Deploy nginx config

  template:
    src: nginx.conf.j2
    dest: /etc/nginx/sites-available/default
    notify: Restart nginx
```

**Example: roles/nginx/handlers/main.yml**

```
yaml
CopyEdit

---
```

```
- name: Restart nginx

  service:
    name: nginx
    state: restarted
```

**Example: roles/nginx/templates/nginx.conf.j2**

```
server {
  listen {{ app_port }};
  server_name {{ server_name }};
  root /var/www/{{ site_folder }};
```

{}

**Example: roles/nginx/defaults/main.yml**

```
app_port: 80
server_name: localhost
site_folder: html
```

**◆ Using a Role in a Playbook**

After creating your role, reference it in your main playbook like this:

---

```
- name: Deploy web server using nginx role
hosts: webservers
become: yes
```

```
roles:
```

```
  - nginx
```

That's it! Ansible will automatically look for the nginx role in roles/nginx/ and run all tasks, handlers, and templates.

**◆ Role Variables: Flexibility and Override**

You can override role defaults by passing variables:

---

```
- name: Use nginx role with custom vars
hosts: webservers
become: yes
```

```
roles:
```

---

- role: nginx

vars:

app\_port: 8080

server\_name: myapp.local

Or use host/group vars as shown earlier.

#### ◆ Sharing Roles via Ansible Galaxy (Optional)

If you want to share your role publicly, you can:

1. Create a GitHub repo for your role
2. Tag and push it
3. Publish it on galaxy.ansible.com

Or install other people's roles:

`ansible-galaxy install geerlingguy.nginx`

#### Summary

In this step, you've:

- Learned what Ansible roles are and why they matter
- Understood the folder structure and responsibilities
- Created your own nginx role with templates and handlers
- Used roles in your playbooks for better organization
- Explored how to customize and reuse them easily

Roles turn your one-off playbooks into **production-grade infrastructure automation**.

## Step 6: Testing & Best Practices

You've built a strong foundation: writing playbooks, managing variables, using templates and handlers, and modularizing everything with roles. Now it's time to **test, validate, and follow best practices** to ensure your infrastructure is reliable, secure, and maintainable.

### ◆ Why Testing Matters in Ansible

Ansible configurations **change infrastructure**, and without testing, a small error can:

- Break your app
- Misconfigure services
- Bring down production

Testing reduces human error and ensures consistent deployments.

### ◆ Types of Testing in Ansible

#### 1. Syntax Checks

Before running a playbook, always do a syntax check:

```
ansible-playbook playbook.yml --syntax-check
```

This catches indentation issues, unknown modules, malformed YAML, etc.

#### 2. Dry Run (Check Mode)

Simulate changes without applying them:

```
ansible-playbook playbook.yml --check
```

It tells you what *would* change, which is great for reviewing changes before pushing to production.

#### 3. Linting (Static Code Analysis)

---

Use ansible-lint to catch common anti-patterns:

`ansible-lint playbook.yml`

You'll get warnings for:

- Deprecated modules
- Hardcoded values
- Missing handlers
- Improper role structures

Install with:

`pip install ansible-lint`

## 4. Molecule for Role Testing

Molecule is a tool designed for **testing Ansible roles** in isolated environments (Docker, Vagrant, etc.).

To install:

`pip install molecule[docker]`

To test a role:

`cd roles/nginx`

`molecule init scenario -r nginx -d docker`

`molecule test`

Molecule runs:

- Linting
- Dependency install
- Converge (provisioning)
- Verify (testing)
- Destroy

Great for **CI/CD pipelines** and local test cycles.

---

◆ Best Practices for Writing Ansible

 1. Use Roles from Day One

Even for small projects, roles promote clean structure and reusability. Avoid long, monolithic playbooks.

 2. Stick to Declarative Patterns

Use state: present/absent, enabled: true, and similar idempotent options. Avoid scripting logic unless necessary.

 3. Use Variables and Defaults

Keep playbooks flexible by extracting values into defaults/, vars/, or external files.

 4. Document Everything

- Use name: for each task (makes output readable)
- Add inline comments for templates and tricky vars
- Version control everything with Git

 5. Limit Hardcoding

Avoid hardcoding values in tasks. Use variables, group/host vars, or encrypted secrets with **Ansible Vault**.

◆ Optional: Ansible Vault for Secrets

To keep sensitive info (passwords, API keys) safe:

**Create a vault file:**

`ansible-vault create secrets.yml`

**Edit an existing vault file:**

`ansible-vault edit secrets.yml`

**Use in playbooks:**

`vars_files:`

`- secrets.yml`

---

**Run playbooks with vault:**

`ansible-playbook playbook.yml --ask-vault-pass`

Or use a vault password file for automation.

◆ **Integrating with CI/CD**

Ansible can be run in pipelines like GitHub Actions, GitLab CI, Jenkins, etc.

A sample GitHub Action might:

- Checkout repo
- Set up Python and Ansible
- Run linting and playbooks
- Use Molecule for testing roles

Automation ensures every change is tested before deployment.

## Summary

In this final step, you've learned to:

- Test playbooks safely using syntax check, dry run, and linting
- Use Molecule to test roles like code
- Protect secrets using Ansible Vault
- Structure your code for readability and reusability
- Integrate Ansible into DevOps workflows using CI/CD

---

## 🏁 Conclusion

Ansible empowers you to automate infrastructure with clarity, consistency, and control. Through this 6-step guide, you've gone from the fundamentals to advanced best practices—transforming your configurations into well-structured, declarative, and reusable automation code.

You started by setting up Ansible and understanding how inventory and playbooks define your infrastructure. Then, you explored the power of variables, templates, and handlers to keep things dynamic and DRY (Don't Repeat Yourself). From there, you modularized your approach using roles, laying the groundwork for scalable, team-friendly automation. Finally, you strengthened your practice with testing strategies, secrets management, and battle-tested best practices to make your playbooks production-grade.

By following these steps, you've not only learned **how to write Ansible code**—you've learned how to write it **the right way**.

Now, whether you're managing a few servers or scaling across data centers and cloud environments, your infrastructure is under version control, reproducible, and secure. You can confidently integrate it into your CI/CD workflows and collaborate with teams across environments.

### 🔧 Next Steps & Recommendations:

- Explore **Ansible Galaxy** for ready-made roles
- Dive into **Ansible Tower / AWX** for GUI-based control
- Start integrating Ansible into **CI/CD pipelines**
- Keep your skills sharp with **real-world use cases**