

TOP 52

# AZURE DATA ENGINEER

Interview Questions



*For Data Engineer*



## \*Disclaimer\*

System Design is the most asked topic in tech interviews.

So, make sure you prepare it thoroughly.

**Take the help of this doc and ace your System Design Interviews.**

## Key Answers:

### QUESTION-1

#### Parameters and Variables in ADF:

- **Parameters:** Used to pass values at runtime to pipelines. They are defined at the pipeline level and cannot be changed during execution.
- **Variables:** Used to store values within the pipeline and can change during execution. Variables are updated using **Set Variable** or **Append Variable** activities.

### QUESTION-2

#### Time Travel in Your Project:

- Time travel is a **Delta Lake** feature that allows querying historical data (snapshots).
- Example: `SELECT * FROM table_name VERSION AS OF 5` or `TIMESTAMP AS OF '2023-01-15'`.
- Use Case: Debugging, auditing, or recreating datasets for ML models.

## QUESTION-3

### Resume Pipeline from Failed Activity:

- Enable **checkpointing** or **activity retry** in ADF. Use a failure path with logic to resume execution by using the **Get Metadata** activity to evaluate where the pipeline failed.

## QUESTION-4

### Pipelines You've Worked With:

- Example: ETL pipelines to ingest and transform raw data from Azure Data Lake using Data Flows and Spark jobs.
- Mention specifics like **copy data activities**, **data validation**, and **orchestration of transformations**.

## QUESTION-5

### Partition vs. Bucketing:

- **Partitioning:** Divides the data into directories based on keys (e.g., year, month).
- **Bucketing:** Hashes data into fixed-sized buckets, optimizing joins and aggregations.



## QUESTION-6

### Medallion Architecture:

- A data architecture that separates data into three layers:
  - **Bronze:** Raw ingested data.
  - **Silver:** Cleaned and transformed data.
  - **Gold:** Business-level aggregates and insights.

## QUESTION-7

### Azure Key Vault:

- Securely stores secrets, keys, and certificates.
- Use **Managed Identity** in ADF to access Key Vault without hardcoding credentials.

## QUESTION-8

### Unity Catalog vs. Hive Metastore:

- **Unity Catalog:** Centralized data governance and access control for all your Databricks workspaces.
- **Hive Metastore:** Manages metadata for Hive and Spark tables, but lacks robust access control.

## QUESTION-9

### Joins in PySpark:

- **Inner Join:** Matches rows from both datasets based on a condition.
- **Left/Right Join:** Keeps all rows from the left/right and matches with the right/left dataset.
- **Full Outer Join:** Includes all rows from both datasets.
- **Cross Join:** Cartesian product of both datasets.

**Want further explanations, detailed examples, or help structuring your answers? 😊**

💡 These are excellent questions to enhance your preparation for Azure Data Engineer interviews. Here's a breakdown of concise, yet impactful answers:

## QUESTION- 10

### How to Implement Parallel Processing in ADF?

- Use **For Each Activity** with the **Batch Count** property set for parallelism.
- Enable **concurrent execution** in pipeline settings.
- Use partitioned datasets for parallel reads/writes to optimize execution.

## QUESTION- 11

### Difference Between Narrow and Wide Transformations:

- **Narrow:** Data is processed within the same partition (e.g., map, filter). Minimal shuffling.
- **Wide:** Data is shuffled across partitions (e.g., groupBy, join). Higher computational cost.

## QUESTION-12

### What is SCD? Explain SCD1, SCD2, SCD3:

- **SCD (Slowly Changing Dimensions)** handles historical changes in dimension data.
- **SCD1:** Overwrites old data with new data.
- **SCD2:** Maintains history by adding new rows for changes (e.g., adding an Effective\_Date ).
- **SCD3:** Adds new columns to store historical data for specific attributes.

## QUESTION-13

### Cluster Options in Databricks:

- **Standard Cluster:** For general-purpose workloads.
- **High-Concurrency Cluster:** Optimized for multiple concurrent users.
- **Single Node Cluster:** For lightweight testing and debugging.
- **Jobs Cluster:** Automatically created for specific jobs and deleted afterward.

## QUESTION-14

### Difference Between Managed and External Tables:

- **Managed Tables:** Databricks manages the data and metadata (stored in default storage).
- **External Tables:** Data is stored outside Databricks, and only metadata is managed in the metastore.

## QUESTION-15

### What is a Surrogate Key?

- A unique identifier for a record, not derived from application data.
- Example: Auto-increment ID in databases.

## QUESTION-16

### Spark Optimization Techniques:

- Cache/persist frequently used data.
- Use broadcast joins for smaller datasets.
- Partition data effectively.
- Enable predicate pushdown for filters.
- Avoid wide transformations where possible.

## QUESTION-17

### Why is Databricks Better Than Dataflow?

- **Flexibility:** Databricks supports more complex workloads (e.g., ML, streaming).
- **Notebook Interface:** Collaborative development environment.
- **Performance:** Databricks uses Apache Spark with optimizations like Delta Lake.
- **Dataflow** is simpler for straightforward ETL use cases.

## QUESTION-18

### Difference Between Data Lake and Delta Lake:

- **Data Lake:** Stores raw, unstructured data. No ACID compliance.
- **Delta Lake:** Built on top of a data lake with ACID transactions, time travel, and schema enforcement.

## QUESTION-19

### Spark Optimization Techniques:

- **Driver:** Coordinates execution, maintains DAG, and schedules tasks.
- **Executors:** Run tasks assigned by the driver. Each executor has its memory and cache.
- **Cluster Manager:** (e.g., YARN, Kubernetes) Allocates resources to the driver and executors.

**Need examples for any of these? Or a deeper dive into any topic? 🚀**

**Here's a solid overview of answers to these questions, tailored to help you shine in interviews! ✨**

## QUESTION- 20

### Difference Between **groupByKey** and **reduceByKey**:

- **groupByKey**: Groups all key-value pairs by key and shuffles all data. More memory-intensive.
- **reduceByKey**: Combines values at the mapper side before shuffling, reducing network traffic. Preferred for better performance.

## QUESTION- 21

### Why is MapReduce Not Widely Used Now? Similarities Between Spark and MapReduce?

- Why not MapReduce:
  - High latency due to disk I/O for intermediate results.
  - Complex to code compared to Spark.

- **Similarities:**

- Both process large-scale data using distributed computing.
- Use key-value pairs for transformations.

- **Spark Advantages:**

- In-memory computation, faster execution, rich APIs (Python, Scala).

## QUESTION- 22

### **What is Delta Lake? Key Features and Creating Delta Tables:**

- **Delta Lake:** A storage layer on top of Data Lake offering ACID compliance and reliability.

- **Key Features:**

- ACID transactions.
- Schema enforcement and evolution.
- Time travel and versioning.

- **Creating Delta Tables:**

```
CREATE TABLE delta_table USING DELTA LOCATION  
'path_to_delta';
```

## QUESTION-23

### Difference Between Serverless Pool and Dedicated SQL Pool:

- **Serverless Pool:**
  - Pay-per-query model.
  - Used for ad-hoc queries on data lakes.
- **Dedicated SQL Pool:**
  - Pre-provisioned resources with fixed cost.
  - Designed for high-performance data warehousing.

## QUESTION-24

### Prerequisites Before Migration:

- Assess source and target environments.
- Ensure schema compatibility.
- Perform data profiling and cleansing.
- Set up network, storage, and permissions.
- Validate data transformation logic.

## QUESTION-25

### What is a Mount Point in Databricks?

- A **mount point** is a shortcut to a storage account, enabling easier access.
- Example: Mounting an Azure Data Lake Gen2 folder using a dbutils.fs.mount command.

## QUESTION-26

### How to Optimize Databricks Performance:

- Enable **Delta Lake optimizations** like Z-ordering and OPTIMIZE.
- Use **Auto-scaling** for clusters.
- Use **broadcast joins** for smaller datasets.
- Optimize shuffling with correct partitioning.
- Persist reusable datasets in memory with `cache()`.

## QUESTION-27

### Difference Between map and flatMap:

- **map:** Transforms each element into another element, 1-to-1 mapping.
- **flatMap:** Can produce 0 or more elements per input, 1-to-n mapping.

## QUESTION-28

### How to Fetch Details from Key Vault:

- Use Azure Key Vault Linked Service in ADF or Databricks.
- In Databricks:

```
secret_value = dbutils.secrets.get(scope="key_vault_scope",  
key="secret_name")
```

## QUESTION-29

### Applying Indexing on a Databricks Table:

- Use Delta Lake Z-order indexing:

```
OPTIMIZE delta_table_name ZORDER BY (column_name);
```

- Helps improve query performance for large datasets.

## QUESTION-30

### Transferring Data to Azure Synapse:

- Use Azure Data Factory for ETL pipelines.
- **COPY INTO** command in Synapse for fast ingestion from Data Lake.
- Databricks-to-Synapse via JDBC connector or PolyBase.

## QUESTION- 31

### What is Incremental Loading? How to Implement It?

- **Definition:** Loading only new or updated data to a target without reloading the entire dataset.
- **Implementation:**
  - **Watermarking:** Use timestamps or surrogate keys to identify changes.
  - **ADF:** Use Lookup + Filter activities.
  - **Delta Lake:** Merge using UPSERT logic:

```
• MERGE INTO target_table AS target
• USING source_table AS source
• ON target.id = source.id
• WHEN MATCHED THEN UPDATE SET target.col = source.col
• WHEN NOT MATCHED THEN INSERT (columns) VALUES (values);
```

**Need any of these elaborated further or some live coding examples? 🚀**

**Here's a breakdown of these advanced Azure Data Engineering topics to keep your prep on point! 🚀**

## QUESTION-32

### How Does Z-Ordering Work?

**Z-Ordering:** A data layout optimization in Delta Lake that reduces I/O by co-locating similar data on disk.

- **How:**

- Applies a multi-dimensional sort algorithm.
- Improves query performance on frequently filtered columns. `OPTIMIZE table_name ZORDER BY (column1, column2);`

## QUESTION-33

### What is Dimension Modeling? Dimension and Fact Tables?

- **Dimension Modeling:** A design technique for data warehouses to optimize query performance using star or snowflake schemas.
- **Fact Tables:** Store numeric measures (e.g., sales amount).
- **Dimension Tables:** Describe the context of facts (e.g., customer, product).

## QUESTION-34

### Difference Between a Data Lake and a Data Warehouse:

- **Data Lake:**

- Stores raw, unstructured data.
- Scalable, cost-effective.
- Example: Azure Data Lake.

- **Data Warehouse:**

- Stores structured, processed data for analytics.
- Schema-on-write.
- Example: Azure Synapse.

## QUESTION-35

### Using Logic Apps in Your Project:

- Automates workflows between services like ADF, Synapse, and notifications.
- Example Use Case:
  - Trigger data pipelines based on events (e.g., file upload).
  - Send failure alerts via email or Teams.

## QUESTION-36

### What is Data Skewness?

- **Definition:** Uneven distribution of data across partitions, leading to performance bottlenecks.
- **Mitigation:**
  - Use salting techniques (adding random keys).
  - Optimize partitioning with balanced keys.

## QUESTION-37

### What is Fault Tolerance and Its Use in Real-Time Applications?

- **Definition:** The ability of a system to recover from failures.
- **Real-Time Use:**
  - Spark achieves fault tolerance by storing lineage and recomputing lost partitions.
  - In ADF, retry policies handle transient failures.

## QUESTION-38

### Converting RDD to DataFrame & Vice Versa:

- **RDD to DataFrame:**

- `from pyspark.sql import SparkSession`

```
df = rdd.toDF(schema=["col1", "col2"])
```

- **DataFrame to RDD:**

```
rdd = df.rdd
```

## QUESTION-39

### Encryption Techniques:

- **At Rest:** Encrypt data in storage using Azure Storage Service Encryption (SSE).
- **In Transit:** Use TLS/SSL for secure data transfer.
- **Column-Level Encryption:** Secure sensitive data fields (e.g., PII).

## QUESTION-40

### How Does Auto Loader Work?

- A feature in Databricks for **incremental file processing** from cloud storage.
- **Working:**
  - Tracks metadata using checkpointing.
  - Processes new files automatically.
- **Example:**
  - `df = spark.readStream.format("cloudFiles") \ .option("cloudFiles.format", "json") \ .load("path")`

## QUESTION-41

### Explain Lazy Evaluation in PySpark:

**Definition:** Transformations are not executed immediately but only when an action (e.g., `count` , `collect` ) is triggered.

- **Benefits:**
  - Optimizes execution by combining transformations into a single stage.
  - Reduces unnecessary computations.

## QUESTION- 42

### What is DAG in Spark?

- **DAG (Directed Acyclic Graph):**
  - A sequence of computations where each node represents a transformation and edges represent dependencies.
  - Spark breaks the execution into stages using DAG, ensuring fault tolerance and optimized execution.
- **Significance:**
  - Tracks lineage for fault recovery.
  - Optimizes execution by combining transformations.

**Want any topic expanded with examples or real-world scenarios? Let me know! 🚀**

**Here's a detailed explanation of these additional Spark and PySpark-related questions:**

## QUESTION-43

### Significance of Catalyst Optimizer in PySpark?

- **What It Is:** A query optimization engine in Spark SQL.
- **Functions:**
  - Converts logical plans into optimized physical plans.
  - Pushes predicates (filter operations) early to minimize I/O.
- **Benefits:** Better performance with optimized execution plans.

## QUESTION-44

### Query to Find the 4th Highest Salary of an Employee:

```
SELECT DISTINCT salary
```

```
FROM employee
```

```
ORDER BY salary DESC
```

```
LIMIT 4 OFFSET 3;
```

- Alternatively, using **ROW\_NUMBER**:
- ```
SELECT salary
```
- ```
FROM (
```
- ```
SELECT salary, ROW_NUMBER() OVER (ORDER BY salary DESC) AS rank
```
- ```
FROM employee
```
- ```
) ranked
```

```
WHERE rank = 4;
```

## QUESTION-45

### PySpark Command to Read Data from a File into a DataFrame:

```
df = spark.read.csv("path/to/file.csv", header=True, inferSchema=True)
```

- Other Formats:

- JSON: spark.read.json("path")
- Parquet: spark.read.parquet("path")

## QUESTION-46

### Handling Nulls and Duplicates in PySpark:

- Drop Nulls: df = df.dropna()
- Fill Nulls: df = df.fillna({'col1': 'default\_value', 'col2': 0})
- Remove Duplicates: df = df.dropDuplicates(['col1', 'col2'])

## QUESTION-47

### Changing the Date Format for a Date Column:

```
from pyspark.sql.functions import date_format  
  
df = df.withColumn("new_date", date_format("date_column", "yyyy-MM-dd"))
```

## QUESTION-48

### What is the Explode Function in PySpark?

- **Explode:** Converts an array or map into multiple rows.
- **Example:**
- from pyspark.sql.functions import explode

```
df = df.withColumn("exploded_col", explode("array_col"))
```

## QUESTION-49

### Code to Read a Parquet File:

```
df = spark.read.parquet("path/to/file.parquet")
```

## QUESTION-50

### Code to Add a Column to a Parquet File:

```
from pyspark.sql.functions import lit  
  
df = spark.read.parquet("path/to/file.parquet")  
  
df = df.withColumn("new_column", lit("value"))  
  
df.write.parquet("path/to/updated_file.parquet")
```

## QUESTION-51

### Different Approaches to Creating RDD in PySpark:

- From a Collection:

```
rdd = spark.sparkContext.parallelize([1, 2, 3, 4])
```

- From a File:

```
rdd = spark.sparkContext.textFile("path/to/file.txt")
```

## QUESTION-49

### Code to Read a Parquet File:

```
df = spark.read.parquet("path/to/file.parquet")
```

## QUESTION-52

### Different Approaches to Creating DataFrame in PySpark:

- **From RDD:**

- ```
from pyspark.sql import Row
```

```
rdd = spark.sparkContext.parallelize([Row(name="Alice", age=25),  
Row(name="Bob", age=30)])
```

```
df = rdd.toDF()
```

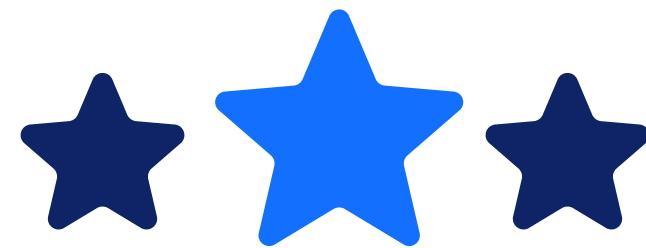
- **From a File:**

```
df = spark.read.csv("path/to/file.csv", header=True, inferSchema=True)
```

- **From a List/Dictionary:**

- ```
data = [("Alice", 25), ("Bob", 30)]
```

```
df = spark.createDataFrame(data, schema=["name", "age"])
```



# WHY BOSSCODER?

 **2200+ Alumni** placed at Top Product-based companies.

 More than **136% hike** for every 2 out of 3 Working Professional.

 Average Package of **24LPA**.

The syllabus is most up-to-date and the list of problems provided covers all important topics.

**Lavanya**  
 Meta



Course is very well structured and streamlined to crack any MAANG company .

**Rahul**  
 Google



[EXPLORE MORE](#)