



**MICROSOFT
CERTIFIED
AZURE DATA
SCIENTIST
ASSOCIATE**
(DP-100)



CloudCertificationStore

Microsoft Certified: Azure Data Scientist Associate (DP-100) - Practice Exam Questions (AZ-DP-100-0010)

© [Cloud Certification Store](https://cloudcertificationstore.com) All rights reserved.

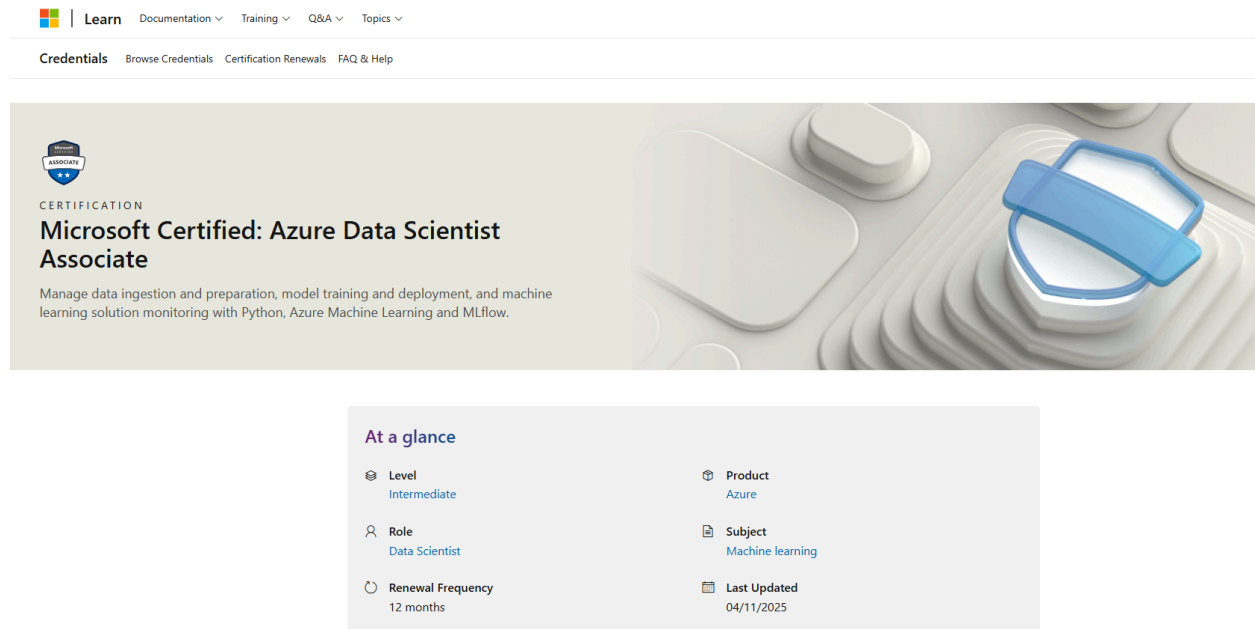
Microsoft Azure is a registered trademark of Microsoft Corporation.

This practice set is an original work for educational use and is NOT endorsed by or affiliated with Microsoft Corporation. “Microsoft,” “Microsoft Azure,” “Microsoft Certified: Azure Solutions Architect Expert,” and related marks are trademarks of Microsoft Corporation, used here for identification only.

DISCLAIMER

- This practice test includes questions **compiled from various exam preparation platforms**.
- Important: **Questions and answers were AI-assisted and human-curated**. Verify accuracy with official documentation before relying on this material.
- **Users are strongly encouraged to double-check all content against official documentation and trusted sources** before using it for exam preparation or making important decisions.
- The creators of this material assume **no responsibility** for any errors, inaccuracies, or outcomes, including exam results, based on the use of this content.
- **Some questions might be duplicated or close** to previous ones, this is done on purpose as a way to re-inforce your learning.
- Single-user licence only
 - Includes one unique Payhip Licence Key per purchase, along with a Product Key.
 - Redistribution, resale, or public posting is prohibited. We can trace any file to the purchaser, with the use of the purchased License Key and Product Key.

Microsoft Certified: Azure Data Scientist Associate (DP-100)



The screenshot shows the Microsoft Learn page for the Microsoft Certified: Azure Data Scientist Associate (DP-100) certification. The page features a navigation bar with links to Learn, Documentation, Training, Q&A, and Topics. Below the navigation bar, there is a section for Credentials, including Browse Credentials, Certification Renewals, and FAQ & Help. The main content area displays the certification logo, the title "Microsoft Certified: Azure Data Scientist Associate", and a brief description: "Manage data ingestion and preparation, model training and deployment, and machine learning solution monitoring with Python, Azure Machine Learning and MLflow." Below this, there is a "At a glance" section with a table of key details.

At a glance	
Level	Intermediate
Product	Azure
Role	Data Scientist
Subject	Machine learning
Renewal Frequency	12 months
Last Updated	04/11/2025

Overview

As a candidate for this certification, you should have subject matter expertise in applying data science and machine learning to implement and run machine learning workloads on Azure. Additionally, you should have knowledge of optimizing language models for AI applications using Azure AI.

Your responsibilities for this role include:

- Designing and creating a suitable working environment for data science workloads.
- Exploring data.
- Training machine learning models.
- Implementing pipelines.
- Running jobs to prepare for production.
- Managing, deploying, and monitoring scalable machine learning solutions.
- Using language models for building AI applications.

As a candidate for this certification, you should have knowledge and experience in data science by using:

- Azure Machine Learning
- MLflow
- Azure AI services, including Azure AI Search
- Azure AI Foundry

Skills earned upon completion

Prepare for the exam

Course

[Designing and implementing a data science solution on Azure](#)

Training in this course

[Explore and configure the Azure Machine Learning workspace](#)

3 hr 37 min

Learning Path

5 modules

[Experiment with Azure Machine Learning](#)

1 hr 10 min

Learning Path

2 modules

[Optimize model training with Azure Machine Learning](#)

2 hr 37 min

Learning Path

4 modules

[Manage and review models in Azure Machine Learning](#)

1 hr 22 min

Learning Path

2 modules

[Deploy and consume models with Azure Machine Learning](#)

1 hr 29 min

Learning Path

2 modules

[Develop generative AI apps in Azure](#)

7 hr 29 min

Learning Path

8 modules

Develop generative AI apps in Azur...

Practice for the exam

Practice Assessment

Assess your knowledge

Practice assessments provide you with an overview of the style, wording, and difficulty of the questions you're likely to experience on the exam. Through these assessments, you're able to assess your readiness, determine where additional preparation is needed, and fill knowledge gaps bringing you one step closer to the likelihood of passing your exam.

<https://learn.microsoft.com/en-us/credentials/certifications/azure-data-scientist/practice/assessment?assessment-type=practice&assessmentId=62&practice-assessment-type=certification>

Exam Sandbox

<https://go.microsoft.com/fwlink/?linkid=2226877>

Video

Exam DP-100 prep videos

Join our experts as they provide tips, tricks, and strategies for preparing for this Microsoft Certification exam.

<https://learn.microsoft.com/en-us/shows/exam-readiness-zone/preparing-for-dp-100-design-and-prepare-a-machine-learning-solution-1-of-4>

Take the exam

You will have 100 minutes to complete this assessment.

Exam policy

This exam will be proctored. You may have interactive components to complete as part of this exam. To learn more about exam duration and experience, visit: [Exam duration and exam experience](#).

If you fail a certification exam, don't worry. You can retake it 24 hours after the first attempt. For subsequent retakes, the amount of time varies. For full details, visit: [Exam retake policy](#).

Assessed on this exam

- Design and prepare a machine learning solution
- Explore data, and run experiments
- Train and deploy models
- Optimize language models for AI applications

Schedule the Exam

<https://learn.microsoft.com/en-us/credentials/certifications/schedule-through-pearson-vue?examUid=exam.DP-100&examUrl=https://learn.microsoft.com/credentials/certifications>

Price

\$165 USD*

Price based on the country or region in which the exam is proctored.

Renew your certification

Do you know that Microsoft role-based and specialty certifications expire unless they are renewed? Learn the latest updates to the technology for your job role and renew your certification at no cost by passing an online assessment on Microsoft Learn.

[Learn more about renewing](#)

Certification resources

[Exam DP-100 study guide](#)

Focus your studies as you prepare for the exam. Review the study guide to learn about the topics the exam covers, updates, and additional resources.

[Certification poster](#)

Check out an overview of fundamentals, role-based, and specialty certifications.

Exam Replay

Boost your odds of success with this great offer.

Support for credentials

Get help through Microsoft Credentials support forums. A forum moderator will respond in one business day, Monday–Friday.

Choose your Microsoft Credential

Microsoft Applied Skills or Microsoft Certifications? Choose the path that fits your career goals, desired skillset, and schedule.

PREVIEW

Study guide for Exam DP-100:

Designing and Implementing a Data Science Solution on Azure

03/12/2025

Purpose of this document

This study guide should help you understand what to expect on the exam and includes a summary of the topics the exam might cover and links to additional resources. The information and materials in this document should help you focus your studies as you prepare for the exam.

Useful links	Description
How to earn the certification	Some certifications only require passing one exam, while others require passing multiple exams.
Certification renewal	Microsoft associate, expert, and specialty certifications expire annually. You can renew by passing a free online assessment on Microsoft Learn.
Your Microsoft Learn profile	Connecting your certification profile to Microsoft Learn allows you to schedule and renew exams and share and print certificates.

Exam scoring and score reports	A score of 700 or greater is required to pass.
Exam sandbox	You can explore the exam environment by visiting our exam sandbox.
Request accommodation s	If you use assistive devices, require extra time, or need modification to any part of the exam experience, you can request an accommodation.
Take a free Practice Assessment	Test your skills with practice questions to help you prepare for the exam.

Updates to the exam

We always update the English language version of the exam first. Some exams are localized into other languages, and those are updated approximately eight weeks after the English version is updated. While Microsoft makes every effort to update localized versions as noted, there may be times when the localized versions of an exam are not updated on this schedule. Other available languages are listed in the Schedule Exam section of the Exam Details webpage. If the exam isn't available in your preferred language, you can request an additional 30 minutes to complete the exam.

Note

The bullets that follow each of the skills measured are intended to illustrate how we are assessing that skill. Related topics may be covered in the exam.

Note

Most questions cover features that are general availability (GA). The exam may contain questions on Preview features if those features are commonly used.

Skills measured as of April 11, 2025

Audience profile

As a candidate for this exam, you should have subject matter expertise in applying data science and machine learning to implement and run machine learning workloads on Azure. Additionally, you should have knowledge of optimizing language models for AI applications using Azure AI.

Your responsibilities for this role include:

- Designing and creating a suitable working environment for data science workloads.
- Exploring data.
- Training machine learning models.
- Implementing pipelines.
- Running jobs to prepare for production.
- Managing, deploying, and monitoring scalable machine learning solutions.
- Using language models for building AI applications.

As a candidate for this exam, you should have knowledge and experience in data science by using:

- Azure Machine Learning
- MLflow

- Azure AI services, including Azure AI Search
- Azure AI Foundry

Skills at a glance

- Design and prepare a machine learning solution (20–25%)
- Explore data, and run experiments (20–25%)
- Train and deploy models (25–30%)
- Optimize language models for AI applications (25–30%)

Design and prepare a machine learning solution (20–25%)

Design a machine learning solution

- Identify the structure and format for datasets
- Determine the compute specifications for machine learning workload
- Select the development approach to train a model

Create and manage resources in an Azure Machine Learning workspace

- Create and manage a workspace
- Create and manage datastores
- Create and manage compute targets
- Set up Git integration for source control

Create and manage assets in an Azure Machine Learning workspace

- Create and manage data assets
- Create and manage environments
- Share assets across workspaces by using registries

Explore data, and run experiments (20–25%)

Use automated machine learning to explore optimal models

- Use automated machine learning for tabular data
- Use automated machine learning for computer vision
- Use automated machine learning for natural language processing
- Select and understand training options, including preprocessing and algorithms
- Evaluate an automated machine learning run, including responsible AI guidelines

Use notebooks for custom model training

- Use the terminal to configure a compute instance
- Access and wrangle data in notebooks
- Wrangle data interactively with attached Synapse Spark pools and serverless Spark compute
- Retrieve features from a feature store to train a model
- Track model training by using MLflow
- Evaluate a model, including responsible AI guidelines

Automate hyperparameter tuning

- Select a sampling method
- Define the search space
- Define the primary metric
- Define early termination options

Train and deploy models (25–30%)

Run model training scripts

- Consume data in a job
- Configure compute for a job run

- Configure an environment for a job run
- Track model training with MLflow in a job run
- Define parameters for a job
- Run a script as a job
- Use logs to troubleshoot job run errors

Implement training pipelines

- Create custom components
- Create a pipeline
- Pass data between steps in a pipeline
- Run and schedule a pipeline
- Monitor and troubleshoot pipeline runs

Manage models

- Define the signature in the MLmodel file
- Package a feature retrieval specification with the model artifact
- Register an MLflow model
- Assess a model by using responsible AI principles

Deploy a model

- Configure settings for online deployment
- Deploy a model to an online endpoint
- Test an online deployed service
- Configure compute for a batch deployment
- Deploy a model to a batch endpoint
- Invoke the batch endpoint to start a batch scoring job

Optimize language models for AI applications (25–30%)

Prepare for model optimization

- Select and deploy a language model from the model catalog
- Compare language models using benchmarks
- Test a deployed language model in the playground
- Select an optimization approach

Optimize through prompt engineering and prompt flow

- Test prompts with manual evaluation
- Define and track prompt variants
- Create prompt templates
- Define chaining logic with the prompt flow SDK
- Use tracing to evaluate your flow

Optimize through Retrieval Augmented Generation (RAG)

- Prepare data for RAG, including cleaning, chunking, and embedding
- Configure a vector store
- Configure an Azure AI Search-based index store
- Evaluate your RAG solution

Optimize through fine-tuning

- Prepare data for fine-tuning
- Select an appropriate base model
- Run a fine-tuning job
- Evaluate your fine-tuned model

Study resources

We recommend that you train and get hands-on experience before you take the exam. We offer self-study options and classroom training as well as links to documentation, community sites, and videos.

Study resources	Links to learning and documentation
Get trained	Choose from self-paced learning paths and modules or take an instructor-led course
Find documentation	Azure Databricks Azure Machine Learning Azure Synapse Analytics MLflow and Azure Machine Learning
Ask a question	Microsoft Q&A Microsoft Docs
Get community support	AI - Machine Learning - Microsoft Tech Community AI - Machine Learning Blog - Microsoft Tech Community
Follow Microsoft Learn	Microsoft Learn - Microsoft Tech Community
Find a video	Microsoft Learn Shows

Change log

The table below summarizes the changes between the current and previous version of the skills measured. The functional groups are in bold typeface followed by the objectives within each group. The table is a comparison between the previous and current version of the exam skills measured and the third column describes the extent of the changes.

Skill area prior to January 16, 2025	Skill area as of January 16, 2025	Change
Audience profile		Minor
Optimize language models for AI applications	Optimize language models for AI applications	No % change
Optimize through prompt engineering and Prompt flow	Optimize through prompt engineering and prompt flow	Minor

Additional resources

Documentation

[Study guide for Exam AZ-500: Microsoft Azure Security Technologies](#)

Study guide for Exam AZ-500: Microsoft Azure Security Technologies | Microsoft Docs

[Study guide for Exam AZ-204: Developing Solutions for Microsoft Azure](#)

Study guide for Exam AZ-204: Developing Solutions for Microsoft Azure | Microsoft Docs

[Study guide for Exam AI-900: Microsoft Azure AI Fundamentals](#)

Study guide for Exam AI-900: Microsoft Azure AI Fundamentals | Microsoft Docs

[Practice Assessment](#)

Practice Assessment

Training

Learning path

[Solution Architect: Design Microsoft Power Platform solutions - Training](#)

Learn how a solution architect designs solutions.

Certification

[Microsoft Certified: Azure Data Scientist Associate - Certifications](#)

Manage data ingestion and preparation, model training and deployment, and machine learning solution monitoring with Python, Azure Machine Learning and MLflow.

This is a PREVIEW, get the full version here

<https://cloudcertificationstore.com/b/hvask>

Practice Questions - PREVIEW 20 out of 500+

Question 1

You must train a machine learning model in Azure Machine Learning using a large dataset stored in Azure Blob Storage. The goal is to minimize data-copy overhead while enabling distributed training with dynamic scaling.

- A. Compute Instance
- B. Compute Cluster
- C. Attached Compute to Synapse Spark Pool
- D. Kubernetes Service

✓ **Correct Answer: C. Attached Compute to Synapse Spark Pool**

📌 Using an attached Synapse Spark pool provides serverless distributed compute, allowing direct access to data in Blob Storage and dynamic scaling for big-data workloads. It eliminates unnecessary data movement and is fully integrated with Azure ML pipelines.

✗ **Why not the other options?**

A. Compute Instance: A single-node environment ideal for experimentation, not for distributed or large-scale training.

B. Compute Cluster: Suitable for parallel training but still requires data copy into compute nodes, adding overhead.

D. Kubernetes Service: Used primarily for deployment and inference, not for distributed training integration with Blob datasets.

Question 2

You need to evaluate a binary classification model and use precision as the primary metric. Which visualization should you choose in Azure Machine Learning Studio?

- A. ROC Curve
- B. Confusion Matrix
- C. Lift Chart
- D. Venn Diagram

✓ **Correct Answer: B. Confusion Matrix**

📌 The confusion matrix provides direct counts of true positives, false positives, and false negatives, which allows for easy computation of precision and recall. It is the standard visualization for assessing classification performance metrics.

❌ **Why not the other options?**

A. ROC Curve: Measures trade-offs between true and false positive rates, not precision specifically.

C. Lift Chart: Evaluates ranking performance in classification, not direct precision calculation.

D. Venn Diagram: A conceptual visualization, not used for quantitative metric evaluation in ML workflows.

Question 3

You plan to train a speech-recognition deep learning model on a Data Science Virtual Machine (DSVM). The model must support GPU acceleration and the latest Python version.

- A. Theano
- B. TensorFlow
- C. Scikit-learn
- D. CNTK

✅ **Correct Answer: B. TensorFlow**

📌 TensorFlow provides native CUDA support for GPUs and full compatibility with the latest Python versions, making it ideal for speech recognition models requiring deep learning on DSVM.

❌ **Why not the other options?**

A. Theano: Deprecated and no longer actively supported by Azure environments.

C. Scikit-learn: Designed for traditional ML, not for large-scale GPU-based deep learning.

D. CNTK: While Microsoft-developed, it's discontinued in favor of TensorFlow and PyTorch integration.

Question 4

A convolutional neural network (CNN) for image classification is showing signs of overfitting. You need to improve model generalization.

- A. Reduce training data size
- B. Add L1/L2 regularization and apply data augmentation
- C. Add more dense layers
- D. Remove dropout

✓ **Correct Answer: B. Add L1/L2 regularization and apply data augmentation**

📌 L1/L2 regularization penalizes large weights, while data augmentation generates diverse samples, both reducing overfitting and improving the model's ability to generalize unseen data.

✗ **Why not the other options?**

A. Reduce training data size: This would worsen overfitting by limiting exposure to variation.

C. Add more dense layers: Increases complexity, exacerbating overfitting.

D. Remove dropout: Dropout is used to combat overfitting; removing it increases model variance.

Question 5

In Azure Machine Learning Designer, you must split a dataset into training and testing subsets. Which module should you use?

- A. Group Categorical Values
- B. Split Data
- C. Clip Values
- D. Edit Metadata

✓ **Correct Answer: B. Split Data**

📌 The Split Data module is designed to partition datasets for training and testing, ensuring balanced sampling and reproducibility in ML experiments.

✗ **Why not the other options?**

A. Group Categorical Values: Used for merging string categories, not data partitioning.

C. Clip Values: Used to restrict numeric ranges, unrelated to data splitting.

D. Edit Metadata: Adjusts data types and field roles, not used for partitioning.

Question 6

You configure k-fold cross-validation for a dataset and must select the most common k value.

- A. 3
- B. 5 or 10
- C. 1
- D. 100

✓ **Correct Answer: B. 5 or 10**

📌 Using $k=5$ or $k=10$ provides a balance between computational cost and robust validation, minimizing bias and variance in model evaluation.

❌ **Why not the other options?**

- A. **3:** Provides weaker statistical confidence in validation.
 - C. **1:** Equivalent to no cross-validation.
 - D. **100:** Computationally excessive and unnecessary for most datasets.
-

Question 7

You use the Clean Missing Data module and choose "Replace with Median." When is this method appropriate?

- A. For categorical features
- B. For continuous numeric features
- C. For text features
- D. Never

✅ **Correct Answer: B. For continuous numeric features**

📌 The median is a robust central tendency measure for numeric data and helps handle outliers while imputing missing values effectively.

❌ **Why not the other options?**

- A. **Categorical features:** Should use mode or most frequent value instead.
 - C. **Text features:** Text cannot be imputed numerically.
 - D. **Never:** Median replacement is a standard numeric imputation technique.
-

Question 8

When configuring hyperparameter tuning in Azure ML, you want to iterate efficiently through parameter combinations with reduced compute cost. Which sweep mode should you select?

- A. Measured grid
- B. Entire grid
- C. Random grid
- D. Selective grid

✅ **Correct Answer: C. Random grid**

📌 Random grid sweep randomly samples from parameter space, achieving high coverage with lower compute demand than exhaustive grid search.

✗ Why not the other options?

- A. Measured grid:** Not a valid tuning mode.
 - B. Entire grid:** Tests all combinations, increasing cost.
 - D. Selective grid:** Not available as a sweep configuration.
-

Question 9

You must import a CSV dataset from a public web URL into Azure ML Designer with the least administrative effort.

- A. Import Data
- B. Dataset
- C. Copy Data
- D. Convert to TXT

✓ Correct Answer: A. Import Data

📌 The Import Data module enables direct ingestion from public URLs or cloud storage with no registration or prior setup, ideal for quick experiments.

✗ Why not the other options?

- B. Dataset:** Requires manual dataset registration.
 - C. Copy Data:** Used for ETL pipelines, not ingestion in Designer.
 - D. Convert to TXT:** Only changes file format, doesn't import data.
-

Question 10

Which compute target allows you to directly connect Azure Machine Learning Designer to a Spark-based big data environment for model training?

- A. Compute Instance
- B. Compute Cluster
- C. Attached Synapse Spark Pool
- D. Inference Cluster

✓ Correct Answer: C. Attached Synapse Spark Pool

📌 An attached Synapse Spark Pool provides seamless Spark integration, enabling large-scale distributed data processing directly from Azure ML Designer.

✗ Why not the other options?

- A. Compute Instance:** Local compute, not distributed.

- B. Compute Cluster:** Suitable for ML jobs but not Spark integration.
 - D. Inference Cluster:** Used for deployment, not training.
-

Question 11

You need to transform categorical features into binary columns suitable for model training. Which module should you use?

- A. Clean Missing Data
- B. Convert to Indicator Values
- C. Edit Metadata
- D. Group Categorical Values

✓ **Correct Answer: B. Convert to Indicator Values**

📌 The Convert to Indicator Values module performs one-hot encoding, converting categorical fields into binary numeric columns for ML algorithms.

✗ **Why not the other options?**

- A. Clean Missing Data:** Handles nulls, not encoding.
 - C. Edit Metadata:** Alters data types, not values.
 - D. Group Categorical Values:** Used for combining similar labels.
-

Question 12

You must visually identify outliers in a dataset. Which method provides this capability?

- A. ROC Curve
- B. Box Plot
- C. Histogram
- D. Line Chart

✓ **Correct Answer: B. Box Plot**

📌 A box plot visualizes the data distribution and highlights outliers through whisker boundaries, making it ideal for anomaly inspection.

✗ **Why not the other options?**

- A. ROC Curve:** Evaluates model performance, not raw data.
- C. Histogram:** Shows frequency but not outlier thresholds.
- D. Line Chart:** Displays trends, not statistical deviations.

Question 13

In a Python experiment, you need to record a metric for the number of rows processed. Which statement logs this metric in Azure ML?

- A. `run.upload_file('row_count', './data.csv')`
- B. `run.log('row_count', rows)`
- C. `run.log_row('row_count', rows)`
- D. `run.tag('row_count', rows)`

✓ **Correct Answer: B. `run.log('row_count', rows)`**

📌 The `run.log()` method records scalar values as metrics, retrievable after the experiment completes for tracking performance.

✗ Why not the other options?

- A. `upload_file`:** Uploads artifacts, not metrics.
 - C. `log_row`:** Used for tabular data, not single values.
 - D. `tag`:** Assigns labels, not numeric logs.
-

Question 14

You are developing a text translation model that must learn language sequences. Which architecture should you choose?

- A. CNN
- B. RNN
- C. GAN
- D. Transformer

✓ **Correct Answer: B. RNN**

📌 RNNs capture temporal dependencies in sequential data such as language, enabling context-aware translation tasks.

✗ Why not the other options?

- A. `CNN`:** Processes spatial data (images).
 - C. `GAN`:** Used for generative image synthesis.
 - D. `Transformer`:** Modern alternative but not explicitly required in this scenario.
-

Question 15

You must increase underrepresented class samples in your dataset. Which Azure ML module should you use?

- A. Join Data
- B. SMOTE
- C. Group Categorical Values
- D. Clip Values

✓ **Correct Answer: B. SMOTE**

📌 SMOTE (Synthetic Minority Oversampling Technique) generates synthetic minority samples, improving class balance for better classification.

✗ **Why not the other options?**

- A. Join Data:** Used for merging datasets.
 - C. Group Categorical Values:** Combines categories, not records.
 - D. Clip Values:** Used for numeric scaling, not sampling.
-

Question 16

You create a batch inference pipeline using Azure ML SDK. Where will the output file be located?

- A. Activity Log
- B. digit_identification.py
- C. parallel_run_step.txt in output folder
- D. Inference Clusters tab

✓ **Correct Answer: C. parallel_run_step.txt in output folder**

📌 Batch pipelines aggregate outputs into a single text file named `parallel_run_step.txt`, stored in the designated output directory.

✗ **Why not the other options?**

- A. Activity Log:** Tracks events, not outputs.
 - B. digit_identification.py:** Script source file, not result storage.
 - D. Inference Clusters:** Used for deployment endpoints, not batch logs.
-

Question 17

During AutoML classification, you must restrict experiments to linear algorithms. What should you configure?

- A. Enable automatic featurization
- B. Disable deep learning
- C. Turn off featurization
- D. Enable forecasting

✅ **Correct Answer: B. Disable deep learning**

📌 Disabling deep learning ensures AutoML evaluates only classical linear algorithms, such as logistic regression or linear SVMs.

❌ **Why not the other options?**

A. Enable automatic featurization: Influences preprocessing, not algorithm choice.

C. Turn off featurization: Reduces model accuracy, not algorithm control.

D. Forecasting: Changes task type, not model selection.

Question 18

When registering an MLflow model in Azure ML, which file defines its signature and dependencies?

- A. model.pkl
- B. MLmodel
- C. conda.yaml
- D. run.json

✅ **Correct Answer: B. MLmodel**

📌 The MLmodel file specifies the model's entry points, environment, and metadata for deployment consistency across platforms.

❌ **Why not the other options?**

A. model.pkl: Stores serialized weights, not metadata.

C. conda.yaml: Lists environment dependencies only.

D. run.json: Contains execution logs, not model definitions.

Question 19

You must select a machine learning environment that supports disconnected training using Caffe2 or Chainer frameworks on personal devices.

- A. Azure ML Service with DSVM
- B. Azure ML Studio
- C. Azure Databricks
- D. Azure Kubernetes Service

✓ **Correct Answer: A. Azure ML Service with DSVM**

📌 DSVM supports both Caffe2 and Chainer frameworks locally and can synchronize pipelines to Azure ML when reconnected.

✗ **Why not the other options?**

- B. ML Studio:** Requires full connectivity.
 - C. Databricks:** Cluster-based, not offline.
 - D. AKS:** For deployment, not training.
-

Question 20

You are optimizing a Retrieval-Augmented Generation (RAG) workflow. Which two Azure components are mandatory?

- A. Azure AI Search index store and vector store
- B. Azure Key Vault and Azure Monitor
- C. Compute Cluster and Batch Endpoint
- D. Prompt Flow and Pipeline Component

✓ **Correct Answer: A. Azure AI Search index store and vector store**

📌 RAG relies on a vector store for embeddings and an Azure AI Search index for retrieval, allowing LLMs to ground responses in enterprise data.

✗ **Why not the other options?**

- B. Key Vault & Monitor:** Handle secrets and metrics, not retrieval.
 - C. Compute Cluster & Batch Endpoint:** Manage compute, not RAG data.
 - D. Prompt Flow:** Useful for orchestration, but not required for retrieval storage.
-

(END OF PREVIEW QUESTIONS)

This is a PREVIEW, get the full version here

<https://cloudcertificationstore.com/b/hvask>

PREVIEW COPY

Final Review Checklist & Exam Readiness Scorecard

✓ How to Use the Final Review Checklist

This section is meant to **validate your hands-on skills and theoretical readiness** across all exam topics.




Step-by-step:

1. **Print it or load it in a note-taking app** (Notion, Google Docs, OneNote, etc.).
2. Go through each checkbox:
 - ✓ Check it if you **fully understand and can implement** the topic without looking up documentation.
 - ✗ Leave it unchecked if you feel unsure or haven't practiced the task.
3. Prioritize unchecked topics by reviewing:
 - Check the official documentation
 - Practice exams
 - Hands-on labs
4. For each **unchecked item**, write a short action plan or resource link next to it.

How to Use the Exam Readiness Scorecard



This part helps you **self-assess your confidence level** and **focus your revision time** wisely.

Instructions:

1. For each domain (e.g., "Hybrid connectivity and routing"), **rate yourself** from 1 to 5:
 -  = No understanding or hands-on practice
 -  = Moderate familiarity, but need review
 -  = Mastered topic and can apply it in real-world use
2. Add **Notes / Action Items** to explain:
 - Why you scored yourself low
 - What resources you'll use to improve (YouTube, whitepapers, exam guides)
 - Practice test scores if relevant
3. Reassess **2–3 days before your exam**, and compare scores to measure improvement.

Bonus Tips

- Do **timed mock exams** and cross-reference errors with checklist topics
- Use the scorecard to **simulate an exam debrief**: where did you fail? What must you strengthen?

Once all checklist items are  and all categories are at **4–5 stars** and you're consistently scoring **85%+** on full practice exams with confidence in scenario-based reasoning, then  you're likely **ready to book the real exam**.

Final Review Checklist

Design and Prepare a Machine Learning Solution (20 – 25%)

- ☐ Identify dataset structures, formats, and compute requirements
- ☐ Choose appropriate ML algorithms and training approaches
- ☐ Create and manage Azure Machine Learning workspaces
- ☐ Create and manage datastores and compute targets
- ☐ Configure Git integration for version control
- ☐ Create and manage data assets and environments
- ☐ Share assets across workspaces using registries

Explore Data and Run Experiments (20 – 25%)

- ☐ Use Automated ML for tabular, computer-vision, and NLP tasks
- ☐ Evaluate AutoML runs with Responsible AI guidelines
- ☐ Use notebooks to access and wrangle data interactively
- ☐ Attach and use Synapse Spark or serverless Spark compute
- ☐ Retrieve features from Feature Store to train models
- ☐ Track experiments and metrics with MLflow

- ☐ Evaluate trained models following Responsible AI principles
- ☐ Perform hyperparameter tuning (sampling, search space, primary metric, early termination)

Train and Deploy Models (25 – 30%)






- ☐ Run training scripts and jobs, define parameters, and troubleshoot logs
- ☐ Configure compute environments for job runs
- ☐ Implement custom components and pipelines
- ☐ Pass data between pipeline steps and schedule runs
- ☐ Monitor and troubleshoot pipeline execution
- ☐ Register and manage MLflow models
- ☐ Define model signatures and package feature retrieval specs
- ☐ Deploy models to online and batch endpoints
- ☐ Test, invoke, and monitor deployed services

Optimize Language Models for AI Applications (25 – 30%)

- ☐ Select, deploy, and benchmark language models from the catalog
- ☐ Test deployed LLMs in the playground and compare results
- ☐ Apply optimization via prompt engineering and Prompt Flow
- ☐ Design prompt templates, track variants, and define chaining logic
- ☐ Use tracing for evaluation of prompt flows

- ☐ Implement Retrieval-Augmented Generation (RAG): cleaning, chunking, embedding
- ☐ Configure vector stores and Azure AI Search-based indexes
- ☐ Evaluate RAG solutions for accuracy and latency
- ☐ Perform fine-tuning: prepare data, choose base model, run jobs, evaluate results

Exam Readiness Scorecard

Domain	Confidence (1-5)	Notes / Action Items
 Design & Prepare ML Solutions (20–25%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Explore Data & Run Experiments (20–25%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Train & Deploy Models (25–30%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Optimize Language Models for AI (25–30%)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	
 Time Management (150-min exam pacing)	<input type="checkbox"/> 1 <input type="checkbox"/> 2 <input type="checkbox"/> 3 <input type="checkbox"/> 4 <input type="checkbox"/> 5	Complete 50-question timed set practice

How to Use

1. Check off each line once you can *perform the task hands-on* in Azure ML Studio or SDK.
2. Rate each domain (1 = weak → 5 = expert) to focus your revision.
3. Prioritize high-weight domains (**Training & Deployment, LLM Optimization**) in your final review week.

4. Re-test using Microsoft's **Practice Assessment** and **sandbox** to simulate the exam interface.



Tips for Use

- Start by marking each checklist item you feel confident about; leave blanks for uncertain areas.
- For each blank or lower-confidence domain, gather relevant Microsoft Learn modules, hands-on labs, or practice problems to reinforce.
- Use timed mock exams to simulate pacing and identify weak spots in your scorecard.
- Revisit the checklist closer to your exam and ensure all critical areas are covered.

● You're ready when all sections score **4 or 5 stars**, and your practice test scores are consistently **above 85%**.

🌟 **Congratulations!! You are on the right path to certification.**


All of our practice exams include **300+ questions**. This one has **over 500**.

Get the full version here <https://cloudcertificationstore.com/b/hvask>

Our writers who have taken the exam recently—and the reviewers who purchased these materials—agree that **over 90 %** of the questions matched what they saw on the live test.

Invest in your future: browse the full catalogue of [Cloud practice exams at our store](https://cloudcertificationstore.com)

Featured Collection

 <p>On Sale</p>	 <p>On Sale</p>	 <p>On Sale</p>	 <p>On Sale</p>	 <p>On Sale</p>
2025 Google Cloud Professional Cloud Architect Exam Questions (GOOG-PCA-0010) \$8.89- \$6.22	2025 Google Cloud Professional Cloud Developer Certification Practice Exam Questions (GOOG-PCD-0010) \$8.89- \$6.22	2025 Google Cloud Professional Machine Learning Engineer Practice Exam Questions (GOOG-PMLE-0010) \$8.99- \$6.29	2025 Google Cloud Professional Data Engineer Certification Practice Exam Questions (GOOG-PDE-0010) \$8.89- \$6.22	2025 Google Cloud Associate Cloud Engineer Exam Questions (GOOG-ACE-0010) \$8.89- \$6.22
 <p>On Sale</p>	 <p>On Sale</p>	 <p>On Sale</p>	 <p>On Sale</p>	 <p>FREE</p>
2025 Google Cloud Certified Professional Cloud Database Engineer Practice Exam Questions (GOOG-PCDE-0010) \$8.89- \$6.22	2025 AWS Certified Cloud Practitioner Practice Exam Questions (AWS-CLF-002-0010) \$8.99- \$6.29	2025 Microsoft Azure Fundamentals AZ-900 Practice Exam Questions (AZ-900-0010) \$8.89- \$6.22	2025 Microsoft Azure Developer Associate AZ-204 Practice Exam Questions (AZ-204-0010) \$8.89- \$6.22	FREE Practice Questions for the Google Cloud Generative AI Leader Exam \$0.00+