



Common Mistakes with GitLAB CI (How to Avoid Them)

BY DEVOPS SHACK

[Click here for DevSecOps & Cloud DevOps Course](#)

DevOps Shack

Common Mistakes with GitLab CI **(How to Avoid Them)**

Table of Content

1. Triggering Pipelines on Every Branch Push

- Overloads runners and deploys unverified code unintentionally.

2. Misusing only and except Instead of rules

- Causes confusing and unreliable job execution conditions.

3. Hardcoding Secrets in .gitlab-ci.yml

- Risks credential exposure and violates security best practices.

4. Not Caching Dependencies Properly

- Slows down builds by reinstalling packages in every job.

5. Using latest Tags for Docker Images

- Leads to non-reproducible builds and hard-to-debug issues.

6. No Separate Stages for Build, Test, and Deploy

- Makes pipelines hard to debug and breaks modularity.

7. Deploying to Production Without Approval Gates

-
- Allows unreviewed changes to go live, increasing risk.

8. Missing before_script and after_script Standardization

- Causes duplicated code and inconsistent environment setup.

9. Ignoring Job and Pipeline Failures

- Creates false positives and hides broken stages.

10. No Monitoring or Notifications on CI/CD Events

- Delays incident response and blinds teams to pipeline health.

☒ 1. Triggering Pipelines on Every Branch Push

The Problem

By default, GitLab CI/CD will trigger pipelines on every push to any branch unless explicitly configured otherwise. In many projects, this leads to **CI jobs being executed on every commit**, including:

- Minor documentation changes
- Feature branches under active development
- Work-in-progress (WIP) commits
- Quick fixes that don't need full pipeline runs

This overloads GitLab runners, slows down the pipeline queue, burns unnecessary compute minutes (especially on shared runners), and — in the worst cases — **triggers deployments from untested or incomplete code**.

Real-World Impact

- **Noise:** Dozens of pipelines running with no intention to deploy
- **Cost:** Runner time consumed needlessly
- **Risk:** A branch push triggered a deployment job and auto-deployed an incomplete feature
- **Delays:** Important pipelines queued behind unimportant ones

The Right Way to Handle This

☒ Use rules: to Define Exactly When a Job Should Run

GitLab CI offers a rules: block to control **when** and **under what conditions** jobs run.

☒ Example 1: Only Run Job on Merge Requests to main

deploy:

stage: deploy

script:

- ./deploy.sh

rules:

- if: '\$CI_COMMIT_BRANCH == "main"'

when: manual

- when: never

☒ Example 2: Run Tests Only on Feature Branches or MRs

test:

stage: test

script:

- npm test

rules:

- if: '\$CI_PIPELINE_SOURCE == "merge_request_event"'

- if: '\$CI_COMMIT_BRANCH =~ /^feature\/.+/'

Key Concepts to Know

- CI_COMMIT_BRANCH → the branch name triggering the pipeline
- CI_PIPELINE_SOURCE → identifies what caused the pipeline (push, merge_request_event, schedule, etc.)
- rules: → supersedes only/except and provides more flexibility
- when: manual → adds a manual approval step for sensitive jobs (like deploy)

Tips to Optimize Triggering Behavior

- Use rules: to **differentiate pipelines** between:
 - Merge requests (for tests & reviews)
 - Pushes to main/release (for deploy)

- Schedules (for backups or cleanup jobs)

- Introduce **manual deployments** that require approval:

deploy_production:

stage: deploy

script: ./deploy-prod.sh

environment: production

rules:

- if: '\$CI_COMMIT_BRANCH == "main"'

when: manual

- Add **commit message conditions** to skip pipelines if needed:

rules:

- if: '\$CI_COMMIT_MESSAGE =~ /^[skip ci\]/'

when: never

☒ What We Learned

1. **Not all pushes need a pipeline** — filter them precisely using rules:
2. Uncontrolled triggers waste resources and may lead to **accidental deployments**
3. Use **branch names**, **pipeline sources**, and **commit messages** to trigger pipelines only when appropriate
4. Manual approval gates and conditional logic improve both **speed** and **safety**

☒ 2. Misusing only and except Instead of rules

The Problem

GitLab CI/CD historically used `only` and `except` to control when jobs run. Many teams still rely on these, but they are now considered **less flexible and harder to maintain** compared to the modern rules: syntax.

Here's an example of older usage:

deploy:

script: ./deploy.sh

only:

- main

At first glance, this seems fine. But `only/except`:

- **Don't allow conditional logic** based on pipeline source
- Can't evaluate **commit messages, file changes, variables, or merge request context**
- Lead to **unexpected job runs** in complex pipelines (especially when multiple triggers exist: push, tag, MR, etc.)

Real-World Impact

- Jobs triggered during `merge_request_event` even when they shouldn't
- CI pipelines ran on tags even when deployment wasn't required
- Developers confused why a job triggered (no context visibility in `only`)
- Multiple jobs skipped unintentionally because `except` masked other branches

The Right Way to Handle This

☒ Use rules: Instead of `only/except`

rules: is **declarative, powerful, and context-aware**. You can use it to:

- Run jobs only on specific branches

- Include/exclude pipelines based on trigger source
- Check for commit message content
- Use custom variables to trigger conditional logic

Migration Example: From only/except to rules:

Before (old style):

deploy:

script: ./deploy.sh

only:

- main

except:

- tags

After (modern style):

deploy:

script: ./deploy.sh

rules:

- if: '\$CI_COMMIT_BRANCH == "main" && \$CI_PIPELINE_SOURCE == "push"'

when: always

- when: never

Useful Conditions for rules:

- `CI_COMMIT_BRANCH == "main"` → for main branch only
- `CI_COMMIT_TAG =~ /^v\d+\.\d+\.\d+$/` → for versioned tags
- `CI_PIPELINE_SOURCE == "merge_request_event"` → for MRs
- `CI_COMMIT_MESSAGE =~ /\[skip-deploy\]/` → to skip on demand
- `$MY_ENVIRONMENT == "prod"` → using custom variables

Safer and Smarter Pipelines

rules: allows for **granular, layered control**, such as:

deploy:

script: ./deploy.sh

rules:

- if: '\$CI_COMMIT_BRANCH == "main"'
when: manual
- if: '\$CI_COMMIT_BRANCH == "release/*"'
when: always
- when: never

This logic handles:

- **Auto deploy on release branches**
- **Manual approval on main**
- **Prevents job from running elsewhere**

What We Learned

1. only/except are deprecated for complex pipelines
2. rules: provide **greater clarity, context awareness**, and **reliability**
3. You can use conditions on **branches, triggers, variables, messages**, and more
4. Migrating to rules: simplifies troubleshooting and improves pipeline maintainability
5. Every job should be **intentionally triggered**, and rules: make that easy

3. Hardcoding Secrets in .gitlab-ci.yml

The Problem

A common and dangerous mistake in GitLab CI/CD is **placing sensitive credentials, tokens, passwords, or keys directly inside the .gitlab-ci.yml file** or hardcoded scripts. For example:

script:

```
- curl -u "admin:SuperSecret123" https://api.example.com/deploy
```

Since .gitlab-ci.yml is typically version-controlled and visible to all project members, this exposes secrets to:

- Internal developers who shouldn't have access
- Accidental commits and Git history
- CI logs and job artifacts
- Potential exfiltration if the repo is public or cloned

This violates **security best practices**, breaks **compliance rules**, and opens the door to **severe breaches**, such as unauthorized access to databases, APIs, cloud resources, or production systems.

Real-World Impact

- Secrets were accidentally committed and **pushed to Git history**, visible to contributors
- An API token printed during a curl command ended up in CI logs and leaked via job artifacts
- A staging token reused in production was discovered by attackers scanning public GitLab instances
- Teams scrambled to **rotate tokens**, revoke credentials, and audit logs after exposure

The Right Way to Handle This

☒ 1. Use GitLab CI/CD Variables for All Secrets

GitLab provides a secure way to manage secrets via **CI/CD Variables**:

- Go to **Settings** → **CI/CD** → **Variables**
- Add secrets as **Masked** and **Protected** (so they're hidden in logs and restricted to protected branches/tags)

Then use them safely in .gitlab-ci.yml:

script:

```
- curl -u "admin:$API_TOKEN" https://api.example.com/deploy
```

variables:

```
API_TOKEN: ${CI_JOB_TOKEN}
```

☒ 2. Never Echo or Print Secrets in Logs

Even if using variables, avoid printing them:

✗ Wrong:

script:

```
- echo $DB_PASSWORD
```

☒ Right:

script:

```
- psql -U user -d db -W <<< "$DB_PASSWORD"
```

Use set +x or disable command echoing in shell scripts to avoid accidental leak.

☒ 3. Use External Secret Managers When Needed

For sensitive or dynamic secrets, integrate with tools like:

- HashiCorp Vault
- AWS Secrets Manager
- Azure Key Vault

You can fetch secrets during the pipeline run:

script:

```
- DB_PASSWORD=$(vault kv get -field=password secret/db)
```

☑ 4. Rotate and Revoke Leaked Credentials Immediately

If a secret was ever committed:

1. **Revoke it**
2. **Rotate it**
3. Use tools like git-filter-repo or BFG to purge it from Git history:

`bfg --delete-files secrets.env`

☑ 5. Scan Your Git Repo for Secrets

Use tools in your pipeline:

- [Gitleaks](#)
- [TruffleHog](#)
- GitLab's own Secret Detection

Example with Gitleaks in .gitlab-ci.yml:

`gitleaks:`

`image: zricethezav/gitleaks`

`script:`

`- gitleaks detect --source=. --exit-code 1`

☑ What We Learned

1. Secrets should **never be stored** in .gitlab-ci.yml or committed to Git
2. Use **GitLab CI/CD variables**, marked as **masked and protected**
3. Prevent secrets from leaking into **logs, artifacts, or echo commands**
4. Use external secret managers for better control and rotation
5. Always **scan your repositories** and pipelines for exposed secrets

☑ 4. Not Caching Dependencies Properly

⚠ The Problem

In many GitLab CI/CD pipelines, developers forget to use **caching** for package dependencies like Node modules, Python wheels, Maven artifacts, Docker layers, etc.

Without caching:

- Dependencies are downloaded and re-installed from scratch on **every pipeline run**
- CI jobs take longer to execute
- Bandwidth is wasted
- Build speed varies inconsistently

This becomes a major bottleneck in medium-to-large projects where **the same dependencies are reused across jobs**.

Real-World Impact

- Pipelines slowed down from **2–3 minutes to 10–15 minutes**
- Redundant downloads increased cost and build time
- Developers lost productivity due to longer feedback loops
- Teams avoided running CI on small changes because it felt “too slow”

The Right Way to Handle This

☒ 1. Use the cache: Keyword Correctly

GitLab’s cache is used to **persist files** (like downloaded dependencies) **between jobs and pipelines**.

Example: Node.js Project

cache:

key: `${CI_COMMIT_REF_SLUG}`

paths:

- `node_modules/`

install_dependencies:

script:

- npm ci

☒ key: defines when to invalidate/reuse the cache. You can use:

- default → shared cache
- \${CI_COMMIT_REF_SLUG} → per-branch cache
- package-lock.json checksum → changes only when dependencies do

☒ 2. Differentiate Between cache: and artifacts:

- cache: → speeds up **future jobs/pipelines** (usually dependencies)
- artifacts: → used to **pass build outputs** (like compiled files) between jobs **in the same pipeline**

Example with Artifacts:

build:

stage: build

script: npm run build

artifacts:

paths:

- dist/

Example with Cache:

cache:

paths:

- node_modules/

☒ 3. Use Language-Specific Cache Directories

Language	Cache Directory
Node.js	node_modules/

Language	Cache Directory
Python	.venv/, pip cache
Java	~/.m2/repository/
Go	GOPATH/pkg/mod/
Rust	~/.cargo/registry/, target/

Python Example:

cache:

paths:

- .venv

☒ 4. Use key: to Scope the Cache Properly

Per-branch cache:

cache:

key: \${CI_COMMIT_REF_SLUG}

Global (shared) cache:

cache:

key: default

Smart cache invalidation:

cache:

key:

files:

- package-lock.json

☒ 5. Avoid Caching Mistakes

⊘ Don't:

- Cache dist/ unless it's reused

- Cache .git directory
- Cache secrets or environment files

☒ Do:

- Limit cache size and expiry if using self-hosted runners
- Regularly review cache hit/miss stats in GitLab pipeline logs

☒ What We Learned

1. Not caching dependencies increases **build time and cost**
2. Use GitLab's cache: keyword to persist packages across jobs
3. Scope your cache with key: to avoid unnecessary rebuilds
4. Know when to use cache: vs artifacts: — they serve different purposes
5. Caching dependencies can **cut pipeline time in half**, improving DevOps velocity

☒ 5. Using latest Tags for Docker Images

The Problem

Using the latest tag in Docker builds and deployments is a common habit—but in CI/CD pipelines, it becomes a **major liability**.

In GitLab CI, developers often write:

```
docker build -t myapp:latest .
```

```
docker push myapp:latest
```

Then reference it in Kubernetes manifests or Docker Compose files:

```
image: registry.example.com/myapp:latest
```

This seems convenient—but the latest tag is **mutable**. Every time it's rebuilt, it **overwrites the previous version**. There's no guarantee of what version is actually running unless it's explicitly tracked.

Real-World Impact

- Deployments pulled the wrong image due to Docker layer cache or pull policy
- CI/CD pipelines broke because latest was updated by another branch
- Rollbacks became impossible — previous versions weren't tagged or stored
- Bugs reappeared due to redeploying an unexpected image under the same tag

The Right Way to Handle This

☒ 1. Use Git SHA or Version Tags Instead of latest

Generate versioned image tags automatically:

variables:

```
GIT_SHA: $CI_COMMIT_SHORT_SHA
```

build:

```
stage: build
```

```
script:
```

- docker build -t registry.example.com/myapp:\$GIT_SHA .
- docker push registry.example.com/myapp:\$GIT_SHA

Update your deployment manifests accordingly:

image: registry.example.com/myapp:\$CI_COMMIT_SHORT_SHA

☒ 2. Use tags in Git to Mark Releases

For production-grade releases, tag your Git commits semantically (v1.3.5) and use that as your Docker image tag:

```
git tag v1.3.5
```

```
git push origin v1.3.5
```

In your GitLab CI:

```
build:
```

```
script:
```

```
- docker build -t myapp:$CI_COMMIT_TAG .
```

☒ 3. Add Image Metadata for Traceability

Add Git metadata during your Docker build:

```
docker build \
```

```
--label "org.opencontainers.image.revision=$CI_COMMIT_SHA" \
```

```
--label "org.opencontainers.image.source=$CI_PROJECT_URL" \
```

```
-t myapp:$CI_COMMIT_SHORT_SHA .
```

☒ 4. Disable Auto Pull of latest in Production

If you still reference latest, be sure to **pin the image digest**:

```
image: myapp@sha256:abc123def456...
```

Or explicitly set imagePullPolicy: IfNotPresent in Kubernetes:

```
imagePullPolicy: IfNotPresent
```

☒ 5. Keep Older Images for Rollback

Avoid deleting old images prematurely. Set retention policies that store at least 5–10 tagged versions, so you can easily roll back:

```
helm rollback myapp 3
```

Or:

```
kubectrl set image deployment/myapp myapp=myapp:<previous-tag>
```

☒ What We Learned

1. The latest tag is **convenient but dangerous** in automated pipelines
2. Always tag images using Git SHA, tags, or semantic versions
3. Use GitLab CI variables like `$CI_COMMIT_SHA`, `$CI_COMMIT_TAG` to generate consistent tags
4. Avoid tag collisions that can **hide bugs, block rollbacks, or break reproducibility**
5. Traceability and rollback capability start with **proper image versioning**

☒ 6. No Separate Stages for Build, Test, and Deploy

The Problem

In many GitLab CI/CD pipelines, especially small or rushed projects, developers combine all tasks (build, test, deploy) into a **single job** or single script: block like this:

```
job:
```

```
script:
```

-
- npm install
 - npm test
 - npm run build
 - ./deploy.sh

This might work initially, but it's an anti-pattern that leads to:

- Poor visibility into which phase failed
- No separation of concerns
- Impossible rollback from failed deploys
- Test failures **after** the build already passed
- Deploys that trigger even if tests break
- Harder reuse of build artifacts or test results

Real-World Impact

- A test failure broke production deployment because the deployment wasn't gated
- Developers wasted time debugging why a job failed — was it the install, test, or deploy?
- Build artifacts weren't reusable across environments
- No way to run only tests or only deployments when needed

The Right Way to Handle This

☒ 1. Use GitLab CI Stages: build, test, deploy

Break your pipeline into **clear, sequential stages**:

stages:

- build
- test
- deploy

☒ 2. Define Jobs per Stage

Each job should handle only **one responsibility**:

build:

stage: build

script:

- npm install
- npm run build

artifacts:

paths:

- dist/

test:

stage: test

script:

- npm test

deploy:

stage: deploy

script:

- ./deploy.sh

dependencies:

- build

☒ 3. Use Artifacts Between Stages

Artifacts let you **pass built files** from one stage to another:

artifacts:

paths:

- dist/

expire_in: 1 hour

This ensures your deploy stage doesn't rebuild code unnecessarily.

☒ 4. Gate Deployments Using Rules or Manual Triggers

Prevent auto-deploys on test failure:

deploy:

stage: deploy

script:

- ./deploy.sh

when: manual

rules:

- if: '\$CI_COMMIT_BRANCH == "main"'

This ensures deploys only happen **after successful build + test**, and optionally with human approval.

☒ 5. Use Stage-Level Timing and Visibility

In the GitLab UI, each stage has its own timeline, logs, and status:

- You can immediately see **what broke and where**
- Retrigger only the **failed stage**
- Cancel pipeline after failed tests (to save time)

☒ What We Learned

1. Mixing build, test, and deploy in one job **hurts clarity, control, and safety**
2. GitLab CI supports clean separation using stages: and artifacts:

3. Each job should be **single-purpose** and **modular**
4. Use rules and manual approvals to **gate deploys**
5. A well-structured pipeline is easier to debug, audit, and extend

7. Deploying to Production Without Approval Gates

The Problem

In many GitLab CI/CD setups, the deploy job runs automatically once the pipeline reaches the deploy stage — **even for critical environments like production**:

deploy_prod:

stage: deploy

script: ./deploy-prod.sh

This means:

- Any push to main or release can **immediately deploy to production**

- Bugs, failed tests, or misconfigurations may **slip through and go live**
- There's no **manual intervention** or approval process
- Teams accidentally **break production with a single commit**

This violates **Change Management** and **Deployment Control** practices — especially in regulated or high-availability environments.

Real-World Impact

- A junior developer pushed code and unintentionally triggered a full production deploy
- An incomplete feature merged into main and deployed during off-hours
- Incident recovery took hours because there was no rollback gate
- Organizations failed audits due to lack of deployment approval workflows

The Right Way to Handle This

☒ 1. Use when: manual for Production Deployments

Gate production jobs with manual approvals:

deploy_prod:

stage: deploy

script: ./deploy-prod.sh

environment:

name: production

url: https://myapp.com

when: manual

only:

- main

This adds a **“Play” button** in the GitLab pipeline UI, so deployment won’t proceed without human confirmation.

☒ 2. Restrict Deployment Jobs to Protected Branches or Tags

Ensure only approved branches can trigger production deployment:

deploy_prod:

rules:

- if: '\$CI_COMMIT_BRANCH == "main"'

when: manual

only:

- tags

Then protect the branch/tag via **GitLab → Settings → Repository → Protected Branches**

Only specific users or maintainers can push or merge.

☒ 3. Use environments: for Better Visibility & Control

Define environments to track deployments in GitLab:

deploy_prod:

environment:

name: production

url: https://app.example.com

Benefits:

- Track deployments by environment
- Rollback from GitLab UI
- Audit logs and change history per environment

☒ 4. Integrate with Slack or Teams for Approval Notifications

Use GitLab's webhook or integrations to notify teams when a production deployment is **waiting for approval**.

☑ 5. Use GitLab's Manual Approval + Merge Request Checks

GitLab Premium supports **MR approval rules** where specific people must approve changes before merge (and deploy).

☑ What We Learned

1. Auto-deploying to production without human approval is **risky and unsafe**
2. Use when: manual + rules: to require human review
3. Limit production deploys to **protected branches**
4. GitLab environments give visibility and rollback capabilities
5. Production pipelines should have **change gates, approvals, and alerting**

☑ 8. Missing `before_script` and `after_script` Standardization

⚠ The Problem

Many GitLab CI/CD pipelines copy the same setup or teardown commands repeatedly in every job. This leads to:

- **Redundant code**
- **Inconsistent behavior** across jobs
- **Hard-to-maintain pipelines**, especially when setup steps change (e.g., auth, env vars, tools)

For example, jobs without `before_script` repeat:

`job1:`

script:

- export ENV=prod
- npm install
- npm run build

job2:

script:

- export ENV=prod
- npm install
- npm test

Or worse, they don't run a proper cleanup — temporary files, dangling containers, and logs are left behind.

Real-World Impact

- Inconsistent environments across jobs (some install tools, others don't)
- Security tools (e.g., Trivy, Gitleaks) not run consistently
- Teams forget to apply updates to all jobs, causing version drift
- Developers waste time debugging jobs with missing dependencies or mismatched config

The Right Way to Handle This

☒ 1. Use Global `before_script:` and `after_script:`

Define shared scripts at the **top level** of `.gitlab-ci.yml`:

`before_script:`

- echo "Setting up CI environment"
- export ENV=prod
- apt-get update && apt-get install -y curl

after_script:

- echo "Cleaning up"
- rm -rf temp/
- docker system prune -f

Now all jobs **inherit** these steps unless overridden.

☒ 2. Override Locally When Needed

Jobs can override global before_script or after_script if they have special requirements:

test:

before_script:

- echo "Custom setup for tests"
- npm ci

script:

- npm test

This **adds flexibility** while keeping core steps centralized.

☒ 3. Use .default-template Jobs for Common Behaviors

For even more structure, create reusable templates using YAML anchors or extends:

.default-job:

before_script:

- export ENV=staging
- npm ci

after_script:

- echo "Job completed."

build:

`extends: .default-job`

`script:`

`- npm run build`

This promotes DRY (Don't Repeat Yourself) pipeline code.

☑ 4. Use `after_script` for Clean-Up and Diagnostics

Typical `after_script`: use cases:

- Cleanup temp files, containers, or cache
- Post-job reporting (e.g., log upload, Slack notification)
- Metrics collection (e.g., upload JUnit or test logs)

`after_script:`

`- curl -X POST -d "Build complete: $CI_JOB_NAME" $SLACK_WEBHOOK`

🔗 Note: `after_script` runs **even if the job fails** — great for cleanup and alerts.

☑ 5. Enforce Common Steps Across Jobs

Examples:

- Ensure security tools like Trivy/Gitleaks are always present
- Standardize docker login steps
- Prepare directories like `/tmp/build`, etc.
- Set env vars like `NODE_ENV`, `JAVA_HOME`, `PYTHONPATH`

☑ What We Learned

1. Without standardized setup/teardown steps, pipelines become **messy and unreliable**
2. `before_script` ensures all jobs start with **clean, consistent environments**
3. `after_script` handles **post-processing**, cleanup, and notification
4. Use `.default-job` or `extends:` to eliminate repetition

-
5. Clean and predictable pipelines reduce maintenance overhead and debugging time

☒ 9. Ignoring Job and Pipeline Failures

The Problem

Many teams fall into the trap of treating CI/CD pipelines as "passive observers." Jobs fail, but no one investigates. This happens due to:

- Job failures being silently ignored in script blocks
- Overuse of `allow_failure: true`
- Lack of alerts or failure notifications
- No owner assigned to broken pipelines
- No policy for failing test cases (pipelines pass even when unit tests fail)

This leads to a dangerous DevOps culture of "**green-looking red pipelines**", where the pipeline looks okay on the surface, but is functionally broken inside.

Real-World Impact

- A pipeline failed during artifact upload, but deployment still happened
- Broken test stage was marked `allow_failure: true` and never fixed
- Staging environment had missing assets because the build job failed silently
- Teams became numb to red pipelines — no one felt responsible to fix them
- A production deploy was based on a failed build, causing an outage

The Right Way to Handle This

☒ 1. Fail Fast and Loud — Don't Suppress Errors

Every job should **fail clearly** if anything goes wrong.

Avoid this:

script:

```
- ./build.sh || true
```

Instead, do this:

yaml

script:

```
- ./build.sh
```

Let the job **fail naturally** and halt the pipeline.

☒ 2. Use `allow_failure: true` Only When Absolutely Needed

Legitimate use cases:

- Experimental or optional jobs
- Security scanners with non-blocking warnings
- Nightly jobs where failure doesn't impact the main workflow

Example:

sast_scan:

script: trivy fs .

allow_failure: true

But NEVER do this on core jobs:

build:

allow_failure: true # ❌ Don't do this!

☒ 3. Assign Job Owners or Set Code Owners

Use GitLab's CODEOWNERS file or team conventions to assign responsibility:

CODEOWNERS

.gitlab-ci.yml @devops-team

/test/ @qa-team

If a job breaks, **someone gets notified** and is responsible to fix it.

☒ 4. Fail the Pipeline When Tests Fail

Avoid constructs that let tests fail silently:

script:

- npm test || true # ❌

Use:

script:

- npm test # ☒ fails on test error

If using test runners, output JUnit reports and make sure they're evaluated.

☒ 5. Enable Notifications for Failures

Set up:

- GitLab email notifications
- Slack/Discord webhooks
- PagerDuty/Alertmanager integration for critical failures

Example:

notify_failure:

stage: notify

script:

```
- curl -X POST -H "Content-Type: application/json" \  
-d '{"text": "Pipeline failed on $CI_COMMIT_REF_NAME"}' \  
$SLACK_WEBHOOK
```

when: on_failure

☒ 6. Monitor Pipeline Health

Use GitLab's **Pipeline Analytics** to track:

- Success/failure trends
- Longest-running or most unstable jobs
- Broken pipelines over time

This helps detect **chronic failures** that are silently ignored.

☒ What We Learned

1. Every job must **fail clearly and stop the pipeline** when something goes wrong
2. Misusing `allow_failure` or suppressing errors hides critical problems
3. Assign job ownership so failures are addressed quickly
4. Never accept test failures — CI must enforce quality
5. Build a culture of **pipeline hygiene** to maintain long-term reliability

10. No Monitoring or Notifications on CI/CD Events

The Problem

Many teams set up complex pipelines but forget one critical piece — **monitoring and alerting**. When builds fail, deployments break, or test coverage drops, there's **no one watching**.

Common issues include:

- No Slack, email, or webhook notifications on pipeline status
- Teams not knowing when a critical stage fails
- No observability into CI/CD trends (e.g., pipeline success rate, test flakiness)
- Delayed reaction to incidents due to lack of visibility
- Developers unaware of broken builds for hours or days

Without feedback loops, CI/CD becomes a **black box** — issues compound silently until they explode.

Real-World Impact

- A staging deploy failed over the weekend, but no one noticed until Monday
- Test coverage dropped by 15%, but no alerts were configured
- A job failed on 7 consecutive MRs, but the team assumed everything was fine
- A production deployment failed silently due to a script typo — no one was notified
- Missed deadlines and SLAs due to untracked pipeline regressions

The Right Way to Handle This

1. Set Up Notifications for Pipeline Events

Use GitLab's built-in notification features:

- **Email Notifications:** Users can configure notifications under their GitLab profile → Preferences → Notifications
- **Slack Integration:**
 - Go to **Settings → Integrations**
 - Enable Slack notifications for push, merge, pipeline events
 - Or use custom webhooks in `.gitlab-ci.yml`:

`notify_slack:`

`stage: notify`

`script:`

```

- curl -X POST -H 'Content-type: application/json' \
  --data '{"text": "🔔 Pipeline *$CI_PIPELINE_ID* failed on
*$CI_COMMIT_BRANCH*"}' \
  $SLACK_WEBHOOK

```

`when: on_failure`

☒ 2. Monitor Pipeline Health with GitLab Analytics

Use GitLab features (Premium/Ultime tiers) or integrate with external tools like:

- Grafana dashboards via GitLab Exporter
- Prometheus metrics for runner usage and job durations
- GitLab CI/CD Reports for pipeline duration, test trends, code quality

Track:

- Pipeline success rate
- Median build time
- Longest-running jobs
- Jobs that fail most frequently

☒ 3. Use when: + Conditional Alerts on Failures

You can run notification or rollback jobs **only on failure**:

notify_failure:

stage: notify

script:

- ./notify.sh "Build failed for \$CI_COMMIT_REF_NAME"

when: on_failure

And on success:

notify_success:

stage: notify

script:

- ./notify.sh "☒ All jobs passed!"

when: on_success

☒ 4. Create Fallback or Auto-Rollback Strategies

Set up conditional jobs that react to failures:

auto_rollback:

stage: deploy

script:

- ./rollback-to-previous.sh

when: on_failure

This allows fast recovery when something breaks.

☒ 5. Make Alerts Actionable and Contextual

Good alerts should answer:

- What failed?
- When did it happen?
- Who committed the change?
- What was the impact?
- What should be done next?

Include metadata in Slack/webhook messages:

```
{  
  "pipeline_id": "12345",  
  "branch": "main",  
  "status": "failed",  
  "commit": "a1b2c3d",  
  "author": "dev@example.com"  
}
```

☒ What We Learned

1. A pipeline without monitoring is a **silent failure waiting to happen**

2. Use GitLab's built-in notification system + Slack/Email integrations
3. Add fallback and notification stages to handle pipeline success/failure
4. Monitor CI/CD health using GitLab analytics or external observability tools
5. The best CI/CD systems **talk back** — they tell you when something needs attention

Conclusion: From Fragile Pipelines to Reliable Automation

GitLab CI/CD is a powerful tool that can accelerate software delivery when used wisely — but like any automation system, it's only as strong as the discipline and structure behind it. The mistakes we've covered — from misconfigured triggers and dependency caching to unsafe deploys and invisible failures — are not just technical oversights. They are process gaps that slowly erode trust in the pipeline and disrupt developer velocity.

The good news? Each of these pitfalls is fixable.

By:

- Breaking down pipelines into logical stages,
- Using semantic versioning over latest,
- Guarding production with manual approvals,
- And monitoring failures with proper notifications,

...you transform CI/CD from a fragile chain of scripts into a **reliable backbone of your DevOps culture**.

Remember, pipelines are not “set it and forget it.” They require **iteration, observation, and ownership**. Treat them as living systems — regularly audited, tested, and improved.

When teams take the time to get CI/CD right, they unlock more than just faster deployments — they build confidence, stability, and resilience into every release.