

KUBERNETES SERVICE

Problem-1: Pods are ephemeral in nature — their IP addresses change whenever a pod restarts or is recreated. Therefore, accessing a pod directly using its IP address is not reliable in the long run.

Problem-2: Multiple users may try to access the application simultaneously, but there is no such mechanism to distribute this traffic across pods by default.

Problem-3: Pods running in K8s cluster are accessible only to those who have access to Kubernetes cluster.


Kubernetes service can resolve this issues:

- Kubernetes Services use **labels** to track pods, enabling them to route traffic even if pod IPs change — a process known as **Service Discovery**.
- Kubernetes Services also provide **load balancing**, distributing incoming traffic evenly across all healthy pods to ensure reliability and high availability.
- Kubernetes service exposes the application to the external world.

In the Below diagram, the user's request goes to the **frontend Service**, which routes it to a healthy frontend pod. The frontend pod forwards the request to the **backend Service**, which sends it to a healthy backend pod. The backend responds to the frontend, allowing the user to interact with the application.


There are 3 types of services:

- 1) Service of ClusterIP Type
- 2) Service of NodePort Type
- 3) Service of Loadbalancer Type.



Types of services:

1) ClusterIP Service:



- A ClusterIP Service is the default Kubernetes Service type that exposes applications internally within the cluster, enabling pod-to-pod communication via an internal IP.

backend.yaml:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: backend-deploy
spec:
  replicas: 3
  selector:
    matchLabels:
      app: backend
  template:
    metadata:
      labels:
        app: backend
    spec:
      containers:
        - name: backend-container
          image: hashicorp/http-echo
          args:
            - "-text=Hello From Backend"

---
apiVersion: v1
kind: Service
metadata:
  name: backend-svc
spec:
  type: ClusterIP
  ports:
    - protocol: TCP
      port: 9090
      targetPort: 5678
  selector:
```

app: backend

Lets apply this backend.yaml and explore the ClusterIP service:

```
kubectl apply -f backend.yaml
```

root@DESKTOP-C6P8EQS:~/kubernetes/4)K8s_Services\$ kubectl get all					
NAME	READY	STATUS	RESTARTS	AGE	
pod/backend-deploy-7ddc9f6f4c-6f5wc	1/1	Running	0	3m55s	
pod/backend-deploy-7ddc9f6f4c-bxlcv	1/1	Running	0	3m55s	
pod/backend-deploy-7ddc9f6f4c-c9qzw	1/1	Running	0	3m55s	

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/backend-svc	ClusterIP	10.96.13.58	<none>	9090/TCP	3m55s
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	22m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/backend-deploy	3/3	3	3	3m55s

NAME	DESIRDED	CURRENT	READY	AGE
replicaset.apps/backend-deploy-7ddc9f6f4c	3	3	3	3m55s

```
kubectl describe svc backend-svc
```

root@DESKTOP-C6P8EQS:~/kubernetes/4)K8s_Services\$ kubectl describe svc backend-svc							
Name:	backend-svc						
Namespace:	default						
Labels:	<none>						
Annotations:	<none>						
Selector:	app=backend						
Type:	ClusterIP						
IP Family Policy:	singlestack						
IP Families:	IPv4						
IP:	10.96.13.58						
IPs:	10.96.13.58						
Port:	<unset> 9090/TCP						
TargetPort:	5678/TCP						
Endpoints:	10.244.2.2:5678,10.244.1.2:5678,10.244.1.3:5678						
Session Affinity:	None						
Internal Traffic Policy:	Cluster						
Events:	<none>						

```
kubectl get pods -o wide
```

root@DESKTOP-C6P8EQS:~/kubernetes/4)K8s_Services\$ kubectl get pods -o wide									
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
backend-deploy-7ddc9f6f4c-6f5wc	1/1	Running	0	8m50s	10.244.2.2	rayeez-cluster-worker2	<none>	<none>	
backend-deploy-7ddc9f6f4c-bxlcv	1/1	Running	0	8m50s	10.244.1.2	rayeez-cluster-worker	<none>	<none>	
backend-deploy-7ddc9f6f4c-c9qzw	1/1	Running	0	8m50s	10.244.1.3	rayeez-cluster-worker	<none>	<none>	

Let's create a test pod and try accessing the backend pods from inside it after logging in.


```
kubectl run -it mynginx --image=nginx -- /bin/bash
```

Pods were accessible through service:


```
curl backend-svc:9090
```

```
root@mynginx:/# curl backend-svc:9090
Hello From Backend
```

2)NodePort Service: It is build on top of ClusterIP Service.



→ A **NodePort Service** exposes pods externally using a node's IP and a fixed port, unlike **ClusterIP**, which is accessible only within the cluster.



→ The user accesses the app via a node's IP and port. The **frontend Service (NodePort)** routes the request to the **frontend Deployment**,

which then forwards it to the **backend Service (ClusterIP)** and finally to the **backend Deployment**.

Note: When a **NodePort** Service is created, a corresponding **ClusterIP** Service is automatically created.

frontend.yaml:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend-deploy
spec:
  replicas: 3
  selector:
    matchLabels:
      app: frontend
  template:
    metadata:
      labels:
        app: frontend
    spec:
      containers:
        - name: frontend-container
          image: nginx
          ports:
            - containerPort: 80
---
apiVersion: v1
kind: Service
metadata:
  name: frontend-svc
spec:
  type: NodePort
  selector:
    app: frontend
  ports:
    - protocol: TCP
```

```

    port: 80          # ClusterIP service port
    targetPort: 80   # Container's port
inside the pod
    nodePort: 31000  # Exposed externally
(must be in 30000-32767)

```

Lets apply frontend.yaml file and explore NodePort service:

```
kubectl apply -f frontend.yaml
```

```

● root@DESKTOP-C6P8EQS:~/kubernetes/4)K8s_Services$ kubectl get all
NAME                               READY   STATUS    RESTARTS   AGE
pod/backend-deploy-7ddc9f6f4c-6f5wc 1/1     Running   0          96m
pod/backend-deploy-7ddc9f6f4c-bxlcv 1/1     Running   0          96m
pod/backend-deploy-7ddc9f6f4c-c9qzw 1/1     Running   0          96m
pod/frontend-deploy-578b67fb7b-599hb 1/1     Running   0          11m
pod/frontend-deploy-578b67fb7b-ds88n 1/1     Running   0          11m
pod/frontend-deploy-578b67fb7b-r1712 1/1     Running   0          11m
pod/mynginx                          1/1     Running   1 (77m ago) 84m

NAME           TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)         AGE
service/backend-svc  ClusterIP   10.96.13.58  <none>        9090/TCP       96m
service/frontend-svc  NodePort    10.96.253.233 <none>        80:31000/TCP   6m52s
service/kubernetes  ClusterIP   10.96.0.1     <none>        443/TCP        115m

NAME                           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/backend-deploy 3/3     3           3           96m
deployment.apps/frontend-deploy 3/3     3           3           14m

NAME           DESIRED  CURRENT  READY   AGE
replicaset.apps/backend-deploy-7ddc9f6f4c  3        3        3      96m
replicaset.apps/frontend-deploy-578b67fb7b  3        3        3      11m
replicaset.apps/frontend-deploy-67447bb9ff  0        0        0      14m

```

→ For frontend deployment, frontend service of type ClusterIP is listening on port 80 and service type NodePort is listening on 31000.

Lets Access the frontend app from browser:

```
http://<NodeIP>:<NodePort>
```


Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.


For online documentation and support please refer to nginx.org. Commercial support is available at nginx.com.

Thank you for using nginx.

3) Load Balancer Service: It is build on top of NodePort service and ClusterIP service.



Remembering the IP addresses of each node is difficult. So when a **LoadBalancer** service is implemented, it automatically provisions a load balancer on the cloud platform, allowing the application to be accessible from the internet (external world).



- When you create a **LoadBalancer service**, Kubernetes requests a **public IP** from the **cloud provider** (e.g., AWS, GCP, Azure).

- This **public IP** is linked to the **LoadBalancer**, which then **routes traffic** to a **NodePort service** within the cluster.
- The **NodePort service** internally uses a **ClusterIP service** to **distribute traffic** to the **pods**.

Service.yaml for **LoadBalancer** mode:

```
apiVersion: v1
kind: Service
metadata:
  name: frontend-svc
spec:
  type: LoadBalancer
  selector:
    app: frontend
  ports:
    - protocol: TCP
      port: 80          # ClusterIP service port
      targetPort: 80   # Container's port
      nodePort: 31000  # Exposed externally
      (must be in 30000-32767)
```


```
root@DESKTOP-C6P8EQS:~/kubernetes/4)K8s_Services$ kubectl get svc frontend-svc
NAME        TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
frontend-svc  LoadBalancer  10.96.253.233  <pending>      80:31000/TCP  40m
```

If it is executed on a cloud platform, It will create a Loadbalancer having public IP address with the help of cloud control manager.

```
kubectl describe svc frontend-svc
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/4)K8s_Services$ kubectl describe svc frontend-svc
Name:           frontend-svc
Namespace:      default
Labels:          <none>
Annotations:    <none>
Selector:       app=frontend
Type:           LoadBalancer
IP Family Policy: SingleStack
IP Families:   IPv4
IP:             10.96.253.233
IPs:            10.96.253.233
Port:           <unset>  80/TCP
TargetPort:     80/TCP
NodePort:       <unset>  31000/TCP
Endpoints:     10.244.1.6:80,10.244.1.5:80,10.244.2.6:80
Session Affinity: None
External Traffic Policy: Cluster
Internal Traffic Policy: Cluster
Events:         <none>
```

4)ExternalName Service:



- Offers a **simple alias** to connect **internal services** to **external resources** using **DNS names**.
- When using **ExternalName**, applications **connect to services using internal DNS names**, and **Kubernetes handles the external redirection**.

- This **separates configuration from application logic**, ensuring **easier maintenance and flexibility**.

Service.yaml:

```
apiVersion: v1
kind: Service
metadata:
  name: cka-db-svc
spec:
  type: ExternalName
  externalName: cka-db-khjukjij.us-east-
1.rds.amazonaws.com #DNS name of the Amazon RDS
Instance
```