


Debugging and Logging Kubernetes Applications

This project focuses on understanding how to **debug, troubleshoot, and analyze** applications running inside a Kubernetes cluster.



Since Kubernetes apps run in multiple Pods and containers, troubleshooting requires checking:

- Pod health
- Logs and events
- Container status
- Service connectivity
- DNS issues
- Network Policies

2. Project Objective

The primary goals of this project were to:

- Debug Kubernetes workloads using kubectl
- Check pod lifecycle, events, and health
- View and analyze container logs
- Perform in-cluster debugging using exec
- Troubleshoot service and connectivity issues
- Verify DNS functionality inside the cluster
- Review NetworkPolicies for traffic restrictions

3. Prerequisites

Before starting the project, the following prerequisites were required:

- A running Kubernetes cluster
 - kubectl installed and configured
 - A sample app deployed for testing
 - Basic knowledge of YAML, Pods, Services, and Deployments
-

4. Implementation Steps

Step 1: Create a Sample Deployment

A deployment was created to run a simple guestbook application.

➤ *kubectl create deployment guestbook --image=ibmcom/guestbook:v1*

```
controlplane:~$ kubectl create deployment guestbook --image=ibmcom/guestbook:v1
deployment.apps/guestbook created
```

This deploys the guestbook container inside a Kubernetes Pod, allowing further debugging actions.

Step 2: Inspecting and Debugging Pods

When a deployment fails or a pod enters an error state, the following inspection methods were used:

a) Retrieve Basic Pod Information

➤ *kubectl get pods*

This outputs pod name, status, restart count, and age.

b) Describe a Pod for Detailed Diagnostics

➤ *kubectl describe pod <pod-name>*

```
controlplane:~$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
guestbook-86cf798756-nm4p9  1/1     Running   0          18s
controlplane:~$ kubectl describe pod guestbook-86cf798756-nm4p9
Name:           guestbook-86cf798756-nm4p9
Namespace:      default
Priority:       0
Service Account: default
Node:           node01/172.30.2.2
Start Time:     Sat, 06 Dec 2025 06:11:51 +0000
Labels:         app=guestbook
                pod-template-hash=86cf798756
Annotations:   cni.projectcalico.org/containerID: fd8b0d20e7ac5c7f0bc3574f7ed72585d9841675ca0aa1ab3e4022e7e4b55cf0
                cni.projectcalico.org/podIP: 192.168.1.4/32
                cni.projectcalico.org/podIPs: 192.168.1.4/32
```

This command provides detailed information including:

- Container states
- Restarts and failure reasons
- Events (Scheduling, ImagePull errors, Probe failures)
- Resource requirements
- Node assignment

c) Check Cluster Events

➤ *kubectl get events -n <namespace>*

controlplane:~\$ kubectl get events				
LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
18d	Normal	Starting	node/controlplane	Starting kubelet.
18d	Normal	NodeHasSufficientMemory	node/controlplane	Node controlplane status is now: NodeHasSufficientMemory
18d	Normal	NodeHasNoDiskPressure	node/controlplane	Node controlplane status is now: NodeHasNoDiskPressure
18d	Normal	NodeHasSufficientPID	node/controlplane	Node controlplane status is now: NodeHasSufficientPID
18d	Normal	NodeAllocatableEnforced	node/controlplane	Updated Node Allocatable limit across pods
18d	Normal	Starting	node/controlplane	Starting kubelet.
18d	Normal	NodeAllocatableEnforced	node/controlplane	Updated Node Allocatable limit across pods
18d	Normal	NodeHasSufficientMemory	node/controlplane	Node controlplane status is now: NodeHasSufficientMemory
18d	Normal	NodeHasNoDiskPressure	node/controlplane	Node controlplane status is now: NodeHasNoDiskPressure
18d	Normal	NodeHasSufficientPID	node/controlplane	Node controlplane status is now: NodeHasSufficientPID
18d	Normal	RegisteredNode	node/controlplane	Node controlplane event: Registered Node controlplane in Cont

Events helped in identifying:

- Image pull issues
- Pod scheduling failures
- Readiness or liveness probe failures
- Networking and configuration errors

Step 3: Viewing Application Logs

Logs were used to analyze runtime behavior inside containers.

a) View Pod Logs

➤ *kubectl logs <pod-name>*

b) View Logs from a Specific Container

➤ *kubectl logs <pod-name> -c <container-name>*

c) View Logs from a Previous Container Instance

➤ *kubectl logs <pod-name> --previous*

This is useful when a container crashed and restarted.

```
controlplane:~$ kubectl logs guestbook-86cf798756-nm4p9
[negroni] listening on :3000
controlplane:~$ kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
guestbook-86cf798756-nm4p9  1/1     Running   0          3m3s
controlplane:~$ kubectl logs guestbook-86cf798756-nm4p9 -c guestbook
[negroni] listening on :3000
controlplane:~$ kubectl logs guestbook-86cf798756-nm4p9 --previous
Error from server (BadRequest): previous terminated container "guestbook" in pod "guestbook-86cf798756-nm4p9" n
controlplane:~$ kubectl logs guestbook-86cf798756-nm4p9 --tail=10
```

Step 4: Debugging Through Shell Access

When logs were insufficient to diagnose issues inside a container, shell access was used.

a) Create a Pod with Shell Support

➤ *kubectl create -f
https://kubernetes.io/examples/application/shell-
demo.yaml*

b) Open Shell Inside the Container

➤ *kubectl exec -it shell-demo -- /bin/bash*

From here, commands like ls, cat, ps, etc., were used for real-time inspection.

```
controlplane:~$ kubectl create -f https://kubernetes.io/examples/application/shell-demo.yaml
pod/shell-demo created
controlplane:~$ kubectl exec -it shell-demo -- /bin/bash
root@node01:/# ls
bin  dev           docker-entrypoint.sh  home  lib64  mnt  proc  run  srv  tmp  var
boot docker-entrypoint.d  etc             lib   media  opt  root  sbin  sys  usr
```

Step 5: Debugging Services

Service-related issues are common in Kubernetes. The following checks were performed:

a) Create a Service if Missing

- *kubectl expose deployment <deployment-name> --type=NodePort --port=3000*

b) List All Services

- *kubectl get svc*

c) Verify Service Configuration

- *kubectl describe svc <service-name>*

```
controlplane:~$ k get deployments.apps
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
guestbook  1/1     1           1           9m40s
controlplane:~$ kubectl expose deployment guestbook --type=NodePort --port=3000
service/guestbook exposed
controlplane:~$ kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
guestbook  NodePort  10.104.10.242  <none>        3000:30560/TCP  9s
kubernetes ClusterIP  10.96.0.1    <none>        443/TCP       18d
controlplane:~$ kubectl describe svc guestbook
Name:           guestbook
Namespace:      default
Labels:         app=guestbook
Annotations:    <none>
Selector:       app=guestbook
Type:          NodePort
IP Family Policy: SingleStack
IP Families:   IPv4
IP:            10.104.10.242
IPs:           10.104.10.242
Port:          <unset>  3000/TCP
TargetPort:    3000/TCP
NodePort:      <unset>  30560/TCP
Endpoints:    192.168.1.4:3000
Session Affinity: None
```

Important aspects checked:

- TargetPort matching container port
- Correct Service type
- Endpoint availability
- Protocol consistency

Common Issue:

If no endpoints are shown, the service cannot forward traffic to pods.

Step 6: Validating DNS Functionality

DNS issues often lead to service communication problems.

a) Test DNS Resolution

➤ *kubectl exec -ti <pod-name> -- nslookup kubernetes.default*

b) Inspect /etc/resolv.conf

➤ *kubectl exec -ti <pod-name> -- cat /etc/resolv.conf*

```
controlplane:~$ kubectl exec -ti guestbook-86cf798756-nm4p9 -- nslookup kubernetes.default
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: kubernetes.default
Address 1: 10.96.0.1 kubernetes.default.svc.cluster.local
controlplane:~$ kubectl exec -ti guestbook-86cf798756-nm4p9 -- cat /etc/resolv.conf
search default.svc.cluster.local svc.cluster.local cluster.local
nameserver 10.96.0.10
options ndots:5
controlplane:~$
```

This ensured that search domains and nameservers were correctly configured.

Step 7: Inspecting NetworkPolicies

NetworkPolicies control inbound and outbound traffic between pods.

Key checks included:

- Confirming if a policy blocks traffic
- Ensuring pods have required labels
- Validating ingress/egress rules
- Testing connectivity using curl / ping

By default, pods accept all traffic unless restricted by NetworkPolicies.

5. Results and Observations

This systematic approach helped in identifying common Kubernetes issues such as:

- Missing services
 - Misconfigured ports
 - Image pull errors
 - CrashLoopBackOff conditions
 - DNS lookup failures
 - Blocked traffic due to NetworkPolicies
-

6. Conclusion

This project provided a comprehensive understanding of Kubernetes debugging and logging techniques. By using various kubectl commands, it was possible to inspect pods, analyze logs, review events, troubleshoot services, and validate networking.



*Thank
You!*