

Kubernetes Outages

8 REAL-WORLD FAILURES AND HOW TO SURVIVE THEM

BY DEVOPS SHACK





Click here for DevSecOps & Cloud DevOps Course

DevOps Shack

Kubernetes Outages: 8 Real-World Failures and How to Survive Them

From minor misconfigurations to catastrophic cluster-wide failures — learn how to detect, diagnose, and recover like a production-grade Kubernetes Admin.

Table of Contents

Each section includes:

- Root Cause
- K How to Diagnose
- Resolution Steps

✓ 1. CrashLoopBackOff Storms

When one pod goes down, then another, then another...

- Root Cause: Misconfigured probes, fatal exceptions, missing configs
- K How to Diagnose: kubectl describe, kubectl logs, check events
- Resolution Steps: Fix env/configs, adjust probe timings, restart with rollback

2. Node-Level Failures or Evictions

Pods suddenly vanish or restart — your app feels flaky

- Root Cause: Node disk full, memory pressure, CPU starvation, kubelet down
- K How to Diagnose: kubectl get nodes, kubectl describe node, dmesg, kubelet logs





• Resolution Steps: Cordon/drain, free space, scale up, tune requests/limits

☑ 3. PersistentVolume Claim (PVC) Issues

Pods get stuck in "Pending" or apps lose data on restart

- Root Cause: StorageClass misconfig, volume zone mismatch, PV unavailable
- K How to Diagnose: kubectl get pvc, describe pvc, look for unbound claims
- Resolution Steps: Fix SC, verify EBS/GCE volumes, use volumeBindingMode: WaitForFirstConsumer

✓ 4. Ingress or LoadBalancer Downtime

The app works, but nobody can reach it

- Root Cause: Broken Ingress rules, SSL cert expired, DNS misrouting, LoadBalancer misprovisioned
- K How to Diagnose: kubectl describe ingress, cert-manager logs, dig/DNS tools
- Resolution Steps: Fix ingress path/host, re-issue certs, re-provision LoadBalancer, check annotations

✓ 5. DNS Resolution Failures Inside Cluster

Pods can't reach services via name — only by IP

- Root Cause: CoreDNS crash, ConfigMap issues, network plugin problem
- Resolution Steps: Restart CoreDNS, validate ConfigMap, upgrade CNI plugin, clear stuck pods



6. Etcd Data Corruption or Unavailability

API server becomes unstable or unresponsive

- Root Cause: Disk pressure, snapshot failure, network split, cluster scaling bug
- K How to Diagnose: Check etcd logs, API latency spikes, use etcdctl for cluster health
- Resolution Steps: Restore from snapshot, clean etcd data dir, scale etcd properly, isolate disk issues

✓ 7. API Server Saturation or Hang

Cluster seems frozen — no pods list, no deployments work

- Root Cause: API server overwhelmed, controller-manager misbehaving, cert expiry
- K How to Diagnose: Check API server metrics, latency, controller logs, Prometheus alerts
- Resolution Steps: Scale control plane, add resource limits, clean long-running watches

✓ 8. Broken Rollouts / Bad Deployments to Prod

Everything was working... until you deployed

- Root Cause: Misconfigured image, failing containers, missing env variables, bad YAML
- How to Diagnose: kubectl rollout status, pod logs, probe failures, image digests
- Resolution Steps: Use rollout undo, implement canary strategy, test configs in staging

✓ 1. CrashLoopBackOff Storms

"Everything was fine... and then pods just started crashing — over and over."



Root Cause

The CrashLoopBackOff status means that a container starts \rightarrow crashes \rightarrow Kubernetes restarts it \rightarrow crashes again — in a loop.

In large clusters, this can quickly become a **storm** if multiple pods, deployments, or namespaces are affected simultaneously.

↑ Common Causes:

1. Invalid or missing environment variables

- App depends on secrets/configs that aren't loaded
- Example: DB connection string is missing or wrong

2. Faulty or misconfigured probes

- Liveness or readiness probe kills the container before it's ready
- Too short initialDelaySeconds, overly strict timeouts

3. Fatal application bugs

- Application throws unhandled exception on boot
- Segfault, null pointer, syntax error in Python, etc.

4. Bad image or command override

- Wrong startup script in Dockerfile or pod spec
- App exits immediately after starting

5. Permission errors

- App tries to write to read-only filesystem
- Non-root user lacks access to required directories

6. OOMKilled (Out Of Memory)

- Pod exceeds memory limits and is killed by kernel
- o Looks like CrashLoopBackOff but root cause is resource starvation







Step 1: kubectl get pods -n <namespace>

NAME READY STATUS RESTARTS AGE

frontend-67c7857bbd-jx97q 0/1 CrashLoopBackOff 5 2m50s

Look at the number of **restarts** — if it's going up rapidly, you're in a loop.

♦ Step 2: Describe the Pod

kubectl describe pod frontend-67c7857bbd-jx97q -n <namespace>

Look for:

- Last State → Terminated
- Reason: OOMKilled / Error
- Events at the bottom, especially related to:
 - Probes failing
 - Image pull success/failure
 - Container exit codes

P Example output:

Last State: Terminated

Reason: Error

Exit Code: 1

Message: fatal: unhandled exception...

♦ Step 3: Get Logs

kubectl logs <pod-name> -n <namespace>

If the pod restarts quickly, use:

kubectl logs <pod-name> -n <namespace> --previous

Look for:

Stack traces



- Config file not found
- Permission denied
- Connection refused
- Syntax errors

♦ Step 4: Exec into the Pod (if it stays alive)

kubectl exec -it <pod> -n <namespace> -- sh

Check:

- File paths
- Permissions
- Missing files or misconfigured environment variables
- Network access to DBs or services

☑ Resolution Steps

Here's how you recover depending on the root cause:

Scenario 1: Missing Config or Secret

- Use kubectl describe deployment <name> and check envFrom, valueFrom, or secretKeyRef fields
- Ensure the Secret or ConfigMap exists

kubectl get secret | grep my-app-secret

kubectl get configmap | grep config



- · Recreate the missing config
- Re-apply the deployment
- Use default values where possible to avoid hard crashes

Scenario 2: Bad Probes

• If probes are too aggressive, the pod may get killed before it's fully ready







• Increase initialDelaySeconds, timeoutSeconds, and failureThreshold

livenessProbe:

httpGet:

path: /health

port: 8080

initialDelaySeconds: 20

periodSeconds: 5

timeoutSeconds: 2

failureThreshold: 5

• Temporarily remove liveness/readiness probes to stabilize, then reintroduce slowly

Scenario 3: App Bug or Crash

- Fix the bug in code or startup script
- Rebuild and push the Docker image
- Redeploy with updated tag (never use :latest)

docker build -t myapp:v1.2.3.

docker push myapp:v1.2.3

kubectl set image deployment/myapp myapp=myapp:v1.2.3

🗞 Scenario 4: OOMKilled

Describe pod and look for:

Last State: Terminated

Reason: OOMKilled

Fix:

• Increase memory limit in spec:

resources:





requests:

memory: "256Mi"

limits:

memory: "512Mi"

- Optimize app memory usage
- Monitor with Prometheus/metrics-server

Scenario 5: Command Error

· Check for command, args, or entrypoint issues

- Fix:
 - Remove unnecessary overrides in the Deployment YAML
 - Test locally using:

docker run -it myapp:v1.2.3 /bin/sh

- Stabilization Tactics
- Scale Down the Broken Deployment Temporarily

kubectl scale deployment <name> --replicas=0

☑ Roll Back to Last Known Good State

kubectl rollout undo deployment <name>

- **✓** Isolate the Issue in a Staging Namespace
 - Use kubectl get deployment -o yaml
 - Modify and apply in a test namespace

Best Practices to Prevent CrashLoopBackOff Storms

Practice	Description
Validate configs before deployment	Use kubectl diff, kubeval, Helm dry-runs





Practice	Description
Use health checks wisely	Tune probes per container startup time
Don't use :latest tags	Use pinned image versions
Automate smoke tests post- deployment	Use jobs, init containers, or canary checks
Always set resource requests & limits	Avoid random crashes and OOMs

Example Incident Timeline

Time	Event
10:00 AM	New image v2.0.1 deployed
10:01 AM	Probes start failing, pods restart
10:02 AM	Team sees CrashLoopBackOff on UI
10:04 AM	kubectl logs shows config missing
10:06 AM	Secret not mounted due to typo
10:08 AM	Secret fixed → redeploy → recovery

□ Wrap-Up

CrashLoopBackOff is **the most common K8s outage** — but also **the most fixable**, if you:

- Read logs carefully
- Use describe/events effectively
- Follow rollback-first, debug-later mindset



2. Node-Level Failures or Evictions

Pods are disappearing. Apps are unstable. And nobody knows why...

Root Cause

When nodes (i.e., Kubernetes worker machines) face issues, Kubernetes begins **evicting pods** to keep the system stable.



This leads to:

- Pods restarting or disappearing unexpectedly
- Services intermittently unavailable
- PVCs failing to mount (due to node change)
- Deployment rollouts getting stuck

These failures often seem "random" — until you realize it's **node pressure**, **kubelet crash**, or **resource exhaustion**.

⚠ Common Node-Level Outage Causes:

1. Memory/CPU pressure

- Node reaches resource limits
- Kubelet begins evicting lower-priority pods

2. Disk pressure / full disk

- Node cannot schedule or run pods
- o Ephemeral storage or image cache uses up space

3. Network partition

- Node becomes NotReady (lost heartbeat)
- Cluster fails to reach kubelet

4. Kubelet crash / misconfiguration

- Broken kubelet config, TLS issues
- No control loop running → node becomes unresponsive

5. Pod misplacement

 Pod is tied to volume/zone and can't be moved easily (e.g., StatefulSets with PVC)

% How to Diagnose





⋄ Step 1: Check Node Status

kubectl get nodes

Look for:

- NotReady → indicates loss of kubelet heartbeat
- SchedulingDisabled → node cordoned
- Taints or labels blocking scheduling

⋄ Step 2: Describe the Node

kubectl describe node < node-name >

Focus on:

- Conditions (MemoryPressure, DiskPressure, PIDPressure)
- **Events** → evictions, kubelet errors
- Allocated resources vs. available

P Example:

Conditions:

MemoryPressure True

DiskPressure False

Ready False

Events:

Type Reason Age From Message

---- -----

Warning EvictionThresholdMet 3m kubelet, ip-10-0-1-29 Attempting to reclaim memory

⋄ Step 3: Pod Events & Evictions

kubectl get events --sort-by='.lastTimestamp' | grep -i evict

You'll see something like:





Warning Evicted pod/frontend-xyz The node was low on resource: memory.

Step 4: Check System-Level Logs (Node OS)

SSH into the node (if possible):

journalctl -u kubelet

dmesg | tail

df -h

Look for:

- Disk full
- OOM events
- Docker/image GC logs
- Kubelet crash messages

☑ Resolution Steps

Scenario 1: Memory or CPU Pressure

Pods are being evicted to maintain node health.

Fix:

- Increase node size or cluster auto-scaler
- Lower pod resource limits
- Prioritize critical workloads with priorityClassName

priorityClassName: high-priority

• Use Vertical Pod Autoscaler (VPA) to right-size containers over time

Scenario 2: Disk Pressure / Full Volume

Node fails due to full disk — usually in:



- /var/lib/docker (image cache)
- /var/log (excessive logs)

Fix:

- SSH into node, clean up /var/log
- Run Docker image prune:

docker system prune -af

• Tune garbage collection thresholds in kubelet config:

evictionHard:

memory.available: "100Mi"

nodefs.available: "10%"

Scenario 3: Kubelet Crash / NotReady

Node enters NotReady state and pods get stuck.

Fix:

Restart kubelet on the node

systemctl restart kubelet

• Verify node joins the cluster again:

kubectl get nodes

If permanent issue → cordon and drain

kubectl cordon <node>

kubectl drain <node> --ignore-daemonsets --delete-emptydir-data

Then reschedule pods elsewhere.

Scenario 4: Pod Gets Evicted and Doesn't Reschedule

PVC stuck or pod with zone-affinity.

Fix:





- Check volumeBindingMode: WaitForFirstConsumer
- Use PodAntiAffinity to distribute pods across zones
- Enable StorageClass with dynamic provisioning
- Use StatefulSet properly to avoid zone/PVC conflicts

Preventive Measures & Best Practices

Task	Why It Helps
Monitor node resource usage	Avoid silent pressure buildup
Set requests and limits for all pods	Prevent noisy neighbor syndrome
Enable Cluster Autoscaler	Adds/removes nodes based on workload
Clean image and log caches	Prevent disk pressure
Prioritize critical services	Evict low-priority pods first (not core APIs)
Spread workloads	Use podAntiAffinity and topologySpreadConstraints

▶ Incident Timeline Example

Time Event

2:00 PM Ingress unavailable for some users

2:01 PM API server shows node ip-10-0-1-29 as NotReady

2:03 PM Events show memory pressure & pod evictions

2:04 PM Team drains node, reroutes workload

2:07 PM Cluster stabilizes after HPA and rescheduling

Summary





Node issues can be stealthy but catastrophic.

Knowing how to detect, drain, debug, and reschedule gives you **real control** as a Kubernetes Admin.

☑ 3. PersistentVolume Claim (PVC) Issues

*My pod is stuck in Pending..." — aka, your app can't store or read any data.

Q Root Cause

A PersistentVolumeClaim (PVC) outage occurs when:

- A pod can't bind to a volume
- A volume is attached to another node
- A pod is stuck in Pending or ContainerCreating





A StatefulSet crashes after restart because data is lost or inaccessible

This affects databases, caches, queues, and any workload that stores state.

↑ Common Causes of PVC-Related Outages

1. StorageClass misconfiguration

- StorageClass name typo
- Non-existent or deprecated provisioner

2. **Zone/Availability Mismatch**

- Pod scheduled in a zone that can't access the volume
- Common in AWS/GCP with zonal storage

3. Volume is stuck in use

- Volume still attached to old pod/node
- Failed to detach during rescheduling

4. Missing or deleted PV

- Underlying EBS/GCE disk deleted manually
- PV deleted while PVC still exists

5. Pod restarted without a PVC

YAML missing volumeClaimTemplates or persistentVolumeClaim

% How to Diagnose

♦ Step 1: Get PVC Status

kubectl get pvc -n <namespace>

Look for:

STATUS: Pending → not bound

STATUS: Lost → bound PV is gone



STATUS: Terminating → stuck deletion

♦ Step 2: Describe PVC

kubectl describe pvc <name> -n <namespace>

Check:

- **Bound Volume** (PV name)
- StorageClass
- Events
 - "waiting for a volume to be created"
 - "waiting for first consumer"

Q	Examp	le:
- 1	LAGITIP	с.

Type	Reason	Age From	Message

Normal Provisioning 2m persistentvolume-controller External provisioner is provisioning volume

Warning ProvisioningFailed 2m persistentvolume-controller StorageClass not found

♦ Step 3: Get Pod Status

kubectl get pods -o wide -n <namespace>

kubectl describe pod <name>

Look for:

- Volume mount issues
- Pod stuck in ContainerCreating
- "Unable to attach or mount volumes"
- Events related to volumeAttachment



◆ Step 4: Inspect the Node (if needed)

SSH into node: Isblk mount df-h Look for volumes still mounted or stale device references (esp. EBS). **✓** Resolution Steps Scenario 1: StorageClass Misconfigured or Missing P Error: Warning ProvisioningFailed StorageClass "ebs-sc" not found Fix: • Create correct StorageClass: apiVersion: storage.k8s.io/v1 kind: StorageClass metadata: name: ebs-sc provisioner: kubernetes.io/aws-ebs volumeBindingMode: WaitForFirstConsumer • Ensure volumeBindingMode matches your cluster zone setup. Scenario 2: Volume Bound but Pod Can't Mount P Error: Failed to attach volume "pvc-xxxx" to node "ip-10-0-1-28": Timeout Fix:





- Volume might still be attached to old node → detach manually via cloud console
- Cordon old node:

kubectl cordon <node>

- Delete stuck pod → allow rescheduling
- Use anti-affinity to prefer same zone

Scenario 3: PVC Stuck in Pending

P Error:

waiting for first consumer

- Fix:
 - Add proper scheduler hints (e.g., nodeSelector, affinity)
 - Set volumeBindingMode: WaitForFirstConsumer (delays provisioning until pod is scheduled)
 - Check if there are enough nodes in the same zone as the volume

Scenario 4: StatefulSet Restart Breaks Volume Attach

- Fix:
 - Use volumeClaimTemplates for each replica
 - Never delete StatefulSet without deleting PVCs manually

kubectl delete pvc data-mongo-0

 Use retentionPolicy: Retain only if you're intentionally persisting PVC after pod delete

Scenario 5: Deleted or Lost PV

P Error:

PVC bound to PV "pv-xyz", but PV does not exist





Fix:

- Recreate the PV with same name, UID (if known)
- Restore from backup
- Update YAML to request a new PVC

Best Practices for PVC Stability

Practice	Why It Helps
Use WaitForFirstConsumer binding mode	Avoid provisioning volumes in wrong zone
Always specify StorageClass	Prevent default mismatch or fallback
Monitor volume events	Detect provisioning delays or failures
Don't hard-delete PVs manually	Avoid PVC stuck/lost state
Tag volumes with pod info (cloud)	Easier to trace and debug in AWS/GCP

- Bonus: How to Avoid Data Loss in Stateful Apps
- ✓ Use backup tools:
 - Velero for cluster + volume snapshots
 - Database-level backups (mysqldump, mongodump)
 - VolumeSnapshot API in Kubernetes 1.17+
- Automate PVC cleanup only after validation
- ✓ Use **Retain reclaimPolicy** on critical volumes:

reclaimPolicy: Retain

Sample Incident Timeline

Time	Event
3:00 PM	MongoDB pod stuck in Pending after node restart



Time	Event
3:02 PM	PVC unable to bind → storageClass missing
3:05 PM	Admin re-applies StorageClass, scales down StatefulSet
3:07 PM	PVC binds successfully → StatefulSet restored

Summary

Storage outages are sneaky — but often stem from simple misconfigurations. The key to recovery is:

- Understand PVC/PV lifecycle
- Know your cloud storage behavior
- Respect zones, bindings, and reclaim policies

✓ 4. Ingress or LoadBalancer Downtime

"The app is working... but nobody can access it."

Q Root Cause

Ingress and LoadBalancer issues result in **external traffic** never reaching your services.

Internally, pods might be healthy, but from the user's perspective — it's down.

⚠ Common Causes of Ingress/LoadBalancer Downtime:

1. Ingress misconfiguration

- Incorrect host/path rules
- Missing or incorrect service backend



Wrong port mapping

2. SSL/TLS certificate issues

- Expired certs
- CertManager not renewing
- Wrong DNS-01 challenge setup

3. Cloud LoadBalancer provisioning failure

- Failed due to wrong annotations
- Service type misconfigured (ClusterIP instead of LoadBalancer)
- LoadBalancer stuck in Pending

4. DNS misrouting

- External domain not pointed to correct LoadBalancer IP
- DNS record TTL caching old IPs

5. Network policy or security group blocking traffic

NSGs/Firewalls deny inbound access on ports (e.g., 80/443)

% How to Diagnose

⋄ Step 1: Check Ingress Resource

kubectl get ingress -n <namespace>

Look for:

- ADDRESS column: Is IP/hostname assigned?
- HOSTS: Does it match expected DNS?
- TLS: Enabled?

♦ Step 2: Describe the Ingress

kubectl describe ingress <name> -n <namespace>

Check for:

Annotations (e.g., nginx.ingress.kubernetes.io/*)



- Backend service mapping
- Events like:

Warning: error resolving service

Step 3: Validate LoadBalancer Service

kubectl get svc -n <namespace>

If EXTERNAL-IP is pending for minutes:

kubectl describe svc <name>

Possible causes:

- Wrong annotations
- No cloud provider integration
- Unsupported service type (must be LoadBalancer)

♦ Step 4: Check Certificate Status (for HTTPS)

If you're using **CertManager**:

kubectl describe certificate -n <namespace>

kubectl describe challenge -n <namespace>

Look for:

- Ready: False
- Failed to complete DNS challenge
- Error creating ACME challenge

⋄ Step 5: DNS Validation

Use dig, nslookup, or curl:

dig yourapp.example.com

Make sure the DNS A/AAAA record points to the **LoadBalancer EXTERNAL-IP**.





Test HTTP/HTTPS:

curl -v https://yourapp.example.com

☑ Resolution Steps

Scenario 1: Ingress Path or Host Mismatch

Symptom:

404 Not Found from ingress-controller

Fix:

Verify path in Ingress spec:

rules:

- host: yourapp.example.com

http:

paths:

- path: /

pathType: Prefix

backend:

service:

name: yourapp-svc

port:

number: 80

- Ensure Service yourapp-svc exists and is ClusterIP
- Ensure app listens on correct port (match service and container port)

Scenario 2: Stuck LoadBalancer / No External IP

Fix:



• Ensure service type is correct:

kind: Service

spec:

type: LoadBalancer

Check cloud provider annotations:

Cloud Example Annotation

AWS service.beta.kubernetes.io/aws-load-balancer-type: nlb

GCP cloud.google.com/load-balancer-type: "External"

Recreate the service if needed:

kubectl delete svc <name>

kubectl apply -f svc.yaml

Scenario 3: SSL Certificate Expired or Failing

Fix:

Re-issue cert:

kubectl delete certificate <name>

kubectl apply -f cert.yaml

- Ensure ClusterIssuer exists and is valid
- Use Let's Encrypt staging environment for testing:

issuerRef:

name: letsencrypt-staging

kind: ClusterIssuer

• Check DNS-01 challenge with CertManager logs

Scenario 4: DNS Not Updated

Fix:





- Update domain provider to point A record to LoadBalancer IP
- Use CNAME for dynamic DNS like AWS ELB
- Flush DNS cache locally:

sudo systemd-resolve --flush-caches

Scenario 5: Firewall/Security Group Blocks



- Open inbound ports (80, 443) in:
 - AWS Security Groups
 - Azure NSGs
 - GCP Firewall Rules

Check NGINX ingress logs:

kubectl logs -n ingress-nginx -l app.kubernetes.io/name=ingress-nginx

Best Practices for Ingress/LoadBalancer Stability

Best Practice	Why It Helps
Use pathType: Prefix	Avoid path match errors
Always monitor cert-manager logs	Avoid silent certificate failures
Use DNS TTLs of 300s or less	Enable fast propagation during cutovers
Tag and track LoadBalancer IPs	Map DNS records accurately
Use external-dns to auto- manage	Automate route53/CloudDNS/GoDaddy records

Sample Outage Timeline



Time	Event
10:00 AM	Service deployed, but Ingress returns 404
10:02 AM	Ingress backend service name was typo (myapp-svcx)
10:03 AM	CertManager fails to renew cert due to DNS challenge
10:06 AM	Fixed service name and DNS TXT record → HTTPS restored
10:09 AM	Domain resolves to LoadBalancer IP → Site up

Summary

Ingress and LoadBalancer outages are common but often:

- Easily diagnosable with describe, logs, dig, and curl
- Caused by simple misconfigurations, not hard bugs
- Fixed by re-checking: ingress → service → cert → DNS → cloud provider

5. DNS Resolution Failures Inside the Cluster

(My pods are running... but they can't talk to each other."

Q Root Cause

In Kubernetes, internal DNS (usually powered by **CoreDNS**) enables service discovery.

Pods use service names like myapp.default.svc.cluster.local instead of IPs.

When DNS resolution fails, it breaks:

- Microservice-to-microservice communication
- Database lookups (db-service.default.svc)
- Init containers waiting on services
- Liveness probes that use hostnames





⚠ Common Causes of DNS Failures in Kubernetes:

- 1. CoreDNS pods crash or become unschedulable
- 2. CoreDNS ConfigMap errors (e.g., malformed forward or stubDomain)
- 3. Network plugin (CNI) breaks DNS routing
- 4. Pods misconfigured or lacking DNSPolicy
- 5. loop plugin missing in CoreDNS config (causes recursive lookups)
- 6. Kubelet config breaking DNS resolution on nodes

K How to Diagnose

♦ Step 1: Check CoreDNS Pod Status

kubectl get pods -n kube-system -l k8s-app=kube-dns

Expected:

coredns-6d8c4cb4d-xxxx 1/1 Running

If status is CrashLoopBackOff, Pending, or Evicted — it's a red flag.

⋄ Step 2: Check CoreDNS Logs

kubectl logs -n kube-system -l k8s-app=kube-dns

Look for:

- plugin/forward: no upstream
- unable to forward request
- loop detected
- cannot resolve...

Step 3: Run DNS Test from Inside a Pod





Start a debugging pod:

kubectl run -i --tty dnsutils --image=busybox --restart=Never -- sh

Inside:

nslookup kube-dns.kube-system.svc.cluster.local

nslookup google.com

- If both fail: cluster DNS is broken
- If only cluster domain fails: **CoreDNS problem**
- Tild google.com fails too: Node DNS config broken

⋄ Step 4: Check CoreDNS ConfigMap

kubectl get configmap coredns -n kube-system -o yaml

Look for:

- Incorrect forward plugin usage
- Recursive DNS loops
- Missing or malformed entries like:

```
::53 {
    errors
    health
    kubernetes cluster.local in-addr.arpa ip6.arpa {
        pods insecure
        fallthrough in-addr.arpa ip6.arpa
    }
    forward . /etc/resolv.conf
    cache 30
    loop
    reload
    loadbalance
```





}

☑ Resolution Steps

Scenario 1: CoreDNS Pod Down or Crashing

Symptom:

CrashLoopBackOff - plugin/forward: no upstream

- Fix:
 - Check if CoreDNS can't reach upstream resolver:
 - /etc/resolv.conf inside CoreDNS container might be misconfigured
 - Restart CoreDNS:

kubectl rollout restart deployment coredns -n kube-system

 If using custom DNS servers, add them to the ConfigMap forward section.

Scenario 2: CoreDNS ConfigMap Errors

Symptom:

plugin/loop: detected recursive DNS query

- Fix:
 - Add loop plugin:

loop

- Always keep this below the forward plugin to prevent recursion
- Reapply ConfigMap and restart CoreDNS:

kubectl apply -f coredns.yaml

kubectl rollout restart deployment coredns -n kube-system

Scenario 3: CNI/Networking Plugin is Blocking DNS





Fix:

- If CoreDNS is running but pods can't resolve, the issue could be **networking (CNI plugin)**:
 - Check iptables, Calico/Weave/Flannel logs
 - Restart CNI pods

kubectl rollout restart ds <cni-daemonset-name> -n kube-system

• Ensure DNS traffic is allowed between namespaces and to kube-system

Scenario 4: Pod DNSPolicy Misconfigured

Fix:

• Pod should have:

dnsPolicy: ClusterFirst

This ensures DNS queries go through CoreDNS first.

• Avoid overriding dnsConfig unless you know what you're doing.

Best Practices to Prevent Cluster DNS Outages

Practice	Why It Matters
Use liveness/readiness probes for CoreDNS	Auto-recovery of DNS service
Always test nslookup and dig during deploys	Validate app-to-app communication
Backup and version your CoreDNS ConfigMap	Prevent accidental edits
Don't delete CoreDNS in dev/test by mistake	Some tools may remove it (like kind delete)
Monitor DNS resolution latency via Prometheus	Catch slow queries before failure





Sample Incident Timeline

Time	Event
9:00 AM	Apps report connection timeout errors internally
9:01 AM	DNS resolution failing for
	mydb.default.svc
9:02 AM	CoreDNS pods seen in CrashLoopBackOff
9:04 AM	Logs show loop plugin missing
9:05 AM	Added loop to ConfigMap, restarted
	CoreDNS
9:07 AM	Internal DNS restored, apps working again

Summary

Internal DNS failures are invisible to external users but cripple microservices.

The fix often comes down to:

- Ensuring CoreDNS is healthy and configured right
- Validating pod-level resolution with tools like nslookup
- Keeping your CNI and CoreDNS plugins updated and monitored



☑ 6. Etcd Data Corruption or Unavailability

(4) "The control plane is choking... nothing responds."

Q Root Cause

Etcd is the **key-value store** behind Kubernetes. If it's down or corrupted:

- API server becomes unresponsive
- Cluster state can't be read or written
- No new pods can be created
- Existing workloads may run, but can't be managed

⚠ Common Triggers

- Disk full on etcd host
- Network split between etcd peers
- Etcd crash or panic due to corruption
- Snapshot restoration gone wrong
- Wrong etcd version or flags during upgrade

% Quick Diagnosis

kubectl get componentstatuses

etcd should be Healthy. If not → investigate immediately.



SSH into master node:

journalctl -u etcd

Check for:

- panic, corrupt, or disk quota exceeded
- Connectivity issues between etcd nodes

Check etcd health manually:

etcdctl endpoint health

✓ Resolution Steps

1. Disk Full / IOPS Bottleneck

- Free up space or resize disk
- Ensure SSD-backed volume with fast IOPS

2. Corrupted DB

- Stop etcd
- Move corrupted DB:

mv /var/lib/etcd /var/lib/etcd.bak

• Restore from snapshot:

etcdctl snapshot restore <snapshot.db>



- Needs majority of nodes (n/2 + 1) to be healthy
- Bring back offline peers or remove stale ones from cluster config

Prevention

Practice	Why It Helps
Schedule etcd snapshots daily	Enables fast recovery
Monitor etcd disk usage	Avoid quota triggers





Practice	Why It Helps
Use dedicated etcd disks	Isolate from OS/app traffic
Never kill etcd without backup	Risk of total cluster loss
Always upgrade etcd with caution	Version mismatch can corrupt

Incident Snapshot

Time	Event
2:00 PM	kubectl hangs — API server slow
2:01 PM	etcd shows disk full, panic in logs
2:05 PM	Snapshot restored, etcd restarted
2:07 PM	API back up, cluster recovered



☑ 7. API Server Saturation or Hang

🔛 "kubectl times out. Dashboards stop loading. Everything feels stuck."

Q Root Cause

The **Kubernetes API server** is the brain of the cluster. If it slows or crashes:

- kubectl commands fail
- CI/CD pipelines hang
- Cluster operations break
- Monitoring tools can't pull data

↑ Common Causes

- Too many requests (e.g., thousands of kubectl/monitoring calls)
- Excessive watch connections (Prometheus, controllers)
- Heavy write load (frequent updates from custom controllers)
- Etcd lagging or unreachable
- Admission webhooks timing out

% Quick Diagnosis

♦ Step 1: Check API server status

kubectl get --raw /healthz

Fails = unhealthy.

♦ Step 2: View metrics (if accessible)



kubectl top pods -n kube-system

High CPU/memory usage in kube-apiserver?

⋄ Step 3: Look into logs

kubectl logs -n kube-system -l component=kube-apiserver

Look for:

timeout, too many open files, etcd timeout, webhook errors

✓ Resolution Steps

1. Reduce Load

- Pause CI/CD pipelines
- Temporarily disable noisy monitoring agents

2. Audit & Remove Excess Watchers

- Reduce Prometheus scrape intervals
- Limit number of kubectl watches from automation

3. Scale the Control Plane

- Increase CPU/memory for API server pods (if self-managed)
- In EKS/GKE → upgrade control plane or contact support

4. Fix Etcd or Webhook Lag

- Etcd timeout = fix etcd first
- Misbehaving webhook = disable via ValidatingWebhookConfiguration

Prevention

Practice	Why It Helps
Set API rate limits via audit/webhook	Prevent request overload
Use HorizontalPodAutoscaler for clients	Avoid spiky loads
Keep admission webhooks lightweight	Prevent processing delays



Practice	Why It Helps
Monitor apiserver_request_duration	Catch early signs of stress

Incident Snapshot

Time	Event
6:00 PM	kubectl get pods hangs
6:02 PM	kube-apiserver logs show timeouts
6:04 PM	Paused CI jobs + throttled Prometheus
6:07 PM	API server load drops → cluster recovers



☑ 8. CrashLoopBackOff Storms

"Pods keep restarting endlessly. Nothing is stable."

Q Root Cause

This happens when pods crash repeatedly and Kubernetes **backs off** from restarting them.

Common symptoms:

- kubectl get pods shows CrashLoopBackOff
- App is never healthy
- Cluster gets overloaded if many pods are affected

↑ Common Causes

- · App code bug or misconfig
- Bad secrets/configMaps/env vars
- Missing volume mount
- Init container fails
- · Liveness/readiness probes misconfigured
- Dependency service (e.g., DB) unavailable

% Quick Diagnosis

kubectl describe pod <pod-name>

kubectl logs <pod-name> -c <container>

Check:

Recent exit codes (137 = OOM, 1 = crash, 126/127 = permission issues)



- Probe failures
- Init container failures

✓ Resolution Steps

1. App Bug or Env Issue

- · Check logs for stack trace
- Fix env vars, image tags, secrets
- Validate command and args

2. Probe Failures

- Comment out livenessProbe and readinessProbe temporarily
- Test basic container startup

3. Missing Volume or PVC

- Check if PVC is Pending or Lost
- Fix volume or storage class

4. Dependency Unavailable

- App waiting on DB, Redis, etc?
- Add initContainers to check connectivity before start

Prevention

Practice	Benefit
Add retry logic in apps	Handles boot-time flakiness
Use startupProbe for slow apps	Avoid early restarts
Delay readiness until ready	Prevents traffic to bad pods
Add BackoffLimit in Jobs	Stops endless retries

Incident Snapshot





Time	Event
11:00 AM	5 microservices stuck in CrashLoop
11:02 AM	Env var typo in secret (DB_HOST)
11:05 AM	Secret fixed → pods stable again





Knowing how to kubectl get pods is just the surface.

Running a Kubernetes cluster in production is about knowing what breaks — and how to fix it fast.

These 8 outage scenarios represent the **most common real-world failures** DevOps engineers face:

- When your app is healthy but unreachable
- When Pods crash endlessly for simple reasons
- When etcd corruption locks the whole cluster
- When DNS silently breaks microservice communication
- When PVCs prevent pods from ever starting
- When Ingress or LoadBalancers misroute live traffic
- When your API server collapses under load
- When the control plane itself stops responding

Want to Stand Out in DevOps Interviews?

X Don't just say: "I deployed apps to Kubernetes."

Say: "I simulated production outages in my home lab and recovered from them—DNS failures, LoadBalancer issues, PVC crashes, API slowdowns, and etcd loss."

That's what real DevOps confidence sounds like.

Your Next Step:

Practice these scenarios in your home lab

Document how you fixed them

Build failure-resilient clusters

Turn these experiences into high-impact stories in interviews and on your resume