

Essential Knowledge for IT Professionals

A Comprehensive Technical Guide

By Mahabir Singh Bisht

November 05, 2025

Table of Contents

1. Linux Fundamentals
2. Networking Essentials
3. Database Management
4. Security Principles
5. Storage Technologies
6. Caching Strategies

[7. Disaster Recovery](#)

[8. Windows Administration](#)

Linux Fundamentals

File Systems

Linux supports multiple file system types, each with specific use cases and characteristics:

Common File System Types:

- **ext4 (Fourth Extended File System)**: The default file system for most Linux distributions. Supports journaling, large file sizes (up to 16 TB), and volumes up to 1 exabyte. Provides excellent performance and reliability for general-purpose use.
- **XFS**: High-performance 64-bit journaling file system, excellent for large files and parallel I/O operations. Commonly used in enterprise environments and with large storage arrays.
- **Btrfs (B-tree File System)**: Modern copy-on-write file system with advanced features like snapshots, compression, and self-healing capabilities. Still evolving but gaining adoption.
- **ZFS**: Originally from Solaris, now available on Linux. Provides data integrity verification, RAID support, snapshots, and excellent scalability.

File System Hierarchy:

- `/` (root): Top-level directory containing all other directories
- `/bin` : Essential user command binaries
- `/boot` : Boot loader files and kernel
- `/dev` : Device files representing hardware
- `/etc` : System configuration files
- `/home` : User home directories
- `/lib` : Essential shared libraries
- `/var` : Variable data (logs, databases, web content)
- `/tmp` : Temporary files

- `/usr` : User programs and data
- `/opt` : Optional third-party software

Key Concepts:

- **Inodes**: Data structures storing metadata about files (permissions, timestamps, ownership)
- **Mounting**: Attaching file systems to the directory tree
- **Journaling**: Logging changes before committing them to prevent corruption

Package Management

Package managers automate software installation, updates, and dependency resolution.

Debian-based Systems (`apt`, `dpkg`):

```
# Update package lists
apt update

# Install a package
apt install package-name

# Remove a package
apt remove package-name

# Search for packages
apt search keyword

# Show package information
apt show package-name

# Upgrade all packages
apt upgrade

# Low-level package installation
dpkg -i package.deb
```

Red Hat-based Systems (yum, dnf, rpm):

```
# Install a package  
dnf install package-name  
  
# Remove a package  
dnf remove package-name  
  
# Update all packages  
dnf update  
  
# Search for packages  
dnf search keyword  
  
# List installed packages  
rpm -qa  
  
# Install from RPM file  
rpm -ivh package.rpm
```

Key Concepts:

- **Repositories:** Centralized locations hosting packages
- **Dependencies:** Required packages for software to function
- **Package signing:** Cryptographic verification ensuring package authenticity

Systemd

Systemd is the modern init system and service manager for Linux, replacing older systems like SysVInit and Upstart.

Core Components:

- **systemd:** The system and service manager

- **journald**: Logging service
- **networkd**: Network management
- **resolved**: DNS resolution
- **timesyncd**: Time synchronization

Service Management:

```
# Start a service
systemctl start service-name

# Stop a service
systemctl stop service-name

# Restart a service
systemctl restart service-name

# Enable service at boot
systemctl enable service-name

# Disable service at boot
systemctl disable service-name

# Check service status
systemctl status service-name

# List all services
systemctl list-units --type=service

# View service logs
journalctl -u service-name
```

Unit Files:

Located in `/etc/systemd/system/` or `/lib/systemd/system/`, unit files define service behavior:

```
[Unit]
Description=My Application
After=network.target

[Service]
Type=simple
User=appuser
ExecStart=/usr/bin/myapp
Restart=on-failure

[Install]
WantedBy=multi-user.target
```

Targets:

Systemd uses targets (similar to runlevels) to define system states:

- `multi-user.target` : Multi-user system without GUI
- `graphical.target` : Multi-user system with GUI
- `rescue.target` : Single-user rescue mode

Permissions

Linux employs a robust permission system controlling access to files and directories.

Permission Types:

- **Read (r/4)**: View file contents or list directory contents
- **Write (w/2)**: Modify file contents or create/delete files in directory
- **Execute (x/1)**: Run file as program or access directory

Permission Categories:

- **User (u)**: File owner
- **Group (g)**: Members of the file's group

- **Others (o):** All other users

Viewing Permissions:

```
ls -l file.txt
# Output: -rw-r--r-- 1 user group 1234 Nov 04 10:00 file.txt
# First character: file type (- = file, d = directory, l = symlink)
# Next 9 characters: permissions (rwx rwx rwx for user, group, others)
```



Modifying Permissions:

```
# Symbolic notation
chmod u+x file.txt          # Add execute for user
chmod g-w file.txt          # Remove write for group
chmod o=r file.txt          # Set others to read-only

# Numeric notation
chmod 755 file.txt          # rwxr-xr-x
chmod 644 file.txt          # rw-r--r--
chmod 600 file.txt          # rw-----
```

Ownership:

```
# Change owner
chown user:group file.txt

# Change group only
chgrp group file.txt

# Recursive ownership change
chown -R user:group /path/to/directory
```

Special Permissions:

- **SUID (Set User ID - 4000)**: Execute file with owner's permissions
- **SGID (Set Group ID - 2000)**: Execute file with group's permissions; new files inherit directory group
- **Sticky Bit (1000)**: Only owner can delete files in directory (used on /tmp)

Logs

Logging is critical for troubleshooting, security auditing, and system monitoring.

Traditional Logging (rsyslog):

- Configuration: /etc/rsyslog.conf
- Log files typically stored in /var/log/

Common Log Files:

- /var/log/syslog or /var/log/messages : General system messages
- /var/log/auth.log : Authentication attempts
- /var/log/kern.log : Kernel messages
- /var/log/dmesg : Boot and hardware messages
- /var/log/apache2/ or /var/log/httpd/ : Web server logs
- /var/log/mail.log : Mail server logs

Systemd Journal (journalctl):

```
# View all logs
journalctl

# Follow logs in real-time
journalctl -f

# Show logs for specific service
```

```
journalctl -u ssh.service

# Show logs since boot
journalctl -b

# Show logs for time range
journalctl --since "2025-11-01" --until "2025-11-04"

# Show only errors
journalctl -p err

# Show kernel messages
journalctl -k

# Limit output
journalctl -n 50
```

Log Rotation:

Managed by `logrotate` to prevent logs from consuming excessive disk space:

- Configuration: `/etc/logrotate.conf` and `/etc/logrotate.d/`
- Automatically compresses, rotates, and removes old logs

Disk Management

Viewing Disk Usage:

```
# Show disk space usage
df -h

# Show directory sizes
du -sh /path/to/directory

# Show largest directories
```

```
du -h /var | sort -rh | head -n 10

# Check inode usage
df -i
```

Partitioning:

```
# List block devices
lsblk

# Partition disk (interactive)
fdisk /dev/sda

# GPT partition table
gdisk /dev/sda

# Partition with parted
parted /dev/sda
```

File System Creation:

```
# Create ext4 file system
mkfs.ext4 /dev/sda1

# Create XFS file system
mkfs.xfs /dev/sda1

# Create file system with label
mkfs.ext4 -L mylabel /dev/sda1
```

Mounting:

```
# Mount file system
mount /dev/sda1 /mnt

# Unmount file system
umount /mnt

# Persistent mounts in /etc/fstab
# Format: device    mount-point    type    options    dump    pass
/dev/sda1    /data    ext4    defaults    0    2

# Remount all filesystems in fstab
mount -a
```

LVM (Logical Volume Manager):

Provides flexible disk management with volume resizing and snapshots:

```
# Create physical volume
pvcreate /dev/sda1

# Create volume group
vgcreate myvg /dev/sda1

# Create logical volume
lvcreate -L 10G -n mylv myvg

# Extend logical volume
lvextend -L +5G /dev/myvg/mylv
resize2fs /dev/myvg/mylv # Resize ext4 filesystem
```

Process Management

Viewing Processes:

```
# List all processes
ps aux

# Process tree
ps auxf
pstree

# Interactive process viewer
top

# Enhanced top
htop

# List processes by resource usage
ps aux --sort=-%cpu | head
ps aux --sort=-%mem | head
```

Process Control:

```
# Run command in background
command &

# List background jobs
jobs

# Bring job to foreground
fg %1

# Send job to background
bg %1

# Kill process by PID
kill 1234
kill -9 1234 # Force kill (SIGKILL)

# Kill process by name
```

```
killall process-name  
pkill process-name  
  
# Change process priority  
nice -n 10 command      # Start with lower priority  
renice -n 5 -p 1234      # Change priority of running process
```

Process Priority:

- Nice values range from -20 (highest priority) to 19 (lowest priority)
- Default nice value is 0
- Only root can set negative nice values

System Resource Monitoring:

```
# CPU and memory usage  
vmstat 1  
  
# I/O statistics  
iostat -x 1  
  
# Network statistics  
netstat -i  
ss -s  
  
# Real-time process monitoring  
watch -n 1 'ps aux --sort=-%cpu | head -n 20'
```

Networking Essentials

TCP/IP

The TCP/IP protocol suite is the foundation of internet communication, consisting of four layers:

TCP/IP Model Layers:

1. **Application Layer:** Provides network services to applications
 - Protocols: HTTP, HTTPS, FTP, SMTP, DNS, SSH, DHCP
2. **Transport Layer:** Provides end-to-end communication
 - **TCP (Transmission Control Protocol):** Connection-oriented, reliable, ordered delivery
 - Three-way handshake (SYN, SYN-ACK, ACK)
 - Flow control and congestion control
 - Used for web browsing, email, file transfer
 - **UDP (User Datagram Protocol):** Connectionless, unreliable, faster
 - No handshake or delivery guarantee
 - Used for streaming, DNS queries, VoIP
3. **Internet Layer:** Handles packet routing
 - **IP (Internet Protocol):** Addressing and routing
 - IPv4: 32-bit addresses (e.g., 192.168.1.1)
 - IPv6: 128-bit addresses (e.g., 2001:0db8:85a3:0000:0000:8a2e:0370:7334)
 - **ICMP:** Error reporting and diagnostics (ping, traceroute)
4. **Network Access Layer:** Physical network interface
 - Ethernet, Wi-Fi, PPP

Port Numbers:

- Well-known ports (0-1023): Reserved for common services
 - 22: SSH
 - 80: HTTP
 - 443: HTTPS
 - 53: DNS
 - 25: SMTP
 - 3306: MySQL
- Registered ports (1024-49151): Assigned to specific services
- Dynamic/private ports (49152-65535): Used for client connections

TCP Three-Way Handshake:

1. Client sends SYN (synchronize) packet
2. Server responds with SYN-ACK (synchronize-acknowledge)

3. Client sends ACK (acknowledge)

4. Connection established

DNS (Domain Name System)

DNS translates human-readable domain names into IP addresses.

DNS Record Types:

- **A**: Maps domain to IPv4 address
- **AAAA**: Maps domain to IPv6 address
- **CNAME**: Canonical name (alias) for another domain
- **MX**: Mail exchange server for domain
- **TXT**: Text information (SPF, DKIM, verification records)
- **NS**: Name server authoritative for domain
- **PTR**: Reverse DNS lookup (IP to domain)
- **SOA**: Start of authority (primary nameserver info)
- **SRV**: Service location record

DNS Hierarchy:

1. Root servers (.)
2. Top-level domain (TLD) servers (.com, .org, .net)
3. Authoritative nameservers (example.com)

DNS Resolution Process:

1. Client queries local DNS resolver
2. If not cached, resolver queries root server
3. Root server responds with TLD server
4. Resolver queries TLD server
5. TLD server responds with authoritative nameserver
6. Resolver queries authoritative nameserver
7. Authoritative nameserver returns IP address
8. Result cached for TTL (Time To Live) period

DNS Commands:

```
# Query DNS records  
nslookup example.com  
  
# Detailed DNS lookup  
dig example.com  
  
# Query specific record type  
dig example.com MX  
dig example.com AAAA  
  
# Reverse DNS lookup  
dig -x 8.8.8.8  
  
# Trace DNS resolution path  
dig +trace example.com
```

HTTP/S

HTTP (HyperText Transfer Protocol):

- Application-layer protocol for web communication
- Stateless protocol (each request independent)
- Uses TCP port 80

HTTP Methods:

- **GET**: Retrieve resource
- **POST**: Submit data to server
- **PUT**: Update/replace resource
- **PATCH**: Partial resource update
- **DELETE**: Remove resource
- **HEAD**: Retrieve headers only

- **OPTIONS:** Query supported methods

HTTP Status Codes:

- **1xx:** Informational (100 Continue)
- **2xx:** Success (200 OK, 201 Created, 204 No Content)
- **3xx:** Redirection (301 Moved Permanently, 302 Found, 304 Not Modified)
- **4xx:** Client errors (400 Bad Request, 401 Unauthorized, 403 Forbidden, 404 Not Found)
- **5xx:** Server errors (500 Internal Server Error, 502 Bad Gateway, 503 Service Unavailable)

HTTPS (HTTP Secure):

- HTTP over TLS/SSL encryption
- Uses TCP port 443
- Provides confidentiality, integrity, and authentication
- Requires SSL/TLS certificate from Certificate Authority (CA)

TLS Handshake:

1. Client sends supported cipher suites
2. Server responds with chosen cipher and certificate
3. Client verifies certificate
4. Key exchange establishes session keys
5. Encrypted communication begins

VPN (Virtual Private Network)

VPNs create secure, encrypted tunnels over public networks.

VPN Types:

1. **Remote Access VPN:** Connects individual users to corporate network
 - Common protocols: OpenVPN, L2TP/IPsec, IKEv2
2. **Site-to-Site VPN:** Connects entire networks
 - IPsec is most common protocol

- Permanent connection between locations

VPN Protocols:

- **IPsec:** Operates at network layer, provides strong encryption
- **OpenVPN:** Open-source, highly configurable, uses SSL/TLS
- **WireGuard:** Modern, lightweight, high-performance
- **L2TP/IPsec:** Layer 2 Tunneling Protocol with IPsec encryption
- **PPTP:** Older, less secure, largely deprecated

VPN Benefits:

- Remote access to internal resources
- Encrypted data transmission
- IP address masking
- Bypass geographic restrictions

Load Balancers

Load balancers distribute network traffic across multiple servers to improve availability and performance.

Load Balancing Algorithms:

1. **Round Robin:** Distributes requests sequentially to each server
2. **Least Connections:** Routes to server with fewest active connections
3. **IP Hash:** Uses client IP to determine server (session persistence)
4. **Weighted Round Robin:** Assigns weights based on server capacity
5. **Least Response Time:** Routes to server with fastest response

Load Balancer Types:

1. Layer 4 (Transport Layer):

- Routes based on IP address and TCP/UDP port
- Fast, simple, protocol-agnostic
- Cannot inspect application data

2. Layer 7 (Application Layer):

- Routes based on HTTP headers, cookies, URL paths
- Content-based routing (e.g., /api → API servers, /images → CDN)
- SSL termination
- More processing overhead

Health Checks:

Load balancers continuously monitor backend server health:

- HTTP status checks
- TCP connection tests
- Custom application checks
- Automatic removal of failed servers

Popular Load Balancers:

- HAProxy
- Nginx
- AWS Elastic Load Balancer (ELB)
- Azure Load Balancer
- Google Cloud Load Balancing

Firewalls

Firewalls control network traffic based on security rules.

Firewall Types:

1. **Packet Filtering:** Examines packet headers (source/destination IP, port, protocol)
2. **Stateful Inspection:** Tracks connection state, allows return traffic
3. **Application Layer:** Inspects application data (e.g., HTTP content)
4. **Next-Generation (NGFW):** Combines features with IPS, malware detection

Linux Firewall Tools:

iptables/nftables:

```
# List rules
iptables -L -n -v

# Allow incoming SSH
iptables -A INPUT -p tcp --dport 22 -j ACCEPT

# Allow established connections
iptables -A INPUT -m state --state ESTABLISHED,RELATED -j ACCEPT

# Block specific IP
iptables -A INPUT -s 192.168.1.100 -j DROP

# Save rules
iptables-save > /etc/iptables/rules.v4
```

firewalld:

```
# List active zones
firewall-cmd --get-active-zones

# Add service to zone
firewall-cmd --zone=public --add-service=http --permanent

# Add port
firewall-cmd --zone=public --add-port=8080/tcp --permanent

# Reload firewall
firewall-cmd --reload
```

UFW (Uncomplicated Firewall):

```
# Enable firewall
ufw enable

# Allow service
ufw allow ssh
ufw allow 80/tcp

# Deny traffic from IP
ufw deny from 192.168.1.100

# Check status
ufw status verbose
```

Network Protocols

Common Application Protocols:

- **SSH (Secure Shell)**: Secure remote access, port 22
- **FTP (File Transfer Protocol)**: File transfer, ports 20-21 (insecure)
- **SFTP/SCP**: Secure file transfer over SSH
- **SMTP (Simple Mail Transfer Protocol)**: Email sending, port 25
- **IMAP/POP3**: Email retrieval, ports 143/110
- **NTP (Network Time Protocol)**: Time synchronization, port 123
- **SNMP (Simple Network Management Protocol)**: Network monitoring, port 161

Network Diagnostic Tools:

```
# Test connectivity
ping example.com

# Trace network path
traceroute example.com

# Port scanning
```

```
nmap -sV example.com

# TCP/UDP connections
netstat -tulpn
ss -tulpn

# Network interface configuration
ip addr show
ip route show

# Packet capture
tcpdump -i eth0 port 80
```

Subnetting

Subnetting divides IP networks into smaller subnetworks for efficient IP address management and security.

IPv4 Address Classes:

- Class A: 0.0.0.0 - 127.255.255.255 (8-bit network, 24-bit host)
- Class B: 128.0.0.0 - 191.255.255.255 (16-bit network, 16-bit host)
- Class C: 192.0.0.0 - 223.255.255.255 (24-bit network, 8-bit host)

CIDR (Classless Inter-Domain Routing) Notation:

- Format: IP address/prefix length
- Example: 192.168.1.0/24
 - Network: 192.168.1.0
 - Subnet mask: 255.255.255.0
 - Usable hosts: 254 (192.168.1.1 - 192.168.1.254)
 - Broadcast: 192.168.1.255

Subnet Mask Reference:

- /24 = 255.255.255.0 (256 addresses, 254 usable)

- $/25 = 255.255.255.128$ (128 addresses, 126 usable)
- $/26 = 255.255.255.192$ (64 addresses, 62 usable)
- $/27 = 255.255.255.224$ (32 addresses, 30 usable)
- $/28 = 255.255.255.240$ (16 addresses, 14 usable)

Private IP Ranges (RFC 1918):

- $10.0.0.0/8$ ($10.0.0.0 - 10.255.255.255$)
- $172.16.0.0/12$ ($172.16.0.0 - 172.31.255.255$)
- $192.168.0.0/16$ ($192.168.0.0 - 192.168.255.255$)

Subnetting Example:

Network: 192.168.1.0/24, need 4 subnets

1. Borrow 2 bits from host portion ($2^2 = 4$ subnets)
2. New subnet mask: $/26$ (255.255.255.192)
3. Subnets:
 - 192.168.1.0/26 (hosts: 1-62)
 - 192.168.1.64/26 (hosts: 65-126)
 - 192.168.1.128/26 (hosts: 129-190)
 - 192.168.1.192/26 (hosts: 193-254)

VLSM (Variable Length Subnet Masking):

Allows different subnet sizes within the same network for efficient IP allocation.

Database Management

SQL vs. NoSQL

SQL (Relational Databases):

Structured data stored in tables with predefined schemas and relationships.

Characteristics:

- Schema-based with strict data types
- ACID compliance (see below)
- Relationships via foreign keys
- SQL query language
- Vertical scaling (increase server resources)

Common SQL Databases:

- PostgreSQL: Advanced features, JSON support, extensible
- MySQL/MariaDB: Popular, fast, widely supported
- Oracle: Enterprise-grade, high performance
- Microsoft SQL Server: Windows integration, business intelligence

Use Cases:

- Financial transactions
- Business applications (ERP, CRM)
- Data requiring complex joins
- Strong consistency requirements

SQL Example:

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    username VARCHAR(50) UNIQUE NOT NULL,
    email VARCHAR(100) NOT NULL,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

```
CREATE TABLE orders (
    id SERIAL PRIMARY KEY,
    user_id INTEGER REFERENCES users(id),
    total DECIMAL(10,2),
```

```
    status VARCHAR(20)
);

SELECT u.username, COUNT(o.id) as order_count
FROM users u
LEFT JOIN orders o ON u.id = o.user_id
GROUP BY u.username;
```

NoSQL (Non-Relational Databases):

Flexible data models optimized for specific access patterns and scalability.

Types:

1. Document Stores (MongoDB, CouchDB):

- JSON-like documents
- Flexible schema
- Nested data structures

2. Key-Value Stores (Redis, DynamoDB):

- Simple key-value pairs
- Extremely fast
- In-memory options

3. Column-Family Stores (Cassandra, HBase):

- Wide-column data model
- Optimized for write-heavy workloads
- Distributed architecture

4. Graph Databases (Neo4j, Amazon Neptune):

- Nodes and relationships
- Complex relationship queries
- Social networks, recommendation engines

Characteristics:

- Schema-less or flexible schema
- Horizontal scaling (add more servers)
- Eventually consistent (often)

- Optimized for specific data patterns
- High throughput

Use Cases:

- Social media platforms
- Real-time analytics
- Content management systems
- IoT data collection
- Caching layers

NoSQL Example (MongoDB):

```
// Insert document
db.users.insertOne({
    username: "john_doe",
    email: "john@example.com",
    profile: {
        age: 30,
        city: "New York"
    },
    interests: ["coding", "music", "travel"]
});

// Query with nested data
db.users.find({ "profile.city": "New York" });
```

ACID Properties

ACID ensures reliable database transactions in SQL databases.

A - Atomicity:

- Transactions are all-or-nothing

- If any operation fails, entire transaction rolls back
- Example: Bank transfer must debit one account and credit another, or neither

```
BEGIN TRANSACTION;  
UPDATE accounts SET balance = balance - 100 WHERE id = 1;  
UPDATE accounts SET balance = balance + 100 WHERE id = 2;  
COMMIT; -- Both succeed or neither
```

C - Consistency:

- Database moves from one valid state to another
- All constraints, triggers, and rules enforced
- Data integrity maintained
- Example: Foreign key constraints prevent orphaned records

I - Isolation:

- Concurrent transactions don't interfere with each other
- Appears as if transactions execute sequentially
- Isolation levels:
 - Read Uncommitted: Allows dirty reads
 - Read Committed: Prevents dirty reads
 - Repeatable Read: Prevents non-repeatable reads
 - Serializable: Complete isolation (slowest)

D - Durability:

- Committed transactions persist even after system failure
- Data written to non-volatile storage
- Write-ahead logging (WAL) ensures recoverability

BASE (Alternative for NoSQL):

- Basically Available: System appears to work most of the time
- Soft state: State may change over time without input

- Eventually consistent: System becomes consistent given enough time

Scalability

Vertical Scaling (Scale Up):

- Increase server resources (CPU, RAM, storage)
- Simpler implementation
- Physical/cost limits
- Single point of failure
- Best for SQL databases

Horizontal Scaling (Scale Out):

- Add more servers
- Better fault tolerance
- Complex implementation
- Theoretically unlimited
- Best for NoSQL databases

Database Scaling Techniques:

1. Replication:

- **Master-Slave:** Master handles writes, slaves handle reads
- **Master-Master:** Multiple masters accept writes
- Improves read performance and availability

2. Sharding (Horizontal Partitioning):

- Split data across multiple servers by key
- Each shard contains subset of data
- Example: Users A-M on server1, N-Z on server2
- Challenges: Cross-shard queries, rebalancing

3. Partitioning:

- **Horizontal:** Split rows across tables
- **Vertical:** Split columns across tables
- Range, hash, or list-based partitioning

4. Caching:

- Store frequently accessed data in memory
- Reduces database load
- See Caching section below

5. Read Replicas:

- Dedicated servers for read operations
- Offload read traffic from primary database
- Asynchronous replication

Connection Pooling:

- Reuse database connections
- Reduces connection overhead
- Configurable pool size based on workload

Data Modeling

Relational Data Modeling:

Normalization:

Process of organizing data to reduce redundancy and improve integrity.

- **1NF (First Normal Form):** Atomic values, no repeating groups
- **2NF (Second Normal Form):** 1NF + no partial dependencies
- **3NF (Third Normal Form):** 2NF + no transitive dependencies

Example:

Unnormalized:

Order (OrderID, CustomerName, CustomerEmail, Items)

3NF:

Customers (CustomerID, Name, Email)

Orders (OrderID, CustomerID, OrderDate)

```
OrderItems (OrderID, ProductID, Quantity)
```

```
Products (ProductID, Name, Price)
```

Denormalization:

Intentionally introduce redundancy for performance (read-heavy workloads).

Entity-Relationship (ER) Modeling:

- **Entities:** Objects or things (Customer, Product)
- **Attributes:** Properties of entities (Name, Price)
- **Relationships:** Connections between entities (Customer places Order)
- **Cardinality:** One-to-One, One-to-Many, Many-to-Many

NoSQL Data Modeling:

Principles:

1. Model for access patterns, not entities
2. Denormalize data for query efficiency
3. Embed related data when accessed together
4. Use references for rarely accessed data

Document Database Example:

```
{
  "_id": "order_123",
  "customer": {
    "name": "John Doe",
    "email": "john@example.com"
  },
  "items": [
    { "product": "Widget", "quantity": 2, "price": 10.00 },
    { "product": "Gadget", "quantity": 1, "price": 25.00 }
  ]
}
```

```
],  
"total": 45.00,  
"created_at": "2025-11-04T10:00:00Z"  
}
```

Indexes:

- Speed up queries by creating lookup structures
- Trade-off: Faster reads, slower writes
- Types: B-tree, Hash, Full-text, Geospatial
- Index frequently queried columns
- Avoid over-indexing

```
CREATE INDEX idx_user_email ON users(email);  
CREATE INDEX idx_order_date ON orders(order_date);  
CREATE INDEX idx_composite ON orders(customer_id, order_date);
```

Security Principles

Encryption

Encryption transforms data into unreadable format to protect confidentiality.

Encryption Types:

1. Symmetric Encryption:

- Same key for encryption and decryption
- Fast, efficient for large data
- Key distribution challenge
- Algorithms: AES (Advanced Encryption Standard), ChaCha20, Blowfish

Example Use Cases:

- File encryption
- Database encryption
- VPN tunnels

2. Asymmetric Encryption:

- Key pair: public key (encrypt) and private key (decrypt)
- Slower than symmetric
- Solves key distribution problem
- Algorithms: RSA, ECC (Elliptic Curve Cryptography), Diffie-Hellman

Example Use Cases:

- SSL/TLS certificates
- Digital signatures
- Secure key exchange

3. Hashing:

- One-way function (cannot be reversed)
- Fixed-length output regardless of input size
- Used for data integrity verification and password storage
- Algorithms: SHA-256, SHA-3, bcrypt, Argon2

Example Use Cases:

- Password storage (with salt)
- File integrity verification
- Digital signatures
- Blockchain

Encryption at Rest vs. In Transit:

At Rest:

- Data stored on disk, database, or backup
- Full-disk encryption (LUKS, BitLocker)
- Database encryption (TDE - Transparent Data Encryption)
- File-level encryption

In Transit:

- Data moving across networks
- TLS/SSL for web traffic
- IPsec for VPNs
- SSH for remote access

Key Management:

- Store keys separately from encrypted data
- Use Hardware Security Modules (HSM) for sensitive keys
- Regular key rotation
- Key escrow and backup procedures

Practical Examples:

```
# Encrypt file with GPG (symmetric)
gpg --symmetric --cipher-algo AES256 file.txt

# Generate SSH key pair (asymmetric)
ssh-keygen -t rsa -b 4096

# Hash password with bcrypt
htpasswd -B -C 12 .htpasswd username

# Encrypt disk partition with LUKS
cryptsetup luksFormat /dev/sdal
```

Authentication

Authentication verifies identity of users or systems.

Authentication Factors:

1. **Something You Know:** Passwords, PINs, security questions
2. **Something You Have:** Security tokens, smart cards, mobile devices
3. **Something You Are:** Biometrics (fingerprint, face, iris)

Multi-Factor Authentication (MFA):

- Combines two or more authentication factors
- Significantly improves security
- Common implementations:
 - SMS/Email codes (less secure)
 - Authenticator apps (TOTP - Time-based One-Time Password)
 - Hardware tokens (YubiKey, FIDO2)
 - Biometric verification

Authentication Protocols:

1. Basic Authentication:

- Username and password in HTTP header (Base64 encoded)
- Not secure without HTTPS
- Simple but limited

2. Token-Based Authentication:

- User receives token after login
- Token included in subsequent requests
- Stateless and scalable

3. OAuth 2.0:

- Authorization framework for third-party access
- Used by Google, Facebook, GitHub login
- Separates authentication from authorization
- Token types: Access tokens, Refresh tokens

4. OpenID Connect:

- Identity layer on top of OAuth 2.0
- Provides user information
- Single Sign-On (SSO) capability

5. SAML (Security Assertion Markup Language):

- XML-based standard for SSO
- Enterprise authentication
- Identity Provider (IdP) and Service Provider (SP)

6. Kerberos:

- Network authentication protocol
- Ticket-based system
- Used in Active Directory
- Mutual authentication

7. LDAP (Lightweight Directory Access Protocol):

- Centralized authentication directory
- Stores user credentials and attributes
- Active Directory uses LDAP

Password Best Practices:

- Minimum length: 12-16 characters

- Complexity requirements (uppercase, lowercase, numbers, symbols)
- Password hashing with salt (bcrypt, Argon2, PBKDF2)
- Never store plaintext passwords
- Implement account lockout after failed attempts
- Password rotation policies
- Check against breached password databases

Session Management:

- Generate secure, random session IDs
- Use HTTPS for session cookies
- Set appropriate cookie flags:
 - `HttpOnly` : Prevents JavaScript access
 - `Secure` : Only sent over HTTPS
 - `SameSite` : CSRF protection
- Implement session timeout
- Regenerate session ID after login

Authorization

Authorization determines what authenticated users can access or do.

Access Control Models:

1. Discretionary Access Control (DAC):

- Resource owners control access
- Unix/Linux file permissions
- Flexible but difficult to manage at scale

2. Mandatory Access Control (MAC):

- System-enforced access based on security labels
- Used in high-security environments
- SELinux, AppArmor

3. Role-Based Access Control (RBAC):

- Permissions assigned to roles
- Users assigned to roles
- Scalable and manageable
- Example: Admin, Editor, Viewer roles

4. Attribute-Based Access Control (ABAC):

- Access based on attributes (user, resource, environment)
- Fine-grained, dynamic policies
- Example: Allow access if user.department == "Finance" AND time < 6PM

Authorization Strategies:

Principle of Least Privilege:

- Grant minimum permissions necessary
- Reduces attack surface
- Regular access reviews

Separation of Duties:

- Critical tasks require multiple people
- Prevents fraud and errors
- No single person has complete control

Implementation:

```
// RBAC example
const roles = {
  admin: ['read', 'write', 'delete', 'manage_users'],
  editor: ['read', 'write'],
  viewer: ['read']
```

```
};

function checkPermission(user, action) {
    return roles[user.role]?.includes(action);
}
```

API Authorization:

- API keys for service identification
- JWT (JSON Web Tokens) for stateless auth
- Scopes to limit permissions
- Rate limiting per client

OWASP Top 10

The Open Web Application Security Project identifies the most critical web application security risks.

OWASP Top 10 (2021):

1. Broken Access Control:

- Users accessing unauthorized functionality or data
- **Prevention:** Implement proper authorization checks, deny by default, log access control failures

2. Cryptographic Failures:

- Sensitive data exposure due to weak encryption
- **Prevention:** Use strong encryption (TLS 1.3, AES-256), encrypt sensitive data at rest and in transit

3. Injection:

- SQL, NoSQL, OS command, LDAP injection
- **Prevention:** Use parameterized queries, input validation, ORM frameworks

```
-- Vulnerable
query = "SELECT * FROM users WHERE id = " + userId;

-- Secure
query = "SELECT * FROM users WHERE id = ?";
preparedStatement.setInt(1, userId);
```

4. Insecure Design:

- Missing or ineffective security controls in design phase
- **Prevention:** Threat modeling, security requirements, secure design patterns

5. Security Misconfiguration:

- Default credentials, unnecessary features, verbose errors
- **Prevention:** Hardening guides, minimal installation, disable unused features, regular patching

6. Vulnerable and Outdated Components:

- Using libraries with known vulnerabilities
- **Prevention:** Inventory components, monitor vulnerabilities (CVE databases), automated scanning, regular updates

7. Identification and Authentication Failures:

- Weak authentication mechanisms
- **Prevention:** MFA, strong password policies, secure session management, rate limiting

8. Software and Data Integrity Failures:

- Untrusted updates, CI/CD pipeline compromise
- **Prevention:** Code signing, integrity checks, review dependencies, secure CI/CD

9. Security Logging and Monitoring Failures:

- Insufficient logging preventing incident detection
- **Prevention:** Log authentication, access control, input validation failures; centralized logging; alerting; incident response plan

10. Server-Side Request Forgery (SSRF):

- Application fetches remote resource without validation
- **Prevention:** Validate and sanitize URLs, deny by default, network segmentation

Security Policies

Information Security Policy:

Defines organization's security objectives and rules.

Key Components:

- Purpose and scope
- Roles and responsibilities
- Acceptable use policy
- Data classification and handling
- Access control requirements
- Incident response procedures
- Compliance requirements

Password Policy:

- Complexity requirements
- Expiration period
- Reuse restrictions
- MFA requirements
- Storage and transmission guidelines

Acceptable Use Policy (AUP):

- Permitted and prohibited uses of systems
- Personal use guidelines
- Monitoring and privacy disclosure
- Consequences of violations

Data Classification Policy:

- **Public:** No harm if disclosed
- **Internal:** Moderate harm if disclosed
- **Confidential:** Significant harm if disclosed
- **Restricted/Sensitive:** Severe harm if disclosed (PII, financial, health)

Incident Response Policy:

- Incident definition and classification
- Reporting procedures
- Response team roles
- Communication protocols
- Post-incident review

Change Management Policy:

- Change request process
- Risk assessment
- Testing requirements
- Rollback procedures
- Documentation standards

Risk Assessment

Systematic process of identifying and evaluating security risks.

Risk Assessment Process:

1. Asset Identification:

- Hardware, software, data, people
- Determine asset value and criticality

2. Threat Identification:

- Natural disasters (floods, earthquakes)
- Human threats (hackers, insiders, terrorists)
- Technical failures (hardware failure, software bugs)

3. Vulnerability Assessment:

- Security weaknesses in systems
- Configuration errors
- Missing patches
- Weak passwords

4. Risk Analysis:

$$\text{Risk} = \text{Likelihood} \times \text{Impact}$$

Likelihood: Probability of threat exploiting vulnerability

Impact: Potential damage if risk materializes

Risk Matrix:

	Low	Medium	High
High	Med	High	Critical
Medium	Low	Med	High
Low	Low	Low	Med

5. Risk Treatment:

- **Avoid:** Eliminate activity causing risk
- **Mitigate:** Implement controls to reduce risk
- **Transfer:** Insurance, outsourcing
- **Accept:** Accept residual risk if low

6. Risk Monitoring:

- Continuous reassessment
- Track control effectiveness
- Update as environment changes

Vulnerability Management:

- Regular vulnerability scanning
- Penetration testing
- Patch management
- Security advisories monitoring
- Bug bounty programs

Common Vulnerability Scoring System (CVSS):

- Standardized vulnerability severity rating
- Score 0-10 (Low: 0-3.9, Medium: 4-6.9, High: 7-8.9, Critical: 9-10)
- Based on exploitability and impact

Compliance Standards

Regulatory frameworks ensuring data protection and privacy.

GDPR (General Data Protection Regulation):

- European Union data protection law
- Applies to any organization processing EU residents' data

- **Key Requirements:**
 - Lawful basis for processing
 - Data minimization
 - Right to access, rectification, erasure
 - Data breach notification (72 hours)
 - Privacy by design
 - Data Protection Impact Assessment (DPIA)
 - Penalties: Up to 4% of annual revenue or €20 million

HIPAA (Health Insurance Portability and Accountability Act):

- US healthcare data protection
- Protects Protected Health Information (PHI)
- **Key Requirements:**
 - Administrative safeguards (policies, training)
 - Physical safeguards (facility access, device security)
 - Technical safeguards (access control, encryption, audit logs)
 - Business Associate Agreements (BAA)
 - Breach notification

PCI DSS (Payment Card Industry Data Security Standard):

- Protects cardholder data
- Applies to organizations handling credit cards
- **12 Requirements:**
 1. Install and maintain firewall
 2. Don't use vendor-supplied defaults
 3. Protect stored cardholder data
 4. Encrypt data transmission
 5. Use and update antivirus
 6. Develop secure systems
 7. Restrict data access
 8. Unique IDs for access
 9. Restrict physical access
 10. Track and monitor access
 11. Test security systems
 12. Maintain information security policy

SOC 2 (Service Organization Control 2):

- Auditing standard for service providers
- **Trust Service Criteria:**
 - Security
 - Availability
 - Processing integrity
 - Confidentiality
 - Privacy

ISO 27001:

- International information security standard
- Information Security Management System (ISMS)
- Risk-based approach
- Continuous improvement cycle

Compliance Best Practices:

- Regular audits and assessments
- Document policies and procedures
- Employee training and awareness
- Data inventory and classification
- Vendor due diligence
- Incident response plan
- Regular testing and monitoring

Storage Technologies

Block Storage

Block storage divides data into fixed-size blocks, each with unique address.

Characteristics:

- Low-level storage, similar to physical hard drive
- High performance, low latency
- Requires file system formatting
- Supports all operating systems
- Can be directly attached to single server
- Supports snapshots and cloning

Use Cases:

- Databases (high IOPS requirements)
- Virtual machine disks
- Boot volumes
- Transactional applications
- High-performance computing

Examples:

- AWS EBS (Elastic Block Store)
- Azure Managed Disks
- Google Persistent Disk
- iSCSI volumes
- Fibre Channel LUNs

Advantages:

- Excellent performance
- Low latency
- Flexible partitioning
- Boot support

Disadvantages:

- Not easily shared between servers
- More expensive than object storage
- Limited metadata capabilities

Object Storage

Object storage manages data as objects with metadata and unique identifiers.

Characteristics:

- Flat address space (no hierarchy)
- HTTP/REST API access
- Extensive metadata support
- Highly scalable (petabytes to exabytes)
- Eventually consistent
- No file system required

Object Components:

- **Data:** The actual file content
- **Metadata:** Attributes describing the object
- **Unique ID:** Global identifier for the object

Use Cases:

- Backup and archival
- Media files (images, videos)
- Static website hosting
- Data lakes and analytics
- Content distribution
- Cloud-native applications

Examples:

- AWS S3 (Simple Storage Service)
- Azure Blob Storage
- Google Cloud Storage
- MinIO (self-hosted)
- Ceph Object Storage

Storage Classes:

- **Hot/Standard:** Frequent access, higher cost
- **Cool/Infrequent Access:** Less frequent, lower cost
- **Archive/Glacier:** Rare access, lowest cost, retrieval delay

Advantages:

- Unlimited scalability
- Cost-effective for large data
- Built-in redundancy
- HTTP access from anywhere
- Rich metadata

Disadvantages:

- Higher latency than block storage
- Cannot be mounted as file system
- Not suitable for databases
- Eventually consistent

File Storage

Network-attached file storage with hierarchical structure.

Characteristics:

- Traditional folder/file hierarchy
- Multiple servers can access simultaneously
- Network File System (NFS) or SMB/CIFS protocols
- File-level locking
- POSIX compliance

Use Cases:

- Shared file systems
- Home directories
- Content management
- Web serving
- Development environments
- Container persistent storage

Examples:

- AWS EFS (Elastic File System)
- Azure Files
- Google Filestore
- NetApp
- Traditional file servers

Advantages:

- Familiar file/folder structure
- Easy sharing between servers
- Standard file operations
- Application compatibility

Disadvantages:

- Performance limited by network
- Scaling complexity
- Higher latency than block storage
- More expensive than object storage

NAS (Network Attached Storage)

Dedicated file-level storage device connected to network.

Characteristics:

- File-based storage over network
- Self-contained unit with own OS
- Multiple protocol support (NFS, SMB, FTP)
- Easy to deploy and manage
- Built-in features (snapshots, replication)

Architecture:

- Ethernet connectivity
- Standard network protocols
- Appears as network share

Use Cases:

- File sharing and collaboration
- Centralized storage for small/medium business
- Media streaming
- Backup destination
- Virtual machine storage (limited)

Popular NAS Systems:

- Synology
- QNAP
- FreeNAS/TrueNAS
- NetApp FAS

Advantages:

- Simple setup
- Cost-effective
- Easy management
- Built-in redundancy (RAID)

Disadvantages:

- Network bandwidth bottleneck
- Limited scalability
- Not ideal for high-performance workloads

SAN (Storage Area Network)

High-speed network providing block-level storage access.

Characteristics:

- Block-level storage access
- Dedicated high-speed network
- Multiple servers access shared storage
- Appears as local disk to servers
- Enterprise-grade performance

Protocols:

- **Fibre Channel (FC):** Traditional, high-performance, expensive
- **iSCSI:** IP-based, lower cost, over Ethernet
- **FCoE (Fibre Channel over Ethernet):** Converged networking

Architecture:

- Storage array with multiple disks
- Switches creating dedicated network
- Host Bus Adapters (HBAs) in servers
- Zoning for access control

Use Cases:

- Enterprise databases
- Virtual machine clusters
- High-availability applications
- Consolidate storage for multiple servers

- Disaster recovery

Advantages:

- Excellent performance and low latency
- Centralized management
- Advanced features (snapshots, replication)
- High availability

Disadvantages:

- Complex setup and management
- Expensive infrastructure
- Requires specialized skills

NAS vs. SAN:

Feature	NAS	SAN
Access Type	File-level	Block-level
Protocol	NFS, SMB	FC, iSCSI
Network	Standard Ethernet	Dedicated high-speed
Cost	Lower	Higher
Performance	Good	Excellent
Complexity	Simple	Complex

SSD vs. HDD

HDD (Hard Disk Drive):

Technology:

- Mechanical spinning platters
- Moving read/write heads
- Magnetic storage

Characteristics:

- Lower cost per GB
- Higher capacity (up to 20TB+)
- Slower performance
- Higher latency (10-15ms)
- Susceptible to physical shock
- Moving parts wear out
- Higher power consumption
- Noise and heat

Performance:

- Read/Write: 100-200 MB/s
- IOPS: 100-200
- Latency: 10-15ms

Use Cases:

- Archival storage
- Cold data storage
- Sequential read/write workloads
- Budget-conscious deployments

SSD (Solid State Drive):

Technology:

- NAND flash memory

- No moving parts
- Electronic storage

Types:

- **SATA SSD:** Uses SATA interface, limited by SATA speed (600 MB/s)
- **NVMe SSD:** Uses PCIe interface, much faster (3000-7000 MB/s)
- **M.2:** Form factor, can be SATA or NVMe

Characteristics:

- Higher cost per GB
- Lower capacity (up to 4-8TB common)
- Excellent performance
- Low latency (<1ms)
- Shock resistant
- No moving parts (more reliable)
- Lower power consumption
- Silent operation

Performance:

- SATA SSD Read/Write: 500-600 MB/s
- NVMe SSD Read/Write: 3000-7000+ MB/s
- IOPS: 10,000-1,000,000+
- Latency: <0.1ms

Use Cases:

- Operating system drives
- Databases
- Virtual machines
- High-performance applications
- Gaming
- Real-time analytics

Hybrid Solutions:

- **SSHD**: Combines HDD with small SSD cache
- **Tiered Storage**: Automatically moves hot data to SSD, cold data to HDD
- **Cache Drives**: SSD as cache for HDD array

Storage Performance Metrics:

- **IOPS**: Input/Output Operations Per Second (random access)
- **Throughput**: MB/s (sequential access)
- **Latency**: Time to complete operation
- **Queue Depth**: Number of pending I/O operations

Caching Strategies

Caching stores frequently accessed data in fast storage to reduce latency and backend load.

In-memory Caches

Redis:

- Open-source in-memory data store
- Data structures: strings, hashes, lists, sets, sorted sets
- Persistence options (RDB snapshots, AOF logs)
- Pub/sub messaging
- Lua scripting
- Cluster support
- Typical latency: <1ms

Use Cases:

- Session storage
- Real-time analytics
- Leaderboards

- Rate limiting
- Message queues

Example:

```
# Set value with expiration
SET user:1000 "John Doe" EX 3600

# Get value
GET user:1000

# Store hash
HSET user:1000 name "John Doe" email "john@example.com"

# Increment counter
INCR page:views

# Add to sorted set (leaderboard)
ZADD leaderboard 1000 "player1"
```

Memcached:

- Distributed memory caching system
- Simple key-value store
- Multi-threaded
- No persistence
- Simpler than Redis
- Lower memory overhead

Use Cases:

- Simple caching scenarios
- Database query results
- API responses
- Session storage (with external session replication)

Redis vs. Memcached:

Feature	Redis	Memcached
Data Structures	Rich (lists, sets, hashes)	Simple key-value
Persistence	Yes	No
Replication	Built-in	No
Clustering	Yes	Client-side
Single-threaded	Yes (with I/O threads)	Multi-threaded

CDN (Content Delivery Network)

Geographically distributed network of servers delivering content to users from nearest location.

How CDN Works:

1. User requests content
2. DNS routes to nearest CDN edge server
3. If cached, edge server returns content
4. If not cached, edge server fetches from origin
5. Edge server caches and returns content
6. Subsequent requests served from cache

Benefits:

- Reduced latency (geographic proximity)
- Lower bandwidth costs
- Improved availability
- DDoS protection
- SSL/TLS termination

Content Types:

- Static assets (images, CSS, JavaScript)
- Video streaming
- Software downloads
- API responses (with appropriate caching headers)

Popular CDN Providers:

- Cloudflare
- Amazon CloudFront
- Akamai
- Fastly
- Azure CDN

CDN Configuration:

```
# Cache-Control headers
Cache-Control: public, max-age=31536000    # Cache for 1 year
Cache-Control: private, no-cache             # Don't cache
Cache-Control: public, s-maxage=3600          # CDN cache for 1 hour
```

Cache Invalidation

"There are only two hard things in Computer Science: cache invalidation and naming things."

Invalidation Strategies:

1. Time-based (TTL - Time To Live):

- Cache expires after fixed duration
- Simple to implement
- May serve stale data until expiration

```
cache.set('key', value, ttl=3600) # Expires in 1 hour
```

2. Event-based:

- Invalidate cache when data changes
- Immediate consistency
- Requires event system

```
def update_user(user_id, data):  
    database.update(user_id, data)  
    cache.delete(f'user:{user_id}')
```

3. Versioning:

- Include version in cache key
- Old versions automatically ignored

```
cache_key = f'user:{user_id}:v{version}'
```

4. Cache Tags:

- Tag related cache entries
- Invalidate by tag

```
cache.set('user:1000', data, tags=['users', 'premium'])  
cache.invalidate_tag('premium')
```

Challenges:

- Distributed cache consistency
- Race conditions
- Thundering herd (many requests when cache expires)

- Cascade failures

Cache Stampede Prevention:

```
def get_with_lock(key):  
    value = cache.get(key)  
    if value is None:  
        with lock(f'lock:{key}'):  
            value = cache.get(key)    # Double-check  
            if value is None:  
                value = database.get(key)  
                cache.set(key, value)  
    return value
```

Write-through vs. Write-back Cache

Write-through Cache:

- Write to cache and database simultaneously
- Data always consistent
- Higher write latency
- Data never lost

```
def write_through(key, value):  
    database.write(key, value)    # Write to database first  
    cache.set(key, value)        # Then update cache  
    return value
```

Advantages:

- Strong consistency
- Simple to implement
- No data loss

Disadvantages:

- Slower writes
- Database always involved

Write-back (Write-behind) Cache:

- Write to cache immediately
- Asynchronously write to database
- Lower write latency
- Risk of data loss if cache fails

```
def write_back(key, value):
    cache.set(key, value)                      # Write to cache
    queue.add({'key': key, 'value': value})    # Queue for async write
    return value

# Background worker
def flush_worker():
    while True:
        item = queue.get()
        database.write(item['key'], item['value'])
```

Advantages:

- Fast writes
- Reduced database load
- Batch writes possible

Disadvantages:

- Eventually consistent
- Data loss risk
- Complex implementation

Write-around Cache:

- Write directly to database
- Bypass cache
- Cache updated on read
- Used for write-heavy, infrequent reads

Cache Hit Ratio

Measures cache effectiveness.

$$\text{Cache Hit Ratio} = \text{Cache Hits} / (\text{Cache Hits} + \text{Cache Misses})$$

Example:

- 900 cache hits, 100 cache misses
- Hit ratio: $900 / (900 + 100) = 0.90$ or 90%

Targets:

- General applications: 80-90%
- Read-heavy applications: 95%+
- Write-heavy applications: 60-70%

Improving Cache Hit Ratio:

- Increase cache size
- Optimize cache key design
- Implement cache warming (pre-populate)
- Use appropriate TTLs
- Monitor and adjust eviction policies

Eviction Policies:

- **LRU (Least Recently Used)**: Remove least recently accessed items
- **LFU (Least Frequently Used)**: Remove least frequently accessed items
- **FIFO (First In, First Out)**: Remove oldest items
- **Random**: Remove random items

Monitoring:

```
# Redis cache statistics
redis-cli INFO stats
```

```
# Key metrics
keyspace_hits
keyspace_misses
used_memory
evicted_keys
```

Caching Best Practices:

- Cache immutable or slowly changing data
- Set appropriate TTLs
- Monitor cache performance
- Handle cache failures gracefully (fallback to database)
- Use compression for large values
- Implement cache warming for critical data
- Consider cache layers (L1: local, L2: distributed)

Disaster Recovery

Disaster Recovery (DR) ensures business continuity after catastrophic events.

Backup and Restore

Backup Types:

1. Full Backup:

- Complete copy of all data
- Slowest backup, fastest restore
- High storage requirements

```
# Example with tar  
tar -czf /backup/full-$ (date +%Y%m%d).tar.gz /data  
  
# Example with rsync  
rsync -av --delete /data/ /backup/full/
```

2. Incremental Backup:

- Only data changed since last backup
- Fast backup, slower restore
- Low storage requirements
- Requires full backup + all incremental backups to restore

3. Differential Backup:

- Data changed since last full backup
- Moderate backup/restore speed
- Moderate storage requirements
- Requires full backup + last differential to restore

Backup Strategy (3-2-1 Rule):

- 3 copies of data
- 2 different media types
- 1 offsite copy

Backup Methods:

- File-level backup

- Block-level backup
- Database dumps
- Snapshot-based backup
- Continuous data protection (CDP)

Database Backup:

```
# MySQL backup  
mysqldump -u root -p database_name > backup.sql  
  
# PostgreSQL backup  
pg_dump database_name > backup.sql  
  
# MongoDB backup  
mongodump --db database_name --out /backup/
```

Testing Restores:

- Regular restore testing (monthly/quarterly)
- Validate backup integrity
- Document restore procedures
- Measure actual restore times

DR Strategies

Pilot Light:

- Minimal version of environment always running
- Core systems replicated
- Scale up when disaster occurs
- Cost-effective for many scenarios

Components:

- Database replication to DR site
- Essential application servers stopped
- Automated scaling scripts ready

Recovery Process:

1. Detect failure
2. Start application servers
3. Update DNS/routing
4. Scale up resources
5. Validate functionality

RTO:

Hours to days

RPO:

Minutes to hours

Cost:

Low to moderate

Warm Standby:

- Scaled-down version of environment running
- All systems operational but reduced capacity
- Faster recovery than pilot light
- Higher ongoing costs

Components:

- All systems running
- Smaller instance sizes

- Active data replication
- Ready to scale up

Recovery Process:

1. Detect failure
2. Scale up resources
3. Update DNS/routing
4. Validate functionality

RTO:

Minutes to hours

RPO:

Minutes

Cost:

Moderate

Multi-site (Hot Site/Active-Active):

- Full production environment in multiple locations
- Active-active or active-passive
- Immediate failover
- Highest cost

Components:

- Multiple fully operational sites
- Load balancing across sites
- Real-time data replication
- Automated failover

Recovery Process:

1. Automatic detection and failover
2. Traffic routed to healthy site
3. No manual intervention

RTO:

Seconds to minutes

RPO:

Near zero (real-time replication)

Cost:

High

Backup and Restore:

- Restore from backups after disaster
- Slowest recovery
- Lowest cost
- Acceptable for non-critical systems

RTO:

Days to weeks

RPO:

Hours to days

Cost:

Very low

RTO and RPO

RTO (Recovery Time Objective):

- Maximum acceptable downtime
- How quickly must systems be restored?
- Influences DR strategy selection

Examples:

- Critical systems: RTO = 1 hour
- Important systems: RTO = 24 hours
- Non-critical systems: RTO = 1 week

RPO (Recovery Point Objective):

- Maximum acceptable data loss
- How much data can be lost?
- Determines backup frequency

Examples:

- Financial transactions: RPO = 0 (no data loss)
- CRM database: RPO = 1 hour
- File server: RPO = 24 hours

Relationship:

Shorter RTO/RPO = Higher Cost

RTO/RPO

Strategy

Cost

Seconds	Active-Active	Very High
Minutes	Warm Standby	High
Hours	Pilot Light	Moderate
Days	Backup/Restore	Low

Determining RTO/RPO:

1. Business impact analysis
2. Calculate cost of downtime
3. Evaluate data criticality
4. Balance cost vs. requirements
5. Document for each system

DR Plan Components:

- System inventory and dependencies
- Recovery procedures (step-by-step)
- Contact information
- Roles and responsibilities
- Communication plan
- Testing schedule

DR Testing:

- **Tabletop Exercise:** Walk through procedures
- **Simulation:** Test without actual failover
- **Full Test:** Complete failover to DR site
- **Frequency:** Annually minimum, quarterly recommended

High Availability vs. Disaster Recovery:

- **HA:** Protects against component failures (99.9%+ uptime)
- **DR:** Protects against site-wide disasters
- Both needed for comprehensive protection

Windows Administration

Windows Internal Tools

Task Manager (taskmgr):

- Real-time performance monitoring
- Process management
- Startup program management
- Service control

Key Tabs:

- **Processes:** Running applications and background processes
- **Performance:** CPU, memory, disk, network utilization
- **Startup:** Programs that start with Windows
- **Services:** Windows services status and control
- **Details:** Detailed process information (PID, CPU, memory)

PowerShell Commands:

```
# Get running processes
Get-Process

# Sort by CPU usage
Get-Process | Sort-Object CPU -Descending | Select-Object -First 10

# Stop process
Stop-Process -Name "notepad"
Stop-Process -Id 1234

# Get services
Get-Service
```

```
# Start/Stop service  
Start-Service -Name "ServiceName"  
Stop-Service -Name "ServiceName"
```

Event Viewer (eventvwr):

- Centralized logging system
- Troubleshooting and auditing

Log Categories:

- **Application:** Application events
- **Security:** Audit events (login attempts, file access)
- **System:** Operating system events
- **Setup:** Installation and update events

PowerShell Event Log Commands:

```
# Get recent errors  
Get-EventLog -LogName System -EntryType Error -Newest 10  
  
# Search for specific event  
Get-EventLog -LogName Security | Where-Object {$_.EventID -eq 4624}  
  
# Export events  
Get-EventLog -LogName Application | Export-Csv events.csv
```

Performance Monitor (perfmon):

- Detailed performance metrics collection
- Custom data collector sets
- Real-time and historical analysis

Key Counters:

- Processor: % Processor Time
- Memory: Available MBytes, Pages/sec
- Disk: % Disk Time, Avg. Disk Queue Length
- Network: Bytes Total/sec

Creating Data Collector Set:

1. Open Performance Monitor
2. Data Collector Sets → User Defined
3. Create New → Data Collector Set
4. Select performance counters
5. Set sample interval
6. Schedule collection

Registry Editor (regedit):

- Windows configuration database
- **WARNING:** Incorrect changes can break system

Registry Hives:

- **HKEY_CLASSES_ROOT (HKCR):** File associations
- **HKEY_CURRENT_USER (HKCU):** Current user settings
- **HKEY_LOCAL_MACHINE (HKLM):** System-wide settings
- **HKEY_USERS (HKU):** All user profiles
- **HKEY_CURRENT_CONFIG:** Current hardware profile

PowerShell Registry Commands:

```
# Read registry value
Get-ItemProperty -Path "HKLM:\SOFTWARE\Microsoft\Windows\CurrentVersion

# Set registry value
```

```
Set-ItemProperty -Path "HKCU:\Software\MyApp" -Name "Setting" -Value "V

# Create registry key
New-Item -Path "HKLM:\SOFTWARE\MyApp"

# Export registry key
reg export "HKLM\SOFTWARE\MyApp" backup.reg
```

Computer Management (compmgmt.msc):

Centralized management console containing:

- Task Scheduler
- Event Viewer
- Shared Folders
- Local Users and Groups
- Performance
- Device Manager
- Disk Management
- Services and Applications

Services (services.msc):

- Manage Windows services
- Configure startup type (Automatic, Manual, Disabled)
- Start, stop, restart services
- View dependencies

Common Services:

- Windows Update
- Windows Defender
- DHCP Client
- DNS Client
- Print Spooler
- Remote Desktop Services

Disk Management (diskmgmt.msc):

- Partition management
- Format drives
- Assign drive letters
- Create volumes
- Extend/shrink partitions

System Information (msinfo32):

- Comprehensive system details
- Hardware resources
- Components inventory
- Software environment
- Network information

Resource Monitor (resmon):

- Real-time resource usage
- More detailed than Task Manager
- Process-specific resource consumption

Tabs:

- **Overview:** Summary of all resources
- **CPU:** Process and service details
- **Memory:** Physical memory usage
- **Disk:** Disk activity by process
- **Network:** Network activity and connections

Group Policy Editor (gpedit.msc):

- Configure system policies (Professional/Enterprise editions)
- User and computer policies
- Security settings
- Software installation

Policy Categories:

- Computer Configuration: System-wide settings
- User Configuration: User-specific settings

Windows Management Instrumentation (WMIC):

Command-line interface for system management:

```
# Get system information  
wmic computersystem get model,manufacturer  
  
# List installed software  
wmic product get name,version  
  
# Get disk information  
wmic logicaldisk get name,size,freespace  
  
# Query processes  
wmic process where name="notepad.exe" get ProcessId,CommandLine
```

System Configuration (msconfig):

- Boot options
- Service startup management
- Startup program control
- Diagnostic startup modes

Windows PowerShell ISE:

- Integrated scripting environment
- Script development and debugging
- IntelliSense and syntax highlighting

VSS (Volume Shadow Copy Service)

VSS enables point-in-time snapshots of volumes for backup and recovery.

How VSS Works:

1. **VSS Coordinator:** Manages snapshot creation
2. **VSS Writers:** Application components that ensure data consistency
3. **VSS Providers:** Create and manage shadow copies
4. **VSS Requesters:** Applications that request shadow copies

VSS Writers:

- SQL Server Writer
- Exchange Writer
- Active Directory Writer
- Registry Writer
- System Writer

Shadow Copy Types:

1. Full Shadow Copy:

- Complete copy of volume at point in time
- High storage requirements

2. Differential Shadow Copy:

- Only stores changes since last full
- Lower storage requirements

VSS Commands (vssadmin):

```
# List shadow copies
vssadmin list shadows

# List providers
vssadmin list providers

# List writers
vssadmin list writers

# Create shadow copy
vssadmin create shadow /for=C:

# Delete shadow copies
vssadmin delete shadows /for=C: /oldest

# Resize shadow storage
vssadmin resize shadowstorage /for=C: /on=C: /maxsize=10GB
```

PowerShell VSS Commands:

```
# Get shadow copies
Get-WmiObject Win32_ShadowCopy

# Create shadow copy
(Get-WmiObject -List Win32_ShadowCopy).Create("C:\", "ClientAccessible")

# Delete shadow copy
$shadow = Get-WmiObject Win32_ShadowCopy | Where-Object {$_.ID -eq "{$Sh
$shadow.Delete()
```

Previous Versions (Shadow Copies):

- Right-click file/folder → Properties → Previous Versions
- Restore from shadow copy snapshots

- Recover accidentally deleted or modified files

Configuring Shadow Copies:

1. Right-click drive → Properties
2. Shadow Copies tab
3. Enable shadow copies
4. Configure schedule and storage limits

Best Practices:

- Allocate sufficient storage (10-20% of volume)
- Regular scheduling (multiple times daily)
- Monitor shadow copy storage
- Test restore procedures
- VSS snapshots are not backups (temporary)
- Use with proper backup solution

VSS Use Cases:

- Application-consistent backups
- File recovery (Previous Versions)
- Database backups (SQL Server, Exchange)
- Bare-metal recovery
- Rapid recovery from ransomware

Resource Monitor

Resource Monitor provides detailed real-time performance information.

Accessing Resource Monitor:

- Task Manager → Performance → Open Resource Monitor
- Run: resmon.exe
- Performance Monitor → Open Resource Monitor

Overview Tab:

- Consolidated view of all resources
- Key processes by resource consumption
- Real-time graphs

CPU Tab:

Information Displayed:

- Processes using CPU
- Services associated with processes
- Associated handles
- Associated modules (DLLs)

Key Columns:

- **Image:** Process name
- **PID:** Process ID
- **Description:** Process description
- **Threads:** Number of threads
- **CPU:** Current CPU usage
- **Average CPU:** Average over time period

Useful For:

- Identifying CPU-intensive processes
- Finding which services run under which processes
- Detecting high CPU usage causes
- Monitoring thread count

Memory Tab:

Information Displayed:

- Physical memory usage
- Committed memory
- Hard faults per second
- Memory composition graph

Memory Categories:

- **In Use:** Active memory
- **Modified:** Changed but not written to disk
- **Standby:** Cached but can be freed
- **Free:** Available memory

Key Metrics:

- **Working Set:** Physical memory used by process
- **Shareable:** Memory that can be shared
- **Private:** Memory exclusive to process
- **Commit:** Reserved virtual memory

Useful For:

- Identifying memory leaks
- Understanding memory allocation
- Diagnosing out-of-memory issues
- Monitoring paging activity

Disk Tab:

Information Displayed:

- Disk activity by process
- Read/write operations
- Response time
- Queue length

Key Columns:

- **File:** File being accessed
- **Read (B/sec):** Bytes read per second
- **Write (B/sec):** Bytes written per second
- **I/O Priority:** Process I/O priority
- **Response Time (ms):** Disk response time

Disk Activity Graphs:

- Disk Queue Length
- Disk Usage (KB/sec)
- Active Time percentage

Useful For:

- Identifying disk bottlenecks
- Finding processes causing high I/O
- Monitoring disk queue length
- Diagnosing slow disk performance

Network Tab:

Information Displayed:

- Network activity by process
- Active connections
- Listening ports
- Sent/received bytes

Key Sections:

1. Processes with Network Activity:

- Process name and PID

- Bytes sent/received per second
- Total bytes sent/received

2. Network Activity:

- Local and remote addresses
- Sent/received bytes
- Packet loss percentage

3. TCP Connections:

- Local address and port
- Remote address and port
- Connection state
- Packet loss

4. Listening Ports:

- Address
- Port
- Protocol
- Firewall status

Useful For:

- Identifying network bandwidth usage
- Finding processes with network connections
- Detecting unexpected network activity
- Monitoring connection states
- Identifying listening services

Filtering and Searching:

- Filter by process name in each tab
- Search for specific handles or modules
- Right-click for additional options:
 - End Process
 - Suspend Process

- Search Online (for unknown processes)

Performance Troubleshooting Workflow:

1. High CPU Usage:

- Open Resource Monitor → CPU tab
- Sort by Average CPU
- Identify top processes
- Check if expected or malicious
- Review services and threads

2. Memory Issues:

- Memory tab → Sort by Working Set
- Check for processes with growing memory
- Monitor Hard Faults (high = memory pressure)
- Review committed memory vs. physical

3. Disk Performance:

- Disk tab → Sort by Total (B/sec)
- Check Response Time (>20ms = slow)
- Identify I/O intensive processes
- Monitor Queue Length (>2 = bottleneck)

4. Network Problems:

- Network tab → Processes with Network Activity
- Check for unexpected connections
- Monitor bandwidth usage
- Review listening ports for security

Resource Monitor vs. Task Manager:

Feature	Task Manager	Resource Monitor
CPU Details	Basic	Detailed (threads, handles)
Memory Info	Summary	Detailed breakdown
Disk Activity	Per process	Per file
Network	Basic	Detailed connections
Real-time Graphs	Yes	Yes, more detailed
Filtering	Limited	Extensive

Performance Counter Integration:

Resource Monitor data available as performance counters:

```
# Get performance counters
Get-Counter -Counter "\Processor(_Total)\% Processor Time"
Get-Counter -Counter "\Memory\Available MBytes"
Get-Counter -Counter "\PhysicalDisk(_Total)\Avg. Disk Queue Length"
Get-Counter -Counter "\Network Interface(*)\Bytes Total/sec"
```

Best Practices:

- Use Resource Monitor for detailed troubleshooting
- Task Manager for quick overview
- Performance Monitor for long-term trending
- Combine tools for comprehensive analysis
- Document baseline performance for comparison
- Regular monitoring prevents surprises

Windows Performance Toolkit:

Windows Performance Recorder (WPR):

- Record detailed performance traces
- Capture system-wide activity
- Analyze with Windows Performance Analyzer

Windows Performance Analyzer (WPA):

- Analyze WPR traces
- Deep dive into performance issues
- CPU sampling and context switches
- Disk I/O analysis
- Memory allocation tracking

Common Windows Performance Issues:

1. High CPU:

- Windows Update
- Antivirus scanning
- Windows Search indexing
- Superfetch/SysMain
- Runtime Broker

2. High Memory:

- Memory leaks in applications
- Insufficient physical RAM
- Many browser tabs
- Background processes

3. Disk Bottleneck:

- 100% disk usage
- Windows Search

- Superfetch
- Antivirus
- Insufficient RAM (paging)

4. Network Issues:

- Bandwidth saturation
- Malware activity
- Background updates
- Cloud sync services

Diagnostic Commands:

```
# System file checker
sfc /scannow

# Check disk
chkdsk C: /f /r

# Memory diagnostic
mdsched.exe

# Network diagnostics
Test-NetConnection -ComputerName google.com -Port 443

# Get system uptime
Get-CimInstance Win32_OperatingSystem | Select-Object LastBootUpTime

# Check Windows version
winver

# System information
systeminfo

# Check for corrupt updates
DISM /Online /Cleanup-Image /RestoreHealth
```

Conclusion

This comprehensive guide covers the essential knowledge areas every IT professional should master. From Linux fundamentals and networking principles to database management, security practices, storage technologies, caching strategies, disaster recovery planning, and Windows administration, these topics form the foundation of modern IT infrastructure.

Key Takeaways:

Linux:

Understanding file systems, package management, systemd, permissions, logging, and resource management is crucial for managing Linux servers effectively.

Networking:

TCP/IP, DNS, HTTP/S, VPNs, load balancers, firewalls, and subnetting are the building blocks of network infrastructure and connectivity.

Databases:

Choosing between SQL and NoSQL, understanding ACID properties, implementing proper data modeling, and planning for scalability are critical for data-driven applications.

Security:

Implementing encryption, authentication, authorization, following OWASP guidelines, establishing security policies, conducting risk assessments, and maintaining compliance protect organizations from threats.

Storage:

Understanding the differences between block, object, and file storage, along with NAS, SAN, SSD, and HDD technologies, enables optimal storage architecture decisions.

Caching:

Leveraging in-memory caches, CDNs, implementing proper invalidation strategies, and monitoring cache hit ratios significantly improve application performance.

Disaster Recovery:

Establishing backup strategies, choosing appropriate DR approaches (pilot light, warm standby, multi-site), and defining RTO/RPO ensure business continuity.

Windows:

Mastering internal tools, VSS for snapshots, and Resource Monitor for troubleshooting provides comprehensive Windows administration capabilities.

Continuous Learning:

Technology evolves rapidly. Stay current through:

- Official documentation
- Online courses and certifications
- Community forums and conferences
- Hands-on practice and labs
- Following industry blogs and publications

Practical Application:

Knowledge becomes valuable through application. Build lab environments, contribute to open-source projects, and practice troubleshooting real-world scenarios to solidify your understanding.

The IT landscape demands both breadth and depth. This guide provides the breadth; developing depth in areas most relevant to your role and interests will set you apart as a skilled IT professional.

© 2025 Mahabir Singh Bisht. All rights reserved.

This guide is for educational purposes. Continuous learning is key in IT.