

Oracle and Java are well positioned in
AI services, tools and platforms

Investments in the Java Platform will benefit
Oracle customers for decades to come

Java Platform

- Java Development Kit (JDK)
- Java Runtime Environment (JRE)
- Java Virtual Machine (JVM)
- Java Language, Java APIs
- Java SE Specification
- Related Oracle Java SE Products



Moving Java forward

30 years of innovation

Java continues to be an **important choice** for enterprises and developers

The Java technology **innovation pipeline** has never been richer

Java 23 planned for Sept 17

#1

Language for today's
Technology Trends

#1

Overall Enterprise/IT
Organizational Use

63B

Active JVMs

41B

Cloud-based JVMs

Choosing what to work on

JDK expectations for compatibility, stability, security, and performance are extreme



Evolving use-cases



Improvements in hardware



New programming paradigms



New trends or hardware architectures



Endeavor to understand problems



Look to identify synergies



Avoid attempting to shoehorn solutions



Avoid the temptation to shortcut

THOUGHTFUL EVOLUTION

“Tip and Tail”

Conservation

Compatibility
Don't alienate users

Innovation

Adapting to change
Fixing mistakes





We continuously
evolve Java to meet
your **future app
development** needs

Java has **30 years** of
experience evolving with the
latest tech trends

Data-centric World

Amber

Continuously **improve developer productivity** through evolution of the Java language.

ZGC

Create a **scalable low-latency** garbage collector capable of handling terabyte heaps.

Cloud-powered World

Loom

Massively scale lightweight threads, **making concurrency simple again.**

Leyden

Improve the start-up time and time to peak performance of cloud applications.

AI-driven World

Babylon

Extend the reach of Java including to machine learning models and GPUs

Panama

Safe and efficient interoperation with native libraries and native Java.

Valhalla

Unify primitives and classes to **improve productivity and performance.**



AI Terminology

AI Compute

Training, Machine learning, Model creation

Computationally intense, large data sets


AI Integration

Making Inferences

Connecting inferences to business logic and business data

Connecting AI Compute processes

Java and AI Today



AI Integration

Several new AI projects in Java

LangChain4J, SpringAI, Semantic Kernel Java, etc

REST calls which perform LLM inference

May also wrap a Runtime and run local but inference computation isn't within Java

Very well served already by Java's Strengths

- Strong typing, good abstractions, good core libraries, memory safety, performance, observability, security, cloud support, web servers, extensive third-party libraries/tools, development speed, developer base, stability and predictability

AI Integration

Several new AI projects in Java

LangChain4J, SpringAI, Semantic Kernel Java, etc

REST calls which perform LLM inference

May also wrap a Runtime and run local but inference computation isn't within Java

Very well served already by Java's Strengths

- Strong typing, good abstractions, good core libraries, memory safety, performance, observability, security, cloud support, web servers, extensive third-party libraries/tools, development speed, developer base, stability and predictability

“Elevate your AI Solutions with Java” – LRN2932 – 3:30pm, Marco Polo 805

Oracle Cloud Infrastructure SDK for Java

Enables you to easily write code to manage OCI Resources

Support for OCI AI Related Services:

- AI Anomaly Detection
- AI Language
- AI Speech
- AI Vision
- Generative AI
- Generative Inference



The screenshot shows the Oracle Cloud Infrastructure Documentation page for the SDK for Java. The header is dark blue with the text "Oracle Cloud Infrastructure Documentation" and a "Try Free Tier" button. Below the header is a breadcrumb trail: "Developer Resources > Developer Guide > SDK Guides >". To the right of the breadcrumb is a link "All Pages". The main content area is white and features the text "Updated 2024-08-27" above the title "SDK for Java". Below the title is a paragraph: "The Oracle Cloud Infrastructure SDK for Java enables you to write code to manage Oracle Cloud Infrastructure resources."



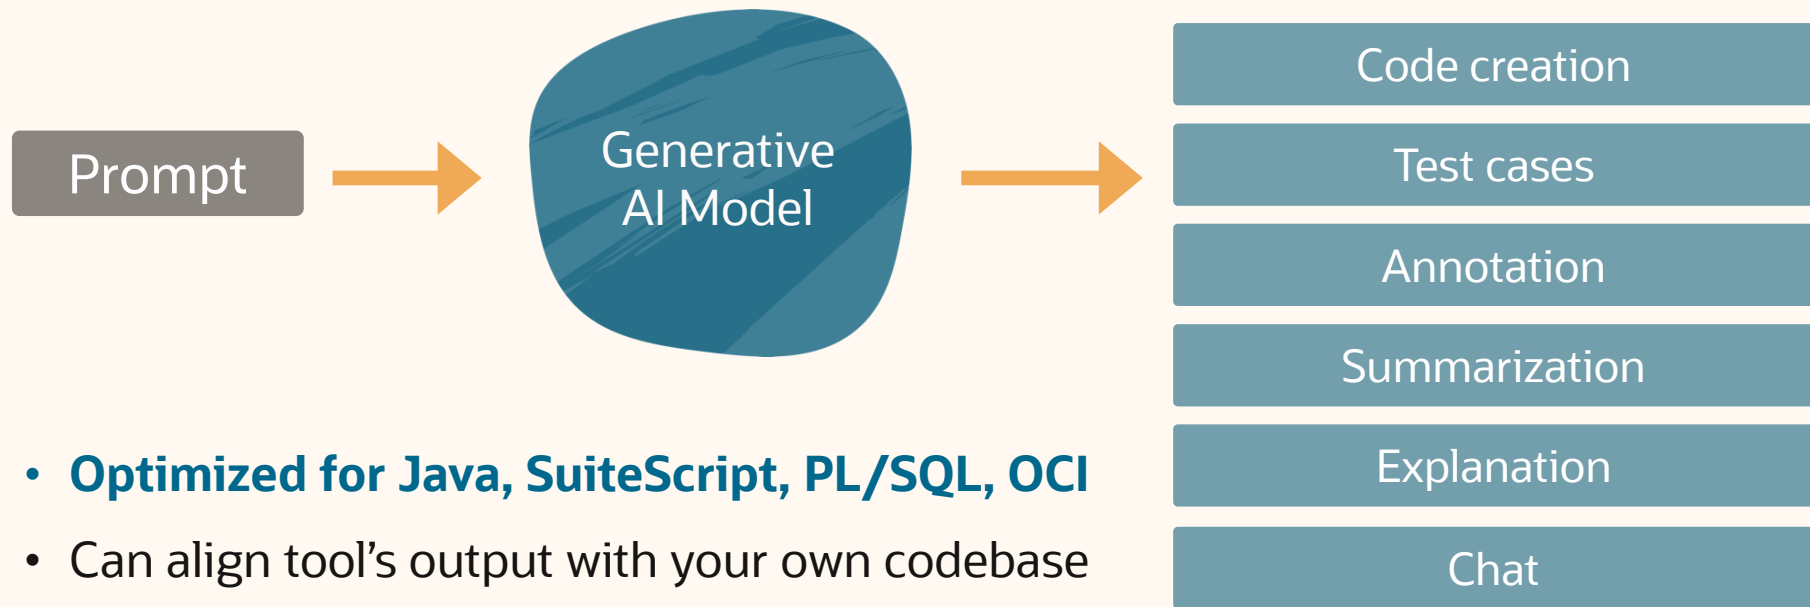
Oracle Code Assist

Build applications and customizations faster with AI

1. Boost developer velocity
2. Enhance code consistency
3. Optimized for Java, SuiteScript, PL/SQL, OCI




Powered by AI models running on OCI



- **Optimized for Java, SuiteScript, PL/SQL, OCI**
- Can align tool's output with your own codebase
- Respect for privacy: customer data is not shared

What is Code Assist and what will it do?

- **Plugin** for JetBrains IntelliJ IDEA and Microsoft Visual Studio Code
- Powered by large language models (LLMs) running on **Oracle Cloud Infrastructure (OCI)**
- Provides **context-specific** suggestions and explanations to help developers stay in flow



What's new with Oracle Code Assist?



- Announcing Oracle Code Assist beta

Now Available!

Optimized for Java initially, but available for C, C++, Go, JavaScript, PL/SQL, Python, Ruby, Rust

- Announcing Oracle Code Assist offering for NetSuite customers

Will help facilitate development when building customizations using SuiteScript

To be available in 2025



AI Compute

Mostly native libraries for GPU, FPGA

CUDA, C++, etc.

AI Compute (currently?) tries to spend as little time as possible in platform languages

Runtimes like from ONNX (Open NN eXchange) provide flexible interface into platforms

Machine learning Java libraries like Tribuo can integrate with ONNX

Platforms must provide

Easy connection to native libraries, functions and data

Access to business logic, data and services

Provide more direct native compute innovations

Java is innovating on both fronts



We continuously
evolve Java to meet
your **future app
development** needs

Java has **30 years** of
experience evolving with the
latest tech trends

Data-centric world

Amber

Continuously improve developer productivity through evolution of the Java language.

ZGC

Create a scalable low-latency garbage collector capable of handling terabyte heaps.

Cloud-powered World

Loom

Massively scale lightweight threads, making concurrency simple again.

Leyden

Improve the start-up time and time to peak performance of cloud applications.

AI-driven World

Babylon

Extend the reach of Java including to machine learning models and GraphQL.

Panama

Safe and efficient interoperability with native libraries and native Java.

Valhalla

Unify primitives and classes to improve productivity and performance.



We continuously
evolve Java to meet
your **future app
development** needs

Java has **30 years** of
experience evolving with the
latest tech trends

Data-centric World

Amber

Continuously **improve developer
productivity** through evolution of the
Java language.

ZGC

Create a **scalable low-latency**
garbage collector capable of
handling terabyte heaps.

Cloud-powered World

Loom

Massively scale lightweight threads,
making concurrency simple again.

Leyden

Improve the start-up time and time
to peak performance of cloud
applications.

AI-driven World

Babylon

Extend the reach of Java
including to machine
learning models and
Graphs.

Panama

**Safe and efficient
interoperation** with
native libraries and
native Java.

Valhalla

Unify primitives and
classes to **improve
productivity and
performance.**



We continuously
evolve Java to meet
your **future app
development** needs

Java has **30 years** of
experience evolving with the
latest tech trends

Data-centric World

Amber

Continuously **improve developer
productivity** through evolution of the
Java language.

ZGC

Create a **scalable low-latency**
garbage collector capable of
handling terabyte heaps.

Cloud-powered World

Loom

Massively scale lightweight threads,
making concurrency simple again.

Leyden

Improve the start-up time and time
to peak performance of cloud
applications.

AI-driven World

Babylon

Extend the reach of Java
including to machine
learning models and
GPUs

Panama

**Safe and efficient
interoperation** with
native libraries and
native Java.

Valhalla

Unify primitives and
classes to **improve
productivity and
performance.**

Project Valhalla – Exec Summary

Biggest refactor of Java language and runtime ever undertaken – 10 years (!) of R&D

Key Motivators

- Flatter and denser memory layouts

- Unify the type system rift between primitives and objects

- Grow the language to allow new numeric types



Project Valhalla – Exec Summary

Biggest refactor of Java language and runtime ever undertaken – 10 years (!) of R&D

Key Motivators

- Flatter and denser memory layouts


- Unify the type system rift between primitives and objects

- Grow the language to allow new numeric types

I contend the benefits to AI Compute here are self evident....

Project Valhalla – Flatter and denser memory layout

For an array of XY point objects, we get this layout:



Over the past 30 years, relative cost of memory fetches vs arithmetic operations has increased by 100-1000x, pointer-heavy layouts beg for expensive cache misses.

Project Valhalla – Flatter and denser memory layout

We would often prefer this layout...

header
x
y
x
y
x
y
x
y



Project Valhalla – Unify the type system

Rift between primitives and objects is “Java’s original sin”

Necessary in 1995 for numeric performance, but dogged us ever since Integer vs int

Puts developers to a bad choice – clear, clean, abstract, safe OO code, or...

Error-prone, unmaintainable primitives to get performance

Valhalla offers the best of both worlds – “Codes like a class, works like an int”

Challenging problem to address:

Classes are user-defined, primitives built in

Polymorphism, Identity, nullability, initialization, layout, atomicity

Project Valhalla – Growing the language

How would we add new numeric types such as half-floats, complex, fixed-point decimals, or unsigned integers, so they feel “built in”?

Wrong way -> add more primitives, bytecodes, conversion rules and API surface

Right way -> as libraries (!)

Focus on making Java extensible via libraries allowing classes to opt into primitive-like semantics, layout, performance, convertibility, operators



Project Valhalla – Examples

```
// Records delivered in JDK 16 – all members final, auto constructors  
// Provided automatically: accessors, toString, equals, hashCode  
record Rectangle(float length, float width) { }  
record Point(int x, int y) { }
```

```
// Records currently in EA release at jdk.java.net  
value record Rectangle(float length, float width) { }  
value record Point(int x, int y) { }  
// ‘value’ indicates this type has no need for identity features  
// This gives the JVM freedom to encode at runtime in ways to optimize memory,  
// locality and GC efficiency.
```

Project Valhalla – What's the potential?

Earlier prototypes (2018-2022) offer great performance results

Benchmarks comparing operations on arrays of objects vs arrays of equivalent value objects:

- 500x500 Complex matrix multiplication: throughput 6x, 1000x less allocation

- Summing array of int pairs 3.5x – 12x depending on factors such as:

 - Mutability, identity, loops vs streams, etc

Gains come from increased locality, reduced allocation, scalarization

Not to mention the benefits of being able to easily create and use new numeric types!

Call To Action

jdk.java.net

Production and Early-Access OpenJDK Builds, from Oracle

Ready for use: JDK 22, JavaFX 22, JMC 9

Early access: JDK 24, JDK 23, JavaFX 24, JavaFX 23, Jextract, Leyden, Loom, & Valhalla



Valhalla - Where Are We? #JVMLS





We continuously
evolve Java to meet
your **future app
development** needs

Java has **30 years** of
experience evolving with the
latest tech trends

Data-centric World

Amber

Continuously **improve developer
productivity** through evolution of the
Java language.

ZGC

Create a **scalable low-latency**
garbage collector capable of
handling terabyte heaps.

Cloud-powered World

Loom

Massively scale lightweight threads,
making concurrency simple again.

Leyden

Improve the start-up time and time
to peak performance of cloud
applications.

AI-driven World

Babylon

Extend the reach of Java
including to machine
learning models and
GPUs

Panama

Safe and efficient
interoperation with
native libraries and
native Java.

Valhalla

Unify primitives and
classes to **improve
productivity and
performance.**

Interconnecting the JVM and native code

Native interop used to be frowned upon – “Pure Java” was the goal

Legacy “Java Native Interface” (JNI) has challenges:

- Native-first programming model, brittle, expensive to maintain

- Passing data is cumbersome and inefficient

While there are many great Java libraries, there are also many important native libraries

- Off-CPU computing (Cuda, OpenCL)

- Machine Learning (Blis, ONNX, Tensorflow)

- Graphics processing (OpenGL, Vulkan, DirectX)

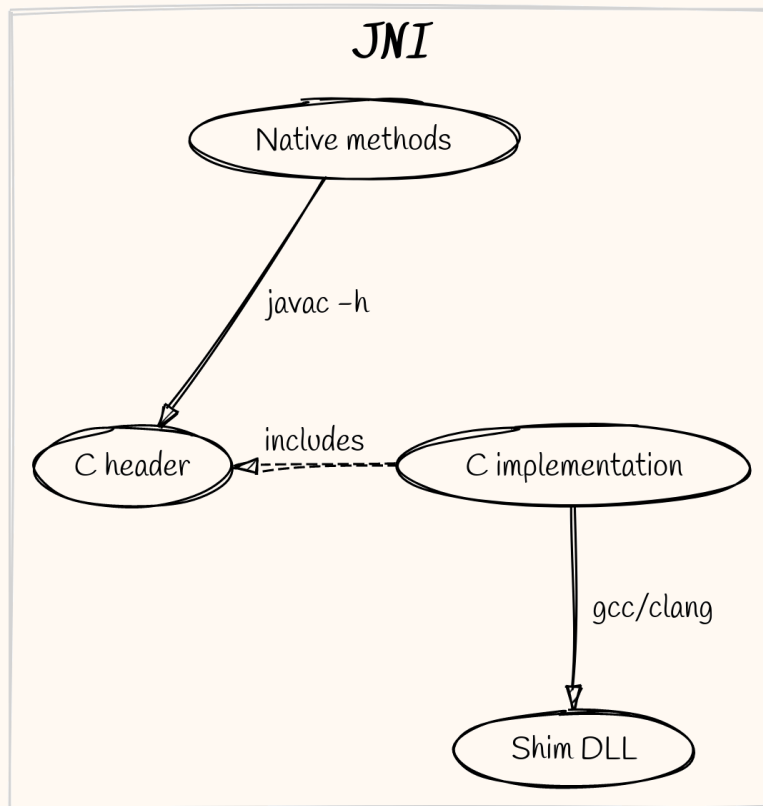

- Others (CRIU, fuse, OpenSSL)

Enter Project Panama


JEP 454: Foreign Function & Memory API

<i>Owner</i>	Maurizio Cimadamore
<i>Type</i>	Feature
<i>Scope</i>	SE
<i>Status</i>	Closed / Delivered
<i>Release</i>	22
<i>Component</i>	core-libs / java.lang.foreign
<i>Discussion</i>	panama dash dev at openjdk dot org
<i>Relates to</i>	JEP 442: Foreign Function & Memory API (Third Preview) JEP 472: Prepare to Restrict the Use of JNI
<i>Reviewed by</i>	Alex Buckley, Jorn Vernee
<i>Endorsed by</i>	Alan Bateman
<i>Created</i>	2023/06/22 09:36
<i>Updated</i>	2024/01/29 21:28
<i>Issue</i>	8310626

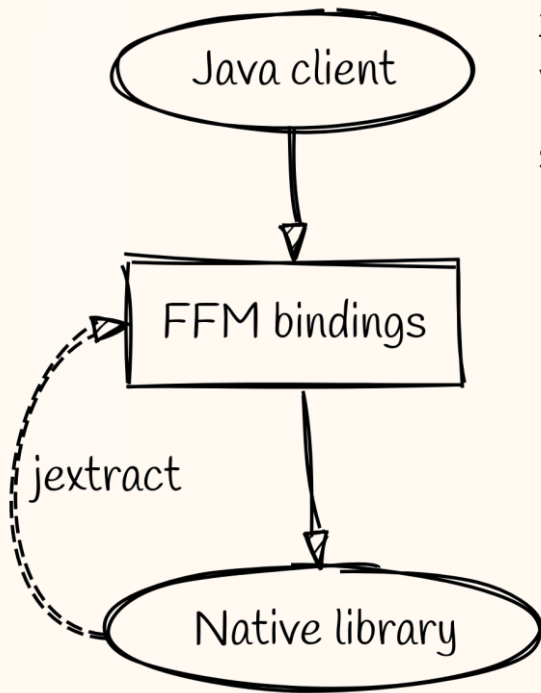
Legacy JNI workflow



FFM API workflow



Enter jextract



```
// stdlib.h
typedef int (*__compar_fn_t) (const void *, const void *);
void qsort (void *__base, size_t __nmemb, size_t __size, __compar_fn_t __compar);
```

```
$ jextract --target-package org.stdlib /usr/include/stdlib.h
```

```
import static org.stdlib.stdlib_h.*;
```


```
...
```

```
try (Arena offHeap = Arena.ofConfined()) {
    MemorySegment array = offHeap.allocateFrom(
        C_INT, 0, 9, 3, 4, 6, 5, 1, 8, 2, 7);

    var compareFunc = __compar_fn_t.allocate((a1, a2) ->
        Integer.compare(a1.get(C_INT, 0), a2.get(C_INT, 0)), offHeap);
    qsort(array, 10L, 4L, compareFunc);
```

```
    int[] sorted = array.toArray(JAVA_INT);
    // [ 0, 1, 2, 3, 4, 5, 6, 7, 8, 9 ]
}
```

Performance



Project Panama - Interconnecting JVM and native code

Accessing memory

Heap Segments (eg, Java array), Native Segments (eg, malloc/mmap)

Size, Lifetime, Confinement (optional)

Arena-based memory management – clients define custom arenas based on need

Global, Automatic, Confined (single-thread), Shared

Strong safety guarantee

Accessing foreign functions

Native linker knows *calling conventions* on the platform where JVM is running

Provides both downcall method handles and upcall stubs to Java

Handles memory layouts for signatures, memory segments, arenas and lifetime

Project Panama – More Information

Search terms:

Panama JVMLS | YouTube Panama Java

Sites:

openjdk.org/projects/panama





We continuously
evolve Java to meet
your **future app
development** needs

Java has **30 years** of
experience evolving with the
latest tech trends

Data-centric World

Amber

Continuously **improve developer
productivity** through evolution of the
Java language.

ZGC

Create a **scalable low-latency**
garbage collector capable of
handling terabyte heaps.

Cloud-powered World

Loom

Massively scale lightweight threads,
making concurrency simple again.

Leyden

Improve the start-up time and time
to peak performance of cloud
applications.

AI-driven World

Babylon

Extend the reach of Java
including to machine
learning models and
GPUs

Panama

**Safe and efficient
interoperation** with
native libraries and
native Java.

Valhalla

Unify primitives and
classes to **improve
productivity and
performance.**

Project Babylon – Exec Summary

“Babylon’s primary goal is to extend the reach of Java to foreign programming models such as SQL, differentiable programming, ***machine learning models, and GPUs***. Babylon will achieve this with an enhancement to reflective programming in Java, called code reflection.”

Developers should not have to:

- Embed snippets of non-Java code

- Write tedious data structures to represent their program

- Use non-standard APIs to access or analyze their program

Project Babylon – GPU Code Example

```
// Developers should write ordinary Java code
static double f(double x, double y) {
    return x * (-Math.sin(x * y) + y) * 4.0d;
}

// Java Developers should not have to write this code, the compiler should do that
var fModel = func("f", methodType(double.class, double.class, double.class))
    .body(entry -> {
        var x = entry.parameters().get(0);
        var y = entry.parameters().get(1);
        var r = entry.op(mul(
            entry.op(mul(
                x,
                entry.op(add(
                    entry.op(neg(
                        entry.op(call(MATH_SIN,
                            entry.op(mul(
                                x,
                                y)))))),
                    y))))), entry.op(constant(DOUBLE, 4.0))));
        entry.op(_return(r));
    });
```


Project Babylon – GPU Code Example

Java Method to differentiate:


```
@CodeReflection
static double f(double x, double y) {
    return x * (-Math.sin(x * y) + y) * 4.0d;
}
```

Serialized Code Model:

```
func @"f" (%0 : double, %1 : double)double -> {
    %2 : Var<double> = var %0 @"x";
    %3 : Var<double> = var %1 @"y";
    %4 : double = var.load %2;
    %5 : double = var.load %2;
    %6 : double = var.load %3;
    %7 : double = mul %5 %6;
    %8 : double = call %7 @"java.lang.Math::sin(double)double";
    %9 : double = neg %8;
    %10 : double = var.load %3;
    %11 : double = add %9 %10;
    %12 : double = mul %4 %11;
    %13 : double = constant @"4.0";
    %14 : double = mul %12 %13;
    return %14;
}
```



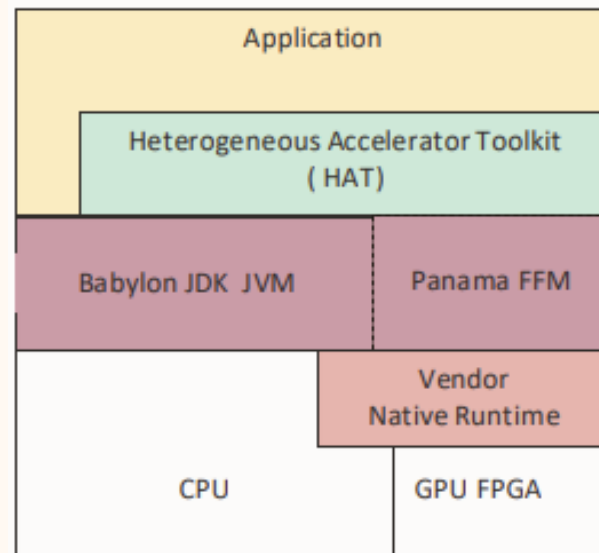
Project Babylon – Spectrum of possibilities



Heterogenous Accelerator Toolkit (HAT)

Leverage programming model from Babylon

Leverage FFM API from Panama





We continuously
evolve Java to meet
your **future app
development** needs

Java has **30 years** of
experience evolving with the
latest tech trends

Data-centric World

Amber

Continuously **improve developer productivity** through evolution of the Java language.

ZGC

Create a **scalable low-latency** garbage collector capable of handling terabyte heaps.

Cloud-powered World

Loom

Massively scale lightweight threads, **making concurrency simple again.**

Leyden

Improve the start-up time and time to peak performance of cloud applications.

AI-driven World

Babylon

Extend the reach of Java including to machine learning models and GPUs

Panama

Safe and efficient interoperation with native libraries and native Java.

Valhalla

Unify primitives and classes to **improve productivity and performance.**





ORACLE CloudWorld

Thank You

Everything you need to know about Artificial
Intelligence and the Java Platform

Please complete attendee survey!