



JAVA GARBAGE COLLECTORS — THE OPTIONS, IN PLAIN ENGLISH

@vvauban 20251021





TL;DR

Java has multiple GCs with different trade-offs:

- **G1** is the *default* general-purpose GC. [inside.java](#)
- **ZGC & Shenandoah** focus on super-low pauses (concurrent, region-based). Generational modes are now mainstream in recent JDKs. [openjdk.org](#)
- **Parallel** maximizes throughput with longer pauses.
- **Serial** is simple/small, OK for tiny apps.
- **Epsilon** does no collection (use for testing/special cases). [openjdk.org](#)
Pick based on **latency vs throughput vs heap size** (and your JDK vendor).



HOW JAVA MEMORY & GC WORKS (QUICK VIEW)

- Objects start in **Young Gen** (Eden → Survivor S0/S1).
Survivors eventually **promote** to **Old Gen**.
- **Minor GC** collects young space; **Major/Old GC** handles old space; some GCs also **compact** to avoid fragmentation.
- **Metaspace** holds class metadata (replaced PermGen since Java 8).



THE COLLECTORS (WHAT, WHEN, WHY)



👉 Serial GC

- **Single-threaded**, stops the world.
- Good for **small heaps**, **single-CPU** or very small container footprints.



Parallel GC (Throughput)

- **Multi-threaded** stop-the-world collector for **max throughput**.
- Accept **longer pauses** to keep overall work high—common for batch jobs.





CMS (Concurrent Mark-Sweep)

- Legacy low-pause GC (removed in modern JDKs).
Prefer **G1** or **low-latency GCs** today.



1 G1 (Garbage-First)

Default GC since Java 9; balances **latency & throughput** on multi-core, large heaps. [inside.java](#)

- Region-based, does **concurrent marking** and **copy/compaction** to limit fragmentation.
- Tuning target: **max pause time goal**.



0 ZGC

- **Ultra-low pause** ($\leq 10\text{ms}$) with **concurrent** mark, relocate, and compact; scales to **very large heaps**.
- **Generational ZGC** is the **default ZGC mode** in recent JDKs (non-generational is deprecated/removed).
Use `-XX:+UseZGC.` openjdk.org
- Minimal tuning: enable it, size the heap, go.





Shenandoah

- **Low-pause** concurrent GC with **concurrent compaction**; great for predictable latencies.
- **Generational Shenandoah** became a **product feature** in JDK 25 (not just experimental). Default mode remains single-gen unless configured. openjdk.org
- Note: **Oracle JDK does not ship Shenandoah**; use vendors like **Red Hat/Temurin/Corretto/Azul** if you need it. wiki.openjdk.org



✖ Epsilon (No-Op)

- **Alloc-only GC** — performs **no reclamation**; JVM exits when heap is exhausted.
- Useful for **performance baselining**, **short-lived apps**, or **GC-free experiments**. It stays experimental on purpose.
openjdk.org





HOW TO CHOOSE (RULE OF THUMB)

- “I want sane defaults” → Start with **G1**. [inside.java](#)
- “I need ultra-low pause times” (APIs, trading, real-timey flows)
→ Try **ZGC**; consider **Shenandoah** if your vendor ships it.
[openjdk.org+1](#)
- “I maximize throughput on batch jobs” → **Parallel GC**.
- “Tiny container or single core” → **Serial GC** can be fine.
- “I’m benchmarking GC impact” → **Epsilon** (carefully).



PRACTICAL TIPS

- **Measure first:** enable GC logs (`-Xlog:gc*`) and track pause P95/P99, allocation rate, survivor promotions, old-gen growth.
- **Start with defaults,** then tune pause targets (G1) or just heap sizing (ZGC/Shenandoah).
- **Match vendor & version** to the GC features you need (e.g., Shenandoah availability).

wiki.openjdk.org



What do you use in prod—G1, ZGC, Shenandoah, or Parallel? Why? Drop your context (heap size, SLA, traffic) in the comments 👇

#Java #JVM #Performance #GarbageCollection #ZGC
#Shenandoah #G1 #ParallelGC #LowLatency #DevOps

