

# Virtual Threads Strengths and Pitfalls



Victor Rentea

👋 I'm **Victor Rentea** 🇷🇴 Java Champion, PhD(CS)

18 years of **Java**, .kt, .js/ts, .scala, .cs, .py ...

10 years of **training** at 150+ companies in 20 countries on:

- Architecture, Refactoring, Unit Testing
- Spring, Hibernate, Performance, Reactive

👥 European Software Crafters **Community** (on Meetup.com)

📺 [YouTube.com/vrentea](https://www.youtube.com/vrentea)

Life += 🧑 + 🧑 + 🧑 + 🐱

in 

@victorrentea  
**victorrentea.ro**

Before Java 21:

## Threads are **blocked** during I/O

keeping a stack of 0.5 MB RAM for the duration of any (long) call

```
public ABC abc(int id) {  
    var a = api.a(id); // REST call  
    var b = api.b(a); // REST call  
    var c = api.c(a, b); // REST call  
    return new ABC(a, b, c);  
}
```

Platform  
Thread

Virtual  
Thread

What's  
wrong  
here?

If using Virtual Threads (Java 21 ★ star feature):

## JVM unmounts the Platform Thread during I/O.

Only a light virtual thread (< 1 KB) blocks in the call.

The OS Platform Thread is used to run other virtual threads.

You keep your code clean and let JVM handle it 🍺.

# Alternative: non-blocking Concurrency

```
// with CompletableFuture since Java 8
public CompletableFuture<ABC> abc(String id) {
    return supplyAsync(() -> api.a(id), executor)
        .thenCompose(a -> api.b(a)
            .thenCompose(b -> api.c(a, b)
                .thenApply(c -> new ABC(a, b, c))));
}
```

```
// with Reactive Programming💀 in Reactor, rxJava,...
public Mono<ABC> abc(int id) {
    return api.a(id)
        .flatMap(a -> api.b(a).flatMap(
            b -> api.c(a, b).map(
                c -> new ABC(a, b, c))));
}
```

Some libs (eg [Jackson](#)) assume threads are reused via a pool and **cache data in ThreadLocals**

No need to reuse VTs via a Thread Pool – they are cheap

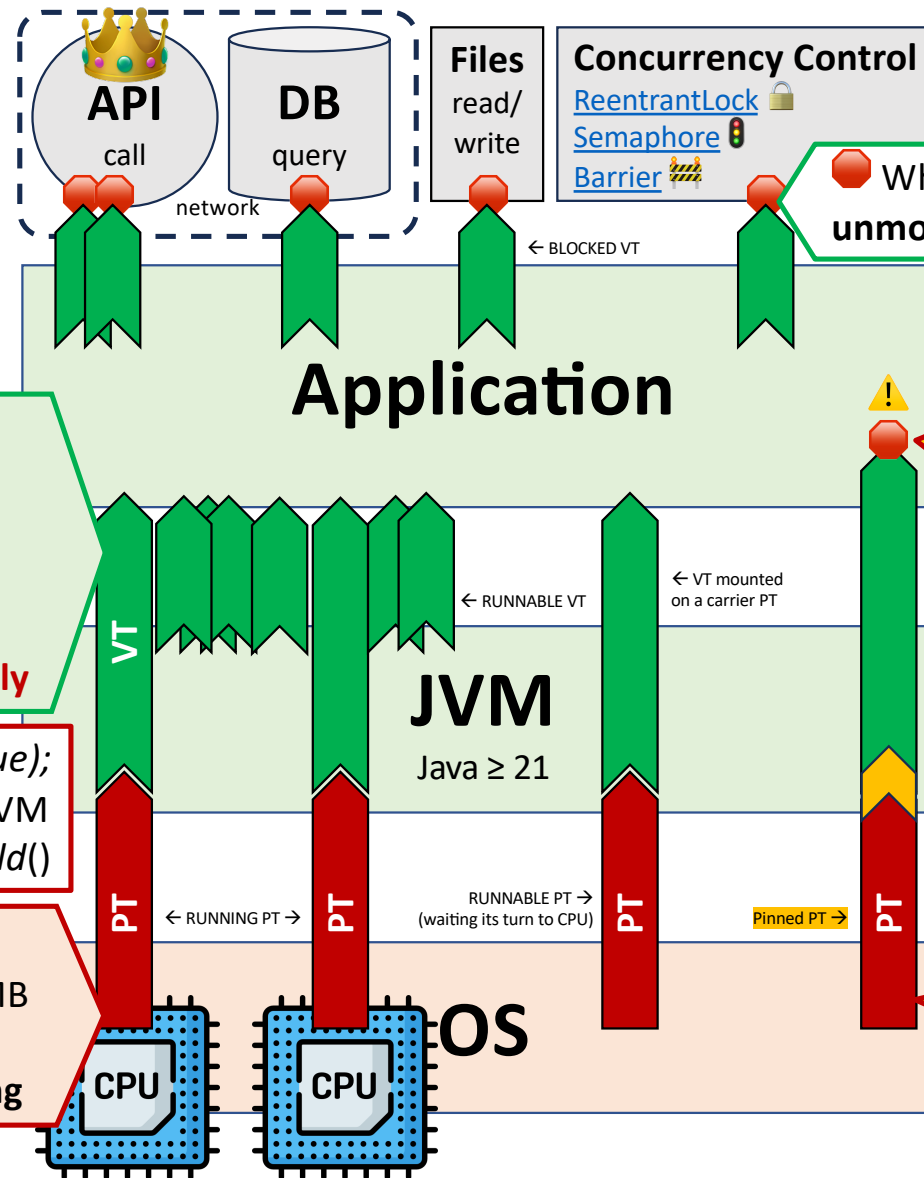
**Virtual Threads (VT)** = cheap

- Java call stack
- **small**, resizable stack: 1-2 KB,
- **many**: 100K .. 1M!
- runs on CPU on a **carrier PT**
- scheduled by JVM **collaboratively**

**CPU Monopolization** eg: `while(true);`  
If a CPU-only task never blocks, JVM cannot interrupt its VT → add `yield()`

**Platform Threads (PT)**

- **heavy**, fixed-size stack: 0.25-1 MB
- **few**: < 1K .. 10K
- scheduled by OS via **time-sharing**



refactor synchronized to

When a VT blocks in Java code, JVM unmounts the PT to run another VT

! **Pinning** of the VT on PT when it blocks in C++ **native** code or **synchronized** block → refactor → trace via: `-Djdk.tracePinnedThreads=full`

! If JVM runs out of PTs, it can:  
→ Let some VT starve → not fair 😞  
→ Use more PT → memory++ 😞  
`-Djdk.virtualThreadScheduler.parallelism=2`  
`-Djdk.virtualThreadScheduler.maxPoolSize=2`

! **Deadlocks**: VT run on an pool of PTs (potentially exhausted?)

# Detecting Thread Pinning using JFR Events

The screenshot displays the JDK Mission Control interface, specifically the Event Browser and Stack Trace views, illustrating thread pinning events.

**Event Browser:**

- Focus:** <No Selection>
- Aspect:** <No Selection>
- Show concurrent:** ☐ Contained ☒ Same threads

**Event Types Tree:**

- pin
- Java Application 303
  - Virtual Thread Pinned 30

**Event Data Table:**

Start Time	Duration	End Time	Event Thread
7/12/24, 12:23:26.412 PM	100.859 ms	7/12/24, 12:23:26.513 PM	
7/12/24, 12:23:26.937 PM	104.329 ms	7/12/24, 12:23:27.041 PM	
7/12/24, 12:23:27.042 PM	105.660 ms	7/12/24, 12:23:27.147 PM	
7/12/24, 12:23:27.148 PM	101.167 ms	7/12/24, 12:23:27.250 PM	
7/12/24, 12:23:27.250 PM	104.826 ms	7/12/24, 12:23:27.355 PM	
7/12/24, 12:23:27.355 PM	101.297 ms	7/12/24, 12:23:27.457 PM	
7/12/24, 12:23:27.457 PM	104.425 ms	7/12/24, 12:23:27.562 PM	
7/12/24, 12:23:28.394 PM	101.126 ms	7/12/24, 12:23:28.495 PM	

**Stack Trace:**

Stack Trace	Samples	Percentage
void java.lang.VirtualThread.parkOnCarrierThread(boolean, long)	1	100
void java.lang.VirtualThread.parkNanos(long)	1	100
void java.lang.VirtualThread.sleepNanos(long)	1	100
void java.lang.Thread.sleep(long)	1	100
void victor.training.java.Util.sleepMillis(int)	1	100
void victor.training.java.virtualthread.experiments.First.synchronizedIsCppCode()	1	100
void victor.training.java.virtualthread.experiments.First.lambda\$main\$0(Map, int, long)	1	100
void victor.training.java.virtualthread.experiments.First\$\$Lambda\$0x000000070011138b8.749282235.run()	1	100
Object java.util.concurrent.Executors\$RunnableAdapter.call()	1	100
void java.util.concurrent.FutureTask.run()	1	100

# Detecting Thread Pinning via Tests

```
@ExtendWith(LoomUnitExtension.class)
public class ExperimentTest {
    @Test
    @ShouldNotPin
    void experiment() throws Exception {
        experiment.execute();
    }
}
```

```
<dependency>
  <groupId>me.escoffier.loom</groupId>
  <artifactId>loom-unit</artifactId>
  <version>0.3.0</version>
  <scope>test</scope>
</dependency>
```

java.lang.AssertionError: The test experiment() was expected to NOT pin the carrier thread, but we collected 30 event(s)



\* Pinning event captured:

```
java.lang.VirtualThread.parkOnCarrierThread(java.lang.VirtualThread.java:675)
java.lang.VirtualThread.parkNanos(java.lang.VirtualThread.java:634)
java.lang.VirtualThread.sleepNanos(java.lang.VirtualThread.java:791)
java.lang.Thread.sleep(java.lang.Thread.java:507)
victor.training.java.Util.sleepMillis(victor.training.java.Util.java:6)
victor.training.java.virtualthread.experiments.Experiment.locks(victor.training.java.virtualthread.experiments.Experiment.java:67)
victor.training.java.virtualthread.experiments.Experiment.lambda$main$0(victor.training.java.virtualthread.experiments.Experiment.java:
```

# Virtual Threads + Resource Pool = Deadlock 💕



```
synchronized (pool) { // pin  
    conn = pool.acquire();  
}  
  
conn.request(); // I/O  
  
pool.release(conn);
```







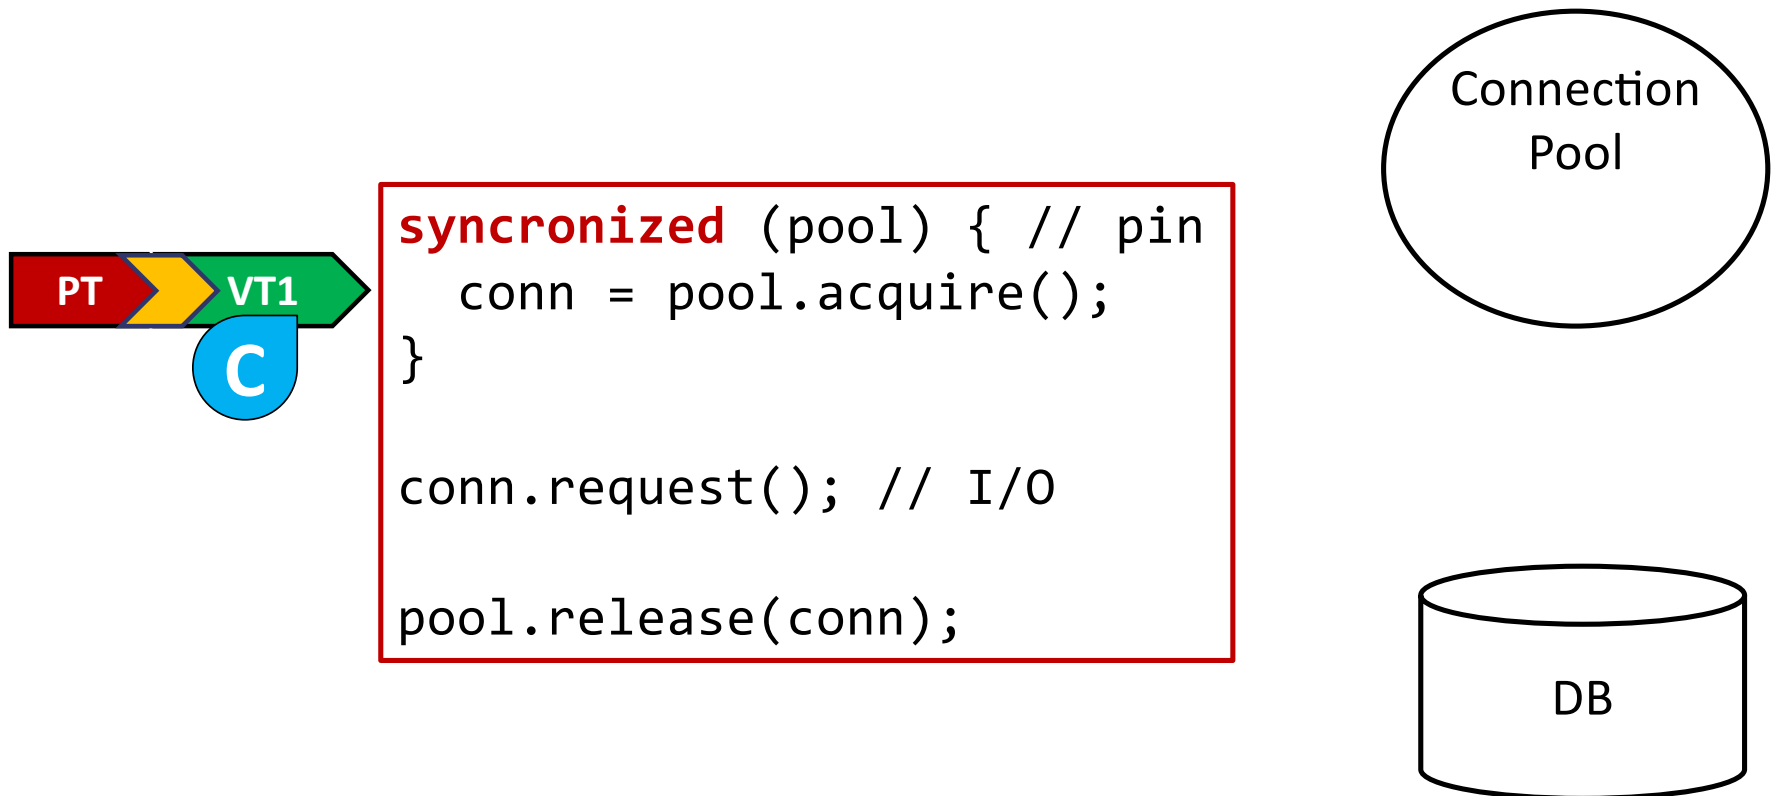
# Virtual Threads + Resource Pool = Deadlock 💕



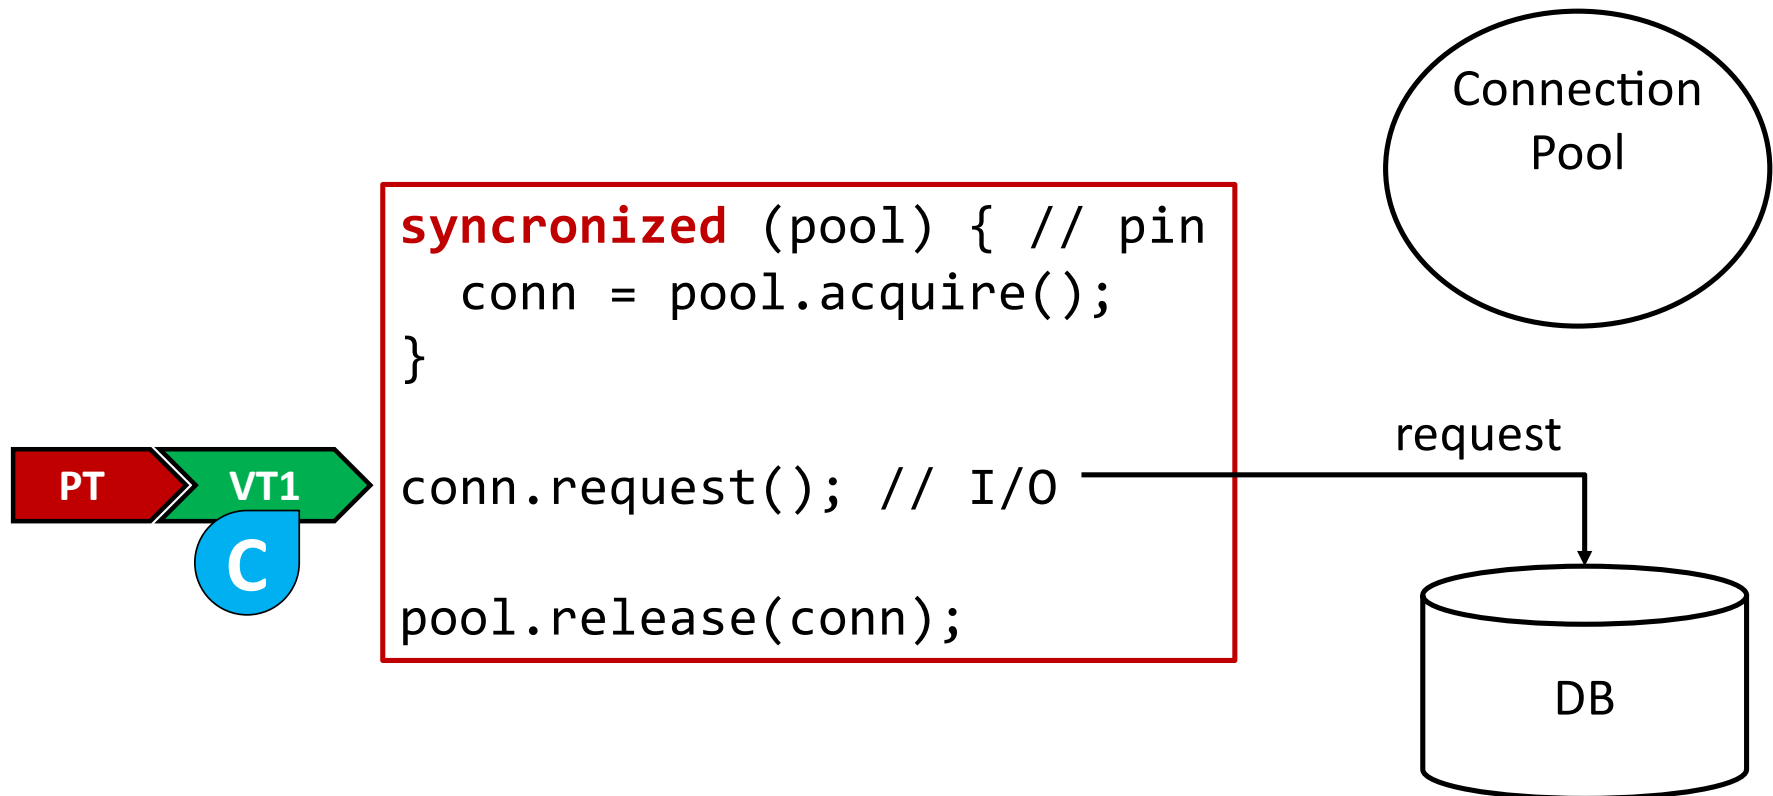
```
synchronized (pool) { // pin  
    conn = pool.acquire();  
}  
  
conn.request(); // I/O  
  
pool.release(conn);
```



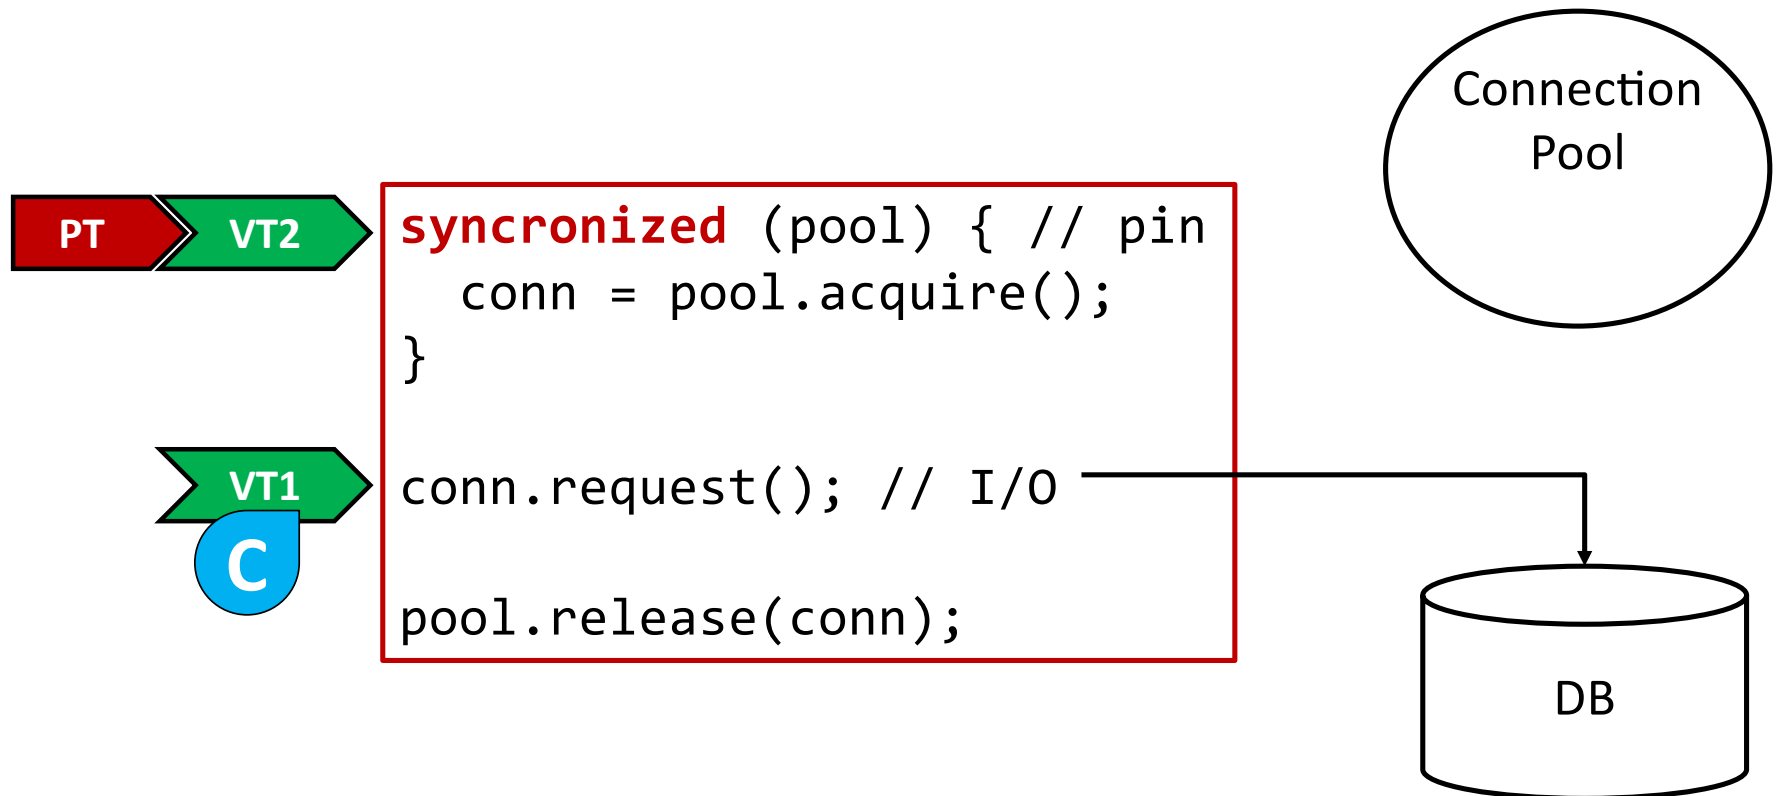
# Virtual Threads + Resource Pool = Deadlock 💕💕



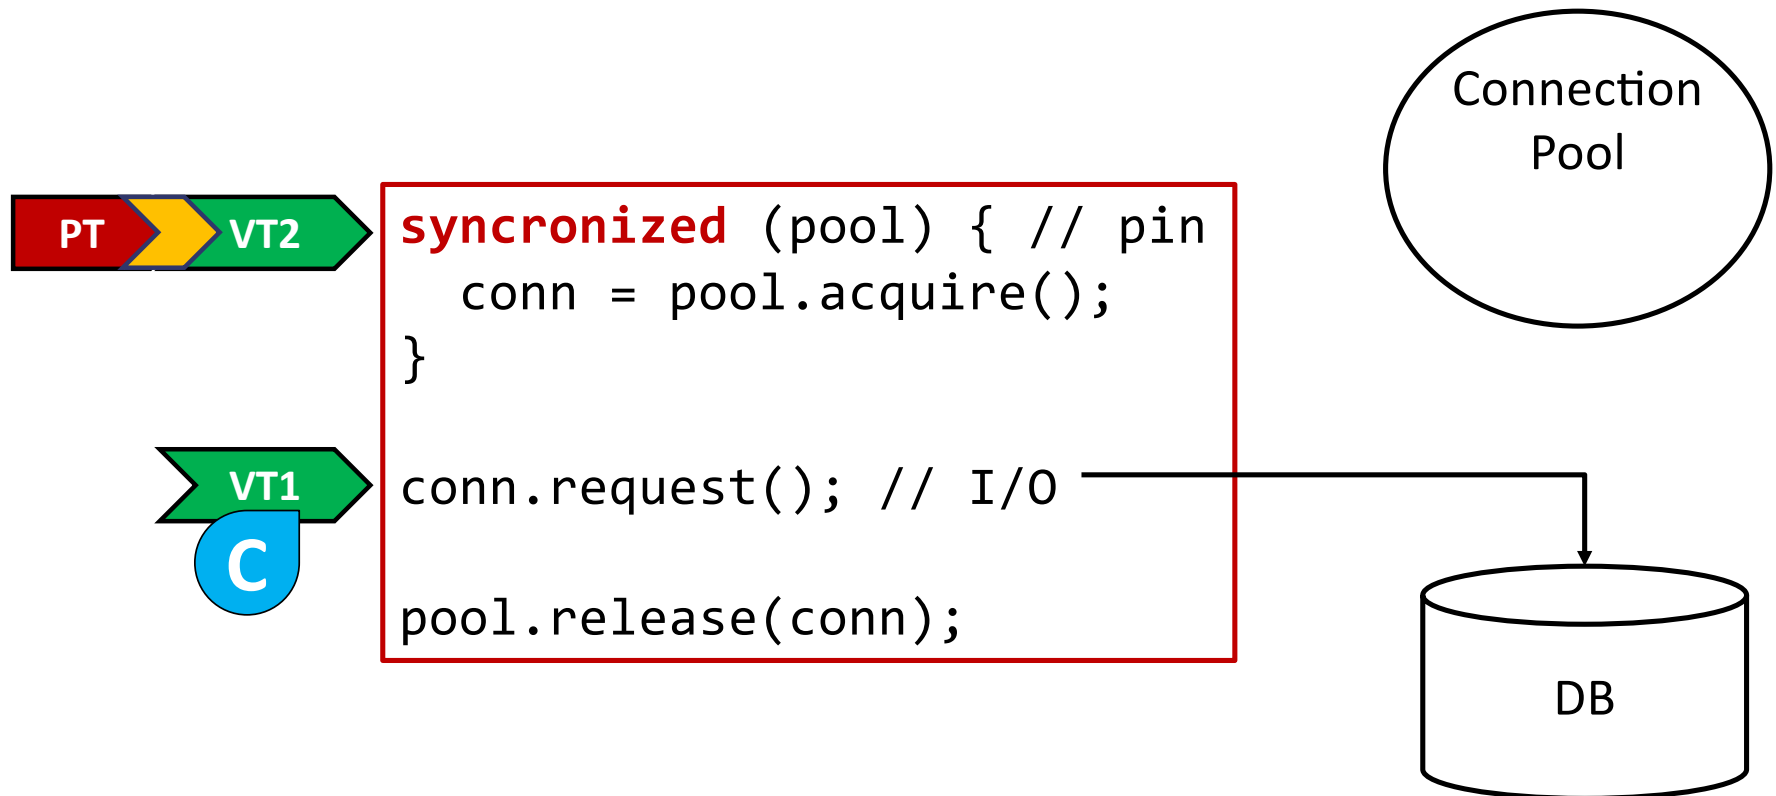
# Virtual Threads + Resource Pool = Deadlock 💕💕



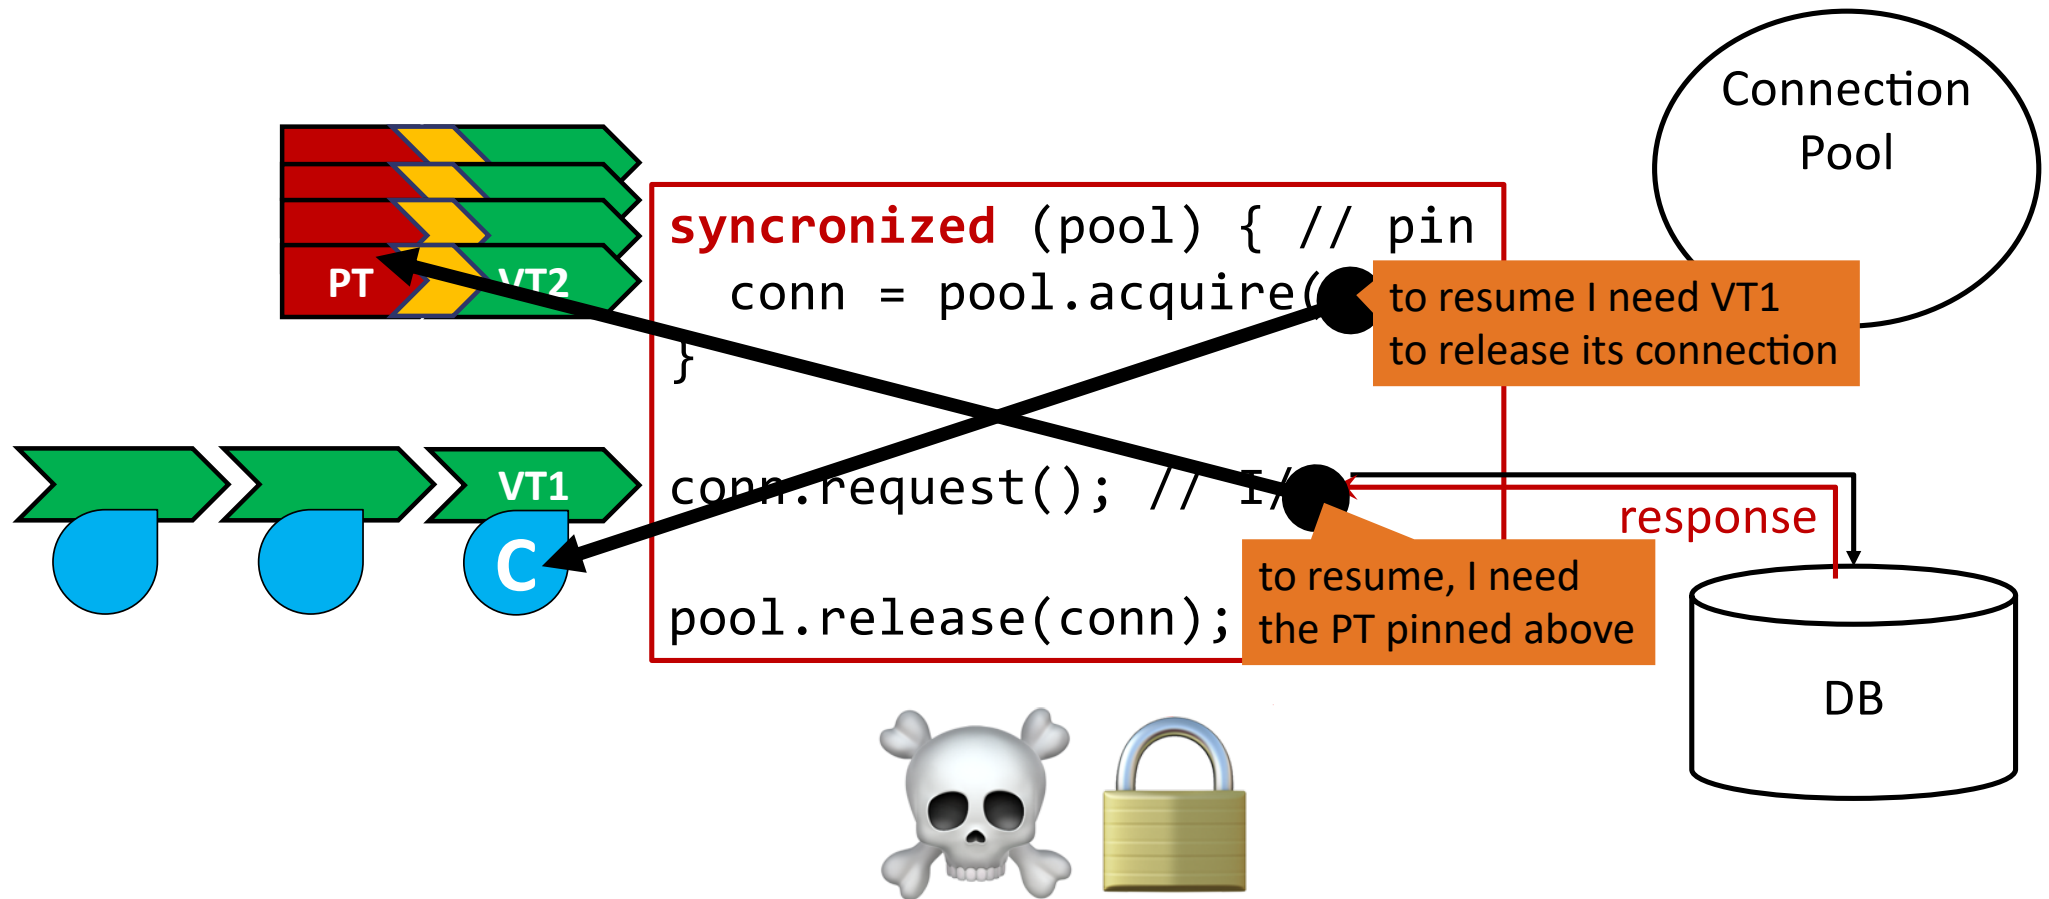
# Virtual Threads + Resource Pool = Deadlock 💕



# Virtual Threads + Resource Pool = Deadlock 💕



# Virtual Threads + Resource Pool = Deadlock 💕



# Virtual Threads + Resource Pool = Deadlock 💕

## *tl;dr – the long story*

- Resource = HTTP/JDBC/Redis... Connection
- Acquire is done via a synchronized block
  - Any incoming VT blocks pinned on the carrier PT
- Using the resource, Java runs an I/O blocking call
  - Example: network read => PT is released
- Last available PT brings a new VT to *acquire* => pin
- No PT is available to resume the VT work after network read

# Old Habbits Die Hard

An old library  
(*Java's HttpClient*)  
creating a new PT  
for each request



ForkJoinPool-1-worker-1		0 ms	(0%)	999 ms
ForkJoinPool-1-worker-10		0 ms	(0%)	999 ms
ForkJoinPool-1-worker-2		0 ms	(0%)	999 ms
ForkJoinPool-1-worker-3		0 ms	(0%)	999 ms
ForkJoinPool-1-worker-4		0 ms	(0%)	999 ms
ForkJoinPool-1-worker-5		0 ms	(0%)	999 ms
ForkJoinPool-1-worker-6		0 ms	(0%)	999 ms
ForkJoinPool-1-worker-7		0 ms	(0%)	999 ms
ForkJoinPool-1-worker-8		0 ms	(0%)	999 ms
ForkJoinPool-1-worker-9		0 ms	(0%)	999 ms
HttpClient-1-SelectorManager		999 ms	(100%)	999 ms
HttpClient-1-Worker-0		0 ms	(0%)	999 ms
HttpClient-1-Worker-1		0 ms	(0%)	999 ms
HttpClient-1-Worker-10		0 ms	(0%)	999 ms
HttpClient-1-Worker-11		0 ms	(0%)	999 ms
HttpClient-1-Worker-12		0 ms	(0%)	999 ms
HttpClient-1-Worker-13		0 ms	(0%)	999 ms
HttpClient-1-Worker-14		0 ms	(0%)	999 ms
HttpClient-1-Worker-15		0 ms	(0%)	999 ms
HttpClient-1-Worker-16		0 ms	(0%)	999 ms
HttpClient-1-Worker-17		0 ms	(0%)	999 ms
HttpClient-1-Worker-2		0 ms	(0%)	999 ms
HttpClient-1-Worker-3		0 ms	(0%)	999 ms
HttpClient-1-Worker-4		0 ms	(0%)	999 ms
HttpClient-1-Worker-5		0 ms	(0%)	999 ms
HttpClient-1-Worker-6		0 ms	(0%)	999 ms
HttpClient-1-Worker-7		0 ms	(0%)	999 ms
HttpClient-1-Worker-8		0 ms	(0%)	999 ms
HttpClient-1-Worker-9		0 ms	(0%)	999 ms



# Old Habbits Die Hard

**Long I/O in  
synchronized  
block**

```
synchronized void bad() {  
    // a network call  
}
```

# Virtual Threads – When / When NOT

- **To accept more parallel requests** with less instances (3, not 7)
  - Especially when calling other APIs a lot
  - Why: less memory blocked / call, VT context switch >>>> faster than PT
- **Won't help CPU-bound systems** (CPU 90-100%)
  - Add occasional `Thread.yield()` for fairness (avoid PT monopolization)
- **Won't help if current bottleneck is outside Java**
  - eg: increasing load on the single shared Database/MQ might be worse

# Coding with Virtual Threads

- Avoid **synchronized** {io | locks} 🙅 use ReentrantLock
- Reduce heap memory used / request = next bottleneck
- Protect remote systems 🙅 Semaphore (not ThreadPools)
- Reduce the size of ThreadLocal data to keep VTs light
- Mind your libraries
  - @ShouldNotPin
  - -Djdk.tracePinnedThreads=short (or full)
  - 'Virtual Thread Pinned' event in JFR recording in [pre]prod

# Legacy Code / Libraries can:

- ... pin VT to PT in **synchronized** blocks, leading to:
  - **Bottlenecks** starving carrier PTs when I/O in synchronized
  - **Deadlocks** with connection pools (JDBC, Redis..)
- ... create **new Platform Threads** (Java HttpClient)
- ... limit **max connections** (Apache HttpClient, JDBC)

# Dig More:




## ■ Understand Virtual Threads

- [Intro - https://blog.rockthejvm.com/ultimate-guide-to-java-virtual-threads](https://blog.rockthejvm.com/ultimate-guide-to-java-virtual-threads)
- [Virtual Threads design explained by Lead of Project Loom@Oracle - https://youtu.be/EO9oMiL1fFo](https://youtu.be/EO9oMiL1fFo)
- [Virtual Threads vs \(Kotlin\) coroutines by Venkat at jPrime'23 - https://youtu.be/uoTyIFvckXA](https://youtu.be/uoTyIFvckXA)

## ■ Industry War Stories

- [medium.com/@phil\\_3582/java-virtual-threads-some-early-gotchas-to-look-out-for-f65df1bad0db](https://medium.com/@phil_3582/java-virtual-threads-some-early-gotchas-to-look-out-for-f65df1bad0db)
- [blog.ydb.tech/how-we-switched-to-java-21-virtual-threads-and-got-deadlock-in-tpc-c-for-postgresql-cca2fe08d70b](https://blog.ydb.tech/how-we-switched-to-java-21-virtual-threads-and-got-deadlock-in-tpc-c-for-postgresql-cca2fe08d70b)
- [blog.ycrash.io/pitfalls-to-avoid-when-switching-to-virtual-threads/](https://blog.ycrash.io/pitfalls-to-avoid-when-switching-to-virtual-threads/)
- <https://www.infoq.com/news/2024/08/netflix-performance-case-study/>

## ■ Library support for Virtual Threads:

- HikariCP: <https://github.com/brettwooldridge/HikariCP/pull/2055>
- Jackson : <https://github.com/FasterXML/jackson-core/issues/919>
- Spring Boot : <https://spring.io/blog/2022/10/11/embracing-virtual-threads>
- Postgres JDBC : <https://jdbc.postgresql.org/changelogs/2023-03-17-42.6.0-release/>