

♦ What is Apache Kafka?

Apache Kafka is an open-source, **distributed streaming platform** built for **real-time data pipelines** and **stream processing**. Originally developed by LinkedIn in 2011, Kafka is now used to **process trillions of messages per day**.

Real-Life Example:

Let's say you use a food delivery app like Swiggy or Zomato.

Here's how Kafka could help:

1. 🍴 **You place an order** — The order info is sent to Kafka.
2. 📄 Kafka stores it in an “Orders” box (called a *Topic*).
3. 🚲 The delivery system checks that box and gets the order (as a *Consumer*).
4. 👨‍🍳 At the same time, the kitchen system also reads the same order!
5. 💳 The payment service reads it too and starts processing.

One event (your order) is shared with **many services at the same time** — all thanks to Kafka.

What Makes Kafka Special?

- Handles **a lot of data** very fast 🚀

- Can store data safely even if systems crash 💾
 - Sends the same data to **many apps at the same time** ↻
 - Helps systems talk to each other without waiting ⌚
-

🧩 Summary:

Kafka is like a smart **delivery boy for data**.

It takes messages from one app, stores them safely, and delivers them to one or more other apps **quickly and reliably**.

♦ Why is Kafka so Popular?

Kafka is the **de facto standard** for stream processing used by over **80% of Fortune 100 companies**. Here's why:

- **High Throughput** – Millions of messages per second
 - **Low Latency** – As low as 2ms
 - **High Scalability** – Thousands of brokers, petabytes of data
 - **Durable Storage** – Distributed commit log ensures reliability
 - **High Availability** – Fault-tolerant across zones and regions
-

♦ Kafka Architecture: How it Works

Kafka is like a **central hub** where data flows in real-time between different systems. It works with the following parts:

- **Producers** – Apps or systems that **send data** into Kafka (e.g., website activity, logs, sensor data).
- **Topics** – Named "folders" inside Kafka that **store the data**.
- **Consumers** – Apps that **read data** from Kafka topics (e.g., dashboards, databases).
- **Brokers** – Kafka servers that **manage and store the data** and help Producers and Consumers talk to Kafka.
- **ZooKeeper / KRaft** – This helps **coordinate** all the Kafka servers and keep the system stable and in sync.

Kafka separates **compute** (processing data) from **storage**, so it can handle real-time data, store it safely, and scale easily.

◆ **Key Kafka Capabilities**

- **Stream Processing** – Kafka Streams lets you **filter, join, and summarize** data while it's still moving.
- **Pub/Sub Messaging** – Producers publish data, and multiple consumers can read it at the same time.
- **Durability & Reliability** – Data is **stored safely** and can be **replayed** if needed.
- **Real-Time Analytics** – You can connect Kafka to tools like **Druid** or **Elasticsearch** to get live insights.
- **Microservices Communication** – Kafka helps microservices **talk to each other** using fast, event-based messages.

♦ **Kafka Use Cases**

- **Data Pipelines** – Move data between systems in real time
- **Stream Processing** – Real-time filtering, transformation, enrichment
- **Streaming ETL** – Extract, transform, and load data continuously
- **Event-Driven Microservices** – Reliable inter-service messaging
- **Real-Time Analytics** – Insights and decisions on live data

♦ **Who Uses Kafka?**

Major companies like **Uber, LinkedIn, Netflix, British Gas**, and many more rely on Kafka for mission-critical streaming infrastructure.

✅ **Apache Kafka Installation Steps (on Ubuntu/Linux)**

1. Update system and install Java

```
sudo apt update
```

```
sudo apt install openjdk-11-jdk -y
```

2. Verify Java installation

```
java -version
```

3. Download Kafka (example: version 3.7.0, Scala 2.13)

```
wget https://downloads.apache.org/kafka/3.7.0/kafka_2.13-3.7.0.tgz
```

4. Extract Kafka archive

```
tar -xvzf kafka_2.13-3.7.0.tgz
```

```
# 5. Navigate to Kafka directory  
cd kafka_2.13-3.7.0
```

Microservices eCommerce Project using Flask, Kafka & MongoDB

Project Overview

This project demonstrates a beginner-friendly **microservices architecture** that simulates the full lifecycle of an online order — from placing an order to shipping it.

- **Flask** is used to build lightweight, RESTful microservices.
 - **Kafka** facilitates **event-driven communication** between services.
 - **MongoDB** is used to persist data at each stage of the order.
-

Why Microservices?

Microservices split large applications into smaller, loosely coupled services that:

- Are easier to maintain and test.
- Can scale independently.

- Improve fault isolation.

Microservices Overview

Each service has a single responsibility and communicates via **Kafka topics**.

- `order_service`: Accepts new orders.
- `inventory_service`: Verifies inventory.
- `payment_service`: Handles payment.
- `shipping_service`: Ships the order.
- All order data is stored and updated in MongoDB.

Project Structure

`ecommerce-kafka-project/`

```
|— docker-compose.yml      # Kafka, Zookeeper, MongoDB, etc.  
|— requirements.txt        # Dependencies for all services  
|— order_service/app.py    # Handles new order placements  
|— inventory_service/app.py # Verifies inventory after order  
|— payment_service/app.py  # Processes payment after inventory
```

|— shipping_service/app.py # Marks order as shipped

Step 1: Docker Setup (docker-compose.yml)

Introduction

This file sets up Kafka, Zookeeper, and MongoDB using Docker containers for a consistent and reproducible environment.

yml

version: '3.8'

services:

zookeeper:

image: confluentinc/cp-zookeeper:latest

environment:

ZOOKEEPER_CLIENT_PORT: 2181

kafka:

image: confluentinc/cp-kafka:latest

ports:

- "9092:9092"

environment:

KAFKA_BROKER_ID: 1

KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181

KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://localhost:9092

KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1

depends_on:

- zookeeper

mongo:

image: mongo

ports:

- "27017:27017"

volumes:

mongo_data:

Step 2: Install Python Dependencies

Introduction

Install the necessary libraries for Flask, Kafka, and MongoDB to interact within microservices.

 requirements.txt:

nginx

flask

kafka-python

pymongo

Install using:

```
pip install -r requirements.txt
```

Step 3: Order Service

Introduction

This service receives incoming orders through an API, saves them in MongoDB, and publishes the event to Kafka.

 order_service/app.py:

```
python
```

```
from flask import Flask, request, jsonify
```

```
from kafka import KafkaProducer
```

```
from pymongo import MongoClient
```

```
import json
```

```
app = Flask(__name__)
```

```
producer = KafkaProducer(
```

```
    bootstrap_servers='localhost:9092',
```

```
    value_serializer=lambda v: json.dumps(v).encode('utf-8')
```

```
)
```

```
mongo = MongoClient("mongodb://localhost:27017/")
```

```
orders = mongo["ecommerce"]["orders"]
```

```
@app.route('/order', methods=['POST'])
```

```
def create_order():
```

```
    data = request.json
```

```
    data['status'] = 'Order Placed'
```

```
orders.insert_one(data)

producer.send('order-created', data)

return jsonify({'message': 'Order placed successfully'}), 200
```


```
if __name__ == '__main__':

    app.run(port=5000)
```

Step 4: Inventory Service

Introduction

This service listens to the order-created topic. It checks inventory and then sends an event to the inventory-checked topic.

 inventory_service/app.py:

```
python
```

```
from kafka import KafkaConsumer, KafkaProducer
```

```
from pymongo import MongoClient
```

```
import json
```

```
consumer = KafkaConsumer(
```

```
'order-created',  
  
bootstrap_servers='localhost:9092',  
  
group_id='inventory-group',  
  
value_deserializer=lambda x: json.loads(x.decode('utf-8'))  
  
)
```

```
producer = KafkaProducer(  
  
    bootstrap_servers='localhost:9092',  
  
    value_serializer=lambda x: json.dumps(x).encode('utf-8')  
  
)
```

```
mongo = MongoClient("mongodb://localhost:27017/")  
  
orders = mongo["ecommerce"]["orders"]
```

for msg in consumer:

```
    order = msg.value  
  
    print(f'[Inventory] Order received: {order}')
```

order['status'] = 'Inventory Checked'


```
orders.update_one({'order_id': order['order_id']}, {'$set': {'status': 'Inventory  
Checked'}})
```

```
producer.send('inventory-checked', order)
```

Step 5: Payment Service

Introduction

This service listens to inventory-checked, processes the payment, and sends an event to the payment-processed topic.

 payment_service/app.py:

```
python
```

```
from kafka import KafkaConsumer, KafkaProducer
```

```
from pymongo import MongoClient
```

```
import json
```

```
consumer = KafkaConsumer(
```

```
    'inventory-checked',
```

```
    bootstrap_servers='localhost:9092',
```

```
    group_id='payment-group',
```

```
    value_deserializer=lambda x: json.loads(x.decode('utf-8'))
```

)

```
producer = KafkaProducer(  
    bootstrap_servers='localhost:9092',  
    value_serializer=lambda x: json.dumps(x).encode('utf-8')  
)
```

```
mongo = MongoClient("mongodb://localhost:27017/")  
orders = mongo["ecommerce"]["orders"]
```


for msg in consumer:

```
    order = msg.value  
  
    print(f"[Payment] Processing: {order}")  
  
    order['status'] = 'Payment Done'  
  
    orders.update_one({'order_id': order['order_id']}, {'$set': {'status': 'Payment  
Done'}})  
  
    producer.send('payment-processed', order)
```

Step 6: Shipping Service

Introduction

The final service listens to payment-processed and marks the order as shipped in MongoDB.

 shipping_service/app.py:

python

```
from kafka import KafkaConsumer
```

```
from pymongo import MongoClient
```

```
import json
```

```
consumer = KafkaConsumer(
```

```
    'payment-processed',
```

```
    bootstrap_servers='localhost:9092',
```

```
    group_id='shipping-group',
```

```
    value_deserializer=lambda x: json.loads(x.decode('utf-8'))
```

```
)
```

```
mongo = MongoClient("mongodb://localhost:27017/")
```

```
orders = mongo["ecommerce"]["orders"]
```

```
for msg in consumer:
```

```
    order = msg.value
```

```
    print(f"[Shipping] Shipped: {order}")
```

```
    orders.update_one({'order_id': order['order_id']}, {'$set': {'status': 'Shipped'}})
```

Step 7: Running the Project

Introduction

You'll now bring everything together: start services, run microservices, and test the end-to-end flow.

Start Docker Services:

```
docker-compose up -d
```

Run Python Services (Each in a separate terminal):

```
python order_service/app.py
```

```
python inventory_service/app.py
```



```
python payment_service/app.py
```

```
python shipping_service/app.py
```

Step 8: Test the Workflow

Introduction

Simulate a real user placing an order and watch it move through all services.

Place an Order via curl:

```
curl -X POST http://localhost:5000/order \  
  
  -H 'Content-Type: application/json' \  
  
  -d '{"order_id": "101", "user": "Alice", "item": "Laptop"}'
```

Step 9: View Orders in MongoDB

Introduction

Verify all updates at each microservice stage from within the MongoDB container.

```
docker exec -it $(docker ps -qf "name=mongo") mongosh
```

Inside the shell:

```
js
```

use ecommerce

```
db.orders.find().pretty()
```

Final Notes

- This project teaches you how to build a **decoupled, event-driven system**.
- Services **only communicate through Kafka**, ensuring flexibility and scalability.
- MongoDB provides **centralized state** and traceability of the order.

:

DevOps Project: Real-Time Monitoring & Alerting System with Kafka

All setups, configuration files, and introductions are included. You can directly run this with **Docker Compose**, visualize logs with **Kibana**, view metrics in **Grafana**, and get alerts when something goes wrong.

Project Name:

Kafka-Powered Monitoring Pipeline

Goal:

To build a real-time monitoring system that:

- Collects logs and metrics from microservices
 - Streams data using **Kafka**
 - Indexes logs with **Elasticsearch**
 - Visualizes logs (Kibana) and metrics (Grafana)
 - Triggers alerts on issues (like CPU > 80%) via Slack/Email
-

Tools Used:

- **Kafka** (message broker)
- **Zookeeper** (Kafka dependency)
- **Filebeat** (log shipper)
- **Logstash** (log processor)
- **Elasticsearch** (log storage & search)
- **Kibana** (log visualization)
- **Prometheus** (metrics collection)
- **Grafana** (dashboard)

- **Alertmanager** (for alerts)
 - **Docker + Docker Compose** (container orchestration)
-

How the System Works:

◆ **Step 1: Spin Up the Core Services with Docker Compose**

Introduction:

We use Docker Compose to set up all services: Kafka, Zookeeper, Prometheus, Grafana, Elasticsearch, Kibana, Filebeat, and Logstash. This allows a complete monitoring system on your local machine.

yaml

```
# docker-compose.yml
version: '3.7'
```

services:

zookeeper:

image: confluentinc/cp-zookeeper:7.2.1

environment:

ZOOKEEPER_CLIENT_PORT: 2181

ports:

- "2181:2181"

kafka:

image: confluentinc/cp-kafka:7.2.1

depends_on:

- zookeeper

ports:

- "9092:9092"

environment:

KAFKA_BROKER_ID: 1

KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181

KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka:9092

KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1

elasticsearch:

image: docker.elastic.co/elasticsearch/elasticsearch:8.6.2

environment:

- discovery.type=single-node
- xpack.security.enabled=false
- bootstrap.memory_lock=true
- ES_JAVA_OPTS=-Xms512m -Xmx512m

ports:

- "9200:9200"

ulimits:

memlock:

soft: -1

hard: -1

volumes:

- esdata:/usr/share/elasticsearch/data

kibana:

image: docker.elastic.co/kibana/kibana:8.6.2

environment:

- ELASTICSEARCH_HOSTS=http://elasticsearch:9200

ports:

- "5601:5601"

depends_on:

- elasticsearch

prometheus:

image: prom/prometheus

volumes:

- ./prometheus.yml:/etc/prometheus/prometheus.yml

ports:
- "9090:9090"

grafana:

image: grafana/grafana
ports:
- "3000:3000"

logstash:

image: docker.elastic.co/logstash/logstash:8.6.2
volumes:
- ./logstash.conf:/usr/share/logstash/pipeline/logstash.conf
depends_on:
- kafka
- elasticsearch
ports:
- "5044:5044"

filebeat:

image: docker.elastic.co/beats/filebeat:8.6.2
volumes:
- ./filebeat.yml:/usr/share/filebeat/filebeat.yml
- ./logs:/logs
depends_on:
- kafka

volumes:
esdata:

♦ **Step 2: Configure Filebeat to Ship Logs to Kafka**



Introduction:

Filebeat monitors log files (like /logs/app.log) and sends them to Kafka.

yaml

```
# filebeat.yml
filebeat.inputs:
  - type: log
    enabled: true
    paths:
      - /logs/*.log
```

```
output.kafka:
  hosts: ["kafka:9092"]
  topic: "logs"
  codec.json:
    pretty: false
```

♦ Step 3: Configure Logstash to Process Logs from Kafka and Send to Elasticsearch

Introduction:

Logstash reads logs from Kafka and forwards them to Elasticsearch for indexing and Kibana visualization.

conf

```
# logstash.conf
input {
  kafka {
    bootstrap_servers => "kafka:9092"
    topics => ["logs"]
    codec => "json"
  }
}

filter {
```

```
# Add transformations or filters here
}

output {
  elasticsearch {
    hosts => ["http://elasticsearch:9200"]
    index => "app-logs-%{+YYYY.MM.dd}"
  }
  stdout { codec => rubydebug }
}
```

♦ Step 4: Create Prometheus Configuration

Introduction:

Prometheus scrapes metrics from your services (e.g., Node Exporter, apps) and stores time-series data.

yaml

```
# prometheus.yml
global:
  scrape_interval: 15s

scrape_configs:
  - job_name: 'node'
    static_configs:
      - targets: ['localhost:9100']
```

♦ Step 5: Launch Grafana and Add Prometheus Data Source

Introduction:

Grafana connects to Prometheus and visualizes your metrics with dashboards.

- Login at <http://localhost:3000>
 - Default credentials: admin / admin
 - Add Prometheus as a data source (<http://prometheus:9090>)
 - Import dashboards or create your own
-

♦ Step 6: Add Alerting Rules

Introduction:

Alertmanager integrated with Prometheus can send alerts when thresholds are breached.

yaml

```
# alerts/alert.rules
groups:
- name: system-alerts
  rules:
  - alert: HighCPUUsage
    expr: node_cpu_seconds_total > 80
    for: 1m
    labels:
      severity: warning
    annotations:
      summary: "High CPU usage detected"
      description: "CPU usage is over 80% for 1 minute"
```

You can plug this into Prometheus config and use Alertmanager to notify via Slack/Email.

♦ Step 7: Create Kafka Topics (Logs, Metrics)

Introduction:

Kafka needs topics where logs and metrics will be published.

```
docker exec -it kafka kafka-topics --create --topic logs --bootstrap-server  
kafka:9092 --replication-factor 1 --partitions 1
```

Repeat this for other topics if needed.

♦ Step 8: View Logs in Kibana

Introduction:

Kibana provides a UI to view, filter, and visualize logs coming from Elasticsearch.

- Access: <http://localhost:5601>
 - Go to “Stack Management” → “Index Patterns” → Add app-logs-*
-



Apache Kafka Commands

Start Zookeeper Server

```
bin/zookeeper-server-start.sh config/zookeeper.properties
```

Start Kafka Broker Server

```
bin/kafka-server-start.sh config/server.properties
```

Create a Kafka Topic

```
bin/kafka-topics.sh --create --topic test-topic --bootstrap-server localhost:9092  
--partitions 1 --replication-factor 1
```

List Kafka Topics

```
bin/kafka-topics.sh --list --bootstrap-server localhost:9092
```

Start Kafka Producer (send messages)

```
bin/kafka-console-producer.sh --topic test-topic --bootstrap-server localhost:9092
```

Start Kafka Consumer (read messages)

```
bin/kafka-console-consumer.sh --topic test-topic --from-beginning  
--bootstrap-server localhost:9092
```

RabbitMQ and Kafka are message brokers, but they are used for different use-cases.

RabbitMQ vs Kafka: Key Differences

Feature	RabbitMQ	Apache Kafka
Type	Message Queue (traditional broker)	Distributed Event Streaming Platform
Message Pattern	Push-based (broker pushes to consumers)	Pull-based (consumers pull from Kafka)

Ordering	Not guaranteed	Guaranteed per partition
Speed	Slower for high throughput	Designed for high throughput (millions/sec)
Storage	Transient (short-lived messages)	Durable (messages stored for days/weeks)
Use-case	Job queues, async processing	Real-time analytics, event streaming
Replay	Not possible after ACK	Yes, you can replay old messages
Built-in Features	Routing, priority queues, TTL, retries	Partitioning, replication, stream replay
Ease of Use	Easier for beginners	Needs deeper understanding (like partitions, offsets)
Dependencies	No Zookeeper	Needs Zookeeper (or KRaft mode)

🔥 Summary: Which is best?

For	Best Choice
Simple task queues	✓ RabbitMQ
Complex routing	✓ RabbitMQ
Real-time streaming	✓ Kafka
High data volume	✓ Kafka
Replay old messages	✓ Kafka

Beginners



RabbitMQ
