



MongoDB

MongoDB Shell Cheat Sheet



Rahul Choudhary

Promoter | Tech Content Creator | Developer | Influencer

MongoDB Shell Cheat Sheet

To get started, install the **MongoDB Shell** (mongosh).

Basic Commands

These basic help commands are available in the MongoDB Shell.

`mongosh`

Open a connection to your local MongoDB instance. All other commands will be run within this mongosh connection.

`db.help()`

Show help for database methods.

`db.<collection>.help()`

Show help on collection methods. The <collection> can be the name of an existing collection or a non-existing collection.

`db.users.help()`

Shows help on methods related to the users collection.

`show dbs`

Print a list of all databases on the server.

`use <db>`

Switch current database to <db>. The mongo shell variable db is set to the current database.

`show collections`

Print a list of all collections for the current database.

`show users`

Print a list of users for the current database.

`show roles`

Print a list of all roles, both user-defined and built-in, for the current database.

`show profile`

Print the five most recent operations that took 1 millisecond or more on databases with profiling enabled.

`show databases` Print a list of all existing databases available to the current user.

`exit` Exit the mongosh session.

Create Operations

Create or insert operations add new documents to a collection. If the collection does not exist, create operations also create the collection.

<code>db.collection.insertOne()</code>	Inserts a single document into a collection. Add a new document with the name of Chris into the users collection
<code>db.collection.insertMany()</code>	Inserts multiple documents into a collection. Add two new documents with the age of 24 and 38 into the users collection

Read Operations

Read operations retrieve documents from a collection; i.e. query a collection for documents.

<code>db.collection.find()</code>	Selects documents in a collection or view and returns a cursor to the selected documents. Returns all users.
<code>db.collection.find(<filterobjec t>)</code> <code>db.users.find({place: "NYC"})</code>	Find all documents that match the filter object Returns all users with the place NYC.
<code>db.collection.find({<field>:1,< field>:1})</code>	Returns all documents that match the query after you explicitly include several fields by setting the <field> to 1 in the projection document.
<code>db.users.find({status:1,item:1})</code>	Returns matching documents only from state field, item field and, by default, the _id field.

<code>db.collection.find({<field>:1,<field>:0, _id:0})</code>	Returns all documents that match the query and removes the <code>_id</code> field from the results by setting it to 0 in the projection.
<code>db.users.find({status:1,item:1,_id:0})</code>	Returns matching documents only from state field and item field. Does not return the <code>_id</code> field.

Update Operations

Update operations modify existing documents in a collection.

<code>db.collection.updateOne()</code>	Updates a single document within the collection based on the filter. Updates all users from the age of 25 to 32.
<code>db.collection.updateMany()</code>	Updates a single document within the collection based on the filter. Updates all users with an age of 27 with an increase of 3.
<code>db.collection.replaceOne()</code>	Replaces a single document within the collection based on the filter. Replaces the first user with the name Kris with a document that has the name Chris in its name field.

Delete Operations

Delete operations remove documents from a collection.

<code>db.collection.deleteOne()</code>	Removes a single document from a collection. Deletes the first user with the age 37.
<code>db.collection.deleteMany()</code>	Removes all documents that match the filter from a collection. Deletes all users with the age less than 18..

Comparison Query Operators

Use the following inside an filter object to make complex queries

\$eq

```
db.users.find({ system: { $eq: "macOS" } })
```

Matches values that are equal to a specified value.

Finds all users with the operating system macOS.

\$gt

```
db.users.deleteMany({ age: { $gt: 99 } })
```

Matches values that are greater than a specified value.

Deletes all users with an age greater than 99.

\$gte

```
db.users.updateMany({ age": { $gte:21 }, {access: "valid"})
```

Matches values that are greater than or equal to a specified value.

Updates all access to "valid" for all users with an age greater than or equal to 21.

\$in

```
db.users.find( { place: { $in: [ "NYC", "SF"] } })
```

Matches any of the values specified in an array.

Find all users with the place field that is either NYC or SF.

\$lt

```
db.users.deleteMany({ "age": { $lt:18 } })
```

Matches values that are less than a specified value.

Deletes all users with the age less than 18..

\$lte

```
db.users.updateMany({ age: { $lte: 17 }, {access: "invalid"})
```

Matches values that are less than or equal to a specified value.

Updates all access to "invalid" for all users with an age less than or equal to 17.

\$ne

```
db.users.find({ "place": { $ne: 'NYC' } })
```

Matches all values that are not equal to a specified value.

Find all users with the place field set to anything other than NYC.

\$nin

```
db.users.find( { place: { $nin: [ "NYC", "SF" ] } } )
```

Matches none of the values specified in an array.

Find all users with the place field that does not equal NYC or SF.

Field Update Operators

Use the following inside an update object to make complex updates

\$inc

```
db.users.updateOne({ age: 22 }, { $inc: { age: 3 } })
```

Increments the value of the field by the specified amount.

Adds 3 to the age of the first user with the age of 22.

\$min

```
db.scores.insertOne( { _id: 1, highScore: 800, lowScore: 200 } )
```

Only updates the field if the specified value is less than the existing field value.

Creates a scores collection and sets the value of highScore to 800 and lowScore to 200.

```
db.scores.updateOne( { _id: 1 }, { $min: { lowScore: 150 } } )
```

\$min compares 200 (the current value of lowScore) to the specified value of 150. Because 150 is less than 200, \$min will update lowScore to 150.

\$max

```
db.scores.updateOne( { _id: 1 }, { $max: { highScore: 1000 } } )
```

Only updates the field if the specified value is greater than the existing field value.

\$max compares 800 (the current value of highScore) to the specified value of 1000. Because 1000 is more than 800, \$max will update highScore to 1000.

\$rename

```
db.scores.updateOne( { $rename: { 'highScore': 'high' } } )
```

Renames a field.

Renames the field 'highScores' to 'high',

\$set

```
db.users.updateOne({ $set: { name:  
  "valid user" } })
```

Sets the value of a field in a document.

Replaces the value of the name field with the specified value valid user.

\$unset

```
db.users.updateOne({ $unset: { name:  
  "" } })
```

Removes the specified field from a document.

Deletes the specified value valid user from the name field.

Read Modifiers

Add any of the following to the end of any read operation

cursor.sort()

```
db.users.find().sort({ name: 1, age:  
 -1 })
```

Orders the elements of an array during a \$push operation.

Sorts all users by name in alphabetical order and then if any names are the same sort by age in reverse order

cursor.limit()

Specifies the maximum number of documents the cursor will return.

cursor.skip()

Controls where MongoDB begins returning results.

cursor.push()

```
db.users.updateMany({}, { $push: {  
 friends: "Chris" } })
```

Appends a specified value to an array.

Add Chris to the friends array for all users

Aggregation Operations

The Aggregation Framework provides a specific language that can be used to execute a set of aggregation operations (processing & computation) against data held in MongoDB.

`db.collection.aggregate()`

A method that provides access to the aggregation pipeline.

```
db.users.aggregate([
  {$match: { access: "valid" } },
  {$group: { _id: "$cust_id",
    total:{$sum: "$amount" } },
  {$sort: { total: -1 } }}])
```

Selects documents in the users collection with accdb.orders.estimatedDocumentCount({})_id field from the sum of the amount field, and sorts the results by the total field in descending order:

Aggregation Operations

Aggregation pipelines consist of one or more stages that process documents and can return results for groups of documents.

`count`

Counts the number of documents in a collection or a view.

`distinct`

Displays the distinct values found for a specified key in a collection or a view.

`mapReduce`

Run map-reduce aggregation operations over a collection

Aggregation Operations

Single Purpose Aggregation Methods aggregate documents from a single collection.

`db.collection.estimatedDocumentCount()`

Returns an approximate count of the documents in a collection or a view.
Retrieves an approximate count of all the documents in the users collection.

`db.users.estimatedDocumentCount({})`

```
db.collection.count()
```

Returns a count of the number of documents in a collection or a view.

```
db.users.count({})
```

Returns the distinct values for the age field from all documents in the users collection.

```
db.collection.distinct()
```

Returns an array of documents that have distinct values for the specified field.

```
db.users.distinct("age")
```

Returns the distinct values for the age field from all documents in the users collection.

Indexing Commands

Indexes support the efficient execution of queries in MongoDB. Indexes are special data structures that store a small portion of the data set in an easy-to-traverse form.

```
db.collection.createIndex()
```

Builds an index on a collection.

```
db.users.createIndex("account  
creation date")
```

Creates the account creation date index in the users collection.

```
db.collection.dropIndex()
```

Removes a specified index on a collection.

```
db.users.dropIndex("account creation  
date")
```

Removes the account creation date index from the users collection.

```
db.collection.dropIndexes()
```

Removes all indexes but the `_id` (no parameters) or a specified set of indexes on a collection.

```
db.users.dropIndexes()
```

Drop all but the `_id` index from a collection.

```
db.users.dropIndex("account creation  
date", "account termination date")
```

Removes the account creation date index and the account termination date index from the users collection.

<code>db.collection.getIndexes()</code>	Returns an array of documents that describe the existing indexes on a collection.
<code>db.users.getIndexes()</code>	Returns an array of documents that hold index information for the users collection.
<code>db.collection.reIndex()</code>	Rebuilds all existing indexes on a collection
<code>db.users.reIndex()</code>	Drops all indexes on the users collection and recreates them.
<code>db.collection.totalIndexSize()</code>	Reports the total size used by the indexes on a collection. Provides a wrapper around the <code>totalIndexSize</code> field of the <code>collStats</code> output.
<code>db.users.totalIndexSize()</code>	Returns the total size of all indexes for the users collection.

Replication Commands

Replication refers to the process of ensuring that the same data is available on more than one MongoDB Server.

<code>rs.add()</code>	Adds a member to a replica set.
<code>rs.add("mongodbd4.example.net:27017")</code>	Adds a new secondary member, <code>mongodbd4.example.net:27017</code> , with default vote and priority settings to a new replica set
<code>rs.conf()</code>	Returns a document that contains the current replica set configuration.
<code>rs.status()</code>	Returns the replica set status from the point of view of the member where the method is run.

<code>rs.stepDown()</code>	Instructs the primary of the replica set to become a secondary. After the primary steps down, eligible secondaries will hold an election for primary.
<code>rs.remove()</code>	Removes the member described by the hostname parameter from the current replica set.
<code>rs.reconfig()</code>	Reconfigures an existing replica set, overwriting the existing replica set configuration.

Sharding Commands

Sharding is a method for distributing or partitioning data across multiple computers. This is done by partitioning the data by key ranges and distributing the data among two or more database instances.

<code>sh.abortReshardCollection()</code> <code>sh.abortReshardCollection("users")</code>	Ends a resharding operation Aborts a running reshard operation on the users collection.
<code>sh.addShard()</code> <code>sh.addShard("cluster"/mongodb3.example.net:27327")</code>	Adds a shard to a sharded cluster. Adds the cluster replica set and specifies one member of the replica set.
<code>sh.commitReshardCollection()</code> <code>sh.commitReshardCollection("records.users")</code>	Forces a resharding operation to block writes and complete. Forces the resharding operation on the records.users to block writes and complete.

<code>sh.disableBalancing()</code>	Disable balancing on a single collection in a sharded database. Does not affect balancing of other collections in a sharded cluster.
<code>sh.disableBalancing("records.users")</code>	Disables the balancer for the specified sharded collection.
<hr/> <code>sh.enableAutoSplit()</code>	Enables auto-splitting for the sharded cluster.
<hr/> <code>sh.disableAutoSplit()</code>	Disables auto-splitting for the sharded cluster.
<hr/> <code>sh.enableSharding()</code>	Creates a database.
<code>sh.enablingSharding("records")</code>	Creates the records database.
<hr/> <code>sh.help()</code>	Returns help text for the sh methods.
<hr/> <code>sh.moveChunk()</code>	Migrates a chunk in a sharded cluster.
<code>sh.moveChunk("records.users", { zipcode: "10003" }, "shardexample")</code>	Finds the chunk that contains the documents with the zipcode field set to 10003 and then moves that chunk to the shard named shardexample.
<hr/> <code>sh.reshardCollection()</code>	Initiates a resharding operation to change the shard key for a collection, changing the distribution of your data.
<code>sh.reshardCollection("records.users", { order_id: 1 })</code>	Reshards the users collection with the new shard key { order_id: 1 }
<hr/> <code>sh.shardCollection()</code>	Enables sharding for a collection.
<code>sh.shardCollection("records.users", { zipcode: 1 })</code>	Shards the users collection by the zipcode field.

<code>sh.splitAt()</code>	Divides an existing chunk into two chunks using a specific value of the shard key as the dividing point.
<code>sh.splitAt("records.users", { x: 70 })</code>	Splits a chunk of the records.users collection at the shard key value x: 70
<code>sh.splitFind()</code>	Divides an existing chunk that contains a document matching a query into two approximately equal chunks.
<code>sh.splitFind("records.users", { x:70 })</code>	Splits, at the median point, a chunk that contains the shard key value x: 70.
<code>sh.status()</code>	Reports on the status of a sharded cluster, as <code>db.printShardingStatus()</code> .
<code>sh.waitForPingChange()</code>	Internal. Waits for a change in ping state from one of the <code>mongos</code> in the sharded cluster.
<code>refineCollectionShardKey</code>	Modifies the collection's shard key by adding new field(s) as a suffix to the existing key.
<code>db.adminCommand({ shardCollection: "test.orders", key: { customer_id: 1 } })</code>	Shard the orders collection in the test database. The operation uses the customer_id field as the initial shard key.
<code>db.getSiblingDB("test").orders.createIndex({ customer_id: 1, order_id: 1 })</code>	Create the index to support the new shard key if the index does not already exist.
<code>db.adminCommand({ refineCollectionShardKey: "test.orders", key: { customer_id: 1, order_id: 1 } })</code>	Run <code>refineCollectionShardKey</code> command to add the order_id field as a suffix

`convertShardKeyToHashed()`

```
use test
db.orders.createIndex( { _id:
"hashed" } )
sh.shardCollection( "test.orders", {
_id : "hashed" } )

{
  _id:
  ObjectId("5b2be413c06d924ab26ff9ca"),
  "item" : "Chocolates",
  "qty" : 25
}

convertShardKeyToHashed(
ObjectId("5b2be413c06d924ab26ff9ca")
)
```

Returns the hashed value for the input.

Consider a sharded collection that uses a hashed shard key.

If the following document exists in the collection, the hashed value of the `_id` field is used to distribute the document:

Determine the hashed value of `_id` field used to distribute the document across the shards,

Database Methods

`db.runCommand()`

Run a command against the current database

`db.adminCommand()`

Provides a helper to run specified database commands against the admin database.

User Management Commands

Make updates to users in the MongoDB Shell.

`db.auth()` Authenticates a user to a database.

`db.changeUserPassword()` Updates a user's password.

`db.createUser()` Creates a new user for the database on which the method is run.

`db.dropUser()`
`db.dropAllUsers()` Removes user/all users from the current database.

`db.getUser()`
`db.getUsers()` Returns information for a specified user/all users in the database.

`db.grantRolesToUser()` Grants a role and its privileges to a user.

`db.removeUser()` Removes the specified username from the database.

`db.revokeRolesFromUser()` Removes one or more roles from a user on the current database.

`db.updateUser()` Updates the user's profile on the database on which you run the method.

`passwordPrompt()` Prompts for the password in mongosh.

Role Management Commands

Make updates to roles in the MongoDB Shell.

`db.createRole()` Authenticates a user to a database.

`db.dropRole()`
`db.dropAllRoles()` Deletes a user-defined role/all user-defined roles associated with a database.

`db.getRole()`
`db.getRoles()` Returns information for the specified role/all the user-defined roles in a database.

`db.grantPrivilegesToRole()` Assigns privileges to a user-defined role.

`db.revokePrivilegesFromRole()` Removes the specified privileges from a user-defined role.

`db.grantRolesToRole()` Specifies roles from which a user-defined role inherits privileges.

`db.revokeRolesFromRole()` Removes inherited roles from a role.

`db.updateRole()` Updates a user-defined role.