


Spring Boot 3.3.2 and Java 21 Performance Benchmark

Web, Reactive,
CDS, AOT, Virtual Threads,
JVM, and Native



@ivanfranchin

Goal

We implemented two **Spring Boot** applications:

- one using **Spring Web** (**Spring MVC** with **Apache Tomcat**);
- the other using **Spring Reactive Web** (**Spring WebFlux** with **Netty**).

We built both **JVM** and **Native Docker images** for different configurations, including options for **virtual threads**, **CDS**, and **AOT** optimizations, and compared them.

Let's understand the strengths and trade-offs of each approach.

Java

Java has added many features over the years to improve performance:

- **Class Data Sharing (CDS):** Introduced in Java 5, it allows classes to be pre-processed into a shared archive file that can be memory-mapped at runtime, reducing startup time and dynamic memory usage when multiple JVMs share the same archive.
- **Application CDS (AppCDS):** Introduced in Java 10 (JEP-310), it extends CDS to include application classes in the shared archive.
- **Virtual Threads:** First previewed in Java 19 and fully supported in Java 21 (JEP-444), these are lightweight threads provided by the JDK instead of the OS.
- **Ahead of Time (AOT) Compilation:** Introduced in Java 9 (JEP-295), it improves startup time by compiling Java classes to native code before launching the JVM.

Spring Boot

Spring Boot has adapted to these Java improvements, making it easier to use new features like **CDS**, **virtual threads**, and **AOT** in Spring apps.

Spring Boot provides two main ways to build web apps:

- **Spring Web**: Uses traditional Spring MVC, a synchronous model with Apache Tomcat as the default web server.
- **Spring WebFlux**: For reactive, non-blocking apps, using Netty.

Docker

When your Java Spring Boot app is ready for production, building and deploying it as a **Docker** image is recommended.

There are two main types of Docker images for Java apps:

- **JVM Docker images:** Use the traditional Java setup with the JVM, offering fast compilation but slower startup and higher memory usage.
- **Native Docker images:** Compiled with GraalVM Native Image, they start faster and use less memory but are more complex to create and have a slower compilation process.

Which one should we choose?

- Should we go for a fully reactive app using Spring Reactive Web?
- If we're not comfortable with reactive programming, what about using Spring Web with Virtual Threads? Will it perform as well as a Spring Reactive Web app?
- No matter which approach we pick, can we make it better with CDS and AOT?
- Will these make it start faster and use less memory?
- Does a native app really use less memory than a JVM one?

Spring Boot Web and Reactive apps

We created two Spring Boot Greetings API applications:

- **spring-boot-greetings-api-web**
- **spring-boot-greetings-api-reactive**

They have a simple business logic that exposes an endpoint **/greetings?name=?**

The endpoint returns a greeting. If no name is provided, it will return "[greeting-word] World!". Otherwise, it will return "[greeting-word] [name]!".

The greeting word is chosen randomly by a service class, which can easily be replaced by a database call or an external API. **To simulate processing time**, we have added a **delay of 1 seconds** to the word selection.

Benchmark and Metrics

To collect important data from our Docker containers, we will use **ivangfr/api-oha-benchmark**. It uses:

- **Testcontainers**: to manage Docker containers.
- **OHA**: to load testing and to obtain metrics such as the slowest, fastest, and average request times in seconds;
- **docker stats**: to collect information like CPU and memory usage.

The benchmark will involve load testing both the JVM and Native Docker images for each configuration.

For each iteration:

1. Start the container
2. Conduct rounds of OHA testing with **100, 300, 900**, and finally **2700** concurrent requests.
3. Shut down the container.


Configuration and Container Name

Configuration	Apps name
JVM Greetings API Web	sb-web-jvm
JVM Greetings API Web with CDS enabled	sb-web-jvm-cds
JVM Greetings API Web with CDS and AOT enabled	sb-web-jvm-cds-aot
Native Greetings API Web	sb-web-native
JVM Greetings API Web with Virtual Threads enabled	sb-web-jvm-vt
JVM Greetings API Web with CDS and Virtual Threads enabled	sb-web-jvm-cds-vt
JVM Greetings API Web with CDS , AOT and Virtual Threads enabled	sb-web-jvm-cds-aot-vt
Native Greetings API Web with Virtual Threads enabled	sb-web-native-vt
JVM Greetings API Reactive Web	sb-reactive-jvm
JVM Greetings API Reactive Web with CDS enabled	sb-reactive-jvm-cds
JVM Greetings API Reactive Web with CDS and AOT enabled	sb-reactive-jvm-cds-aot
Native Greetings API Reactive Web	sb-reactive-native

@ivanfranchin


Benchmark Results

Average Time per Request




Benchmark Results

Startup Time




Benchmark Results

Maximum Memory Usage




Benchmark Results

Maximum CPU Usage



That's all



You can read more about Java Concurrency in the **Medium** article:

"Spring Boot 3.3.2 Benchmark: Web, Reactive, CDS, AOT, Virtual Threads, JVM, and Native"

Let's connect:  *@ivanfranchin*

Follow me:    *@ivangfr*